

1 Modular GPU Programming with Typed Perspectives

2 ANONYMOUS AUTHOR(S)

3 To achieve peak performance on modern GPUs, one must balance two frames of mind: issuing instructions to
4 individual threads to control their behavior, while simultaneously tracking the convergence of many threads
5 acting in concert to perform *collective operations* like Tensor Core instructions. The tension between these two
6 mindsets makes modular programming error prone. Functions that encapsulate collective operations, despite
7 being called per-thread, must be executed cooperatively by groups of threads. In this work, we introduce Prism,
8 a new GPU language that restores modularity while still giving programmers the low-level control over collective
9 operations necessary for high performance. Our core idea is *typed perspectives*, which materialize, at the type
10 level, the granularity at which the programmer is controlling the behavior of threads. We describe the design
11 of Prism, implement a compiler for it, and lay its theoretical foundations in a core calculus called Bundl. We
12 implement state-of-the-art GPU kernels in Prism and find that it offers programmers the safety guarantees
13 needed to confidently write modular code without sacrificing performance.

15 1 Introduction

16 CUDA [32] is a low-level, imperative programming language for NVIDIA GPUs. These GPUs are
17 organized into a hierarchy of compute resources. *Threads* are the basic unit of sequential execution,
18 *blocks* are groups of threads that can cooperate through shared scratchpad memory, and the *grid* is
19 the full collection of blocks launched for a computation. A GPU program executes in parallel across
20 this hierarchy, but is written *from the perspective of a single thread*.

21 While operations are specified per-thread, some are only valid when executed *collectively* by a
22 group of threads. The `__syncthreads()` intrinsic, for instance, synchronizes all threads within a
23 block, and it will cause a deadlock if executed by only some of those threads. There are many such
24 collective intrinsics, including Tensor Core instructions [34–36], warp shuffle operations [17], and
25 other kinds of synchronization primitives [16]. These operations require programmers to carefully
26 marshal compute resources to coordinate which threads execute which lines of code. As a result,
27 such collective operations break the illusion of threads executing independently.

28 The conceptual clash between regular statements (executed by a single thread) and collective
29 operations (executed by a group of threads) impacts not only hardware intrinsics, but also *user-defined*
30 collective operations, like functions. While functions are invoked individually by threads, they may
31 encapsulate collective behavior, creating a contradiction between the per-thread syntax of their
32 invocation and the cooperative semantics of their execution.

33 This contradiction puts modularity at odds with correctness and is apparent even in widely used
34 CUDA libraries that package common functionality via function interfaces [11, 31, 43]. Consider the
35 following snippet of documentation taken directly from CUB [31], a library of parallel primitives.
36 It describes the **BlockReduce** function [9], in which all threads in a block collaboratively apply a
37 reduction operation, such as a maximum or prefix sum, over an array:

38 Computes a block-wide reduction for `thread0` using the specified binary reduction functor.

- 40 • The return value is undefined in threads other than `thread0`.
- 41 • A subsequent `__syncthreads()` threadblock barrier should be invoked after calling this method if the collective's
42 temporary storage (e.g., `temp_storage`) is to be reused or repurposed.

43 This passage attempts to convey several assumptions implicit in **BlockReduce**'s implementation.
44 First, the function is only well defined when invoked by all threads in a block. Second, because it
45 accesses memory shared among threads, any subsequent reuse of that storage requires synchroniza-
46 tion to ensure all threads have completed their accesses. In effect, CUB attempts to retrofit CUDA
47 with information about the compute and memory requirements of collective operations. By prefixing
48 functions with identifiers such as **Block**, CUB creates ad-hoc "namespaces" for different functions

50 that assume similar invariants. However, without a type system to statically enforce these invariants,
 51 correct usage of this function—and collective operations in general—depends on carefully reading
 52 and interpreting the documentation.

53 In this work, we ask: can we provide low-level control over collective operations while statically
 54 guaranteeing that they execute with the necessary compute resources? By reifying configurations
 55 of compute resources as type-level *perspectives*—so named because they describe the view of GPU
 56 resources against which each statement is written—we find that we can. Our key insight is that GPU
 57 programmers naturally map computations onto different compute resources, and, unlike CUDA,
 58 which obscures this mapping, we can track it with a type system. This tracking enforces that collective
 59 operations are executed with the necessary resources, while still allowing users to access low-level
 60 intrinsics. Further, it allows users to define and compose their own collective operations by specifying
 61 the perspective with which their code must be run.

62 Our approach departs from previous work, which attempts to resolve the mismatch between
 63 per-thread syntax and collective execution by restricting access to low-level operations. Tile-based
 64 languages like Triton [46], Helion [37], and Tilos [21], limit the user’s perspective to the block
 65 level, which prevents them from writing the highest-performance kernels. A variety of functional
 66 languages, [7, 22, 26] provide compile-time guarantees through their type systems but do not expose
 67 low-level control over hardware. Some other efforts, like Descend [29], aim to provide a low-level,
 68 memory-safe GPU systems programming language, but lack support for modern GPU features (like
 69 Tensor Cores or asynchrony) entirely. As a result, despite its flaws, CUDA remains the de-facto
 70 standard for writing high-performance kernels on modern GPUs [18].

71 We introduce Prism, a new low-level GPU programming language that enforces that operations are
 72 only executed with the correct view of hardware resources, enabling fearless composition. Inspired by
 73 dependency calculi [1], Prism statically tracks the configurations of compute and memory resources
 74 with type-level perspectives. We also develop Bundl, a core calculus underpinning Prism, provide
 75 formal rules for manipulating the perspectives of both code and data, and prove *type-and-perspective*
 76 *safety*. This ensures that Bundl is sound, and that code always has the right perspective to execute
 77 operations at run time. A parallel goal, alongside safe composition, is performance. We incorporate
 78 modern GPU features such as Tensor Cores and asynchronous data movement into Prism, and
 79 demonstrate that Prism can achieve the same performance as hand-written, highly optimized code
 80 on an H100 and a 4070 Ti Super. Our contributions are:

- 81 (1) Prism, a low-level GPU language that tracks the logical grouping of compute and memory
 resources with type-level perspectives, empowering users to write modular code (Section 3);
- 82 (2) Bundl, a formal model of Prism that tracks perspectives in its type system and operational se-
 mantics, along with a soundness theorem guaranteeing that programs execute operations only
 when they have been statically proven to possess the necessary perspective (Section 4); and
- 83 (3) An implementation of Prism (Section 5) that demonstrates that it can support modern GPU
 features like Tensor Cores and asynchronous data movement, achieving performance com-
 parable to hand-optimized CUDA implementations (Section 6).

90 Section 7 discusses related work, and Section 8 discusses the limitations of our approach and
 91 outlines future work.

93 2 Background & Motivation

94 Before diving into the design of Prism, we begin with an overview of the GPU’s compute and memory
 95 hierarchies and outline the challenges posed by reasoning about them collectively. We also discuss
 96 how Prism can help solve these problems.

```

99  1 int tid = threadIdx.x;
100 2 if (tid >= 0 && tid < 32){
101 3   float A[4], B[2];
102 4   float C[4] = { 0 };
103 5   // Populate A with unique values.
104 6   for (int i = 0; i < 4; i++)
105 7     A[i] = tid * 4 + i;
106 8   // Populate B with unique values.
107 9   for (int i = 0; i < 2; i++)
108 10    B[i] = tid * 4 + i;
109 11   // Issue a warp-level Tensor Core
110 12   // operation: D = A * B + C
111 13   // (eliding some typecasts).
112 14   asm("mma.sync.aligned.m16n8k8..."
```

Fig. 1. Warp-level Tensor Core instruction in CUDA.

```

1 tid : int @ thread[1] = id();
2 with group(thread[32]):
3   A : float[4] @ thread[1]
4   B : float[2] @ thread[1]
5   C : float[4] @ thread[1]
6
7   # Populate A with unique values
8   for i in range(0, 4, 1):
9     A[i] = tid * 4 + 1
10
11   # Populate B with unique values
12   for i in range(0, 2, 1):
13     B[i] = tid * 4 + 1
14
15   # Issue Tensor Core op.
16   intrinsic.mma(
17     A[0], A[1], A[2], A[3],
18     B[0], B[1],
19     C[0], C[1], C[2], C[3],
20     out=[A[0], A[1], A[2], A[3]])
```

Fig. 2. Warp-level Tensor Core instruction in Prism.

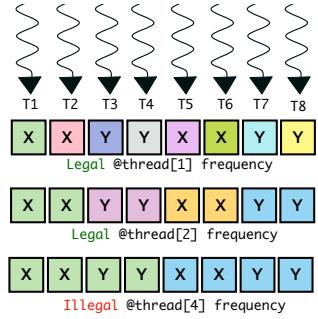


Fig. 3. The `@` syntax represents the frequency at which a variable varies across threads T1–T8. Each `@` denotes a coloring of the variable; threads that “view” the variable with the same color must observe the same value.

2.1 Compute Hierarchy

In CUDA, programmers launch computations that run on thousands of threads. These threads are organized into a *compute hierarchy* that defines how work is distributed and scheduled on the GPU. At the top of this hierarchy is the *grid*, representing all threads launched as part of a single computation. The grid is divided into *blocks*, each containing a user-specified number of threads. *Threads* are the finest unit of execution and the machine’s basic unit of sequential control.

Users define the behavior of a computation by writing a single program which is replicated across all threads. This program controls the behavior of individual threads by reading built-in identifiers like `threadIdx.x`—to determine a thread’s position within a block—and `blockIdx.x`—to determine the block’s position within the grid—at runtime. By making control flow decisions based on these two variables, users can assign each thread to its share of the full computation.

A natural and tempting way to interpret these built-in variables is to think of them as the indices of implicit “parallel-for loops” surrounding the program, where each iteration executes simultaneously. While this view is sufficient to understand CUDA’s programming model when threads do independent work, it quickly breaks down in the presence of collective operations.

Unlike most instructions, which are executed by a single thread, collective operations must be executed collaboratively by a group of threads. For example, on line 14 of Figure 1 we perform a Tensor Core operation [36], which is only meaningful when invoked by a collection of 32 threads—a *warp*—acting together. For the call, each thread sets up its portion of the operands, and the operation is performed *once* for the entire warp, with the results scattered across participating threads. In this example, the first 32 threads in a block are tasked with this operation; programmers mentally group these 32 threads into a collective unit, and then reflect that grouping in the program via the `if`-statement on line 2. If the condition on line 2 were instead `tid >= 0 && tid < 30`, making fewer than 32 threads reach line 14, the result would be undefined. To make matters worse, the restriction is not only on the number of threads executing the operation, but also on their alignment. For threads to form a warp, the starting ID of the group must be aligned to a multiple of 32. So, if the condition on line 2 were instead `tid >= 1 && tid < 33`, the result would still be invalid, even though 32 threads would execute it.

Collective operations make reasoning about CUDA programs challenging because they force programmers to track the *convergence behavior* of threads. Specifically, programmers must reason both about how many threads reach each in the program and how those threads are arranged, accounting for alignment. This difficulty is further amplified by two factors. First, programs may have multiple points of convergence, requiring programmers to mentally track the relative ID of a thread within

148 a logical group as that group evolves over the course of a program’s execution. Second, threads may
 149 be participating in multiple levels of convergence within the same program. In our example, we
 150 only considered the warp-level Tensor Core operation, but there are other collective operations that
 151 require convergence at different granularities like *warpgroup*-level Tensor Core operations—which
 152 must be issued collectively by four warps—or the block-level `__syncthreads()` primitive—which
 153 must be executed by all threads within a block.

154 Reasoning about collective operations is already error prone within the context of a single function,
 155 but becomes even more difficult when reasoning across functions. A callee may assume a certain
 156 configuration of threads and may structure its computation around that assumption. However, such
 157 assumptions are not visible in the callee’s function interface, which only exposes the input and
 158 output types. To invoke the function correctly, users must read its documentation—or worse, its
 159 implementation—to understand its assumptions, breaking modular reasoning.

160 In Prism, we materialize the programmer’s mental grouping of the compute hierarchy explicitly
 161 in the program’s source. Consider the example in Figure 2, which shows an equivalent rewrite of the
 162 CUDA program in Figure 1 using Prism. In Figure 2, the call to the `mma` intrinsic is valid only because
 163 it is executed within a **group** of 32 threads, made explicit on line 2. Since Prism’s **group** construct
 164 enforces both the size and alignment of participating threads, the validity of the `mma` operation is
 165 guaranteed by construction.

166 However, it is not sufficient to only consider the compute hierarchy when reasoning about con-
 167 vergence; we must consider the memory hierarchy as well. Whenever threads see different data,
 168 they can use it to induce divergent behavior. We’ve already seen an example of this: the `tid` variable
 169 in Figure 1. In general, the possibility of divergence lurks anywhere threads branch on data.

170 2.2 Data and the Memory Hierarchy

171 Data on the GPU is organized into a *memory hierarchy*, mirroring the compute hierarchy. All threads
 172 in a grid can read and write from *global memory*, where operands reside at the start of a computation
 173 and where results are eventually written. Each block has access to a limited amount of *shared memory*,
 174 a programmer-managed scratchpad typically used to stage repeatedly-used data. Finally, each thread
 175 maintains its own state in *registers* and *local memory*. These are private to each thread, while shared
 176 and global memory are accessible by multiple threads.

177 *Thread-Local Data.* Registers and local memory are owned by individual threads, so variables with
 178 the same name can have different values on different threads at run time. In general, CUDA offers
 179 no support for dealing with the divergence that arises as a consequence.

180 A key challenge for Prism is managing this divergence, ensuring that all threads organized by
 181 other language constructs like **group** remain logically unified, even in the face of data-dependent
 182 control flow. To do this, Prism tracks the *frequency* with which values vary in space. The `@` syntax
 183 attached to each declaration denotes this frequency. For example, `thread[1]` variables can vary for
 184 every thread, while `block[1]` variables can vary for every block, but not for threads within that block.
 185 Intuitively, the `@` construct “colors” a variable across threads, and any threads that share a color are
 186 required to agree on that variable’s value. Figure 3 illustrates valid and invalid colorings of a variable.

187 Using this information, Prism enforces rules for reading from and writing to variables. So, for
 188 example, if we introduce a condition based on `tid` within the **group**’s scope in Figure 2, Prism will
 189 reject that program; code that diverges at low frequency cannot read from variables that change at
 190 higher frequency. We explain the rules of Prism programs in detail in Section 3.

191 *Thread-Shared Data.* Shared and global memory, on the other hand, is not replicated per-thread;
 192 instead, all threads in a block have the same view into shared memory, while all threads on the grid
 193 have the same view into global memory.

CUDA does not explicitly model shared and global memory spaces, nor does it track whether allocations in a given memory space exceed device limits. Prism, on the other hand, distinguishes these spaces and restricts shared memory to static allocations, throwing an error if an allocation exceeds device limits. We will discuss these aspects in detail in Section 3.4. For now, however, we focus on how views of memory converge and diverge in tandem with the compute hierarchy.

As an example, when launching a kernel on the GPU, a global memory pointer initially belongs to the grid because every thread sees the same pointer value at the start of the computation. To write to that memory, each thread computes an offset from the original base address. In doing so, the pointer effectively diverges across threads so each can write independently. After these writes, the view on the memory must reconverge, ensuring that all threads have completed their updates before it can return to its original logical owner.

Let us reconsider the Tensor Core example from Figure 1, this time initializing the operands of the Tensor Core operation from pointers into global memory. In Figure 4, **A** and **B** are now populated from **a_mem** and **b_mem**. After the Tensor Core operation completes, the result is written back to **c_mem**, also in global memory. This variable **c_mem** is logically accessed from two different levels of the compute hierarchy. At the beginning, each thread in a group of 32 sees the same value for **c_mem**. Then, they each locate an offset within **c_mem** that they write to (lines 20–23). To restore **c_mem** back to its 32 thread-level ownership, all 32 threads must synchronize (line 25) to ensure all per-thread writes have completed. This is similar to the requirement we saw documented in CUB in Section 1.

Similarly to the compute hierarchy, CUDA programmers are responsible for tracking the logical owner of a view of memory as it evolves over the course of a program, and for ensuring that appropriate synchronization occurs.

Prism, by contrast, makes the evolution of memory’s view explicit in the syntax and automatically inserts the synchronization required to restore that view to its original owner. In Figure 5, we show how memory is lowered through the compute hierarchy for the same Tensor Core operation from Figure 4. The lowering of **c_mem** is required in this program as Prism only permits writes from the view of a single thread. To lower **c_mem**, we use Prism’s **partition** operation on line 28. The operation takes a pointer to partition, **c_mem**, and lowers it to a single thread, assigning it a new name, **c_thrd**; Prism will not allow use of the old variable **c_mem** within the **partition**’s scope. The **partition** also takes an indexing function, and each time **c_thrd** is accessed, this function is implicitly applied. Once the **partition**’s scope ends, Prism inserts a synchronization barrier before the next use of the original variable, so that all per-thread writes have completed. In this way, at the end of the **partition**’s scope, the original variable represents a convergent view of the data once again.

```

1 int x = threadIdx.x;
2 if (x >= 0 && x < 32){
3     float A[4], B[2];
4     float C[4] = { 0 };
5
6     // k is a_mem's stride.
7     A[0] = a_mem[(x/4)*k+(x%4)];
8     A[1] = a_mem[(x/4)*k+(x%4)+1];
9     A[2] = a_mem[(x/4+8)*k+(x%4)];
10    A[3] = a_mem[(x/4+8)*k+(x%4)+1];
11
12    // n is b_mem's stride.
13    B[0] = b_mem[(x%4)*n+(x/4)];
14    B[1] = b_mem[(x%4+4)*n+(x/4)];
15
16    asm("mma.sync...");
17
18    // Write back into c_mem.
19    // n is c_mem's stride.
20    c_mem[(x/4)*n+2*(x%4)] = C[0];
21    c_mem[(x/4)*n+2*(x%4)+1] = C[1];
22    c_mem[(x/4+8)*n+2*(x%4)] = C[2];
23    c_mem[(x/4+8)*n+2*(x%4)+1] = C[3];
24
25    __syncwarp();
}
```

Fig. 4. Invoking a warp-level Tensor Core instruction in CUDA after loading data from memory.

```

1 @cuda("device")
2 #requires(thread[32])
3 def simple_mma(
4     a: ptr(const(float)) @ thread[32],
5     b: ptr(const(float)) @ thread[32],
6     c: ptr(float) @ thread[32],
7     x: int @ thread[] - id()
8     with group(thread[32]):
9         A : float[4] @ thread[1]
10        B : float[2] @ thread[1]
11        C : float[4] @ thread[1]
12        # Reads do not need to be lowered.
13        A[0] = a_mem[(x/4)*k+(x%4)];
14        A[1] = a_mem[(x/4)*k+(x%4)+1];
15        A[2] = a_mem[(x/4+8)*k+(x%4)];
16        A[3] = a_mem[(x/4+8)*k+(x%4)+1];
17        B[0] = b_mem[(x%4)*n+(x/4)];
18        B[1] = b_mem[(x%4+4)*n+(x/4)];
19        # Skipping initialize C to ...
20        intrinsic.mma(
21            A[0], A[1], A[2], A[3],
22            B[0], B[1],
23            C[0], C[1], C[2], C[3],
24            out=[C[0], C[1], C[2], C[3]])
25        # Must be at thread[1] to write.
26        idx = \
27            lambda ro, co: (x/4+ro)*n+2*(x%4)+co
28        with partition(c_mem, p=thread[32], f=idx)
29            as c_thrd:
30                c_thrd[0, 0] = C[0]
31                c_thrd[0, 1] = C[1]
32                c_thrd[8, 0] = C[2]
33                c_thrd[8, 1]] = C[3]
34
35        # --- Sync point ---
```

Fig. 5. Invoking a warp-level Tensor Core instruction in Prism after loading data from memory.

```

246 1  %%cud("global")
247 2  # Top-level perspective bounds and shared memory usage.
248 3  #requires(grid[1], block[1], thread[32], smem=1280)
249 4  def mmaTF32NaiveKernel(A: ptr(const(float)) @ grid[1],
250 5   B: ptr(const(float)) @ grid[1],
251 6   C: ptr(float) @ grid[1],
252 7   M: int @ grid[1],
253 8   N: int @ grid[1],
254 9   K: int @ grid[1]):
255 10
256 11  # Starts out with grid[1] perspective.
257 12  with group(grid[1]):
258 13  # @ grid[1] inferred from current perspective
259 14  # Each block computes an 16 x 8 tile
260 15  num_blocks_n : const(int) = (N + 8 - 1) / 8
261 16
262 17  # id() function returns the block id
263 18  # inferred from @ block[1].
264 19  blk_row : const(int) @ block[1] = (id() / num_blocks_n) * 16
265 20  blk_col : const(int) @ block[1] = (id() % num_blocks_n) * 8
266 21
267 22  # Give each block an offset into C
268 23  offset = lambda x: blk_row * N + blk_col + x
269 24  with partition(C, p=block[1], f=offset) as Cs_blk:
270 25  with group(block[1]):
271 26  # SHMEM declarations are only allowed
272 27  # with a block[1] perspective.
273 28  A_smem : shared(float[16 * 8]) @ block[1]
274 29  B_smem : shared(float[8 * 8]) @ block[1]
275 30  C_smem : shared(float[16 * 8]) @ block[1]
276 31  # Now, id() returns the thread id
277 32  idx = lambda x: x * 4
278 33  # To write to C_smem, drop to thread[1] perspective
279 34  with partition(C_smem, p=thread[1], f=idx) as C_th:
280 35  for i in range(0, 4, 1):
281 36  with group(thread[1]):
282 37  C_th[i] = 0

1 1  for i in range(0, K_tiles, 1):
2 2
3 3  # --- Sync point --- (backedge from for loop)
4 4  for j in range(0, 4, 1):
5 5  global_row : int @ thread[1] = blk_row + row
6 6  global_col : int @ thread[1] = i * 8 + col
7 7  with partition(A_smem, p=thread[1], f=..) as A_th:
8 8  with group(thread[1]):
9 9  A_thrd[0] = A[global_row * K + global_col]

10 10
11 11  # --- Sync point --- (backedge from for loop)
12 12  for j in range(0, 2, 1):
13 13  # Similar to write into A_smem ...
14 14  with partition(B_smem, p=thread[1], f=..) as B_th:
15 15  with group(thread[1]):
16 16  B_th[0] = B[global_row * b * N + global_col_b]

17 17
18 18  # Give each warp an offset into C_smem.
19 19  # --- Sync point --- (backedge from for loop)
20 20  with claim(C_smem, p=thread[32]) as Cs_warp:
21 21  match split(thread): # Masks off other threads
22 22  case 32:
23 23  # Call function that performs a
24 24  # Tensor Core instruction.
25 25  simple_mma(A_smem, B_smem, Cs_warp)

26 26
27 27  # Write back final result to C.
28 28  for j in range(0, 4, 1):
29 29  flat_idx_c : int @ thread[1] = id() * 4 + j
30 30  row_c : int @ thread[1] = flat_idx_c / MMA_K
31 31  col_c : int @ thread[1] = flat_idx_c % MMA_K
32 32  idx = lambda x: row_c * N + col_c + x
33 33  with partition(C_blk, p=thread[1], f=idx) as C_th:
34 34  with group(thread[1]):
35 35  C_th[0] = C_smem[row_c * MMA_N + col_c]

36 36
37 37  return

```

Fig. 6. Tensor Float 32 matrix Multiplication in Prism (full program can be found in).

3 The Prism Language

Prism is an imperative, low-level language designed at a level of abstraction comparable to that of CUDA. Unlike CUDA, however, Prism’s syntax materializes the mapping of computations onto the compute and memory hierarchy explicitly in the program source. Using this information, Prism enforces that programs only execute collective operations with sufficient resources.

To guide our discussion about Prism’s design, we use the program in Figure 6 as a running example. It computes a matrix multiplication between two float arrays, **A** and **B**, to produce an output matrix **C**. In this program, each block computes an independent 16×8 tile of the output. To do so, blocks first locate the tile index assigned to them (line 19-20). Next, threads in each block load corresponding rows and columns from **A** and **B** (line 41-53) into shared memory. Finally, the program invokes a warp-level, Tensor Core instruction to compute the output (line 62), requiring threads in a warp to converge. We encapsulate this Tensor Core instruction in a function, demonstrating how function composition works in Prism. This function is the same as the one in Figure 5.

3.1 Levels

Prism models the machine's compute hierarchy through *levels*. There are three levels in Prism—**grid**, **block** and **thread**—which are organized as expected: a **grid** consists of multiple **blocks**, each of which consist of multiple **threads**. Levels are ordered, with **thread < block < grid**.

There are two key differences between Prism’s levels and those in CUDA. First, Prism does not model CUDA’s three-dimensional grid or block structure. Second, there are two commonly-used “levels”, warp and warpgroup, absent from our hierarchy. On the hardware, the units of each level are arranged in a linear order, and the three-dimensional structures of CUDA are simply interpretations of this ordering, not distinct hardware resources. Similarly, warps and warpgroups are organizational constructs defined in terms of existing levels. Namely, a warp is a group of 32 threads whose first

295 thread ID is aligned to 32. A warpgroup, which was introduced with the release of the Hopper
 296 architecture [8], consists of 4 consecutive warps.

297 Rather than baking these interpretations into Prism by adding new dimensions and levels, we let
 298 users express multi-dimensional structures and define groupings of custom sizes.

299 300 3.2 Perspectives

301 Perspectives are the central concept of Prism, representing the view of the hierarchy from which
 302 the programmer is defining the machine's behavior. Perspectives allow Prism to determine which
 303 compute resources the programmer is controlling, whether they are available in the program's
 304 context, and if those resources are sufficient for a given operation.

305 A perspective is a level—**grid**, **block**, or **thread**—paired with a static constant **n**, specifying the
 306 number of units at that level. For example, **thread[2]** denotes a perspective of two threads, **block[4]**
 307 denotes a perspective of four blocks, and so on. Perspectives also carry *alignment information*: a
 308 perspective of size **n** is aligned to **n**. In this way, a warp is simply a desugaring of **thread[32]**, and a
 309 warpgroup is a desugaring of **thread[128]**.

310 Finally, perspectives are partially ordered. We say that **level₂[n₂]** is *broader* than **level₁[n₁]**, or
 311 that **level₁[n₁]** is *narrower* than **level₂[n₂]**, if either:

- 312 (1) **level₁** < **level₂**; or,
 313 (2) **level₁** = **level₂** and **n₁** divides **n₂**.

314 315 3.3 Perspectives on Code

316 Code is associated with a set of perspectives, called a *perspective bound*, which corresponds to the
 317 compute resources whose behavior it defines. At every point in the program, the perspective bound
 318 for that point indicates which layer of the hierarchy is being programmed, and how that layer can be
 319 destructed into narrower perspectives. For example, a line of code with a **{block[1], thread[4]}**
 320 perspective bound tells Prism that the current line of code is being programmed at the **block[1]**
 321 perspective and that the **block[1]** perspective can be destructed into a some number of **thread[4]**
 322 perspectives. As a shorthand, when we refer to a code's perspective, we mean the broadest perspective
 323 available in its perspective bound (in this example, **block[1]**).

324 Functions begin with a top-level perspective bound. In Figure 6, the perspective bound is defined
 325 on line 2, using the notation **@requires(grid[n₁], block[n₂], thread[n₃])**. Programmers can
 326 then shape the program's current perspective using two constructs: **group** and **split**.

327 **Group**. The **group** construct lets programmers shift from a broader perspective to some number
 328 of narrower perspectives contained in it. Operationally, the **group** construct does this by replicating
 329 code written from the narrower perspective across the broader one. In effect, **group** forks many
 330 copies of a narrower perspective.

331 Let's consider this in context of our example. In Figure 6, execution begins at **grid[1]** on line 12.
 332 At that point, a programmer controls the whole grid's behavior. To produce different output tiles, the
 333 programmer shifts their perspective to **block[1]** on line 25. The code within the **group(block[1])**
 334 defines the behavior of a single block, and is replicated across all blocks in the grid.

335 Not all uses of **group** are valid: the examples in Figure 7 have no
 336 meaning on the hardware. On line 4, the program tries to broaden
 337 its perspective to **block[1]** from **thread[2]**, which is illegal. On
 338 the other hand, in the second example, the program tries to narrow
 339 its perspective to **block[5]** from **block[6]**. While 5 is indeed less
 340 than 6, Prism cannot evenly replicate **block[5]** inside a **block[6]**
 341 perspective, and so rejects this program. Recall, from Section 3.2,

```

1 # Example 1
2 with group(thread[2]):
3     # Illegal because block > thread.
4     with group(block[1]):
5         pass
6 # Example 2
7 with group(block[6]):
8     # Illegal because 6 % 5 != 0.
9     with group([block[5]]):
10        pass
  
```

Fig. 7. Illegal uses of **group**.

344 that the condition on whether one perspective is narrower than
 345 another is a condition on divisibility, not just size.

346 To eliminate such cases, Prism only allows an invocation of `group(level[n])` if the current
 347 perspective bound contains a perspective broader than `level[n]`. Once `group(level[n])` is invoked,
 348 it modifies the current perspective bound in two ways. First, it removes all perspectives broader than
 349 `level[n]` from it. Second, it sets the broadest perspective within the `group` to be `level[n]`.

350 **`Split`**. Unlike `group`, which is used for replication into equally-sized, narrower perspectives,
 351 `split` is used for sharding the current perspective unequally. For example, Figure 8 shows a `split`
 352 from `thread[4]` into one branch with `thread[2]` perspective and two with `thread[1]`. The three
 353 arms of the `split` execute independently in parallel¹, as a form of unordered composition. When
 354 `split(level)` is invoked, the perspectives of its branches diverge. At the end of the `split`, they
 355 reconverge, and continue execution with the original perspective.
 356

357 Use of `split` is necessary to write warp-specialized [4] code, a pro-
 358 gramming pattern used in high-performance kernels. Another important
 359 use of `split`—masking off threads—can be found in our running example.
 360 Line 58 in Figure 6 shows a `split` that requests the first warp in the block,
 361 narrowing from `block[1]` into a single `thread[32]`. This warp is later
 362 used to execute a Tensor Core operation.

363 Because `split` corresponds to unordered composition, it must provide
 364 each of its branches their requested perspectives simultaneously. Prism
 365 thus checks that the sum of the perspectives requested by all branches
 366 of the `split` can be satisfied. For example, the program in Figure 9 does
 367 not type check. Finally, because perspectives enforce alignment, every
 368 branch of the `split` must also be aligned; not all `splits` whose sizes are
 369 at most the available units are valid. Figure 10 shows an example of this
 370 constraint being violated: the second branch of the split is not aligned
 371 to 2.

372 Once `split` is invoked, for each branch that requests `n` units, all per-
 373 spectives broader than `level[n]` are removed from its perspective bound,
 374 and the available units for `level` are set to exactly `n`.

375 3.4 Perspectives on Data

376 Section 3.2 described how programmers can control different layers of the hierarchy by changing
 377 their perspectives on code through `group` and `split`. To ensure these operations remain meaningful,
 378 Prism must ensure that threads inside a perspective remain logically grouped, even when they
 379 encounter `for`, `while`, and `if` statements. As we saw in Section 2.2, making this guarantee requires
 380 Prism to track how data varies across threads.
 381

382 **3.4.1 Thread-Local Data**. In Prism, each local variable has a perspective, which indicates the fre-
 383 quency at which its values change in space. This frequency remains constant for the duration of a
 384 program, and it tells Prism that a `level[n]` variable is always indistinguishable to threads within
 385 that perspective. For example, `blk_row @ block[1]` on line 19 in Figure 6 has the same value across
 386 all threads in a block.

387 Programmers specify the perspective that each variable lives at in its declaration. A variable `v` of
 388 type `int` is declared at `thread[1]` perspective using the syntax `v : int @ thread[1]`². To enforce

```
1 with group(thread[4]):
2   match split(thread):
3     case 2:
4       ...
5     case 1:
6       ...
7     case 1:
8       ...
```

Fig. 8. Example `split`.

```
1 with group(thread[4]):
2   match split(thread):
3     case 4:
4       ...
5     case 1:
6       ...
```

Fig. 9. Illegal split exceeds the available perspectives.

```
1 with group(thread[3]):
2   match split(thread):
3     case 1:
4       ...
5     case 2:
6       ...
```

Fig. 10. Illegal split violates alignment.

390 ¹Despite the use of the `match` syntax, *all* branches of the `split` execute.

391 ²If not explicitly annotated, Prism infers a variable's perspective to be the perspective of the code where it was declared.

393 this frequency invariant, Prism restricts reads from and writes to variables based on their perspectives.
 394 The rule can be summarized as follows: Prism allows programs to “read up” from broader perspectives
 395 and “write down” to narrower ones.

396 **Read Up.** Variables can only be read if their perspective is at
 397 least as broad as the current code perspective. Figure 11 gives an
 398 example of an illegal read that would violate this constraint. While
 399 `__syncthreads()` should always be safe in `block[1]`, branching on
 400 the variable `flag`—which may take different values across threads
 401 in a block—can cause only some of those threads to reach the
 402 `__syncthreads()`, violating its collective invariant.

```
1 flag : bool __thread[1] = ...
2 with group(block[1]):
3   if (flag)
4     __syncthreads();
```

Fig. 11. Illegal read of `thread[1]` variable.

403 **Write Down.** Writes are dually constrained. Only values that live at broader perspectives can
 404 be written into variables that live at narrower ones. For example, a `block[1]` variable can be used
 405 to write into a `thread[1]` variable, but not vice versa. An example of an illegal write is shown in
 406 Figure 12, where writing from a `thread[1]` variable into a `block[1]` one would cause a deadlock.

407 Together, the “read up” and “write down” rules ensure that in-
 408 formation only flows from broader perspectives to narrower ones.
 409 In Figure 6, we can see the “read up” rule in action on lines 42 and 43.
 410 Meanwhile, lines 19 and 20 are instances of the “write down” rule.

```
1 x : bool __thread[1] = ...
2 y : bool __block[1] = ...
3 with group(block[1]):
4   y = x
5   if(y):
6     __syncthreads();
```

Fig. 12. Illegal write of `thread[1]` variable into a `block[1]` variable.

411 3.4.2 *Thread-Shared Data.* Unlike thread-local data, which is liter-
 412 ally replicated across threads and is backed by distinct physical
 413 storage, thread-shared data consists of pointers into shared and
 414 global memory that are visible to some collection of threads. As a
 415 result, the perspective such data inhabits can evolve as the program executes.

416 In Prism, there are three mechanisms for obtaining thread-shared data. The first is to directly
 417 allocate data residing in shared memory³. The second and third are to obtain offsets into existing
 418 data by using the `partition` or `claim` operations. These constructs mirror `group` and `split`, and
 419 are used to temporarily narrow the perspective associated with a pointer.

420 Prism simplifies shared memory allocations by requiring them to have a static size. The `@requires`
 421 annotation specifies the amount of shared memory a function expects to have. Prism uses this
 422 information to ensure that a function’s allocations do not exceed its declared limit, and checks
 423 whether there is enough shared memory available at its call sites. The function in Figure 6 declares
 424 the amount of shared memory it will require on 2. We use standard techniques [28] to statically
 425 bound memory usage with Prism’s type system.

426 Additionally, since shared memory is only visible to threads within the same block, such allocations
 427 are only permitted at `block[1]`. An example of such an allocation is shown on lines 28-30 of Figure 6.
 428 During compilation, Prism will automatically handle the necessary pointer arithmetic to assign each
 429 allocation an appropriate offset within the shared memory space.

430 **Partition.** The `partition` operation plays the same role for memory that `group` plays for code.
 431 It is used to refine the perspective of a pointer by computing offsets for each narrowed perspective.
 432 Concretely, the `partition` operation takes a pointer variable `x`, an indexing function `f`, and produces
 433 a new variable `y` at a narrower perspective `level[n]` to be used inside the `partition`. After the

434

 435 ³Prism assumes that all global memory allocations have been made before launching the CUDA kernel, as is typically the case
 436 with CUDA programs.

partition ends, the threads participating in the **partition** must synchronize to restore the original pointer **x** to its perspective⁴.

Inside the scope of the partition, the original variable **x** cannot be used, and the pointer can only be accessed using **y**. Each use of **y** applies the indexing function **f** to compute the true offset of the access. For example, the use of **partition** on line 34 in Figure 6 takes a pointer **C_smem** at **block[1]**, gives it a new name **C_th**, and distributes it across **thread[1]** perspectives by transforming each occurrence of **C_th[i]** in the body into **C_smem[i*4]**.

At this point, it is necessary to consider how different indexing functions affect the possibility of data races. If the indexing function is injective, each narrower perspective receives a distinct offset into the underlying array, and the resulting partition is free of data races. When the indexing function is not injective, multiple threads may race on the same location, introducing a potential data race.

Preventing data races is not one of Prism's goals. In Prism, data races *can* occur within a **partition**; however, since the data is eventually synchronized before it is reused, the last writer wins. Out-of-bounds accesses are considered undefined behavior. Prior work, in particular Descend [29], describes a type system for data-race free GPU programs, and we believe a similar approach can be combined with Prism's. Prism instead focuses on the interaction between the compute and memory hierarchies and on reasoning about them simultaneously to ensure that operations are executed only with sufficient compute resources, a guarantee that Descend cannot provide.

Out-of-bounds accesses are undefined behavior.

Claim. The **claim** operation lets programmers narrow a pointer's perspective by giving it to only one collection of threads with that narrower perspective. For example, line 57 of Figure 6 shows an example of a **claim**, where the pointer **C_smem @ block[1]** is only available to a single warp in the block and narrowed to **Cs_warp @ thread[32]** to call a Tensor Core operation.

A **claim** takes the original pointer **x**, the target perspective to narrow it to, and a new name **y** to assign to the narrowed memory. In all **splits** within the **claim**, the new pointer **y** is accessible only within one branch; sibling branches are not permitted to read or write from this memory. As with **partition**, the original variable **x** is not accessible within the **claim**.

3.5 The **id()** Function

As opposed to exposing users to special hardware variables like **blockIdx.x**, and **threadIdx.x**, Prism provides an **id()** function instead. The **id** function returns the relative index of a narrow perspective within a broader perspective. That is, the interpretation of the **id** function depends on both the perspective of the variable it is being written to and the perspective of the code invoking it. In Figure 6, we use a call to **id** with **grid[1]** perspective on lines 19 and 20 to locate the tile that each block is in charge of computing. On line 32, however, a write of **id** into a **thread[1]** variable at **block[1]** perspective will return the relative ID of the thread within the block, not in the grid.

3.6 Collective Operations

There are two types of collective operations in Prism.

Function Calls. Each function carries a *perspective signature*, which consists of its top-level perspective bound, shared memory usage, and its arguments' perspectives. An example of such a signature can be seen in lines 3-9 in Figure 6. At call sites, Prism ensures that the callee's perspective signature can be satisfied by the caller. Since the perspective bound describes a minimum requirement, functions can be called with a broader code perspective than necessary, accounting for alignment. When checking

⁴Strictly speaking, this synchronization is only required at the next point the memory is to be *re-used*. Prism has a mechanism for optimizing the placement of these synchronization points, which is described in Section 5.

491 function arguments at call sites, Prism distinguishes between primitive data types and pointers. For
 492 primitive data types, like **int** or **bool**, arguments can have broader perspectives than specified in the
 493 function signature. On the other hand, pointers that live at broader perspectives can only be passed if
 494 the pointer is marked as **const**, indicating that it will only be used for reading. Otherwise, we require
 495 that the pointer's perspective exactly match that of the function's perspective signature.

496

497 *Intrinsics.* The compiler provides a pre-defined set of collective intrinsics –like the Tensor Core
 498 instruction discussed in Section 2.1—each of which declares its perspective signature. Programmers
 499 may also add to this set using **unsafe**. From within **unsafe** code, programmers can inline assembly
 500 instructions and wrap them in a Prism function, specifying its perspective signature. Prism checks
 501 these call sites like those of any other function.

502

503 3.7 Asynchrony

504 Asynchronous data movement works similarly to other memory operations. Users can mark storage as asynchronous with
 505 the **async** construct. As with **partition** and **claim**, **async** hides the old variable and exposes a fresh one that is only accessible
 506 within the **async** statement. Inside, Prism ensures the new variable is only used by async data-movement intrinsics.

507 Prism includes two such intrinsics: bulk [14] and non-bulk [15] asynchronous data-movement
 508 instructions. We show an example of the former in Figure 13. As with the data operations in Section 3.4,
 509 Prism inserts the necessary synchronization before the program’s next use of the original variable,
 510 ensuring that the asynchronous transfer has completed.

511

512 4 Formalization in Bundl

513 Having introduced the full Prism language, we now describe Bundl, a core calculus that formalizes
 514 its most fundamental aspects by statically tracking perspectives on code and data. We use Bundl to
 515 argue that well-typed Prism programs are not only type-safe, but will also never get stuck trying to
 516 execute operations for which they lack the correct perspective.

517 In this section, we describe Bundl’s type system and operational semantics—in particular how it
 518 manages compute and data perspectives—and build up to a formal proof of type-and-perspective
 519 safety. We instrument Bundl’s operational semantics with runtime perspective enforcement: its rules
 520 will get stuck if they encounter an operation for which they have the wrong perspective. This runtime
 521 enforcement means that our safety theorem can guarantee that dynamically-realized perspectives
 522 match the ones inferred by the type system.

523

524 4.1 Bundl Type System

525 The core idea in Bundl is to track, at the type-level, the program’s perspective on code and data. To
 526 achieve this, we borrow techniques from the literature on *dependency tracking* [1]; in particular, the
 527 code perspective is tracked on the typing judgment, which has the form $\Gamma \vdash^\pi e : \tau$ for expressions and
 528 $\Gamma \vdash^\pi s$ for statements. The π over the \vdash is the code perspective on e and s , and is comprised of a level
 529 and a size, the same structure as a perspective in Prism.

530 The typing context also tracks the perspective at which each variable lives; data can only be read
 531 or written from a variable when its perspective is compatible with that of the code interacting with it.
 532 This requirement is made manifest in the **T-VAR** rule, which can be found in Figure 14. Observe that
 533 the π in the variable rule must match exactly between data and code; principles like “read up” and
 534

```
1 with async(old_name) as new_name:
2   match split(thread):
3     case 1:
4       cp_async(smem, new_name, 16);
```

Fig. 13. Example of an **cp.async** instruction in Prism.

535

$\boxed{\Gamma \vdash^\pi e : \tau}$ $\frac{x : \pi \in \Gamma}{\Gamma \vdash^\pi x : \tau} \text{ T-VAR}$	$(Expression \ typing)$ $\frac{\Gamma \vdash^{\pi'} e_1 : \tau[] \quad \Gamma \vdash^\pi e_2 : \mathbf{int} \quad \pi \leq \pi'}{\Gamma \vdash^\pi e_1[e_2] : \tau} \text{ T-ARR-ACCESS}$
$\boxed{\Gamma \vdash^\pi s}$	$(Statement \ typing)$
$\frac{\Gamma \vdash^{(h,n_1)} s_1 \quad \Gamma \vdash^{(h,n_2)} s_2 \quad n_1, n_2 \text{ align to } n}{\Gamma \vdash^{(h,n)} \mathbf{split}(n_1, n_2)\{s_1\}\{s_2\}} \text{ T-SPLIT}$	$\text{where } n_1, n_2 \text{ align to } n ::= (n_1 + n_2 \leq n) \text{ and } (n_1 n) \text{ and } (n_2 n) \text{ and } (n_2 n_1 + n)$
$\frac{\Gamma \vdash^\pi s}{\Gamma \vdash^{q \cdot \pi} \mathbf{group} \ q \ s} \text{ T-GROUP}$	$\frac{\Gamma, y : \downarrow^\pi \tau[] \vdash^\pi s \quad l \neq \mathbf{Local}}{\Gamma, x : \pi \tau[] \vdash^\pi \mathbf{lower} \ x \ \mathbf{into} \ y \ \mathbf{in} \ s} \text{ T-LOWER}$
$\frac{\Gamma, y : (h, n/c) \tau[] \vdash^\pi s \quad c n \quad l \neq \mathbf{Local}}{\Gamma, x : (h, n) \tau[] \vdash^\pi \mathbf{partition} \ x \ \mathbf{into} \ y \ \mathbf{by} \ c \ \mathbf{in} \ s} \text{ T-PARTITION}$	$\frac{\Gamma \vdash^\pi s}{\Gamma \vdash^\pi \mathbf{destruct} \ \mathbf{in} \ s} \text{ T-DESTRUCT}$
$\frac{\Gamma, y : (h, n') \tau[] \vdash^\pi s \quad n' \leq n \quad l \neq \mathbf{Local}}{\Gamma, x : (h, n) \tau[] \vdash^\pi \mathbf{claim} \ x \ \mathbf{into} \ y \ \mathbf{at} \ n' \ \mathbf{in} \ s} \text{ T-CLAIM}$	

Fig. 14. Core typing rules of Bundl. The typing rules presented here are a simplified selection of the full rules, which can be found in Appendix A.2.

“write down” are instead encoded directly in the rules for reading and writing, like **T-ARR-ACCESS**, which views the array being read with broader perspective than the current code perspective.

Figure 14 also shows other key rules, and these rules fall into two main categories: rules for managing the code perspective and rules for managing data perspective.

4.1.1 Managing Perspectives on Code. In Prism, there are two operations which manage code perspectives: **group** and **split**. In Bundl, to better model the details of how perspectives shift throughout programs, we introduce a third construct, **destruct**.

Bundl’s **group** statement directly corresponds to Prism’s, and is checked by the **T-GROUP** rule. Given some perspective that can be divided into q equally sized narrower π s, the statement **group** qs will check with perspective $q \cdot \pi$ provided that s itself checks with perspective π . The starting perspective is of the form $q \cdot \pi$, and encodes Prism’s alignment requirement.

The **split** statement, meanwhile, is checked by the **T-SPLIT** rule, and functions like a binary version of the n-ary **split** construct in Prism. It enforces the same divisibility requirements to ensure that the perspectives on code and data remain properly aligned, and then checks the two sub-statements s_1 and s_2 with the divided, narrower perspectives.

The final rule for managing code perspective is **T-DESTRUCT**, which shifts perspective down the GPU hierarchy and makes explicit in Bundl programs exactly where such shifts occur. The \downarrow operation on π s “destructs” the perspective into many narrower perspectives at a lower level; $\downarrow(\mathbf{Grid}, 1) = (\mathbf{Block}, B)$ and $\downarrow(\mathbf{Block}, 1) = (\mathbf{Thread}, T)$, where B and T are parameters to a particular instantiation of Bundl to describe the number of blocks per grid and threads per block.⁵ Because \downarrow

⁵Bundl is abstracted over these B and T values, so instead of tracking perspective bounds the way that Prism does, it only tracks the top-level perspective described in Section 3.2.

589 $\frac{L(t), S(b), \Sigma, t, b, 0 \vdash^{(\text{Grid},1)} s \rightsquigarrow s' \dashv \eta', \sigma', \Sigma' \quad P(t, b) = s}{L, S, \Sigma, P \rightsquigarrow L[t \mapsto \eta'], S[b \mapsto \sigma'], \Sigma', P[(t, b) \mapsto s']} \text{ S-PROGRAM}$

590
591
592
593 $\frac{p < n_1 \quad n_1, n_2 \text{ align to } n \quad \eta, \sigma, \Sigma, t, b, p \vdash^{(h, n_1)} s_1 \rightsquigarrow s'_1 \dashv \eta', \sigma', \Sigma'}{\eta, \sigma, \Sigma, t, b, p \vdash^{(h, n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\} \rightsquigarrow \text{split}(n_1, n_2)\{s'_1\}\{s_2\} \dashv \eta', \sigma', \Sigma'} \text{ S-SPLIT-LEFT}$

594
595
596 $\frac{p \geq n_1 \quad p < n_1 + n_2 \quad n_1, n_2 \text{ align to } n \quad \eta, \sigma, \Sigma, t, b, p - n_1, \vdash^{(h, n_2)} s_2 \rightsquigarrow s'_2 \dashv \eta', \sigma', \Sigma'}{\eta, \sigma, \Sigma, t, b, p \vdash^{(h, n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\} \rightsquigarrow \text{split}(n_1, n_2)\{s_1\}\{s'_2\} \dashv \eta', \sigma', \Sigma'} \text{ S-SPLIT-RIGHT}$

597
598
599
600 $\frac{\eta, \sigma, \Sigma, t, b, p \text{ mod } n \vdash^{(h, n)} s \rightsquigarrow s' \dashv \eta', \sigma', \Sigma'}{\eta, \sigma, \Sigma, t, b, p \vdash^{(h, q \cdot n)} \text{group } q \ s \rightsquigarrow \text{group } q \ s' \dashv \eta', \sigma', \Sigma'} \text{ S-GROUP}$

601
602
603 $\frac{\eta, \sigma, \Sigma, t, b, t \text{ mod } T \vdash^{(\text{Thread}, T)} s \rightsquigarrow s' \dashv \eta', \sigma', \Sigma'}{\eta, \sigma, \Sigma, t, b, 0 \vdash^{(\text{Block}, 1)} \text{destruct in } s \rightsquigarrow \text{destruct in } s' \dashv \eta', \sigma', \Sigma'} \text{ S-DESTRUCT-BLOCK}$

604
605
606
607 $\frac{}{\eta, \sigma, \Sigma, t, b, p \vdash^{(\text{Block}, 1)} x := \text{alloc Shared } \tau \ n \ \text{in } s \rightsquigarrow s \dashv \eta, \sigma[x \mapsto^{\pi} \langle x, n \rangle], \Sigma} \text{ S-ALLOC-SHARED}$

608

609 Fig. 15. Core semantic rules of Bundl. As with the typing rules, we present only a simplified selection of the full
610 rules, which can be found in Appendix A.3.

611
612
613
614
615 is only defined on **(Grid,1)** and **(Block,1)**, the rule enforces that one can only **destruct** their code
616 perspective with exactly one grid or block.

617
618 4.1.2 *Managing Perspectives on Data*. The mechanism for managing data perspective mirrors that
619 of code perspective, with each operation for data corresponding to an operation for code.

620 The **partition** operation is analogous to **group-ing** a code perspective. The typing rule for this
621 operation, **T-PARTITION**, requires that the data perspective on the variable to be partitioned, x , is the
622 same as the current perspective on code. After partitioning, a fresh variable y is introduced with
623 a new perspective π/c , which we use as a shorthand to denote a division of a perspective π into c
624 equally-sized narrower perspectives. Within the body of the **partition**, we disallow any references
625 to the original variable x and continue checking with the original code perspective π ; the **partition**
626 has no effect on the current code perspective.

627 Unlike **partition**, which divides up a piece of data equally among narrower perspectives, the
628 **claim** operation views the claimed data with exactly one narrower perspective. Accordingly, Bundl
629 needs to ensure that only one branch of a **split** operation, with the appropriate π , can refer to
630 the claimed variable. To ensure that this is the case, the **T-CLAIM** rule links the data perspective of
631 the variable to the compute perspective of the code claiming it by changing *both* at the same time.
632 This represents a minor difference from Prism, which uses additional static analysis to ensure that a
633 claimed variable is only accessed in a single **split** branch.

634 Lastly, the **T-LOWER** rule mirrors the **T-DESTRUCT** rule; it uses the \downarrow operator to move a variable
635 from one level of the hierarchy to another, distributing it equally among all the narrower perspectives
636 at that level in the same manner as **T-PARTITION**.

638 4.2 Bundl Semantics

639 Having explained the key rules of the type system, we can move on to discuss Bundl’s operational
 640 semantics. To reflect the fact that GPU programs execute in parallel across numerous threads, we
 641 model the semantics of Bundl in the style of Turon et al. [49], using a two-level small step judgment.
 642 We present the key rules of this semantics in Figure 15.

643 The top level (i.e., device-level) judgment has just one rule: **S-PROGRAM**. This rule acts as a “frame”
 644 for the lower level (i.e., thread-level) judgment, and steps a collection of thread-ID-indexed local
 645 memories (L), a collection of block-ID-indexed shared memories (S), a global memory (Σ), and a *thread*
 646 *pool* to an updated collection of memories and thread pool. The thread pool maps thread and block IDs
 647 to code, intuitively representing the program being executed by each thread at the current moment.
 648 The **S-PROGRAM** non-deterministically chooses a thread ID and block ID and steps it according to
 649 the thread-level judgment. This allows the semantics to model the full range of non-deterministic
 650 behavior arising from the GPU’s thread scheduler.⁶

651 The thread-level judgment has the shape $\eta, \sigma, \Sigma, t, b, p \vdash^\pi s \rightsquigarrow s' \dashv \eta', \sigma', \Sigma'$, where

- 653 • η denotes thread-local memory,
- 654 • σ denotes shared memory,
- 655 • Σ denotes global memory,
- 656 • t denotes the thread’s ID,
- 657 • b denotes the ID of the block in which the thread lives, and
- 658 • p denotes the relative position of the thread within π (the perspective ID).

659 Critically, notice that a π also appears on the thread-level judgment just as it does on the typing
 660 judgment. This is because the thread semantics *dynamically tracks and enforces* perspectives. The
 661 same way evaluation of a program “gets stuck” if a value does not have the right type, the semantics of
 662 Bundl also get stuck if code attempts to access data or invoke commands with the wrong perspective.
 663 As an example, observe the **S-ALLOC-SHARED** rule in Figure 15, which requires a **(Block,1)** perspective
 664 and will fail to step if encountered with a different one. This runtime perspective is present in Bundl,
 665 but is erased by Prism during compilation; we will later use it to prove that well-typed programs will
 666 always execute with the same perspective that the type system viewed them with.

667 The semantic rules for perspectives involve manipulating p to track which threads take which code
 668 paths when perspectives are **split** or **grouped**. Notice that in **S-PROGRAM**, the thread-level judgment
 669 always begins with perspective **(Grid,1)**: all the perspective management rules are congruences,
 670 narrowing the perspective of further evaluation as determined by the particular rule used.

671 These rules take great care to ensure that p always describes the relative position of a compute
 672 resource within its perspective; the payoff is that Bundl’s semantics can later use this p value
 673 to model the way that Prism automatically adjusts indices into data when partitioning a data
 674 perspective. Beyond these key rules for managing perspective on code, we have modeled many of
 675 the other features of Prism, such as asynchronous operations and thread synchronization, in Bundl.
 676 To handle features like these we equip the operational semantics with additional structure, including
 677 sets of semaphores [20] for synchronization and a stack of effect handlers for modeling deferred
 678 asynchronous computations inspired by Ahman and Pretnar [3]. We have elided these details here
 679 for simplicity, but interested readers can see further details in Appendix A.3.

680
 681
 682
 683
 684 ⁶In reality, the GPU’s thread scheduler guarantees that the threads of a warp execute in lockstep, but modeling every thread
 685 as completely independent is both simpler and a conservative overestimate of the nondeterministic behavior of the GPU.

687 4.3 Soundness Theorem

688 Together, the type system and operational semantics allow us to prove the following syntactic
 689 soundness theorem, which says that Bundl programs are type safe and do not get stuck trying to
 690 execute operations for which they lack the required perspective:

691
 692 THEOREM 4.1. (*Type-and-Perspective Safety*). For any program s such that $\Gamma \vdash^\pi s$, either:

- 693 (1) s is **skip**, or
 694 (2) for any well-typed environments η, σ , and Σ , there is an s', η', σ' , and Σ' such that
 695 $\eta, \sigma, \Sigma, t, b, p \vdash^\pi s \rightsquigarrow^* s' \dashv \eta', \sigma', \Sigma'$ such that $\Gamma' \vdash^\pi s'$, where Γ' is an extension of Γ , and η', σ' , and
 696 Σ' are well-typed with respect to Γ' .

697 PROOF. Via the usual progress and preservation lemmas, available in Appendix A.4. \square

698 It is worth noting that this soundness theorem guarantees a syntactic safety property, not a liveness
 699 property: it does not guarantee that all threads sharing perspective π that *can* reach a program point
 700 typed with π *will* eventually do so. Indeed, in the presence of nontermination, liveness does not
 701 hold—some of the threads could **split** off and loop forever. While we believe the liveness version of
 702 this theorem holds for a terminating fragment of Bundl, it is not provable with syntactic methods; the
 703 proof would require semantic techniques that are notoriously challenging and would be a research
 704 contribution [5, 23, 49] in and of itself. We hope to tackle this proof for Bundl in future work.

705 5 Implementation

706 Prism is implemented as an embedded language in Python. Once a program type checks, Prism lowers
 707 it to a CUDA file. All perspective information is erased during this step, and the generated CUDA
 708 contains no run-time checks. The file can then be compiled by **nvcc**, NVIDIA’s closed-source compiler,
 709 to produce an executable. Because Prism operates at roughly the same abstraction level as CUDA,
 710 there is a one-to-one mapping between most language constructs and their CUDA counterparts. A
 711 notable change is the addition of three parameters to each device function: the thread’s relative ID,
 712 the block’s ID, and an offset into shared memory that it can use for allocations.

713
 714 *Inserting Synchronization.* Most of Prism’s implementation is straightforward,
 715 but inserting synchronization points is more involved. As described in Section 3.4,
 716 once data has been partitioned, Prism is responsible for inserting the appropriate
 717 synchronization points when the partition’s scope ends.

718 To determine where these points need to be inserted, Prism constructs a *data-control-flow graph* from the program. Each node corresponds to a partition, and each
 719 edge captures program-order precedence: the parent partition must complete before
 720 the child begins. The graph can have backedges introduced through loops. In this
 721 graph, each partition is categorized as a *read* or a *write* partition. Synchronization
 722 points are inserted according to the following scheme:

- 723 (1) If the parent partition is a write partition, a synchronization point is inserted
 724 to ensure that the current partition observes the most recent data; or
 725 (2) If the current partition is a write partition, a synchronization point is inserted
 726 to ensure that all preceding reads on the same memory have completed.

727 Figure 16 shows an example graph with the synchronization points derived from
 728 these two conditions. The inferred synchronization points in Figure 6 have also
 729 been marked on lines **TODO**.

730 Using this information, Prism emits *wait* operations before a partition begins (to
 731 wait for its parent) and *arrive* operations when a partition ends (to signal completion).

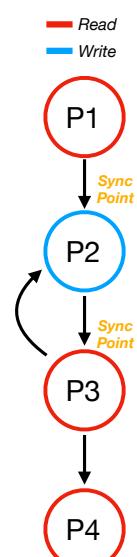


Fig. 16. An example data-control-flow graph lifted by Prism with its synchronization points.

CUDA exposes a general *split-barrier* primitive that we use to implement these waits and arrives. For special cases, such as synchronizing an entire block or warp, Prism instead uses primitives like `__syncthreads()` or `__syncwarp()`.

Synchronization for asynchronous data movement is handled in exactly the same way and uses the same underlying graph. CUDA allows asynchronous load operations to be associated with a split barrier, so Prism binds each asynchronous transfer to the appropriate wait–arrive pair inferred from the graph. Certain features, such as commit-group–style synchronization, require additional reasoning, and Prism performs further static analysis to insert the necessary synchronization around these operations.

It is worth noting that naively inserting synchronization immediately after each partition would be correct but prohibitively slow. To avoid this, Prism applies two optimization passes. A wait-motion pass pushes waits downward toward the first use of the partitioned name, and an arrive-motion pass pulls arrives upward toward its last use, reducing unnecessary stalls while preserving correctness.

6 Evaluation

Having explained the design of Prism, we now evaluate it in the context of four main questions:

RQ1 Can Prism express a diversity of composable CUDA programs?

RQ2 Can Prism express programs that use advanced GPU features?

RQ3 Can Prism match the performance of existing, speed-of-light CUDA code?

To perform this evaluation, we use two GPUs. The first is the NVIDIA H100 SXM5, a server-grade chip that supports Tensor Core operations and a dedicated hardware copy engine, the Tensor Memory Accelerator. Notably, the H100 introduces a new logical level called the *warpgroup* (collection of 4-aligned warps), and we show that our programming model can accommodate this new level. Moreover, because the H100 has historically served as the primary GPU for large-scale AI training, many CUDA kernels are already highly optimized for this hardware and achieve near-speed-of-light performance, providing a rigorous baseline for comparison. To ensure our results generalize beyond the H100, we additionally test programs on a second GPU, the NVIDIA 4070 SuperTi GPU, a consumer-grade chip.

Like mentioned in Section 5, when name typechecks a program, it produces a CUDA file. We compile the CUDA file with `nvcc` version 12.3 [33] with flags `-O3 -use_fast_math`. We initialize all inputs using a random number generator. We report the average runtime sampled over 10 iterations, following a warm-up phase of 5 iterations. All aggregate results reported in this section are geometric means. All programs in the evaluation have been reproduced in Section B.

6.1 RQ1: Can Prism express a diversity of CUDA programs?

We evaluate Prism’s expressivity by writing programs that have fundamentally different patterns of convergence, and writing a library modeled after CUB that contains composable pieces.

6.1.1 Programs with Different Convergence Behavior.

Matrix Multiplication. We chose matrix multiplication as our first benchmark for two main reasons. First, there exist several implementations that achieve near-peak performance, providing a strong baseline. Second, matrix multiplication allows for a range of increasingly sophisticated implementations that stress different parts of the language, making it ideal for evaluating expressiveness.

To undertake this exploration, we adapt the codebase from [6] to implement matrix multiplication on the `float` datatype, commonly referred to as `sgemm`. Concretely, we compute $C = \alpha * A \times B + \beta$ where A , B and C are matrices with appropriate sizes and α and β are scalars. As this is a `float`

benchmark, it *does not need* to use advanced GPU features like asynchrony or Tensor Cores to achieve speed-of-light performance. We will discuss these advanced features in Section 6.2.

We implement five variants of the **sgemm** benchmark: (1) a naive version that follows the traditional single-program multiple-data pattern, written from the perspective of a thread; (2) a version that exploits memory coalescing, still expressed at the thread level; (3) a version that builds on (2) by introducing 2D tiling and staging data in shared memory, which requires shifting first to the perspective of a block and then to the perspective of a thread, while also requiring block-level synchronization; (4) a version that applies 2D tiling with vectorized loads, again written from the perspective of a block; and (5) a version that combines the optimizations from (2) through (4) while adding an additional level of tiling from the perspective of a warp.

We can express all five variants cleanly, and the resulting programs remain close to their CUDA counterparts, which adopt a separation between perspectives through disciplined style. Prism, on the other hand, requires programmers to follow this discipline, and enforces it at compile time. The only notable difference is that some expressions must be hoisted earlier in the program so that their logical introduction corresponds to the required perspective structure. We evaluate the performance of these variants in Section 6.3.

Single-pass Parallel Prefix Scan with Decoupled Look-Back. We also implement scan, a widely used parallel primitive, in Prism. We focus on the prefix-sum scan, which computes, for each position in an array, the sum of all elements up to that position. Prefix-sum sits in a different corner of the GPU design space from matrix multiplication: it is memory intensive, requires careful attention to the convergence behavior of different threads, and traditionally requires multiple passes over data.

We implement the single-pass parallel prefix scan with decoupled look-back, introduced by Merrill and Garland [30], an elegant algorithm that does not require multiple passes over the input data. In this algorithm, each block computes a local prefix sum over its assigned region of the array. Once finished, it writes the final element of its region into a global array. Each block then *looks back* to accumulate the contributions of prior blocks, allowing them to independently determine the global prefix without a full sweep over memory. This decoupled look-back mechanism lets blocks progress at different speeds while still producing correct global results.

The algorithm involves several distinct points of convergence. Within each block, work is decomposed into fine-grained thread-level and warp-level scans. After producing the local result, blocks publish their partial prefix to global memory. Finally, each block waits until enough prefix information from earlier blocks becomes available, at which point it accumulates the value and completes its section of the global scan.

We implement this full strategy in Prism. Our implementation uses the unsafe feature to implement a global-memory spinlock that lets blocks check when the previous block’s data is ready.

6.1.2 Case Study: Can Prism help build composable functions? Over the course of our evaluation, we found ourselves developing a small library of functions—similar in spirit to the CUB library—that we could call to execute operations. In this section, we would like to qualitatively study how Prism can help programmers design libraries that they can compose with confidence.

As mentioned in Section 1, CUB occupies a unique design point in the GPU library ecosystem. Unlike many other libraries such as cuBLAS [10], cuDNN [12], and cuSPARSE [13], which provide global interfaces that users can call, CUB provides a device-side library organized into different levels. It makes these levels apparent by prefixing each of its functions with **Device**, **Block**, and **Thread**. The single-scan prefix sum in used variants of functions we translated into Prism already.

Currently, we have a library of functions that perform loads functions across block, warp and thread; store functions across block, warp, thread, and scans across a block and a warp. Let’s turn

our attention to a particular CUB function—the store function—and examine how it is equivalently expressed in Prism, and how Prism’s type system reifies the assumptions implicit in CUB.

In CUB, the store function is implemented as a class, as shown below:

```
838 1 template<typename T, int BlockDimX,
839   int ItemsPerThread, BlockLoadAlgorithm Algorithm = BLOCK_LOAD_DIRECT,
840   int BlockDimY = 1, int BlockDimZ = 1>
841 class BlockStore:
```

To use it, users must first instantiate an object of this class, and then call it with shared memory.

```
841 1 using BlockLoad = cub::BlockLoad<int, 128, 4, BLOCK_LOAD_DIRECT>;
842 2 // Allocate shared memory for BlockLoad
843 3 __shared__ typename BlockLoad::TempStorage temp_storage;
844 4 int thread_data[4]; // Thread local data
845 5 BlockLoad(temp_storage).Load(d_data, thread_data);
```

CUB exposes a leaky abstraction, where information about the number of threads, block sizes, and other details seeps through:

- (1) The CUB documentation needs to specify the number of threads that the function can assume to be available, because within the function, each thread must locate itself in the computation and use its `threadIdx.x` accordingly. If the starting `threadIdx.x` is not 0, the function must compute its relative ID internally.
- (2) The “item per thread” design is interesting and serves two purposes. The first is related to performance: if loops have constant bounds, they can be unrolled, enabling downstream optimizations. The second is related to correctness: the function relies on the assumption that all threads call the function with an equal number of values to load.
- (3) The CUB documentation also needs to specify that `thread_data` can be data local to each thread.
- (4) Finally, the CUB documentation specifies that if shared memory is being overwritten, a `__syncthreads()` call must be made to ensure that all reads have completed.

On the other hand, the same function has a completely different interface in Prism:

```
864 1 @cuda("device")
865 2 @requires(block[1], thread[32])
866 3 def block_load(input : ptr(const(int)) @ block[1], output : ptr(int) @ thread[1], items_per_thread : int @ block[1]):
```

In Prism, we are able to encapsulate code effectively, reducing the need to communicate numerous implementation details through documentation:

- (1) We do not need to pass in the number of threads at all. Whenever Prism calls a function, it threads the relative ID through, so each function can be written locally as if it were running alone, rather than having to determine where the thread resides in the global array.
- (2) We do not need to make `item_per_thread` a template argument for *correctness*. Its frequency is set at the function signature, so Prism will never allow a function to be called with a value living at a narrower perspective than its signature requires.
- (3) In our interface, `thread_data` is explicitly set to a thread-local value. Since it is not marked as `const`, Prism conservatively assumes it may be written to, and enforces at compile time that only `thread[1]` values are passed in.
- (4) Finally, using our synchronization pass outlined in Section 5, a `__syncthreads()` call will be inserted automatically if `src` is going to be used for writing.

Moreover, our CUB functions is built as a composition of 2 other functions:

```

883 1 @cuda("device")
884 2 requires(thread[32])
885 3 # block_load calls into warp_load
886 4 def warp_load(input : ptr(const int) @ thread[32], output : ptr(int) @ thread[1], items_per_thread : int @ thread[32]):
887 5
888 6 # warp_load calls into thread_load
889 7 @cuda("device")
890 8 requires(thread[1])
891 9 def thread_load(input : ptr(const int) @ thread[1], output : ptr(int) @ thread[1], items_per_thread : int @ thread[1]):
```

892

If we had mistakenly attempted to call **warp_load** from inside **thread_load**, Prism would return an error, not silently compute fail to meet the contract set out up the function.

893

6.2 RQ2: Can Prism express programs that use advanced GPU features?

To test whether Prism can express programs that use advanced features of modern GPUs, we write a matrix multiplication for the **bf16** datatype for the H100.

This benchmark is an acid test of our language, as the H100 **bf16** matrix multiplication pushes several language features to the extreme. To write a matrix multiplication that can hit peak throughput on an H100, we need to write a warp-specialized kernel that uses the Tensor Memory Accelerator (TMA), an asynchronous hardware copy engine that can move tiles of data at a time, and uses the warpgroup-level Tensor Core instructions, or **wmma**, specifically introduced for the Hopper architecture. A high-performance kernel for this matrix-multiplication overlaps computation with data movement by pipelining loads.

The implementation in Prism looks different from normal CUDA code, particularly in how pipelining is expressed. Since Prism uses *names* and subsequent **partition** or **claim** constructs to determine the synchronization each region requires, when pipelining, we cannot dynamically change the pipeline slot simply by maintaining an index variable that wraps around based on the pipeline length. Instead, each pipeline slot must be given a separate name so that Prism can track them independently and actually overlap compute with data-movement. This leads to pipeline slots that must be individually named and forces the load logic to be effectively “unrolled,” since we can no longer iterate over a pipeline-slot index. In turn, this forces that all pipelines in Prism be statically sized. In practice, these pipelines are statically sized in CUDA anyway because they occupy shared memory, which is a small, finite resource that must be explicitly managed.

Notably, for this benchmark, in addition to **wmma** and TMA, the program needs to dynamically reallocate registers between producer and consumer warpgroups, an instruction only available for the Hopper. This redistribution is a warpgroup-level collective operation, and Prism can check it like any other collective operation. Moreover, getting **wmma** to work did not require introducing a new perspective into the language; **thread[128]** was sufficient. We did, however, need to add a dedicated TMA-style asynchronous data-movement construct, since Prism must eventually insert the appropriate synchronization for these transfers.

921

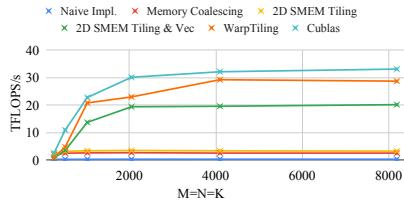
6.3 RQ3: Can Prism match the performance of existing, speed-of-light CUDA code?

In Section 6.1 and Section 6.2, we examined programs that expressed the same computation in multiple ways, expressed operations that rely on multiple points of convergence, and used advanced GPU features. We now discuss the performance of these programs.

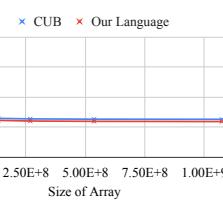
The performance of the matrix multiplication benchmarks is shown in , where we demonstrate that Prism is competitive with cuBLAS [10].

For the prefix scan, we compare our performance to , the library introduced earlier in Section 1, and show that we can reach approximately % of CUB’s peak bandwidth.

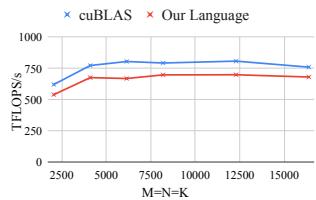
931



932 Fig. 17. Performance of `sgemm` on
933 square matrices as matrix dimension
934 $M=N=K$ increases.
935



936 Fig. 18. Performance for prefix
937 sum as input array size increases.
938



939 Fig. 19. Performance of `hgemm` on
940 square matrices as matrix dimension
941 $M=N=K$ increases.
942

943 Finally, we evaluate our H100 implementation and show that it is competitive with cuBLAS on
944 square sizes, falling between %. We emphasize that this is an exacting benchmark, and achieving this
945 level of performance requires writing large kernels with precise control over low-level features.

7 Related Work

946 Prism builds on a rich tradition for systems languages for GPU programming and theoretical foundations
947 for reasoning about parallel programs. However, Prism differs from these systems in a crucial
948 way: it treats collective operations, and therefore composability, as first-class concerns.
949

950 **Imperative Languages for GPUs.** CUDA [32], ROCm [2], and OpenCL [45] are imperative,
951 C-style programming languages that expose low-level access to GPU hardware. These languages
952 do not model perspectives for code or data, statically or otherwise; instead, programmers must
953 orchestrate computation on the machine by describing how individual threads execute code. The
954 subtle and often silent failures permitted by this model are the motivation for Prism.
955

956 Descend [29], a new, Rust-inspired language, uses a type system to track aspects of the compute
957 and memory hierarchy, and, is thus the closest to our work. However, Descend’s main concern is
958 preventing data races, while ours is ensuring that collective operations execute in valid contexts.
959 Moreover, Descend lacks support for many collective operations like Tensor Cores; in fact, because
960 Descend allows threads to read their own IDs and thus induce data-dependent control flow, adding
961 such support would require major changes to the language. It is also worth noting that while Descend
962 does formalize a type system, it does not attempt to prove any properties about it.
963

964 **Functional Languages for GPUs.** Futhark [27], Accelerate [7] and Vertigo [22] are functional
965 array languages with compilers targeting CUDA. These languages expose a high level interface,
966 abstracting away details of the hardware entirely in exchange for stronger safety guarantees. Prism,
967 however, makes an intentional decision to expose low-level details of the hardware, and thus does
968 not trade performance for safety.
969

970 **Tile-Based Kernel DSLs.** More recently, tile-based GPU languages—Triton [47], Pallas [44],
971 Tilus [21], Helion [37]—offer a middle ground between high-level abstraction and low-level control.
972 However, these languages sidestep the question of composability entirely because they restrict
973 programmers to a single vantage point in the GPU compute hierarchy: a block (Triton, Tilus, Helion),
974 or a warpgroup (Pallas). Gluon [48] and is a new low-level, tile-based language. However, it does not
975 check and cannot enforce that collective operations are executed at the correct layer of the hierarchy.
976

977 **Task-Based and Scheduling Languages.** Languages like Cypress [50], Halide [38], Fireiron [24]
978 and RISE/ELEVATE [25] expose a mapping specification or scheduling language that lets users modify
979 an existing reference program through external commands. Because they operate by transforming
980

981 a fixed source-of-truth rather than generating the program directly, they expose a fundamentally
 982 different programming model than Prism.

983
 984 **Libraries Built on CUDA.** Many GPU libraries, such as CUTLASS [43], CUB [31], and ThunderKittens [39], expose device functions that operate at various levels of the compute hierarchy.
 985 These functions are typically organized via C++ namespaces to reflect their intended usage (e.g.,
 986 warp-level, block-level). However, this organization is purely conventional: it encodes hierarchy
 987 through naming rather than enforcing it with a compiler. As a result, correct use requires careful
 988 discipline from both the library implementer (to uphold naming and usage invariants) and the user
 989 (to correctly interpret them). Any mismatch or subtle misunderstanding between the intended and
 990 actual use of these functions goes unchecked by the compiler, leaving room for fragile errors.
 991

992
 993 **Theoretical Foundations.** The design of Bundl is heavily inspired by existing work on dependency
 994 tracking [1]. Dependency tracking calculi allow type systems to track how data and code depend on
 995 each other, and have commonly been used to implement secure information flow analyses [19]. In
 996 Bundl, data that lives at a narrow perspective is unable to flow into data living at a broader perspective,
 997 and we use dependency tracking to capture this restriction in Bundl’s type system.

998 8 Conclusion, Limitations, and Future Work

1000 We have presented Prism, a new, low-level GPU language that statically guarantees safe usage of
 1001 compute resources by construction, without sacrificing low-level control. Prism introduces a new
 1002 mental model for writing GPU code, which we are excited to make more expressive and ergonomic.

1003 Currently, Prism does not natively support explicit references to pointers; programmers must rely
 1004 on built-in constructs to interact with memory. Prior work such as Descend [29], which builds on
 1005 Rust’s ownership model [42], provides a promising direction for extending Prism with more flexible
 1006 pointer semantics.

1007 We also aim to improve the ergonomics of programming with Prism, particularly by enhancing
 1008 the pipelining experience. As discussed in Section 6.2, Prism currently requires pipeline slots to be
 1009 explicitly name; we can improve this by adding language support for generating pipeline-style code.

1010 Prism can be extended to more architectures; in particular, it can accommodate newer GPUs—such
 1011 as Blackwell—that support coarser-grained Tensor Core operations. More broadly, we hope to
 1012 generalize the model presented here to hierarchical compute environments, including distributed
 1013 systems.

1014 On the theoretical side, we plan to explore a terminating fragment of Bundl and prove the liveness
 1015 property discussed in Section 4.3: that all threads sharing perspective π eventually reach the parts
 1016 of a program viewed at that π . This amounts to showing that threads sharing the same perspective
 1017 execute the same code and observe the same data, which we hope to prove using logical relations,
 1018 following Turon et al. [49] and Spies et al. [40, 41].

1019 We believe that Prism is a promising low-level substrate that enables confident composition and
 1020 can serve as a foundation for building higher-level libraries and abstractions.

1021

1022 References

- 1023 [1] Martín Abadi, Anindya Banerjee, Nevin Heintze, and Jon G. Riecke. 1999. A core calculus of dependency. In *Proceedings*
 1024 *of the 26th ACM SIGPLAN-SIGACT symposium on Principles of programming languages* (San Antonio Texas USA, 1999-01).
 1025 ACM, 147–160. [doi:10.1145/292540.292555](https://doi.org/10.1145/292540.292555)
- 1026 [2] Advanced Micro Devices, Inc. 2024. AMD ROCm™ Software. <https://www.amd.com/en/products/software/rocm.html>.
 Accessed: 2025-11-10.
- 1027 [3] Danel Ahman and Matija Pretnar. 2021. Asynchronous effects. *Proceedings of the ACM on Programming Languages* 5,
 1028 POPL (Jan. 2021), 1–28. [doi:10.1145/3434305](https://doi.org/10.1145/3434305)

1029

- [4] Michael Bauer, Sean Treichler, and Alex Aiken. 2014. Singe: leveraging warp specialization for high performance on GPUs. *SIGPLAN Not.* 49, 8 (Feb. 2014), 119–130. [doi:10.1145/2692916.2555258](https://doi.org/10.1145/2692916.2555258)
- [5] Lars Birkedal, Filip Sieczkowski, and Jacob Thamsborg. 2012. A Concurrent Logical Relation. In *Computer Science Logic (CSL'12) - 26th International Workshop/21st Annual Conference of the EACSL (2012)*. Schloss Dagstuhl – Leibniz-Zentrum für Informatik, 107–121. [doi:10.4230/LIPIcs.CSL.2012.107](https://doi.org/10.4230/LIPIcs.CSL.2012.107)
- [6] Simon Boehm. 2022. SGEMM_CUDA: Fast CUDA matrix multiplication from scratch. https://github.com/siboehm/SGEMM_CUDA. GitHub repository. Accessed: 2025-11-10.
- [7] Manuel M.T. Chakravarty, Gabriele Keller, Sean Lee, Trevor L. McDonell, and Vinod Grover. 2011. Accelerating Haskell array codes with multicore GPUs. In *Proceedings of the Sixth Workshop on Declarative Aspects of Multicore Programming* (Austin, Texas, USA) (DAMP '11). Association for Computing Machinery, New York, NY, USA, 3–14. [doi:10.1145/1926354.1926358](https://doi.org/10.1145/1926354.1926358)
- [8] NVIDIA Corporation. 2022. NVIDIA H100 GPU Whitepaper. (2022). <https://resources.nvidia.com/en-us-hopper-architecture/nvidia-h100-tensor-c> Accessed: 2025-11-12.
- [9] NVIDIA Corporation. 2025. cub::BlockReduce. https://nvidia.github.io/cccl/cub/api/classcub_1_1BlockReduce.html Accessed: 2025-11-09.
- [10] NVIDIA Corporation. 2025. cuBLAS. <https://docs.nvidia.com/cuda/cublas/index.html> Accessed: 2025-11-09.
- [11] NVIDIA Corporation. 2025. cuBLASDx. <https://docs.nvidia.com/cuda/cublasdx/#nvidia-cublasdx> Accessed: 2025-11-11.
- [12] NVIDIA Corporation. 2025. cuDNN. <https://docs.nvidia.com/deeplearning/cudnn/latest/> Accessed: 2025-11-09.
- [13] NVIDIA Corporation. 2025. cuSPARSE. <https://docs.nvidia.com/cuda/cusparse/index.html> Accessed: 2025-11-09.
- [14] NVIDIA Corporation. 2025. Data Movement and Conversion Instructions: Bulk Copy. <https://docs.nvidia.com/cuda/parallel-thread-execution/#data-movement-and-conversion-instructions-bulk-copy> Accessed: 2025-11-12.
- [15] NVIDIA Corporation. 2025. Data Movement and Conversion Instructions: Non-bulk Copy. <https://docs.nvidia.com/cuda/parallel-thread-execution/#data-movement-and-conversion-instructions-non-bulk-copy> Accessed: 2025-11-12.
- [16] NVIDIA Corporation. 2025. Synchronization Functions. <https://docs.nvidia.com/cuda/cuda-c-programming-guide/#synchronization-functions> Accessed: 2025-11-11.
- [17] NVIDIA Corporation. 2025. Warp Shuffle Functions. <https://docs.nvidia.com/cuda/cuda-c-programming-guide/#warp-shuffle-functions> Accessed: 2025-11-10.
- [18] DeepSeek-AI, Aixin Liu, Bei Feng, Bing Xue, Bingxuan Wang, Bochao Wu, Chengda Lu, Chenggang Zhao, Chengqi Deng, Chenyu Zhang, Chong Ruan, Damai Dai, Daya Guo, Dejian Yang, Deli Chen, Dongjie Ji, Erhang Li, Fangyun Lin, Fucong Dai, Fuli Luo, Guangbo Hao, Guanting Chen, Guowei Li, H. Zhang, Han Bao, Hanwei Xu, Haocheng Wang, Haowei Zhang, Honghui Ding, Huajian Xin, Huazuo Gao, Hui Li, Hui Qu, J. L. Cai, Jian Liang, Jianzhong Guo, Jiaqi Ni, Jiashi Li, Jiawei Wang, Jin Chen, Jingchang Chen, Jingyang Yuan, Junjie Qiu, Junlong Li, Junxiao Song, Kai Dong, Kai Hu, Kaige Gao, Kang Guan, Kexin Huang, Kuai Yu, Lean Wang, Lecong Zhang, Lei Xu, Leyi Xia, Liang Zhao, Litong Wang, Liyue Zhang, Meng Li, Miaojun Wang, Mingchuan Zhang, Minghua Zhang, Minghui Tang, Mingming Li, Ning Tian, Panpan Huang, Peiyi Wang, Peng Zhang, Qiancheng Wang, Qihao Zhu, Qinyu Chen, Qiushi Du, R. J. Chen, R. L. Jin, Ruiqi Ge, Ruisong Zhang, Ruizhe Pan, Runji Wang, Runxin Xu, Ruoyu Zhang, Ruyi Chen, S. S. Li, Shanghao Lu, Shangyan Zhou, Shanhua Chen, Shaoqing Wu, Shengfeng Ye, Shengfeng Ye, Shirong Ma, Shiyu Wang, Shuang Zhou, Shuiping Yu, Shunfeng Zhou, Shuting Pan, T. Wang, Tao Yun, Tian Pei, Tianyu Sun, W. L. Xiao, Wangding Zeng, Wanjia Zhao, Wei An, Wen Liu, Wenfeng Liang, Wenjun Gao, Wenqin Yu, Wentao Zhang, X. Q. Li, Xiangyu Jin, Xianzu Wang, Xiao Bi, Xiaodong Liu, Xiaohan Wang, Xiaojin Shen, Xiaokang Chen, Xiaokang Zhang, Xiaosha Chen, Xiaotao Nie, Xiaowen Sun, Xiaoxiang Wang, Xin Cheng, Xin Liu, Xin Xie, Xingchao Liu, Xingkai Yu, Xinnan Song, Xinxia Shan, Xinyi Zhou, Xinyu Yang, Xinyuan Li, Xuecheng Su, Xuheng Lin, Y. K. Li, Y. Q. Wang, Y. X. Wei, Y. X. Zhu, Yang Zhang, Yanhong Xu, Yanhong Xu, Yanping Huang, Yao Li, Yao Zhao, Yaofeng Sun, Yaohui Li, Yaohui Wang, Yi Yu, Yi Zheng, Yichao Zhang, Yifan Shi, Yiliang Xiong, Ying He, Ying Tang, Yishi Piao, Yisong Wang, Yixuan Tan, Yiyang Ma, Yiyuan Liu, Yongqiang Guo, Yu Wu, Yuan Ou, Yuchen Zhu, Yuduan Wang, Yue Gong, Yuheng Zou, Yujia He, Yukun Zhu, Yunfan Xiong, Yunxian Ma, Yuting Yan, Yuxiang Luo, Yuxiang You, Yuxuan Liu, Yuyang Zhou, Z. F. Wu, Z. Z. Ren, Zehui Ren, Zhangli Sha, Zhe Fu, Zhean Xu, Zhen Huang, Zhen Zhang, Zhenda Xie, Zhengyan Zhang, Zhewen Hao, Zhibin Gou, Zhicheng Ma, Zhiqiang Yan, Zhihong Shao, Zhipeng Xu, Zhiyu Wu, Zhongyu Zhang, Zhushu Li, Zihui Gu, Zijia Zhu, Zijun Liu, Zilin Li, Ziwei Xie, Ziyang Song, Ziyi Gao, and Zizheng Pan. 2025. DeepSeek-V3 Technical Report. [doi:10.48550/arXiv.2412.19437](https://doi.org/10.48550/arXiv.2412.19437) arXiv:2412.19437.
- [19] Dorothy E. Denning and Peter J. Denning. 1977. Certification of programs for secure information flow. *Commun. ACM* 20, 7 (July 1977), 504–513. [doi:10.1145/359636.359712](https://doi.org/10.1145/359636.359712)
- [20] Edsger Wybe Dijkstra. 1963. Over de sequentialeit van procesbeschrijvingen. (1963). <https://training-ir7.tdl.org/handle/123456789/594>
- [21] Yaoyao Ding, Bohan Hou, Xiao Zhang, Allan Lin, Tianqi Chen, Cody Yu Hao, Yida Wang, and Gennady Pekhimenko. 2025. Tilitus: A Tile-Level GPGPU Programming Language for Low-Precision Computation. arXiv:2504.12984 [cs.LG] <https://arxiv.org/abs/2504.12984>

- 1079 [22] C. Elliott. 2005. Vertigo – GPU Compiler & Embedded Language for Graphics Processors. <http://conal.net/Vertigo/>.
 1080 Accessed: 2025-11-10.
- 1081 [23] Azadeh Farzan, Zachary Kincaid, and Andreas Podelski. 2016. Proving Liveness of Parameterized Programs. In *Proceedings of the 31st Annual ACM/IEEE Symposium on Logic in Computer Science*. ACM, New York NY USA, 185–196.
 1082 doi:[10.1145/2933575.2935310](https://doi.org/10.1145/2933575.2935310)
- 1083 [24] Bastian Hagedorn, Archibald Samuel Elliott, Henrik Barthels, Rastislav Bodik, and Vinod Grover. 2020. Fireiron: A
 1084 Data-Movement-Aware Scheduling Language for GPUs. In *Proceedings of the ACM International Conference on Parallel
 1085 Architectures and Compilation Techniques* (Virtual Event, GA, USA) (PACT ’20). Association for Computing Machinery,
 1086 New York, NY, USA, 71–82. doi:[10.1145/3410463.3414632](https://doi.org/10.1145/3410463.3414632)
- 1087 [25] Bastian Hagedorn, Johannes Lenfers, Thomas Kundefinedhler, Xueying Qin, Sergei Gorlatch, and Michel Steuwer. 2020.
 1088 Achieving high-performance the functional way: a functional pearl on expressing high-performance optimizations as
 1089 rewrite strategies. *Proc. ACM Program. Lang.* 4, ICFP, Article 92 (Aug. 2020), 29 pages. doi:[10.1145/3408974](https://doi.org/10.1145/3408974)
- 1090 [26] Troels Henriksen, Niels G. W. Serup, Martin Elsman, Fritz Henglein, and Cosmin E. Oancea. 2017. Futhark: purely
 1091 functional GPU-programming with nested parallelism and in-place array updates. In *Proceedings of the 38th ACM
 1092 SIGPLAN Conference on Programming Language Design and Implementation*. ACM, Barcelona Spain, 556–571. doi:[10.1145/3062341.3062354](https://doi.org/10.1145/3062341.3062354)
- 1093 [27] Troels Henriksen, Niels G. W. Serup, Martin Elsman, Fritz Henglein, and Cosmin E. Oancea. 2017. Futhark: purely
 1094 functional GPU-programming with nested parallelism and in-place array updates. *SIGPLAN Not.* 52, 6 (June 2017),
 1095 556–571. doi:[10.1145/3140587.3062354](https://doi.org/10.1145/3140587.3062354)
- 1096 [28] Jan Hoffmann. 2011. *Types with potential: polynomial resource bounds via automatic amortized analysis*. Ph. D. Dissertation.
 1097 lmu.
- 1098 [29] Bastian Köpcke, Sergei Gorlatch, and Michel Steuwer. 2024. Descend: A Safe GPU Systems Programming Language. 8
 1099 (2024), 841–864. Issue PLDI. doi:[10.1145/3656411](https://doi.org/10.1145/3656411)
- 1100 [30] Duane Merrill and Michael Garland. 2016. *Single-pass Parallel Prefix Scan with Decoupled Look-back*. Technical Report
 1101 NVR-2016-002. NVIDIA Corporation. Accessed: 2025-11-10.
- 1102 [31] NVIDIA Corporation. 2025. *CUB: CUDA Unbound*. <https://docs.nvidia.com/cuda/cub/index.html> CUDA Toolkit
 Documentation.
- 1103 [32] NVIDIA Corporation. 2025. *CUDA C++ Programming Guide*. NVIDIA. <https://docs.nvidia.com/cuda/cuda-c-programming-guide/> Accessed: 2025-11-06.
- 1104 [33] NVIDIA Corporation. 2025. CUDA Toolkit 12.3 Downloads (Archive). <https://developer.nvidia.com/cuda-12-3-0-download-archive>. Accessed: 2025-11-11.
- 1105 [34] NVIDIA Corporation. 2025. Parallel Thread Execution (PTX) ISA – Asynchronous Warpgroup Level Matrix Multiply-
 1106 Accumulate Instructions. <https://docs.nvidia.com/cuda/parallel-thread-execution/index.html?highlight=tngen05%2520cp#asynchronous-warpgroup-level-matrix-instructions>. Accessed: 2025-11-10.
- 1107 [35] NVIDIA Corporation. 2025. Parallel Thread Execution (PTX) ISA – TensorCore 5th Generation Instructions (tc-
 1108 gen05). <https://docs.nvidia.com/cuda/parallel-thread-execution/index.html?highlight=tngen05%2520cp#tensorcore-5th-generation-instructions>. Accessed: 2025-11-10.
- 1109 [36] NVIDIA Corporation. 2025. Parallel Thread Execution (PTX) ISA – Warp Level Matrix Instructions. <https://docs.nvidia.com/cuda/parallel-thread-execution/index.html?highlight=tngen05%2520cp#warp-level-matrix-instructions>. Accessed: 2025-11-10.
- 1110 [37] PyTorch Team at Meta. 2025. Helion: A High-Level DSL for Performant and Portable ML Kernels. <https://pytorch.org/blog/helion/>. Accessed: 2025-11-10.
- 1111 [38] Jonathan Ragan-Kelley, Connelly Barnes, Andrew Adams, Sylvain Paris, Frédo Durand, and Saman Amarasinghe. 2013.
 1112 Halide: a language and compiler for optimizing parallelism, locality, and recomputation in image processing pipelines.
SIGPLAN Not. 48, 6 (June 2013), 519–530. doi:[10.1145/2499370.2462176](https://doi.org/10.1145/2499370.2462176)
- 1113 [39] Benjamin Frederick Spector, Simran Arora, Aaryan Singhal, Arjun Parthasarathy, Daniel Y Fu, and Christopher Re. 2025.
 1114 ThunderKittens: Simple, Fast, and \$textit{adorable}\$\$ Kernels. In *The Thirteenth International Conference on Learning
 1115 Representations*. <https://openreview.net/forum?id=0fjfVOSUra>
- 1116 [40] Simon Spies, Lennard Gäher, Daniel Gratzer, Joseph Tassarotti, Robbert Krebbers, Derek Dreyer, and Lars Birkedal.
 1117 2021. Transfinite Iris: resolving an existential dilemma of step-indexed separation logic. In *Proceedings of the 42nd ACM
 1118 SIGPLAN International Conference on Programming Language Design and Implementation*. ACM, Virtual Canada, 80–95.
 1119 doi:[10.1145/3453483.3454031](https://doi.org/10.1145/3453483.3454031)
- 1120 [41] Simon Spies, Neel Krishnaswami, and Derek Dreyer. 2021. Transfinite step-indexing for termination. *Proceedings of the
 1121 ACM on Programming Languages* 5, POPL (Jan. 2021), 1–29. doi:[10.1145/3434294](https://doi.org/10.1145/3434294)
- 1122 [42] Rust Team. 2025. Rust Programming Language. <https://rust-lang.org/>
- 1123 [43] Vijay Thakkar, Pradeep Ramani, Cris Cecka, Aniket Shivam, Honghao Lu, Ethan Yan, Jack Kosaian, Mark Hoemmen,
 1124 Haicheng Wu, Andrew Kerr, Matt Nicely, Duane Merrill, Dustyn Blasig, Aditya Atluri, Fengqi Qiao, Piotr Majcher, Paul
 1125

- 1128 Springer, Markus Hohnerbach, Jin Wang, and Manish Gupta. 2023. *CUTLASS*. <https://github.com/NVIDIA/cutlass>
- 1129 [44] The JAX Authors. 2024. Pallas:Mosaic GPU – A JAX Kernel Language for GPUs. <https://docs.jax.dev/en/latest/pallas/gpu/index.html>. Accessed: 2025-11-10.
- 1130 [45] The Khronos Group Inc. 2025. OpenCL™ – The Open Standard for Parallel Programming of Heterogeneous Systems. <https://www.khronos.org/opencl/>. Accessed: 2025-11-10.
- 1131 [46] Philippe Tillet. 2025. Introducing Triton: Open-source GPU programming for neural networks. <https://openai.com/index/triton/>
- 1132 [47] Philippe Tillet, H. T. Kung, and David Cox. 2019. Triton: an intermediate language and compiler for tiled neural network computations. In *Proceedings of the 3rd ACM SIGPLAN International Workshop on Machine Learning and Programming Languages* (Phoenix, AZ, USA) (MAPL 2019). Association for Computing Machinery, New York, NY, USA, 10–19. doi:[10.1145/3315508.3329973](https://doi.org/10.1145/3315508.3329973)
- 1133 [48] triton-lang. 2025. “01-intro.py” – Introduction to Gluon tutorial in Triton. <https://github.com/triton-lang/triton/blob/main/python/tutorials/gluon/01-intro.py>. Accessed: 2025-11-10.
- 1134 [49] Aaron J. Turon, Jacob Thamsborg, Amal Ahmed, Lars Birkedal, and Derek Dreyer. 2013. Logical relations for fine-grained concurrency. In *Proceedings of the 40th annual ACM SIGPLAN-SIGACT symposium on Principles of programming languages* (Rome Italy, 2013-01-23). ACM, 343–356. doi:[10.1145/2429069.2429111](https://doi.org/10.1145/2429069.2429111)
- 1135 [50] Rohan Yadav, Michael Garland, Alex Aiken, and Michael Bauer. 2025. Task-Based Tensor Computations on Modern GPUs. *Proc. ACM Program. Lang.* 9, PLDI, Article 163 (June 2025), 25 pages. doi:[10.1145/3729262](https://doi.org/10.1145/3729262)
- 1136
- 1137
- 1138
- 1139
- 1140
- 1141
- 1142
- 1143
- 1144
- 1145
- 1146
- 1147
- 1148
- 1149
- 1150
- 1151
- 1152
- 1153
- 1154
- 1155
- 1156
- 1157
- 1158
- 1159
- 1160
- 1161
- 1162
- 1163
- 1164
- 1165
- 1166
- 1167
- 1168
- 1169
- 1170
- 1171
- 1172
- 1173
- 1174
- 1175
- 1176

1177 **A Complete Bundl Type System, Semantics, and Syntactic Soundness Proofs**

1178 **A.1 Basic Definitions**

1180 Hierarchy Levels $h ::= \text{Grid} \mid \text{Block} \mid \text{Thread}$
 1181 Memory Kinds $l ::= \text{Local} \mid \text{Shared} \mid \text{Global}$
 1182 Perspectives $\pi : h \times \mathbb{N}$
 1183 Base Types $b ::= \text{bool} \mid \text{int} \mid \text{float}$
 1184 Types $\tau ::= b \mid b[]^l \mid \text{Fun}(\Gamma, \pi, m) \mid \text{async } \tau$
 1185 Contexts $\Gamma ::= \cdot \mid \Gamma, x : \pi \tau$

1186 Shared Memory Remaining $m : \mathbb{N}$

1189
 1190
 1191 Perspectives π are part of an algebra parameterized over some constant values T (the number of
 1192 threads per block) and B (the number of blocks per grid). With these values, we have two isomorphisms:

$$(\text{Block}, 1) \cong (\text{Thread}, T)$$

1193 and

$$(\text{Grid}, 1) \cong (\text{Block}, B)$$

1194
 1195 The **group** and **split** operations of Prism allows us to move along this isomorphism, from left
 1196 to right. For clarity, in Bundl we split these operations into three: a **split** operation that can split
 1197 perspectives into multiple narrower ones, and a **destruct** operation that directly moves us along
 1198 the isomorphism, and a **group** that divides our current perspective into equal sized parts.

1199 Perspectives π are also lexicographically ordered in the obvious way. hs are ordered such that
 1200 $\text{Thread} \leq \text{Block} \leq \text{Grid}$, and $(h_1, n_1) \leq (h_2, n_2)$ iff $n_1 | n_2$ and $h_1 \leq h_2$.

1201 We define scalar multiplication $i \times h$ of natural numbers with hs : $i \times (h, n) = (h, in)$.

1202 We also define division of perspectives and hierarchy levels of type $\pi \times \pi \rightarrow \mathbb{N}$. $\text{Grid}/\text{Block} = B$
 1203 and $\text{Block}/\text{Thread} = T$. We lift this to perspectives like so: $(h_1, n_1) / (h_2, n_2) = ((h_1/h_2) \cdot n_1) / n_2$.

1204 Lastly we define a partial \downarrow operator on hs such that $\downarrow \text{Block} = \text{Thread}$ and $\downarrow \text{Grid} = \text{Block}$. Note
 1205 that $\downarrow \text{Thread}$ is undefined. This operator lifts to π s whose second component is 1 and encodes the
 1206 leftward component of the isomorphism above: $\downarrow (\text{Block}, 1) = (\text{Thread}, T)$ and $\downarrow (\text{Grid}, 1) = (\text{Block}, B)$.

1207 Note also that the presentation of these rules in the main body of the paper elide the m portion,
 1208 which tracks the maximum amount of memory a given computation is allowed to use. In the full
 1209 system presented here, both the typing rules and the operational semantics carry an additional piece
 1210 of information tracking allocated memory.

1214 **A.2 Complete Bundl Typing Rules**

1215 **A.2.1 Expressions.**

$$\frac{x : \pi \tau \in \Gamma}{\Gamma \vdash^\pi x : \tau} \text{-Var} \quad \frac{}{\Gamma \vdash^\pi n : \text{int}} \text{-Int} \quad \frac{}{\Gamma \vdash^\pi f : \text{float}} \text{-Float}$$

$$\frac{}{\Gamma \vdash^\pi b : \text{bool}} \text{-Bool} \quad \frac{\pi < (\text{Grid}, 1)}{\Gamma \vdash^\pi \text{partition_id} : \text{int}} \text{-Partition-Id}$$

$$\frac{\Gamma \vdash^\pi' e_1 : \tau[]^l \quad \Gamma \vdash^\pi e_2 : \text{int} \quad l = \text{Global or } l = \text{Local} \quad \pi \leq \pi'}{\Gamma \vdash^\pi e_1[e_2] : \tau} \text{-Arr-Access}$$

1226	
1227	$\frac{\Gamma \vdash^{\pi'} e_1 : \tau[]^{\text{Shared}} \quad \Gamma \vdash^{\pi} e_2 : \text{int} \quad \pi \leq (\text{Block}, 1) \quad \pi \leq \pi'}{\Gamma \vdash^{\pi} e_1[e_2] : \tau} \text{ T-Arr-Access-Shared}$
1228	
1229	
1230	$\frac{\Gamma \vdash^{\pi} e_1 : \text{int} \quad \Gamma \vdash^{\pi} e_2 : \text{int}}{\Gamma \vdash^{\pi} e_1 \text{ bop } e_2 : \text{int}} \text{ T-Bop} \quad \frac{\Gamma \vdash^{\pi} e_1 : \text{int} \quad \Gamma \vdash^{\pi} e_2 : \text{int}}{\Gamma \vdash^{\pi} e_1 \text{ cmp } e_2 : \text{bool}} \text{ T-Cmp}$
1231	
1232	A.2.2 <i>Statements</i> .
1233	$\frac{f : \pi' \text{Fun}(x_i : \tau_i, \pi, m') \in \Gamma \quad \Gamma \vdash^{\pi} e_i : \tau_i \quad m' \leq m}{\Gamma \vdash_m^{\pi} f(e_1, \dots, e_n)} \text{ T-Function-Call}$
1234	
1235	
1236	$\frac{\Gamma \vdash_m^{(h, n_1)} s_1 \quad \Gamma \vdash_m^{(h, n_2)} s_2 \quad n_1, n_2 \text{ align to } n}{\Gamma \vdash_m^{(h, n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\}} \text{ T-Split} \quad \frac{\Gamma \vdash_m^{\downarrow \pi} s}{\Gamma \vdash_m^{\pi} \text{destruct in } s} \text{ T-Destruct}$
1237	
1238	
1239	
1240	$\frac{\Gamma \vdash_m^{(h, n)} s}{\Gamma \vdash_m^{(h, q, n)} \text{group } q \ s} \text{ T-Group} \quad \frac{x : \pi' \tau \in \Gamma \quad \Gamma \vdash^{\pi} e : \tau \quad \pi' \leq \pi}{\Gamma \vdash_m^{\pi} x = e} \text{ T-Assn}$
1241	
1242	
1243	$\frac{\Gamma \vdash_m^{\pi} \text{init}_{\psi}}{\Gamma \vdash_m^{\pi} \text{dec}_{\psi}} \text{ T-Sync-Init} \quad \frac{\Gamma \vdash_m^{\pi} \text{dec}_{\psi}}{\Gamma \vdash_m^{\pi} \text{dec}_{\psi}} \text{ T-Sync-Dec}$
1244	
1245	$\frac{}{\Gamma \vdash_m^{\pi} \text{wait}_{\psi}} \text{ T-Sync-Wait} \quad \frac{}{\Gamma \vdash_m^{\pi} \text{skip}} \text{ T-Skip} \quad \frac{n \leq m}{\Gamma \vdash_m^{\pi} \text{free } n} \text{ T-Free}$
1246	
1247	
1248	$\frac{\Gamma \vdash^{\pi} e : \tau \quad \Gamma, x : \pi' \tau \vdash_m^{\pi} s \quad \pi' \leq \pi \quad \tau \text{ not an array type}}{\Gamma \vdash_m^{\pi} x : \tau @ \pi' = e \text{ in } s} \text{ T-Decl}$
1249	
1250	
1251	$\frac{\Gamma \vdash^{\pi'} e_1 : \tau[]^l \quad \Gamma \vdash^{\pi} e_2 : \text{int} \quad \Gamma \vdash^{\pi} e_3 : \tau \quad l = \text{Global or } l = \text{Local} \quad \pi' \leq \pi}{\Gamma \vdash_m^{\pi} e_1[e_2] = e_3} \text{ T-Arr-Assn}$
1252	
1253	
1254	$\frac{\Gamma \vdash^{\pi'} e_1 : \tau[]^{\text{Shared}} \quad \Gamma \vdash^{\pi} e_2 : \text{int} \quad \Gamma \vdash^{\pi} e_3 : \tau \quad \pi \leq (\text{Block}, 1) \quad \pi' \leq \pi}{\Gamma \vdash_m^{\pi} e_1[e_2] = e_3} \text{ T-Arr-Assn-Shared}$
1255	
1256	
1257	$\frac{\Gamma \vdash^{\pi} e : \text{bool} \quad \Gamma \vdash_m^{\pi} s_1 \quad \Gamma \vdash_m^{\pi} s_2}{\Gamma \vdash_{\max(m_1, m_2)}^{\pi} \text{if } e \text{ then } s_1 \text{ else } s_2} \text{ T-If}$
1258	
1259	
1260	$\frac{\Gamma \vdash^{\pi} e : \text{bool} \quad \Gamma \vdash_m^{\pi} s}{\Gamma \vdash_m^{\pi} \text{while } e \text{ do } s} \text{ T-While} \quad \frac{\Gamma \vdash_m^{\pi} s_1 \quad \Gamma \vdash_m^{\pi} s_2}{\Gamma \vdash_{\max(m_1, m_2)}^{\pi} s_1 ; s_2} \text{ T-Seq}$
1261	
1262	
1263	$\frac{\Gamma, x : \pi \tau[]^l \vdash_m^{\pi} s \quad l = \text{Global or } l = \text{Local}}{\Gamma \vdash_{m+n \cdot \text{size}(\tau)}^{\pi} x := \text{alloc } l \ \tau \ n \text{ in } s} \text{ T-Alloc}$
1264	
1265	
1266	$\frac{\Gamma, x : (\text{Block}, 1) \tau[]^{\text{Shared}} \vdash_m^{\pi} s}{\Gamma \vdash_{m+n \cdot \text{size}(\tau)}^{\pi} x := \text{alloc Shared } \tau \ n \text{ in } s} \text{ T-Alloc-Shared}$
1267	
1268	
1269	$\frac{\Gamma, y : (h, n/c) \tau[]^l \vdash_m^{(h, n)} s \quad c n \quad l \neq \text{Local}}{\Gamma, x : (h, n) \tau[]^l \vdash_m^{(h, n)} \text{partition}_{\psi} x \text{ into } y \text{ by } c \text{ in } s} \text{ T-Partition}$
1270	
1271	
1272	$\frac{\Gamma, y : (h, n') \tau[]^l \vdash_m^{(h, n')} s \quad n' \leq n \quad l \neq \text{Local}}{\Gamma, x : (h, n) \tau[]^l \vdash_m^{(h, n)} \text{claim}_{\psi} x \text{ into } y \text{ at } n' \text{ in } s} \text{ T-Claim}$
1273	
1274	

$$\begin{array}{c}
 \frac{\Gamma, y : \downarrow^\pi \tau[]^l \vdash_m^\pi s \quad l \neq \text{Local}}{\Gamma, x : \pi \tau[]^l \vdash_m^\pi \text{lower}_\psi x \text{ into } y \text{ in } s} \text{ T-Lower} \\
 \\
 \frac{\Gamma, y : (\text{Thread}, 1) \text{ async } \tau[]^l \vdash_m^{(\text{Thread}, 1)} s}{\Gamma, x : (\text{Thread}, 1) \tau[]^l \vdash_m^{(\text{Thread}, 1)} \text{async_partition}_\phi x \text{ into } y \text{ in } s} \text{ T-Async-Partition} \\
 \\
 \frac{}{\Gamma, x : \pi \tau[]^l, y : \pi \tau[]^{l'} \vdash_m^\pi \text{async_memcpy}(x, y)} \text{ T-Async-Memcpy} \\
 \\
 \frac{}{\Gamma, x : \pi \tau[]^l, y : \pi \tau[]^{l'} \vdash_m^\pi \text{memcpy}(x, y)} \text{ T-Memcpy}
 \end{array}$$

A.3 Complete Bundl Semantics

A.3.1 Definitions.

Global Memory $\Sigma ::= \cdot | \Sigma, n \mapsto^\pi v$

Shared Memory $\sigma ::= \cdot | \sigma, n \mapsto^\pi v$

Local Memory $\eta ::= \cdot | \eta, n \mapsto^\pi v$

Block Memory Map $S ::= \forall n \in B, n \mapsto \sigma$

Thread Memory Map $L ::= \forall n \in T, n \mapsto \eta$

Synchronization Map $\Psi : \psi \rightarrow \mathbb{N} \rightarrow \mathbb{N}$

Deferred Computations Map $\Phi : \phi \rightarrow \{s\}$

In real GPUs, thread IDs are only unique within their block. However, in this calculus for simplicity we assume thread IDs are global. One can convert back and forth between this abstracted notion of a thread ID and a block-unique ID via addition modulo T . That is, $t_{\text{real}} = t_{\text{simplified}} \bmod T$ and $t_{\text{simplified}} = t_{\text{real}} + b \cdot T$.

By convention the names for local and shared and global memory do not conflict, as on the GPU they will be separate pointer spaces. Additionally, we freely interchange between using names for variables and integer locations.

In the main body of the paper, for simplicity we elide the synchronization map and the deferred computation map from the operational semantics, as our theorems do not make any guarantees about non-interference. However, as they are part of the full semantics, we include them here for completeness. By convention the synchronization and deferred computation maps are a total functions, initialized to map to $\lambda_. 0$ for ψ s and $\lambda_. \{s\}$ for ϕ s not explicitly initialize.

The shape of the judgment for a single thread is $\eta, \sigma, \Sigma, t, b, p, \Psi \vdash_m^\pi s \rightsquigarrow s' \dashv_m \eta', \sigma', \Sigma', \Psi'$. The t here is the thread ID, the b is the block ID, and the p is the perspective ID. The last of these three is modified and managed by the rules for **split**, **group** and **destruct**, and tracks the relative position of the thread within a group. This semantics is in a small step style.

The shape of the judgment for expressions is $\eta, \sigma, \Sigma \vdash_\pi^\pi e \Downarrow v$. The two π s represent the ambient compute context (i.e., the context in which resources are being read), while π' , represents the target compute context (i.e., the compute context of the variable into which the result of the expression is going to be written. This is relevant for computing the value of **partition_id**, which divides the two contexts. As a shorthand, we can divide a perspective by a scalar value like so: $(h, n)/c = (h, n)/(h, c)$.

1324 The overall evaluation of a program is expressed as

$$L, S, \Sigma, P, \Psi, \Phi \rightsquigarrow L', S', \Sigma', P', \Psi', \Phi'.$$

1325
1326 In this judgment P serves as a *thread pool*, mapping pairs of thread and block IDs (which don't change) to statements and memory (which can be updated by stepping). One can think of P as tracking which
1327 program is running on each thread. This steps according to the following rule:

$$\frac{L(t), S(b), \Sigma, t, b, 0, \Psi, \Phi \vdash_m^{(\text{Grid}, 1)} s \rightsquigarrow s' \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi' \quad P(t, b) = (s, m)}{L, S, \Sigma, P, \Psi, \Phi \rightsquigarrow L[t \mapsto \eta'], S[b \mapsto \sigma'], \Sigma', P[(t, b) \mapsto (s', m')], \Psi', \Phi'} \text{ S-Program}$$

1330
1331 For simplicity of notation, we define an **update** operation that searches the three environments
1332 for the one that contains the variable being used (by convention, there is no conflict between the
1333 environments, as in reality they exist in three separate address spaces). We also define a similar **get**
1334 operation that retrieves a variable from memory, and a **rename** operation that remaps a variable
1335 with the same value but under a different name.
1336
1337

$$\text{update}(\eta, \sigma, \Sigma, x, v) = (\eta[x \mapsto^\pi v], \sigma, \Sigma) \text{ when } x \in^\pi \eta$$

$$\text{update}(\eta, \sigma, \Sigma, x, v) = (\eta, \sigma[x \mapsto^\pi v], \Sigma) \text{ when } x \in^\pi \sigma$$

$$\text{update}(\eta, \sigma, \Sigma, x, v) = (\eta, \sigma, \Sigma[x \mapsto^\pi v]) \text{ when } x \in^\pi \Sigma$$

$$\text{get}(\eta, \sigma, \Sigma, x) = \eta(x) \text{ when } x \in^\pi \eta$$

$$\text{get}(\eta, \sigma, \Sigma, x) = \sigma(x) \text{ when } x \in^\pi \sigma$$

$$\text{get}(\eta, \sigma, \Sigma, x) = \Sigma(x) \text{ when } x \in^\pi \Sigma$$

$$\text{rename}(\eta, \sigma, \Sigma, x, y, \pi') = (\eta[y \mapsto^\pi \eta(x)], \sigma, \Sigma) \text{ when } x \in^\pi \eta$$

$$\text{rename}(\eta, \sigma, \Sigma, x, y, \pi') = (\eta, \sigma[y \mapsto^\pi \sigma(x)], \Sigma) \text{ when } x \in^\pi \sigma$$

$$\text{rename}(\eta, \sigma, \Sigma, x, y, \pi') = (\eta, \sigma, \Sigma[y \mapsto^\pi \Sigma(x)]) \text{ when } x \in^\pi \Sigma$$

1354 A.3.2 Perspective Management Rules.

$$\frac{p < n_1 \quad n_1, n_2 \text{ align to } n \quad \eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n_1)} s_1 \rightsquigarrow s'_1 \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi'}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\} \rightsquigarrow \text{split}(n_1, n_2)\{s'_1\}\{s_2\} \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi'} \text{ S-Split-Left}$$

$$\frac{p < n_1 \quad n_1, n_2 \text{ align to } n}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n)} \text{split}(n_1, n_2)\{\text{skip}\}\{s_2\} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-Split-Left-Done}$$

$$\frac{p \geq n_1 \quad p < n_1 + n_2 \quad n_1, n_2 \text{ align to } n \quad \eta, \sigma, \Sigma, t, b, p - n_1, \Psi, \Phi \vdash_m^{(h, n_2)} s_2 \rightsquigarrow s'_2 \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi'}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\} \rightsquigarrow \text{split}(n_1, n_2)\{s_1\}\{s'_2\} \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi'} \text{ S-Split-Right}$$

$$\frac{p \geq n_1 \quad p < n_1 + n_2 \quad n_1, n_2 \text{ align to } n}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n)} \text{split}(n_1, n_2)\{s_1\}\{\text{skip}\} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-Split-Right-Done}$$

$$\frac{p \geq n_1 + n_2 \quad n_1, n_2 \text{ align to } n}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-Split-None}$$

1373			
1374	$\eta, \sigma, \Sigma, t, b, t \text{ mod } T, \Psi, \Phi \vdash_m^{(\text{Thread}, T)} s \rightsquigarrow s' \dashv_m \eta', \sigma', \Sigma', \Psi', \Phi'$		
1375		S-Destruct-Block	
1376	$\eta, \sigma, \Sigma, t, b, 0, \Psi, \Phi \vdash_m^{(\text{Block}, 1)} \text{destruct in } s \rightsquigarrow \text{destruct in } s' \dashv_m \eta', \sigma', \Sigma', \Psi', \Phi'$		
1377			
1378	$\eta, \sigma, \Sigma, t, b, b \text{ mod } B, \Psi, \Phi \vdash_m^{(\text{Block}, B)} s \rightsquigarrow s' \dashv_m \eta', \sigma', \Sigma', \Psi', \Phi'$		S-Destruct-Grid
1379	$\eta, \sigma, \Sigma, t, b, 0, \Psi, \Phi \vdash_m^{(\text{Grid}, 1)} \text{destruct in } s \rightsquigarrow \text{destruct in } s' \dashv_m \eta', \sigma', \Sigma', \Psi', \Phi'$		
1380			
1381		S-Destruct-Done	
1382	$\eta, \sigma, \Sigma, t, b, 0, \Psi, \Phi \vdash_m^\pi \text{destruct in skip} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi$		
1383			
1384	$\eta, \sigma, \Sigma, t, b, p \text{ mod } n, \Psi, \Phi \vdash_m^{(h, n)} s \rightsquigarrow s' \dashv_m \eta', \sigma', \Sigma', \Psi', \Phi'$		S-Group
1385	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, q \cdot n)} \text{group } q \ s \rightsquigarrow \text{group } q \ s'; \dashv_m \eta', \sigma', \Sigma', \Psi', \Phi'$		
1386			
1387		S-Group-Done	
1388	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{group } q \ \text{skip} \rightsquigarrow \text{skip}; \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi$		

A.3.3 *Thread Synchronization.* We define a **size** operation on perspectives to compute the size of a perspective (the number of individual compute resources sharing it). The operation is defined thusly:

1391			
1392	$\text{size}(\text{Thread}, n) = n$		
1393		S-Sync-Wait-Done	
1394	$\text{size}(\text{Block}, n) = n \cdot T$		
1395		S-Sync-Wait-Spin	
1396			
1397			
1398	$\Psi(\psi)(p) = 0$		
1399	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{wait}_\psi \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi$	S-Sync-Dec	
1400			
1401	$\Psi(\psi)(p) \neq 0$		
1402	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{wait}_\psi \rightsquigarrow \text{wait}_\psi \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi$	S-Sync-Init-Zero	
1403			
1404	$\Psi' = \Psi(\psi)[p \mapsto \Psi(\psi)(p) - 1]$		
1405	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{dec}_\psi \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi$	S-Sync-Init-Nonzero	
1406			
1407	$\Psi(\psi)(p) = 0 \quad \Psi' = \Psi(\psi)[p \mapsto \text{size}(\pi)]$		
1408	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{init}_\psi \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi$	S-Async-Partition-Congr	
1409			
1410	$\Psi(\psi)(p) \neq 0$		
1411	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{init}_\psi \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi$	S-Async-Partition-Unwind	

A.3.4 *Asynchrony.*

1413	$\text{rename}(\eta, \sigma, \Sigma, x, y, (\text{Thread}, 1)), t, b, p, \Psi, \Phi \vdash_m^{(\text{Thread}, 1)} s \rightsquigarrow s' \dashv_m \eta', \sigma', \Sigma', \Psi', \Phi'$		
1414		S-Async-Partition-Congr	
1415	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(\text{Thread}, 1)} \text{async_partition}_\phi x \text{ into } y \text{ in } s \rightsquigarrow$		
1416	$\text{async_partition}_\phi x \text{ into } y \text{ in } s' \dashv_m \eta', \sigma', \Sigma', \Psi', \Phi'$		
1417			
1418	$\Phi = \Phi'[\phi \mapsto \Phi'(\phi) \cup \{s\}]$		S-Async-Partition-Unwind
1419	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(\text{Thread}, 1)} \text{async_partition}_\phi x \text{ into } y \text{ in skip} \rightsquigarrow$		
1420	$(\text{async_partition}_\phi x \text{ into } y \text{ in } s) \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi'$		
1421			

1422		$\Phi(\phi) = \emptyset$	
1423			S-Async-Partition-Done
1424	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(\text{Thread}, 1)} \text{async_partition}_\phi x \text{ into } y \text{ in } \text{skip} \rightsquigarrow$	$\text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi$	
1425			
1426			
1427			S-Async-Memcpy
1428	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(\text{Thread}, 1)} \text{async_memcpy}(x, y) \rightsquigarrow$	$\text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi[\phi \mapsto \Phi(\phi) \cup \{\text{memcpy}(x, y)\}]$	
1429			
1430			
1431	$(\eta', \sigma', \Sigma') = \text{update}(\eta, \sigma, \Sigma, x, \text{get}(\eta, \sigma, \Sigma, y))$		S-Memcpy
1432	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{memcpy}(x, y) \rightsquigarrow \text{skip} \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi$		
1433			
1434	A.3.5 <i>Variables and Memory.</i>		
1435	$\eta, \sigma, \Sigma \vdash_{\pi'}^\pi e \Downarrow v \quad \pi' \leq \pi$		S-Decl
1436	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi x : \tau @ \pi' := e \text{ in } s \rightsquigarrow s \dashv_m \eta[x \mapsto^{\pi'} v], \sigma, \Sigma, \Psi, \Phi$		
1437			
1438			S-Free
1439	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{free } n \dashv_{m-n} \eta, \sigma, \Sigma, \Psi, \Phi$		
1440			
1441			S-Alloc-Local
1442	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi x := \text{alloc Local } \tau n \text{ in } s \rightsquigarrow$		
1443	$s; \text{free } (n \cdot \text{size}(\tau)) \dashv_{m+n \cdot \text{size}(\tau)} \eta[x \mapsto^\pi \langle x, n \rangle], \sigma, \Sigma, \Psi, \Phi$		
1444			S-Alloc-Shared
1445	$\pi = (\text{Block}, 1)$		
1446	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi x := \text{alloc Shared } \tau n \text{ in } s \rightsquigarrow$		
1447	$s; \text{free } (n \cdot \text{size}(\tau)) \dashv_{m+n \cdot \text{size}(\tau)} \eta, \sigma[x \mapsto^\pi \langle x, n \rangle], \Sigma, \Psi, \Phi$		
1448			S-Alloc-Global
1449	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi x := \text{alloc Global } \tau n \text{ in } s \rightsquigarrow$		
1450	$s; \text{free } (n \cdot \text{size}(\tau)) \dashv_{m+n \cdot \text{size}(\tau)} \eta, \sigma[x \mapsto^\pi \langle x, n \rangle], \Psi, \Phi$		
1451			
1452	$x \in \pi' \quad \eta, \sigma, \Sigma \quad \eta, \sigma, \Sigma \vdash_{\pi'}^\pi e \Downarrow v \quad (\eta', \sigma', \Sigma') = \text{update}(\eta, \sigma, \Sigma, v) \quad \pi' \leq \pi$		S-Assn
1453	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi x = e \rightsquigarrow \text{skip} \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi$		
1454			
1455	$\eta, \sigma, \Sigma, t, b, p \vdash_\pi^\pi e_1 \Downarrow \langle l, n \rangle \quad i < n \quad \pi' \leq \pi$		S-Arr-Assn
1456	$\eta, \sigma, \Sigma, t, b, p \vdash_\pi^\pi e_2 \Downarrow i \quad x \in \pi' \quad \eta, \sigma, \Sigma$		
1457	$\eta, \sigma, \Sigma \vdash_{\pi'}^\pi e_3 \Downarrow v \quad (\eta', \sigma', \Sigma') = \text{update}(\eta, \sigma, l+i, v)$		
1458	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi e_1[e_2] = e_3 \rightsquigarrow \text{skip} \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi$		
1459			
1460			S-Partition
1461	$s' = s[(y + c \cdot p)/y] \quad (\eta', \sigma', \Sigma') = \text{rename}(\eta, \sigma, \Sigma, x, y, \pi/c)$		
1462	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{partition}_\psi x y c s \rightsquigarrow \text{init}_\psi; s'; \text{dec}_\psi; \text{wait}_\psi \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi$		
1463			S-Claim
1464	$(\eta', \sigma', \Sigma') = \text{rename}(\eta, \sigma, \Sigma, x, y, (h, n_1))$		
1465	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n_1+n_2)} \text{claim}_\psi x \text{ into } y \text{ at } n_1 \text{ in } s \rightsquigarrow$		
1466	$\text{init}_\psi; \text{split}(n_1, n_2) \{s'\} \{\text{skip}\}; \text{dec}_\psi; \text{wait}_\psi \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi$		
1467			
1468	$(\eta', \sigma', \Sigma') = \text{rename}(\eta, \sigma, \Sigma, x, y, \downarrow \pi)$		S-Lower
1469	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{lower}_\psi x \text{ into } y \text{ in } s \rightsquigarrow \text{init}_\psi; s; \text{dec}_\psi; \text{wait}_\psi \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi$		
1470			

1471 A.3.6 Control Flow.

$$\begin{array}{c}
 \frac{\eta, \sigma, \Sigma \vdash_{\pi}^{\pi} e \Downarrow \text{true}}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{if } e \text{ then } s_1 \text{ else } s_2 \rightsquigarrow s_1 \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-If-True} \\
 \\
 \frac{\eta, \sigma, \Sigma \vdash_{\pi}^{\pi} e \Downarrow \text{false}}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{if } e \text{ then } s_1 \text{ else } s_2 \rightsquigarrow s_2 \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-If-False} \\
 \\
 \frac{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{while } e \text{ do } s \rightsquigarrow \text{if } e \text{ then } (s; \text{while } e \text{ do } s) \text{ else skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} s_1 \rightsquigarrow s'_1 \dashv_{m'}^{\pi'} \eta', \sigma', \Sigma', \Psi', \Phi'} \text{ S-While} \\
 \\
 \frac{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} s_1 \rightsquigarrow s'_1 \dashv_{m'}^{\pi'} \eta', \sigma', \Sigma', \Psi', \Phi'}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} s_1; s_2 \rightsquigarrow s'_1; s_2 \dashv_{m'}^{\pi'} \eta', \sigma', \Sigma', \Psi', \Phi'} \text{ S-Seq-First} \\
 \\
 \frac{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{skip}; s_2 \rightsquigarrow s_2 \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{skip}; s_2 \rightsquigarrow s_2 \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-Seq-Done} \\
 \\
 \frac{\Sigma(f) = \{[x_1 : \tau_1, \dots, x_n : \tau_n], s\} \quad \sigma, \Sigma \vdash_{\pi}^{\pi} e_i \Downarrow v_i \quad m' \leq m}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} f(e_1, \dots, e_n) \rightsquigarrow s \dashv_m \eta[x_i \mapsto v_i], \sigma, \Sigma, \Psi, \Phi} \text{ S-Function-Call}
 \end{array}$$

1486 A.3.7 Expressions.

$$\begin{array}{c}
 \frac{\pi < (\text{Grid}, 1) \quad \pi' \leq \pi}{\eta, \sigma, \Sigma \vdash_{\pi}^{\pi} \text{partition_id} \Downarrow \pi / \pi' - 1} \text{ E-Partition-Id} \\
 \\
 \frac{}{\eta, \sigma, \Sigma \vdash_{\pi}^{\pi} x \Downarrow \text{get}(\eta, \sigma, \Sigma, x)} \text{ E-Var} \\
 \\
 \frac{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi} e_1 \Downarrow \langle l, n \rangle \quad \eta, \sigma, \Sigma \vdash_{\pi'}^{\pi} e_2 \Downarrow i \quad i < n \quad \pi' \leq \pi}{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi} e_1[e_2] \Downarrow \text{get}(\eta, \sigma, \Sigma, l + i)} \text{ E-Arr-Access} \\
 \\
 \frac{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi} \text{int} \Downarrow n}{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi} n \Downarrow n} \text{ E-Int} \quad \frac{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi} b \Downarrow b}{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi} b \Downarrow b} \text{ E-Bool} \\
 \\
 \frac{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi} e_1 \Downarrow v_1 \quad \eta, \sigma, \Sigma \vdash_{\pi'}^{\pi} e_2 \Downarrow v_2 \quad v = v_1 \text{ bop } v_2}{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi} e_1 \text{ bop } e_2 \Downarrow v} \text{ E-Bop} \\
 \\
 \frac{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi} e_1 \Downarrow v_1 \quad \eta, \sigma, \Sigma \vdash_{\pi'}^{\pi} e_2 \Downarrow v_2 \quad v = v_1 \text{ cmp } v_2}{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi} e_1 \text{ cmp } e_2 \Downarrow v} \text{ E-Cmp}
 \end{array}$$

1505 A.4 Theorems and Proofs

1506 Note that in this section we assume no out of bounds array accesses. In general Prism (and by
 1507 extension Bundl) makes no guarantees about array out of bounds.
 1508

1509 A.4.1 More definitions. As a premise to our type safety theorems, we need to assume we have a
 1510 well-typed environment, written $\Gamma \vdash \eta, \sigma, \Sigma$. We define what this means inductively
 1511

$$\begin{array}{c}
 \frac{}{\eta, \sigma, \Sigma \vdash n : \text{int}} \text{ V-Int} \quad \frac{}{\eta, \sigma, \Sigma \vdash b : \text{bool}} \text{ V-Bool} \quad \frac{}{\eta, \sigma, \Sigma \vdash f : \text{float}} \text{ V-Float} \\
 \\
 \frac{\forall i < n, \eta, \sigma, \Sigma \vdash \text{get}(\eta, \sigma, \Sigma, x + i) : \tau}{\eta, \sigma, \Sigma \vdash \langle x, n \rangle : \tau[]^l} \text{ V-Array} \quad \frac{\Gamma, x_i : \pi \tau_i \vdash_m^{\pi} s \quad \Gamma \vdash \cdot, \cdot, \text{fns} \Sigma}{\eta, \sigma, \Sigma \vdash \{x_i : \tau_i, s\} : \text{Fun}(x_i : \tau_i, \pi, m)} \text{ V-Function} \\
 \\
 \frac{}{\cdot \vdash \eta, \sigma, \Sigma} \text{ G-Empty} \quad \frac{\eta(x) =^{\pi} v \quad \eta, \sigma, \Sigma \vdash v : \text{int}}{\Gamma, x : \pi \text{ int} \vdash \eta, \sigma, \Sigma} \text{ G-Int}
 \end{array}$$

$$\begin{array}{c}
 \frac{\eta(x) =^\pi v \quad \eta, \sigma, \Sigma \vdash v : \text{bool}}{\Gamma, x :^\pi \text{bool} \vdash \eta, \sigma, \Sigma} \text{ G-Bool} \quad \frac{\eta(x) =^\pi v \quad \eta, \sigma, \Sigma \vdash v : \text{float}}{\Gamma, x :^\pi \text{float} \vdash \eta, \sigma, \Sigma} \text{ G-Float} \\
 \\
 \frac{\eta(x) =^\pi v \quad \eta, \sigma, \Sigma \vdash v : \tau[]}{\Gamma, x :^\pi \tau[] \text{Local} \vdash \eta, \sigma, \Sigma} \text{ G-Local} \quad \frac{\sigma(x) =^\pi v \quad \eta, \sigma, \Sigma \vdash v : \tau[]}{\Gamma, x :^\pi \tau[] \text{Shared} \vdash \eta, \sigma, \Sigma} \text{ G-Shared} \\
 \\
 \frac{\Sigma(x) =^\pi v \quad \eta, \sigma, \Sigma \vdash v : \tau[]}{\Gamma, x :^\pi \tau[] \text{Global} \vdash \eta, \sigma, \Sigma} \text{ G-Global} \quad \frac{\Sigma(x) =^\pi v \quad \eta, \sigma, \Sigma \vdash v : \text{Fun}(\Gamma', \pi, m)}{\Gamma, x :^\pi \text{Fun}(\Gamma', \pi, m) \vdash \eta, \sigma, \Sigma} \text{ G-Function}
 \end{array}$$

We can prove a couple simple lemmas about well-typed environments under operations like **rename**, **update**, and **get**.

LEMMA A.1. (*Well-typed get*) If $\Gamma \vdash \eta, \sigma, \Sigma$ and $x :^\pi \tau \in \Gamma$ then $\eta, \sigma, \Sigma \vdash \text{get}(\eta, \sigma, \Sigma, x) : \tau$.

LEMMA A.2. (*Well-typed rename*) If $\Gamma, x :^\pi \tau \vdash \eta, \sigma, \Sigma$ then $\Gamma, x :^\pi \tau, y :^\pi \tau' \vdash \text{rename}(\eta, \sigma, \Sigma, x, y, \tau')$.

LEMMA A.3. (*Well-typed update*) If $\Gamma \vdash \eta, \sigma, \Sigma$ and $\eta, \sigma, \Sigma \vdash v : \tau$ then $\Gamma, x :^\pi \tau \vdash \text{update}(\eta, \sigma, \Sigma, x, v)$.

We also define a well-formedness precondition on p with respect to π :

$$(h, n) \vdash p ::= p < n$$

We also define well-formedness for the async stack:

$$\Gamma \vdash \Phi ::= \forall \phi, s \in \Phi(\phi), \Gamma \vdash_m^{(\text{Thread}, 1)} s$$

A.5 Proofs

LEMMA A.4. (*Expression Safety*) If $\Gamma \vdash^\pi e : \tau$ and $\Gamma \vdash \eta, \sigma, \Sigma$ and $\pi \vdash p$ and $\pi' \leq \pi$, then there is some v such that $\eta, \sigma, \Sigma \vdash_{\pi'} e \Downarrow v$ and $v : \tau$.

PROOF. This proof proceeds by induction on the typing relation for expressions. Despite the fact that this property implies termination, we do not need a logical relation to prove it because the expression language is very simple.

The **T-Int**, **T-Float**, and **T-Bool** cases are trivial, using the rules **E-Int**, **E-Bool** and **E-Float** to compute values. In the case for **T-Partition-Id**, the π premises of the typing rules and our assumption that $\pi' \leq \pi$ match the premises of the evaluation rule, so this rule is simple as well.

The cases for **T-Bop** and **T-Cmp** follow directly from the inductive hypotheses, assuming a valid and correctly implemented set of binary operators and comparators.

The only interesting cases are **T-Arr-Access** and **T-Arr-Access-Shared**.

In both cases our inductive hypotheses and inversion give us that $\pi' \leq \pi$, and e_1 evaluates to a $\langle x, n \rangle$, and that all the values between x and $x + n$ in the appropriate environment are typed at τ . We also know that e_2 evaluates to an integer i . We assume that all array accesses are in bounds, so $i < l$, which is sufficient to use the **E-Arr-Access** rule to complete this case, and the proof. \square

LEMMA A.5. (*Expression Determinism*) If $\eta, \sigma, \Sigma, t, b, p \vdash_{\pi'}^\pi e \Downarrow v_1$ and $\eta, \sigma, \Sigma, t, b, p \vdash_{\pi'}^\pi e \Downarrow v_2$ then $v_1 = v_2$.

PROOF. Straightforward by induction on the semantic derivation. \square

LEMMA A.6. (*Expression Well-Typedness*) If $\Gamma \vdash_{\pi'}^\pi e : \tau$ and $\Gamma \vdash \eta, \sigma, \Sigma$ and $\pi \vdash p$, and $\eta, \sigma, \Sigma, t, b, p \vdash_{\pi'}^\pi e \Downarrow v$, then $v : \tau$.

PROOF. By our expression safety lemma our well-typed expression must evaluate to a well-typed value v' . By our determinism lemma v' must be the same as v , so v is well-typed. \square

LEMMA A.7. (*Statement Progress*) If $\Gamma \vdash \eta, \sigma, \Sigma$ and $\Gamma \vdash_m^\pi s$ and $\pi \vdash p$, then either s is **skip** or there is some s' such that $\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi s' \rightsquigarrow s' \vdash_m^{\pi'} \eta', \sigma', \Sigma, \Psi', \Phi'$

PROOF. This proceeds by induction on the typing derivation.

Perspective Management Rules.

- Case **T-Split**:

In this case we have by our assumption that $\pi \vdash p$ that $p < n$. We also have that n_1, n_2 align to n , so $n_1 + n_2 \leq n$. There are three cases to consider, then: when $p < n_1$, when $p \geq n_1$ and $p < n_1 + n_2$, and when $p \geq n_1 + n_2$.

In the first case, we have by our inductive hypothesis that s is either **skip** or that it can step in an (h, n_1) context. In the former case we can use the **S-Split-Left-Done** rule and in the latter we can use the **S-Split-Left rule**.

The second case is almost symmetric. The only additional work we have to do is to argue that $(h, n_2) \vdash p - n_1$, or equivalently that $p - n_1 < n_2$. This, however, is immediate from our assumption that $p < n_1 + n_2$.

In the last case, we just use the **S-Split-None** rule to step to **skip**.

- Case **T-Destruct**

By our inductive hypothesis, we know that s can step at $\downarrow \pi$ if $\downarrow \pi \vdash p$. The \downarrow operation is only defined at **(Block, 1)** or **(Grid, 1)**, so we only need to consider the cases where π is one of those.

In the former case p becomes $t \text{ mod } T$ while $\downarrow \pi$ is **(Thread, T)**. $t \text{ mod } T < T$ for any t so this satisfies the requirement that $\pi \vdash p$, which lets us use our inductive hypothesis: s is either **skip** or can step. If it can step, we can use this to satisfy the premise of **S-Destruct-Block** to step in this case. If it is **skip**, then we use the rule **S-Destruct-Done** to step instead.

The latter case is the same, except using the fact that $b \text{ mod } B < B$ and the **S-Destruct-Grid** rule.

- Case **T-Group**

In this case we have by assumption that $(h, q \cdot n) \vdash p$, i.e., that $p < q \cdot n$.

In this case we have our IH that if $(h, n) \vdash p'$ for some p' , then s is either **skip** or steps with (h, n) perspective with p' as our perspective ID.

We choose p' to be $p \text{ mod } n$. This is always $< n$, so $(h, n) \vdash p'$. This lets us use our IH to get that s is either **skip** (in which case we can use the **S-Group-Done** rule to step) or itself steps, which lets us use the **S-Group** rule to step.

Thread Synchronization Rules.

- Case **T-Sync-Wait**

$\Psi(\psi)(p)$ is either zero or it is not. In the former case we use the **S-Sync-Wait-Done** rule and in the latter we use **S-Sync-Wait-Spin**.

- Case **T-Sync-Dec**

We use the **S-Sync-Dec** rule to step.

- Case **T-Sync-Init**

We use the **S-Sync-Init-Zero** or **S-Sync-Init-Nonzero** rules depending on whether $\Psi(\psi)(p)$ is zero or not.

Asynchrony Rules.

- Case **T-Async-Partition**

In this case, we have via our IH that if $\Gamma, y : (\text{Thread}, 1) \text{ async } \tau[]^l \vdash \eta', \sigma', \Sigma'$, then we can either step s with **(Thread, 1)** perspective under environments η', σ' , and Σ' , or s is **skip**.

1618 We have by assumption that $\Gamma, x : (\text{Thread}, 1) \text{ async } \tau[]^l \vdash \eta, \sigma, \Sigma$. By our environment renaming
 1619 lemma, this gives us what we need to use our IH with $(\eta', \sigma', \Sigma')$ as **rename** $(\eta, \sigma, \Sigma, x, y, (\text{Thread}, 1))$.
 1620 Thus s either steps or is **skip**. In the former case we can step with **S-Async-Partition-Congr**,
 1621 and in the latter we can use either **S-Async-Partition-Unwind** or **S-Async-Partition-Done**
 1622 depending on whether $\Phi(\phi)$ is empty or not.

- Case **T-Async-Memcpy**

1623 Immediate via use of the **S-Async-Memcpy** rule.

- Case **T-Memcpy**

1624 Immediate via use of the **S-Memcpy** rule.

1625 *Memory Rules.*

- Case **T-Decl**

1626 Via our lemma about expression type safety and our hypothesis that e is well-typed, we obtain
 1627 the premises necessary to use the **S-Decl** rule to step.

- Case **T-Arr-Assn**

1628 Each of e_1, e_2 and e_3 must evaluate to a well-typed value by the expression type safety lemma.
 1629 In particular, both e_1 evaluates to some $\langle l, n \rangle$ and e_2 evaluates to some i . We assume all array
 1630 accesses are in bounds, so this is sufficient to use the **S-Arr-Assn** rule to step.

- Case **T-Arr-Assn-Shared**

1631 Same as previous case.

- Case **T-Free**

1632 Trivial via the **S-Free** rule.

- Case **T-Partition**

1633 Trivial via the **S-Partition** rule.

- Case **T-Claim**

1634 Trivial via the **S-Claim** rule.

- Case **T-Lower**

1635 Trivial via the **S-Lower** rule.

- Case **T-Alloc**

1636 We assume that l is not **Shared**, so we can use the **S-Alloc-Local** or **S-Alloc-Global** rule,
 1637 depending on whether l is **Local** or **Global**.

- Case **T-Alloc-Shared**

1638 Trivial via the **S-Alloc-Shared** rule.

1639 *Control Rules.*

- Case **T-Skip**

1640 Trivial

- Case **T-While**

1641 Trivial, all while loops step via the **S-While** rule

- Case **T-If**

1642 By our proof of expression type safety, the expression e steps to either the boolean value true
 1643 or false. We can thus use either the **S-If-True** or **S-If-False** rules to step.

- Case **T-Seq**

1644 In this case we know by our IH that s_1 is either **skip** or can step. In the former case we use
 1645 the **S-Seq-Done** rule and in the latter we use the **S-Seq-First** rule.

- Case **T-Function-Call**

1646 In this case we know by our expression safety lemma that each of the arguments will evaluate to
 1647 a well-typed value. We also have by assumption that f has a function type, which by inversion

on the **V-Function** rule tells us that it is a closure type. Additionally our assumption that $\Gamma \vdash \eta, \sigma, \Sigma$ tell us that Σ contains f at the same type that Γ does. These premises are sufficient to use the **S-Function-Call** rule.

□

LEMMA A.8. (Statement Preservation) If $\Gamma \vdash \eta, \sigma, \Sigma$ and $\Gamma \vdash_m^\pi s$ and $\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_{m'}^\pi s \rightsquigarrow s' \dashv_{m''} \eta', \sigma', \Sigma, \Psi', \Phi'$ and $\pi \vdash p$ and $m \geq m'$ and $\Gamma \vdash \Phi$, then there is some Γ' such that $\Gamma \subseteq \Gamma'$ and $\Gamma' \vdash_m^\pi s'$ and $\Gamma' \vdash \eta', \sigma', \Sigma'$ and $m \geq m''$ and $\Gamma' \vdash \Phi'$.

PROOF. We proceed by induction on the derivation of $\Gamma \vdash_m^\pi s$.

Perspective Management Rules.

- Case **T-Split**

In this case we have by assumption that $(h, n) \vdash p$, $\Gamma \vdash \eta, \sigma, \Sigma$, and $\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_{m'}^\pi \text{split}(n_1, n_2)\{s_1\}\{s_2\} \rightsquigarrow s' \dashv_{m''} \eta', \sigma', \Sigma, \Psi', \Phi'$. Our inductive hypotheses give us that if for any p , if $(h, n_1) \vdash p$ and s_1 steps with perspective ID p , or if $(h, n_2) \vdash p$ and s_2 steps with perspective ID p then their results are well typed.

By inversion on our semantic derivation, we are in one of 5 cases.

In the **S-Split-Left** case $p < n_1$ and s_1 steps to s'_1 . This is sufficient to tell us that s'_1 is well-typed and the output environments of that relation η', σ' , and Σ' are all well-typed by $\Gamma' \supseteq \Gamma$, and that the memory is properly bounded by the typing rules.

We can thus use the **T-Split-Left** rule to conclude that the result of this case is well-typed.

The **S-Split-Left-Done** rule is trivial via the **T-Skip** rule.

The **Right** cases are symmetric, with the observation that when $p \geq n_1$ and $p < n_1 + n_2$ then $p - n_1 < n_2$.

The last **S-Split-None** rule is trivial via the **T-Skip** rule.

- Case **T-Destruct**

In this case we have by assumption that $\downarrow \pi$ is defined, so π is either **(Block,1)** or **(Grid,1)**. We also assume that $\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_{m'}^\pi \text{destruct in } s \rightsquigarrow s' \dashv_{m''} \eta', \sigma', \Sigma, \Psi', \Phi'$. We also have by our inductive hypothesis that for any p such that $\downarrow \pi \vdash p$ and s'' such that s steps to s'' at p , then that step preserved well-typedness.

By inversion on the step relation, we are in one of three cases.

If the rule used **S-Destruct-Block**, then we know that s steps to s'' and p is $t \bmod T$. $(\text{Thread}, T) \vdash t \bmod T$ for any t , so we can use our inductive hypothesis to conclude that the s'' stepped to by s is well-typed, as are its environments and memory usage. The **T-Destruct** rule then gives us our desired goal.

The **S-Destruct-Grid** case proceeds similarly, while the **S-Destruct-Done** case is trivial.

- Case **T-Group**

In this case we have by assumption that

$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_{m'}^{(h, q \cdot n)} \text{group } q \rightsquigarrow s' \dashv_{m''} \eta', \sigma', \Sigma, \Psi', \Phi'$. We have by our inductive hypothesis that s steps to s'' at some perspective ID p' and $(h, n) \vdash p'$, then s'' is well typed, as are the other outputs of that step.

By inversion on the step relation, we are in one of two cases. The **S-Group-Done** case is trivial, so we shall focus on the **S-Group** case. In this case we have that s steps to s'' at perspective ID $p \bmod n$. It is always the case that $(h, n) \vdash p \bmod n$ for any p , so we can use our inductive hypothesis to conclude that s'' is well-typed, as are its output environments and memory usage. From there, it is a simple application of the **T-Group** rule to conclude that **group** q s' is well-typed, and to finish the case.

Thread Synchronization Rules. These rules are all trivial: with one exception all thread synchronization primitives step to skip without changing environment or memory, and are thus obviously well-typed. **S-Sync-Wait-Spin** does not produce skip, but it steps to the same statement as we already assumed typechecks in the premise of the lemma, so is straightforward nonetheless.

If we wanted to say something about deadlock freedom we'd have more work here, but we aren't doing that, so these rules are easy.

Asynchrony Rules.

- Case **T-Async-Partition** In this case we have that s is well-typed in a context where x has been renamed into y , with **(Thread,1)** perspective. We also have that $\Gamma, x:(\text{Thread},1) \vdash \tau[]^l \vdash \eta, \sigma, \Sigma$ and $\Gamma, x:(\text{Thread},1) \vdash \tau[]^l \Phi$.

By inversion, we are in one of three cases.

In the **S-Async-Partition-Done** case, we are done.

In the **S-Async-Partition-Congr** case, our inductive hypothesis gives us that there is some Γ' such that $\Gamma, y:(\text{Thread},1) \text{ async } \tau[]^l \subseteq \Gamma'$ and $\Gamma' \vdash \Phi'$ and $\Gamma' \vdash \text{rename}(\eta, \sigma, \Sigma, x, y, (\text{Thread},1))$ via our well-typed renaming lemma. This lets us use the **T-Async-Partition** rule to check this case, with a choice of Γ' as $\Gamma, x:(\text{Thread},1) \vdash \tau[]^l$.

In the **S-Async-Partition-Unwind** case, our assumption that Φ is well-typed tells us that $\Gamma, y:(\text{Thread},1) \vdash \{\ (\text{Thread},1)_m s$. Thus, we can use the **T-Async-Partition** rule to type this case.

- Case **T-Async-Memcpy**

In this case the statement and environment typing are trivial, we need only to show that the `async` stack remains well typed.

In this case we have that x and y have the same type at **(Thread,1)**. This is sufficient for us to check $x = y$ at **(Thread,1)**, meaning that adding that instruction to the stack maintains its well-typedness.

- Case **T-Memcpy**

Immediate via use of the **S-Memcpy** rule. We just need to show that the environment remains well typed, which we know via our lemmas about **update** and **get**.

Memory Rules.

- Case **T-Decl**

By inversion, we are using **S-Decl** rule for evaluation. Our well-typed expression lemma gives us that v is well-typed, so it follows from our assumptions and our lemmas about extending environments that the extended η and s are well-typed by $\Gamma, x:\pi \tau$.

- Case **T-Free**

Trivial.

- Case **T-Alloc**

By inversion we are either in the **S-Alloc-Local** or **S-Alloc-Global** rules. In either case, we assume that $\Gamma, x:\pi \tau[]^l$ checks s , meaning we can use our extended environment lemmas and the **T-Seq** and **T-Free** rules to check these cases.

- Case **T-Alloc-Shared**

Same as previous case.

- Case **T-Partition**

In this case we have by assumption that $\Gamma, y:(h,n/c) \tau[]^l$'s and l is not local and c divides n . By inversion on our step relation we must be in the **S-Partition** case, so we have $\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash \pi_m \text{ partition} \psi \text{ y c s} \sim \text{init}_\psi; s'; \text{dec}_\psi; \text{wait}_\psi \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi$ where $s' = s[(y + c \cdot p)/y]$ and $(\eta', \sigma', \Sigma') = \text{rename}(\eta, \sigma, \Sigma, x, y, \pi/c)$. Via our well-typed renaming lemma contexts we know that we can check the renamed environments in the extended environment $\Gamma, y:\pi/c \tau[]^l, x:\pi$

$\tau[]^l$, and this context is also sufficient to check s' (via a substitution-preserves-typing lemma that is obvious). Using the **T-Skip** rule this is exactly what we need to show to complete this case, as the thread sync primitives check trivially via their typing rules.

- Case **T-Claim**

Essentially the same as **T-Partition**.

- Case **T-Lower**

Essentially the same as **T-Partition**.

Control Rules.

- Case **T-If**

We have by assumption that **if** e **then** s_1 **else** s_2 steps to some s' , and by inversion we know that either e evaluates to **true** and s' is s_1 , or e evaluates to **false** and s' is s_2 .

In either case, our inductive hypotheses is sufficient to tell us that these are well-typed. In particular, in both cases our IHs tell us that the amount of memory used by stepping each branch of the **if** is less than the amount of memory computed by the type system for each branch. Because the whole **if** expression checks using the greater of the memory usage of m_1 or m_2 (i.e., the memory usage on each branch), the resulting usage for the whole conditional is also bounded by the type system.

- Case **T-Skip**

Trivial: **skip** does not step

- Case **T-Seq**

In this case we have that s_1 and s_2 are both well-typed (with m_1 memory and m_2 memory respectively), and $m = \mathbf{max}(m_1, m_2)$. We also have by inversion that s_1 either steps to **skip** or s'_1 , and our inductive hypothesis tells us that s'_1 is well-typed.

In the former case we can use the **S-Seq-Done** rule to trivially finish the case. In the latter, our IH allows us to finish the case, since m_1 is always $\leq \mathbf{max}(m_1, m_2)$

- Case **T-While**

By inversion, we have that

$$\begin{aligned} & \eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \mathbf{while } e \mathbf{ do } s \\ & \rightsquigarrow \mathbf{if } e \mathbf{ then } (s; \mathbf{while } e \mathbf{ do } s) \mathbf{ else skip } \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi. \end{aligned}$$

We also have by assumption that e and s are well-typed. With this information, through use of the **T-If**, **T-Seq**, **T-While**, and **T-Skip** rules, we can conclude that the result of this rule is also well-typed.

- Case **T-Function-Call**

By inversion we have that

$$\begin{aligned} & \eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi f(e_1, \dots, e_n) \\ & \rightsquigarrow \mathbf{call } s \mathbf{ with } (x_i : \tau_i \mapsto v_i) @ \{\eta, \sigma, \Sigma, m'\} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi, \end{aligned}$$

and also that $\Sigma(f) = \{[x_1 : \tau_1, \dots, x_n : \tau_n], s\}$, and $\eta, \sigma, \Sigma, t, b, p \vdash^\pi e_i \Downarrow v_i$, and $m' \leq m$.

We also have from the premises of our case that $f : \pi \mathbf{Fun}(x_1 : \tau_1, \dots, x_n : \tau_n, \pi, m') \in \Gamma$, and $\Gamma \vdash^\pi e_i : \tau_i$, and $m' \leq m$.

Our lemma for expression well-typedness tells us that that each v_i is a well-typed value, and our assumption that $\Gamma \vdash \eta, \sigma, \Sigma$ tells us that $\Sigma(f)$ is a well-typed function and thus that $\mathbf{fns } \Gamma, x_i : \tau_i \vdash_m^\pi s$. We can take the union of this with Γ to produce $\Gamma, x_i : \tau_i \vdash_m^\pi s$, which clearly checks s and the output environments.

□

```

1814
1815     @ccuda("global")
1816     @requires(grid[1], block[1], thread[1], smem=49000)
1817     def my_sgemm_kernel_1(M : int @ grid[1],
1818                           N : int @ grid[1],
1819                           K : int @ grid[1],
1820                           alpha : float @ grid[1],
1821                           A: ptr(const(float)) @ grid[1],
1822                           B: ptr(const(float)) @ grid[1],
1823                           beta : float @ grid[1],
1824                           C: ptr(float) @ grid[1],
1825                           block_size : int @ grid[1]):
1826         with group(grid[1]):
1827             bid : int @ block[1] = id()
1828             with partition(C, dimension=block[1], offset=0) as C_blk:
1829                 with group(block[1]):
1830                     tid : int @ thread[1] = bid * block_size + id()
1831                     with partition(C_blk, dimension=thread[1], offset=0) as C_thrd:
1832                         x : const(int) @ thread[1] = tid / N
1833                         y : const(int) @ thread[1] = tid % N
1834                         with group(thread[1]):
1835                             tmp : float @ thread[1] = 0
1836                             for i in range(0, K, 1):
1837                                 tmp += A[x * K + i] * B[i * N + y]
1838                         C_thrd[x * N + y] = alpha * tmp + beta * C[x * N + y]
1839
1840     return

```

1829

1830

1831

```

1832     @ccuda("global")
1833     @requires(grid[1], block[1], thread[1], smem=49000)
1834     def my_sgemm_kernel_2(M : int @ grid[1],
1835                           N : int @ grid[1],
1836                           K : int @ grid[1],
1837                           alpha : float @ grid[1],
1838                           A: ptr(const(float)) @ grid[1],
1839                           B: ptr(const(float)) @ grid[1],
1840                           beta : float @ grid[1],
1841                           C: ptr(float) @ grid[1],
1842                           tile_size : int @ grid[1]):
1843         with group(grid[1]):
1844             global_tid: int @ thread[1] = id()
1845             total_elements : const(int) @ grid[1] = M * N
1846             num_tiles_per_row : const(int) @ grid[1] = (N + tile_size - 1) / tile_size
1847             tile_id : const(int) @ thread[1] = global_tid / (tile_size * tile_size)
1848             tile_row : const(int) @ thread[1] = tile_id / num_tiles_per_row
1849             tile_col : const(int) @ thread[1] = tile_id % num_tiles_per_row
1850             local_id : const(int) @ thread[1] = global_tid % (tile_size * tile_size)
1851             local_row : const(int) @ thread[1] = local_id / tile_size
1852             local_col : const(int) @ thread[1] = local_id % tile_size
1853             cRow : const(int) @ thread[1] = tile_row * tile_size + local_row
1854             cCol : const(int) @ thread[1] = tile_col * tile_size + local_col
1855             with partition(C, dimension=thread[1], offset=cRow * N + cCol) as C_thrd:
1856                 with group(thread[1]):
1857                     if global_tid < total_elements:
1858                         if cRow < M and cCol < N:
1859                             tmp : float @ thread[1] = 0
1860                             for i in range(0, K, 1):
1861                                 tmp += A[cRow * K + i] * B[i * N + cCol]
1862                         C_thrd[0] = alpha * tmp + beta * C_thrd[0]
1863
1864     return

```

1851

1852

1853

B Programs from Evaluation

We reproduce programs that we wrote for Section 6 in this section.

1856

1857

1858

1859

1860

1861

1862


```

1912
1913 1 __cuda("global")
1914 2 __attribute__((__launch_bounds__(((BM * BN) / (TM * TN), 1))
1915 3 @template([("BM", c_int),
1916 4   ("BN", c_int),
1917 5   ("BK", c_int),
1918 6   ("TM", c_int),
1919 7   ("TN", c_int)])
1920 8 requires(grid[1], block[1], thread[1], smem=49000)
1921 9 def my_sgemm_kernel_4C(M : int @ grid[1],
1922 10   N : int @ grid[1],
1923 11   K : int @ grid[1],
1924 12   alpha : float @ grid[1],
1925 13   A: ptr(float) @ grid[1],
1926 14   B: ptr(float) @ grid[1],
1927 15   beta : float @ grid[1],
1928 16   C: ptr(float) @ grid[1],
1929 17   num_blocks_N : int @ grid[1]):
1930 18   with group(grid[1]):
1931 19     cRow : const(int) @ block[1] = id() / num_blocks_N
1932 20     cCol : const(int) @ block[1] = id() % num_blocks_N
1933 21     totalResultsBlocktile: const(uint) @ grid[1] = BM * BN
1934 22     numThreadsBlocktile: const(uint) @ grid[1] = totalResultsBlocktile / (TM * TN)
1935 23     with partition(C, dimension=block[1], offset = cRow * BM * N + cCol * BN) as C_blk:
1936 24       with partition(A, dimension=block[1], offset = cRow * BM * K) as A_blk:
1937 25         with partition(B, dimension=block[1], offset = cCol * BN) as B_blk:
1938 26           with group(block[1]):
1939 27             As : shared(float[128 * 8]) @ block[1]
1940 28             Bs : shared(float[128 * 8]) @ block[1]
1941 29             threadRow : const(int) @ thread[1] = id() / (BN / TN)
1942 30             threadCol : const(int) @ thread[1] = id() % (BN / TN)
1943 31
1944 32             innerRowA : const(int) @ thread[1] = id() / (BK / 4)
1945 33             innerColA : const(int) @ thread[1] = id() % (BK / 4)
1946 34             strideA : const(int) @ thread[1] = numThreadsBlocktile / BK
1947 35
1948 36             innerRowB : const(int) @ thread[1] = id() / (BN / 4)
1949 37             innerColB : const(int) @ thread[1] = id() % (BN / 4)
1950 38             strideB : const(int) @ thread[1] = numThreadsBlocktile / BN
1951 39
1952 40             threadResults : float[64] @ thread[1] = { 0 }
1953 41             regM : float[8] @ thread[1] = { 0 } # Need to go add the {}
1954 42             regN : float[8] @ thread[1] = { 0 }
1955 43
1956 44             i : int @ thread[1] = 0
1957 45
1958 46             for bkIdx in range(0, K, BK):
1959
1960               with partition(As, dimension=thread[1], offset = 0) as As_thrd:
1961                 with partition(Bs, dimension=thread[1], offset = 0) as Bs_thrd:
1962                   with group(thread[1]):
1963
1964                     tmp : float4 @ thread[1] = float4_cast(A_blk[i * BK + innerRowA * K + innerColA * 4])
1965                     As_thrd[(innerColA * 4 + 0) * BM + innerRowA] = tmp.x
1966                     As_thrd[(innerColA * 4 + 1) * BM + innerRowA] = tmp.y
1967                     As_thrd[(innerColA * 4 + 2) * BM + innerRowA] = tmp.z
1968                     As_thrd[(innerColA * 4 + 3) * BM + innerRowA] = tmp.w
1969
1970                     i += 1
1971
1972                     for dotIdx in range(0, BK, 1):
1973                       for i in range(0, TN, 1):
1974                         regM[i] = As[dotIdx * BM + threadRow * TM + i]
1975                         for i in range(0, TN, 1):
1976                           regN[i] = Bs[dotIdx * BN + threadCol * TN + i]
1977                           for resIdxM in range(0, TM, 1):
1978                             for resIdxN in range(0, TN, 1):
1979                               threadResults[resIdxM * TN + resIdxN] += regM[resIdxM] * regN[resIdxN]
1980
1981                     for resIdxM in range(0, TM, 1):
1982                       for resIdxN in range(0, TN, 4):
1983                         with partition(C_blk, dimension=thread[1], offset = (threadRow * TM + resIdxM) * N + threadCol * TN +
1984                           ↪ resIdxN) as C_thrd:
1985                           with group(thread[1]):
1986                             temp : float4 = float4_cast(C_thrd[0])
1987                             temp.x = alpha * threadResults[resIdxM * TN + resIdxN] + beta * temp.x
1988                             temp.y = alpha * threadResults[resIdxM * TN + resIdxN + 1] + beta * temp.y
1989                             temp.z = alpha * threadResults[resIdxM * TN + resIdxN + 2] + beta * temp.z
1990                             temp.w = alpha * threadResults[resIdxM * TN + resIdxN + 3] + beta * temp.w
1991                             float4_cast[C_thrd[0]] = temp
1992
1993             return

```

```

1961 1
1962
1963
1964
1965 1 @ccuda("device")
1966 2 @requires(block[1], thread[32])
1967 3 def block_load(input : ptr(const(int)) @ block[1],
1968 4           output : ptr(int) @ thread[1],
1969 5           items_per_thread : int @ block[1]):
1970 6     with group(block[1]):
1971 7         tid : int @ thread[32] = id()
1972 8         with partition(input, dimension=thread[32], offset = tid * items_per_thread) as input_thrd:
1973 9             with group(thread[32]):
1974 10                warp_load(input_thrd, output, items_per_thread)
1975 11
1976 12 return
1977
1978
1979
1980
1981
1982
1983
1984 1 @ccuda("device")
1985 2 @requires(thread[32])
1986 3 def warp_load(input : ptr(const(int)) @ thread[32],
1987 4           output : ptr(int) @ thread[1],
1988 5           items_per_thread : int @ thread[32]):
1989 6     with group(thread[32]):
1990 7         tid : int @ thread[1] = id()
1991 8         with partition(input, dimension=thread[1], offset = tid * items_per_thread) as input_thrd:
1992 9             with group(thread[1]):
1993 10                thread_load(input_thrd, output, items_per_thread)
1994 11
1995
1996
1997
1998
1999
2000
2001
2002
2003
2004
2005
2006
2007
2008
2009

```

```
2010
2011 1 @ccuda("device")
2012 2 #requires(thread[1])
2013 3 def thread_store(input : ptr(const(int)) @ thread[1],
2014 4         output : ptr(int) @ thread[1],
2015 5         items_per_thread : int @ thread[1]):
2016 6     with group(thread[1]):
2017 7         for i in range(0, items_per_thread, 1):
2018 8             output[i] = input[i]
2019 9
2020
2021
2022
2023
2024
2025
2026
2027
2028
2029
2030
2031
2032
2033
2034
2035
2036
2037
2038
2039
2040
2041
2042
2043
2044
2045
2046
2047
2048
2049
2050
2051
2052
2053
2054
2055
2056
2057
2058
```

```

259
260     1 #include("global")
261     2 #attr __launch_bounds__((128*3))
262     3 #requires(grid[1], block[1], thread[1], warp[12], smem=227000)
263     4 def cutlass_my_h100_match_no_tail(
264     5     M: int @ grid[1],
265     6     N: int @ grid[1],
266     7     K: int @ grid[1],
267     8     C: ptr(bf16) @ grid[1],
268     9     tensorMapA: const(CutensorMap) @ grid[1],
269    10    tensorMapB: const(CutensorMap) @ grid[1]):
270
271     11
272     12     with group(grid[1]):
273     13         BM : constexpr(int) @ grid[1] = 128
274     14         BN : constexpr(int) @ grid[1] = 256
275     15         BK : constexpr(int) @ grid[1] = 64
276     16         NUM_THREADS : constexpr(int) @ grid[1] = 128*3
277     17         QSIZE : constexpr(int) @ grid[1] = 4
278     18         NUM_SM : constexpr(int) @ grid[1] = 128
279
280     19         WGMMA_M : constexpr(int) @ grid[1] = 64
281     20         WGMMA_K : constexpr(int) @ grid[1] = 16
282     21         WGMMA_N : constexpr(int) @ grid[1] = BN
283     22         num_consumers : constexpr(int) @ grid[1] = (NUM_THREADS / 128) - 1
284     23         B_WG_M : constexpr(int) @ grid[1] = BM / num_consumers
285
286     24
287     25         TM : constexpr(int) @ grid[1] = 16
288     26         TN : constexpr(int) @ grid[1] = 8
289     27         schedule_block : int @ block[1] = id()
290
291     28     with partition(C, offset=0, dimension=block[1]) as block_C:
292     29         with group(block[1]):
293
294     30             # Allocate SA
295             # SA[QSIZE][BK*BK]
296             sA_slot0 : shared(bf16[64* 128], align=128) @ block[1]
297             sA_slot1 : shared(bf16[64* 128]) @ block[1]
298             sA_slot2 : shared(bf16[64* 128]) @ block[1]
299             sA_slot3 : shared(bf16[64* 128]) @ block[1]
300
301             # Allocate SB
302             # sB[QSIZE][BK*BK]
303             sB_slot0 : shared(bf16[64* 256], align=128) @ block[1]
304             sB_slot1 : shared(bf16[64* 256]) @ block[1]
305             sB_slot2 : shared(bf16[64* 256]) @ block[1]
306             sB_slot3 : shared(bf16[64* 256]) @ block[1]
307
308
309             num_blocks_k : const(int) @ block[1] = K / BK
310             wg_idx : int @ warp[4] = id()
311             blk_thrd_id : int @ thread[1] = id()
312
313
314             tid : int @ thread[1] = id() % 128
315             is_producer : bool @ warp[4] = wg_idx == 0
316
317             num_block_m : int @ block[1] = 0
318             num_block_n : int @ block[1] = 0
319
320             with group(warp[4]):
321                 if (is_producer):
322                     pass
323                 else:
324                     wg_idx = wg_idx - 1
325
326             schedule_it : int @ block[1] = 0
327             total_blocks_m : int @ block[1] = (((M + BM) - 1) / BM)
328             total_blocks_n : int @ block[1] = (((N + BN) - 1) / BN)
329             unsafe("assert(CEIL_DIV(M, BM)%TM == 0 && total_blocks_n%TN == 0);")
330
331             while(true):
332                 num : int @ block[1] = ((schedule_it * NUM_SM) + schedule_block)
333                 if (num >= (total_blocks_m * total_blocks_n)):
334                     break
335                 cur_tile : int @ block[1] = (num / (TM * TN))
336                 cur_tile_pos : int @ block[1] = (num % (TM * TN))
337                 num_block_m = (TM * (cur_tile / (total_blocks_n / TN)))
338                 num_block_n = (TN * (cur_tile % (total_blocks_n / TN)))
339                 num_block_m += (cur_tile_pos / TN)
340                 num_block_n += (cur_tile_pos % TN)
341                 schedule_it += 1
342                 d : c_float[1][16][8] @ thread[1]
343
344                 with group(thread[1]):
345                     d_ptr : ptr(float) @ thread[1]
346                     # TODO: ANNOYING TO NOT HAVE TO SET EACH INDIVIDUAL TO 0.
347                     unsafe("d_ptr = (float *)0;")
348                     memset_wrapper(d_ptr, 0, 1 * 16 * 8 * 4)
349
350
351                 # Now remove these registers...
352                 with group(warp[4]):
353                     if (is_producer):
354                         num_regs : constexpr(int) @ warp[4] = 24 if (num_consumers <= 2) else 32

```