

1 Modular GPU Programming with Typed Privileges

2 ANONYMOUS AUTHOR(S)

3 Abstract

4 1 Introduction

5 CUDA [11] is a low-level, imperative programming language for NVIDIA GPUs. These GPUs are organized into a hierarchy of compute resources: threads, blocks, and a grid. *Threads* are the basic units of sequential execution; *blocks* are groups of threads that can cooperate through shared scratchpad memory; and the *grid* is the full collection of blocks launched for a GPU kernel. CUDA programs are written from the perspective of a single thread.

6 Although programs are written from the perspective of a single thread, some operations are only semantically valid when executed *collectively* by a group of threads. For example, the `__syncthreads()` primitive, which synchronizes all threads within a block, causes the machine to deadlock if executed by a subset of threads within that block. Similarly, tensor core instructions, which are hardware-accelerated matrix-multiply operations, are only well-defined when invoked by a group of 32-threads whose starting index is 32-aligned. Therefore, these operations require programmers to carefully marshal compute resources and coordinate which threads execute which lines of code. As a result, collective operations *break the illusion* of threads executing independently.

7 The conceptual clash between regular statements (executed by a single thread), and collective operations (executed by a group of threads) hinders compositional reasoning. In CUDA, function calls, like all other statements, are per-thread operations. The functions, however, may contain code that executes a collective operation, creating a semantic gap between the per-thread syntax for a function's invocation and the underlying semantics of the function's execution.

8 This gap between the syntax and semantics of a CUDA program creates tension between abstraction and correctness. This tension is apparent even in widely used CUDA libraries that package common functionality via function interfaces. Consider the snippet of documentation presented in Figure 1 taken directly from CUB [10], a library from NVIDIA that provides wrappers around parallel primitives like scan, sort, and histogram, among others. The documentation describes a function called `BlockReduce`, in which all threads in a block collaboratively apply a reduction operation, such as a maximum or prefix sum, over an array.

9 The CUB documentation attempts to make several implicit assumptions of `BlockReduce`'s implementation clear to the programmer. First, the function is only well-defined when invoked by an entire block's worth of threads. Second, the function accesses shared memory, a scratchpad space visible to all threads in a block. Users wishing to reuse this storage must first ensure all the block's threads have finished using it by inserting a `__syncthreads()` barrier. Whether this synchronization is required cannot be determined from the function alone; it depends on how the storage is used afterward.

10 In effect, CUB attempts to informally retrofit CUDA with information about the compute and memory requirements of collective operations. By prefixing functions with identifiers such as `Block`, CUB effectively creates "namespaces" for different kinds of these operations. Even then, this organization is purely conventional: it encodes hierarchy through naming and cannot be enforced by the compiler. Ultimately, correctness depends not only on the user carefully reading and interpreting the documentation, but also on the library implementer upholding the invariants of each namespace.

11 Previous attempts to resolve the mismatch between the syntax and semantics of CUDA programs have generally done so by raising the level of abstraction. Triton [15], a tile-based, multi-dimensional array language, circumvents the challenges of compositional reasoning by restricting users to a single level of the hierarchy, namely the block level. While this makes for a simplified programming model,

```

50
51 //include <cub/cub.cuh> // or equivalently <cub/block/block_reduce.cuh>
52
53 _global_ void ExampleKernel(...)

54 { // Specialize BlockReduce for a 1D block of 128 threads of type int
55     using BlockReduce = cub::BlockReduce<int, 128>;
56     // Allocate shared memory for BlockReduce
57     _shared_ typename BlockReduce::TempStorage temp_storage;
58     // Obtain a segment of consecutive items that are blocked across threads
59     int thread_data[4];
60     ...
61     // Compute the block-wide max for thread0
62     int aggregate = BlockReduce(temp_storage).Reduce(thread_data, cuda::maximum<>());
63 }
64
65 
```

Computes a block-wide reduction for `thread0` using the specified binary reduction functor. The first `num_valid` threads each contribute one input element.

- The return value is undefined in threads other than `thread0`.
- For multi-dimensional blocks, threads are linearly ranked in row-major order.
- A subsequent `__syncthreads()` threadblock barrier should be invoked after calling this method if the collective's temporary storage (e.g., `temp_storage`) is to be reused or repurposed.

Fig. 1. CUB documentation for the BlockReduce function

the lack of fine-grained control prevents GPU programmers from writing the highest-performance kernels. A variety of functional languages, such as Futhark [8], provide compile-time guarantees through their type systems but do not expose low-level control over the hardware. Some other efforts, like Descend [9], aim to provide a low-level, safe GPU systems programming language, but lack support for collective operations entirely. As a result, CUDA remains the de-facto standard for writing high-performance kernels on modern GPUs.

Rather than sacrificing low-level control for safety, we aim to make fine-grained GPU programming both safe and compositional by design. Our key insight is that GPU programmers naturally map computations onto different compute resources, and unlike CUDA, which obscures this mapping, we can expose it explicitly using the surface syntax of the language. Using this information, we can check at compile time whether the usage of collective operations contradicts the resources available in that context.

We introduce Snap, a new low-level GPU programming language that guarantees correct usage of hardware resources by construction. Inspired by coeffect systems [12] and dependency calculi [1], Snap statically tracks the configurations of compute and memory resources at the type level.

In this work, we reify the notion of compute resources via *privileges*, so named because they dictate which operations are permitted in which contexts. Snap uses these privileges to verify that every collective operation is executed with the necessary resources. To enable modular reasoning, function interfaces include *privilege signatures* that declare the privilege each function requires for correct execution.

We implement Snap and to validate our design, we build Bundl, a core calculus model of Snap that tracks privileges in both its type system and operational semantics. We provide formal rules for manipulating the privilege of both code and data, and prove *type-and-privilege safety*, ensuring that Bundl is sound, and that all threads are statically privileged to run the operations they attempt at run time.

A parallel goal, in addition to correctness, is performance. To that end, we incorporate modern GPU features such as tensor cores and asynchronous data movement into Snap, and demonstrate that Snap can achieve the same performance as hand-written, highly optimized code on an H100 and a 4070 Ti Super. Our contributions are:

- 99 (1) Snap, a low-level GPU programming language with type-level privileges that track compute
 100 and memory resources (Section 3);
 101 (2) Bundl, a calculus for Snap that tracks privileges in both its type system and operational
 102 semantics, ultimately helping us prove a soundness theorem that states that threads only
 103 execute operations for which they are statically privileged (Section 4); and
 104 (3) An implementation of Snap (Section 5) that demonstrates that Snap can expose modern
 105 GPU features like tensor cores and asynchronous data-movement, and achieve performance
 106 comparable to hand-optimized CUDA implementations (Section 6).

2 Background & Motivation

109 Before diving into Snap’s design, we begin with an overview of the GPU’s compute and memory
 110 hierarchies and outline the challenges posed by reasoning about them collectively.

2.1 Compute Hierarchy

113 In CUDA, programmers launch computations that run on thousands of threads. These threads are
 114 organized into a *compute hierarchy* that defines how work is distributed and scheduled on the GPU. At
 115 the top of this hierarchy is the *grid*, representing all threads launched as part of a single computation.
 116 The grid is divided into *blocks*, each containing a user-specified number of threads. A *thread*, in turn,
 117 is the finest unit of execution, and it is the machine’s basic unit of sequential control.

118 Users define the behavior across different levels of the compute hierarchy by defining the behavior of individual threads that constitute
 119 those levels. To do so, CUDA allows programs to read built-in identifiers like `threadIdx.x`—to determine a thread’s position within
 120 a block—and `blockIdx.x`—to determine the block’s position within
 121 the grid—at runtime. Combining these two IDs, threads can locate
 122 themselves in their share of the full computation.,

123 A natural and tempting way to interpret these built-in identifiers
 124 is to think of them as indices of implicit “parallel for-loops” where
 125 each iteration is executing simultaneously. While this view is suf-
 126 ficient to understand CUDA’s programming model when threads
 127 do independent work, it quickly breaks down in the presence of
 128 collective operations.

129 Unlike most instructions, which are executed by a single thread,
 130 collective operations must be executed collaboratively by a group
 131 of threads. For example, a “warp-level” tensor-core operation is only
 132 meaningful when invoked by a collection of 32 threads—a *warp*—acting together. Consider Figure 2
 133 which shows a tensor-core operation being invoked by a warp on line 14. For the call, each thread
 134 sets up its portion of the operands, and the operation is performed once for the entire warp, with
 135 the results scattered across participating threads. If the condition on line 2 were instead `tid > 0`
 136 `&& tid < 30`, making fewer than 32 threads reach line 14, the instruction would be undefined. To
 137 make matters worse, the restriction is not only on the *number* of threads executing the operation,
 138 but also on their *alignment*. In this case, the starting ID of the group of 32 participating threads must
 139 be aligned to a multiple of 32. So, if the condition on line 2 were instead `tid > 1 && tid < 33`, the
 140 instruction would still be invalid, even though 32 threads would execute it.

141 Collective operations make reasoning about CUDA programs challenging because they force
 142 programmers to track the convergence behavior of threads. Specifically, programmers must reason
 143 about how many threads reach a particular point in the program, and how those threads are arranged,
 144 accounting for alignment. This difficulty is further amplified due to two reasons. First, programs may

```

1 int tid = threadIdx.x;
2 if (tid >= 0 && tid < 31) {
3     float A[4];
4     float B[2];
5     float C[4] = { 0 };
6     # Populate A with unique values
7     for (int i = 0; i < 4; i++)
8         A[i] = tid * 4 + i;
9     # Populate B with unique values
10    for (int i = 0; i < 2; i++)
11        B[i] = tid * 4 + i;
12    # Issue a warp-level tensor-core
13    # operation. D = A * B + C
14    asm("mma.sync.aligned.m16n8k8... "
15        "%0, %1, %2, %3, " /*D*/
16        "%4, %5, %6, %7, " /*A*/
17        "%8, %9, " /*B*/
18        "%10, %11, %12, %13;" /*C*/
19        : "=r"(C[0]), "=r"(C[1]), ...
20        : "r"(A[0]), "r"(A[1]), ...
21        "r"(B[0]), "r"(B[1]),
22        "r"(C[0]), "r"(C[1]), ...);
```

Fig. 2. Invoking a warp-level tensor-core instruction in CUDA.

have *multiple points of convergence*, requiring programmers to mentally track the relative ID of a thread within a logical group as it changes over the course of a program’s execution. Second, threads may be participating in *multiple levels of convergence* within the same program. In our example, we only considered the warp-level tensor-core operation, but there are other collective operations that require convergence at different granularities like the *warp-group*-level tensor-core operations, which must be issued collectively by four warps, or the block-level `__syncthreads()` synchronization primitive, which must be executed by all threads within a block.

Reasoning about collective operations is already error-prone within the context of a single function, but becomes even more difficult when reasoning interprocedurally *across* functions. A callee may assume a certain number of threads or blocks and may structure its computation, including convergence behavior, around that assumption. However, such assumptions are not visible in the callee’s function interface, which only exposes the input and output types. To invoke that function correctly, users must read documentation, or worse, read the callee’s implementation to understand its assumptions, breaking modular reasoning.

In this work, we make the programmer’s implicit assumptions about a program’s convergence behavior explicit in the program’s source. Consider the example in Figure 3 which shows an equivalent rewrite of the CUDA program in Figure 2 using Snap. In Figure 3, the call to the `mma` intrinsic is legal only because it is executed within a group of 32 threads, as made explicit on line 2. Since Snap’s `group` construct already enforces both the size and alignment of participating threads, the legality of the `mma` operation is guaranteed by construction.

Further, Snap also tracks the *frequency* at which values diverge in *space*. The `@` syntax attached to each variable deceleration denotes the rate at which a value changes across compute resources on the machine. For example, a `@ thread[1]` value can diverge per-thread, while a `@ block[1]` can diverge for every block, but not for threads within that block. Using this information, Snap enforces rules for reading from and writing to data. So, if we had attempted to introduce a condition based on `tid` *within* the group’s scope, Snap would reject that program at compile time. This is because code executing at a lower frequency cannot read variables that change at a higher frequency than it. We will explain the syntax and rules of legal Snap programs in greater detail in Section 3.

2.2 Memory Hierarchy

GPUs also have a *memory hierarchy*, which, like the compute hierarchy, is organized into three levels. All threads in a grid can read and write from *global memory*, where operands typically reside at the start of a computation and where results are eventually written. Each block has access to a limited amount of *shared memory*, a programmer-managed scratchpad typically used to stage data that will be repeatedly used. Finally, each thread maintains its private state in *registers* and *local memory*, which is used for storing the thread’s stack.

This memory hierarchy has its own notion of convergence and divergence, mirroring that of compute. For example, when launching a kernel on the GPU, a pointer may initially belong to the grid because every thread sees the same pointer value at the start of the computation. To write to that memory, each thread computes an offset from the original base address. In doing so, the pointer *diverges* across threads within the larger memory region so each

```

1  tid : int @ thread[] = id();
2  with group(thread[32]):
3      A : float[4] @ thread[1]
4      B : float[2] @ thread[1]
5      C : float[4] @ thread[1]
6      # Populate A with unique values
7      for i in range(0, 4, 1):
8          A[i] = tid * 4 + 1
9          C[i] = 0
10     # Populate B with unique values
11     for i in range(0, 2, 1):
12         B[i] = tid * 4 + 1
13     # Issue a warp-level tensor-core
14     intrinsic.mma(
15         A[0], A[1], A[2], A[3],
16         B[0], B[1],
17         C[0], C[1], C[2], C[3],
18         out=[A[0], A[1], A[2], A[3]])
19

```

Fig. 3. Invoking a warp-level tensor-core instruction in Snap.

```

1  int x = threadIdx.x;
2  if (x >= 0 && x < 31){
3      float A[4];
4      float B[2];
5      float C[4] = { 0 };
6
7      // k is A's stride
8      A[0] = a_mem[(x/4)*k+(x%4)];
9      A[1] = a_mem[(x/4)*k+(x%4)+4];
10     A[2] = a_mem[(x/4)*k+(x%4)+1];
11     A[3] = a_mem[(x/4)*k+(x%4)+3];
12
13     // n is B's stride
14     B[0] = b_mem[(x%4)*n+(x/4)];
15     B[1] = b_mem[(x%4+4)*n+(x/4)];
16
17     asm("mma.sync...");
18
19     // Write back into c_mem

```

197 thread can write independently. After these writes, the memory
 198 must *reconverge*, ensuring that all threads have completed their
 199 updates before the memory can safely be returned to its original
 200 logical owner in the hierarchy.

201 CUDA does not explicitly model these memory spaces, nor does it
 202 track whether allocations in a given memory space, such as shared
 203 memory, exceed device limits. Snap, on the other hand, reifies these
 204 spaces in its type system by explicitly decorating each variable
 205 with its associated memory space and restricting shared memory to
 206 static allocations, throwing an error if an allocation exceeds device
 207 limits. We will discuss these aspects in detail in Section 3.5. For now,
 208 however, we focus on how memory diverges and reconverges in
 209 tandem with the compute hierarchy.

210 Let us reconsider the tensor-core example in Figure 2, this time initializing the operands of the
 211 tensor-core operation from pointers into memory. In Figure 4, A and B are now populated from data
 212 in global memory, a_mem and b_mem. After the tensor-core operation completes, the result is written
 213 back to c_mem, also in global memory.

214 In this program, c_mem is logically accessed from two different levels of the compute hierarchy.
 215 When entering this code, each thread in a group of 32 threads sees the same value for c_mem. Then,
 216 each thread within these 32 locates an offset within c_mem that it can write to (lines 21-24). To
 217 restore c_mem back to its 32 thread-level ownership, all 32 threads must first synchronize to ensure all
 218 per-thread writes have completed. This is similar to the requirement we saw documented in Figure 1.

219 Similarly to the compute hierarchy (Section 2.1), CUDA programmers are responsible for tracking
 220 the logical owner of a piece of memory as it evolves over the course of a program, and for ensuring
 221 that appropriate synchronization occurs whenever that ownership changes.

222 Snap, by contrast, makes the evolution of a memory object's
 223 logical ownership explicit in the type system and *automatically*
 224 inserts the synchronization required to restore memory to its
 225 original ownership . In Figure 5, we show how memory is
 226 lowered through the compute hierarchy for the same tensor-
 227 core operation introduced in Figure 4. The lowering of c_mem
 228 is required in this program as Snap only permits writes when
 229 data is owned by a single thread. To lower c_mem, we use Snap's
 230 partition operation on line 31. The operation takes the name
 231 of the memory to partition, c_mem, and lowers it to a compute
 232 hierarchy below the original one, a single thread, and assigns
 233 it a new name, c_thrd; Snap will not allow references to the
 234 old name c_mem within the partition's scope. The partition
 235 function also takes an indexing function, and each time c_thrd
 236 is accessed, this indexing function is implicitly applied. Once
 237 the partition's scope ends, Snap will insert a synchronization
 238 before the first use of the original name, so that all per-thread
 239 writes have completed.

240 We would like to emphasize that preventing data races is not
 241 an one of Snap's goals. In Snap, data races *can* occur; however,
 242 since the data is eventually synchronized before it is reused,
 243 the last-writer wins. Out-of-bounds accesses are considered

```

1 //cuda("device")
2 //requires(thread[32])
3 def simple_mma(
4     a: ptr(const float) @ thread[32],
5     b: ptr(const float) @ thread[32],
6     c: ptr(float) @ thread[32],
7     x : int @ thread[] = id()
8     with group(thread[32]):
9         A : float[4] @ thread[][]
10        B : float[2] @ thread[][]
11        C : float[4] @ thread[][]
12
13 # Reads do not need to be lowered.
14 A[0] = a_mem[(x/4)*k+(x%4)];
15 A[1] = a_mem[(x/4)*k+(x%4)+1];
16 A[2] = a_mem[(x/4+8)*k+(x%4)];
17 A[3] = a_mem[(x/4+8)*k+(x%4)+1];
18
19 B[0] = b_mem[(x%4)*n+(x/4)];
20 B[1] = b_mem[(x%4+4)*n+(x/4)];
21 # Skipping initialize C to ...
22 intrinsic.mma(
23     A[0], A[1], A[2], A[3],
24     B[0], B[1],
25     C[0], C[1], C[2], C[3],
26     out=[C[0], C[1], C[2], C[3]]
27 )
28 # Must be at thread[1] to write.
29 idx : int @ thread[] =
30     lambda ro, co: (x/4+ro)*n+2*(x%4)+co
31     with partition(c_mem, p=thread[32],
32     ↪ f=idx) as c_thrd:
33         c_thrd[0, 0] = C[0]
34         c_thrd[0, 1] = C[1]
35         c_thrd[0, 2] = C[2]
36         c_thrd[0, 3] = C[3]
37
38 # <-- Sync point inferred by name

```

Fig. 5. Invoking a warp-level tensor-core instruction in Snap after loading data from memory.

246 undefined behavior. We believe that prior work, such as Descend [9], lays the foundation for reasoning about data-race
247 free programs, and that a Descend-like memory model can be
248 adapted to Snap. Snap is instead focused on the *interaction* be-
249 tween the compute and memory hierarchies, and on reasoning
250 about them simultaneously to ensure that an operation is executed only when sufficient compute
251 resources are available, a guarantee that prior work like Descend cannot provide.
252

253

254

255

256

257

258

259

260

261

262

263

264

265

266

267

268

269

270

271

272

273

274

275

276

277

278

279

280

281

282

283

284

285

286

287

288

289

290

291

292

293

294

3 The Snap Language

Snap is an imperative, low-level programming language designed at a level of abstraction comparable to CUDA. Snap's surface syntax explicitly represents how compute (grid, block, threads) and data (global, shared and local memory) resources are used and distributed over the course of its execution. Using this explicit representation, Snap can enforce safety properties that CUDA cannot. Namely, Snap rejects programs that attempt to execute operations with insufficient compute and memory resources.

In this section, we introduce Snap's core constructs through a running example, a tensor-core matrix multiplication, shown in Figure 6. This is a standard computation a GPU programmer might write while being sufficiently complex to reveal the interesting challenges in GPU programming.

The program in Figure 6 computes a matrix multiplication between two float arrays, A and B to produce an output matrix C. In this kernel, each block computes an independent tile of the output C. To perform the computation, each block first locates its assigned tile index, then loads the corresponding rows and columns from A and B (line 20-21). Since matrix multiplication offers ample data reuse, to amortize the cost of loading data, we stage A and B in shared memory (lines 46-59). At this point, each thread within the block diverges and loads different elements of A and B. Finally, the program invokes an Amere warp-level tensor-core instruction to compute the output, requiring threads in a warp to converge (line 71). We encapsulate this tensor-core instruction in a function, demonstrating how function composition works in Snap. This function's implementation is identical to the one shown in Figure 5.

3.1 Levels

Snap models the machine's compute hierarchy through *levels*. There are three levels in Snap—grid, block and thread—which are organized as expected: a grid consists of multiple units of the block level, which in turn consist of multiple units of the thread level. Levels have a strict

```

1  @ccuda("global")
2  # Top-level privilege class and required shared memory usage.
3  @requires(grid[1], block[1], thread[32], smem=1280)
4  def mmaTF32NaiveKernel(A: ptr(const(float)) @ grid[1],
5                           B: ptr(const(float)) @ grid[1],
6                           C: ptr(float) @ grid[1],
7                           M : int @ grid[1],
8                           N : int @ grid[1],
9                           K : int @ grid[1]):
# Replicate this code across 1 grid.
with group(grid[1]):
    # @ grid[1] inferred from compute privilege
    MMA_N : constexpr(int) = 8
    ...
    num_blocks_n : const(int) = (N + MMA_N - 1) / MMA_N
    # id() function returns the block id, inferred from
    # @ block[1].
    block_row : const(int) @ block[1] = id() / num_blocks_n
    block_col : const(int) @ block[1] = id() % num_blocks_n

    warp_row: const(int) @ block[1] = block_row * MMA_M
    warp_col : const(int) @ block[1] = block_col * MMA_N
    # Give each block an offset into C
    offset = lambda x: warp_row * N + warp_col + x
with partition(C, p=block[1], f=offset) as C_blk:
    with group(block[1]):
        # SEM declarations only allowed with block[1] privilege.
        A_smem : shared(float[16 * 8]) @ block[1]
        B_smem : shared(float[8 * 8]) @ block[1]
        C_smem : shared(float[16 * 8]) @ block[1]
        # Now, id() returns the thread id
        idx : int @ thread[1] = id() * 4
        # To write to C_smem, must drop to thread[1] privilege
        with partition(C_smem, p[thread[1]], offset=idx) as C_thrd:
            for i in range(0, 4, 1):
                with group(thread[1]):
                    C_thrd[i] = 0

for i in range(0, K_tiles, 1):
    a_idx : int @ thread[1] = id() * 4
    # <-- Name will insert a sync point
    # because this partition is a parent partition
    # or itself (back edge from for loop)
    for j in range(0, 4, 1):
        global_row : int @ thread[1] = warp_row + row
        global_col: int @ thread[1] = i * MMA_K + col
        with partition(A_smem, p[thread[1]], f...) as As_thrd:
            with group(thread[1]):
                As_thrd[0] = A[global_row * K + global_col]

        b_idx : int @ thread[1] = id() * 2
        # <-- Name will insert a sync point
        # (back edge from for loop)
        for j in range(0, 2, 1):
            # Similar to write into C_smem ...
            with partition(B_smem, p[thread[1]], f...) as Bs_thrd:
                with group(thread[1]):
                    Bs_thrd[0] = B[global_row_b * N + global_col_b]

        # <-- Name will insert a sync point
        # (back edge from for loop)
        # Give each warp an offset into C_smem
        with claim(C_smem, p[thread[32]]) as Cs_warp:
            match split(thread):
                case 32:
                    # Can pass A_smem and B_smem
                    # as arguments to simple_mma for
                    # those two are declared const.
                    # "read up" is okay.
                    simple_mma(A_smem, B_smem, Cs_warp)

# Write back final result to C
for j in range(0, 4, 1):
    flat_idx_c : int @ thread[1] = id() * 4 + j
    row_c : int @ thread[1] = flat_idx_c / MMA_K
    col_c : int @ thread[1] = flat_idx_c % MMA_K
    offset = lambda x: row_c * N + col_c + x
    with partition(C_blk, p[thread[1]], f=offset) as C_thrd:
        with group(thread[1]):
            C_thrd[0] = C_smem[row_c * MMA_N + col_c]
return

```

Fig. 6. Tensor Float 32 matrix Multiplication in Snap.

344 order defined over them. In particular, thread
 345 <block <grid.

346 There are two key differences between Snap's
 347 levels and those in CUDA. First, unlike CUDA,
 348 Snap does not model a a three-dimensional grid
 349 or block structure, say through different levels for grid_x, grid_y, or grid_z. Second, there are two
 350 "levels", warp and warp group, noticeably absent from our hierarchy.

351 These differences are a deliberate design choice. On the hardware, the units of each level are
 352 arranged in a single linear order, and the three-dimensional structures are simply *interpretations* of this
 353 linear ordering, not distinct hardware resources. Similarly, warp and warp group are organizational
 354 constructs defined in terms of existing levels . Namely, a warp is a group of 32 threads whose first
 355 thread ID is aligned to 32. A warp group, which only became a meaningful construct with the release
 356 of the Hopper architecture , consists of 128 threads aligned to 128.

357 Rather than embedding these interpretations into the language as special cases by introducing
 358 new levels, Snap provides constructs that let users group units of a given level . With these constructs,
 359 users can express multi-dimensional structures and define custom groupings with specific quantities
 360 and alignments.

361 The primary mechanism for organizing units at each level is a *privilege*.

362

363 3.2 Privileges & Privilege Classes

364 Privileges are the central concept in Snap, allowing it to determine which compute resources are
 365 being requested by the programmer, whether they are available in the program's context, and, once
 366 provided, if those resources are sufficient for a given operation.

367 A *privilege* is a level—grid, block, or thread—parametrized by a static constant n, specifying the
 368 number of units at that level. A privilege is written as level[n]. For example, thread[2] denotes a
 369 privilege of two threads, block[4] denotes a scope of four blocks, and so on.

370 Privileges also carry *alignment information*: a privilege parameterized by n is aligned to n. In this
 371 way, a warp is simply a desugaring of thread[32], and a warp group is a desugaring of thread[128].

372 Finally, privileges have a *partial order* defined on them. We say that level₂[n₂] is a *higher privilege*
 373 than level₁[n₁] iff either level₁ < level₂, or, if level₁ = level₂, then n₂ % n₁ = 0. The condition n₂
 374 % n₁ = 0 may appear somewhat mysterious at first, but we will see why it is required in Section 3.2.2.

375 In Snap, both code and data are bound to privileges . Code is bound to a set of privileges, called a
 376 privilege class meanwhile data is bound to a single privilege.

377 3.2.1 Privileges for Code, or Privilege Classes.

378 Privilege classes capture what compute resources a program has access to. Each function starts
 379 out with a top-level privilege class that defines the total available resources that it *assumes* are
 380 available to it. As a shorthand, when we refer to a code's privilege, we mean the highest privilege
 381 available to it in its class. In Figure 6, the privilege class is declared on line 2 using the notation
 382 @requires(grid=..,block=..) Starting from this top-level declaration, programmers shape the
 383 current privilege class in the function's body through two constructs, group and split.

384 **Group.** The group construct bundles privileges available in the current privilege class into "groups"
 385 of a particular granularity. Intuitively, group narrows the current privilege class and replicates it
 386 across collections of a finer granularity. For example, Figure 6, line 26 shows a grouping of a single
 387 block, while line 22 in Figure 5 shows a grouping of 32 warps , that will ultimately invoke a tensor-core
 388 operation.

389

393 An invocation of `group(level[n])` is allowed if and only if the
 394 privilege class contains privilege higher than `level[n]`. For example,
 395 `group(warp[32])` on 22 in Figure 5 is allowed because the privilege
 396 class of the function declares that the top level on line 22 starts out
 397 with at least 32 threads.

398 Once `group(level[n])` is invoked, it modifies the current privi-
 399 lege class in two key ways. First, it removes all parents of `level[n]`
 400 from it. Second, the number of available units for `level` is set to
 401 exactly `n`. Combined, these rules eliminate malformed programs that would result in unsatisfiable
 402 requirements. We shown some examples of illegal group statements in Figure 7.

```
1 # Example 1
2 with group(thread[2]):
3     # Illegal because block > thread
4     with group(block[1]):
5         pass
6 # Example 2
7 with group(block[1]):
8     # Illegal because 1 % 2 != 0
9     with group([block[2]]):
10        pass
```

Fig. 7. Illegal uses of group.

404 **Split.** Unlike `group`, which is used for replication over smaller, equal sized resource sets, `split` is
 405 used for sharding currently available resources into unequal sized sets. For example, line 63 in Figure 6
 406 shows a `split` of a `thread[32]` privilege that is used to execute a warp-level tensor-core operation.

407 The `split` operation reflects unordered composition that is not
 408 mere replication, and all branches of the `split` execute in parallel.
 409 Once `split(level)` is invoked, the privilege classes of its branches
 410 diverge. A `split` construct will eventually compile to an `if`-statement
 411 during code generation.

412 Therefore, to satisfy a `split` operation, Snap checks that the
 413 *cumulative sum* of the privileges requested by each branch of the
 414 `split` can be satisfied. For example, if we were to change the program
 415 in Figure 6 slightly (on line 63-69), to the one shown in Figure 8,
 416 Snap would throw a compile-time error. Finally, because privileges
 417 assume alignment, every branch of the `split` must also be aligned.
 418 In other words, not all splits whose sizes sum to less than or equal to
 419 the available units are necessarily valid. Figure 9 shows an example
 420 of this constraint. This alignment constraint is also why we compare
 421 privileges by checking whether one quantity is a multiple of the
 422 other's.

```
1 # Example 1
2 with group(thread[32]):
3     match split(thread):
4         case 32:
5             simple_mma(..)
6             # Illegal! 32 + 1 > 33
7         case 1:
8             pass
```

Fig. 8. Illegal split that exceeds available privileges.

```
1 # Example 2
2 with group(thread[33]):
3     match split(thread):
4         case 1:
5             pass
6             # Illegal! 32 + 1 <= 33,
7             # but the second branch
8             # is not aligned by 32.
9         case 32:
10            simple_mma(..)
```

Fig. 9. Illegal split that violates alignment.

423 Inside each branch that requests `n` units, all parent levels of
 424 `level[n]` are removed from its privilege class, and the number of available units for `level` is
 425 set to exactly `n`, and, within that scope, all resource requirements are checked against that privilege
 426 class.

427 3.2.2 Privileges for Data.

428 Like code, data also has a privilege attached to it that dictates the set of privileges that must be held
 429 for reading and writing to it. Ultimately, it is this interaction between privilege classes for code and
 430 privileges for data that let us make meaningful statements about the set of privileges available at
 431 runtime. The precise details of this interaction are formalized in Section 4.

432 In this section, we focus tracking thread-local stack allocations and local variables. We defer
 433 discussion of pointers to memory to Section 3.5, where we will show how shared views of memory
 434 are lowered and restored through the compute hierarchy.

435 For thread-local data, its privilege is specified in its declaration using the `@ level[n]` syntax.
 436 For example, an integer variable `v` is declared with `thread[1]` privilege with the syntax `v : int @`
 437 `thread[1]`. If not explicitly annotated, Snap infers a variable's privilege to be the privilege of the
 438 code where it was declared. For example, `MMA_N`, `num_blocks_n` are all inferred to be at `grid[1]`
 439 scope on lines 13 and 15 in Figure 6.

442 The rules for reading and writing to data are asymmetric: higher-privilege data can flow into lower-
 443 privilege data, but not the other way around. In particular, a value coming from a higher-privilege
 444 variable can be written into lower-privilege data, but a value coming from a lower-privilege variable
 445 cannot be written into higher-privilege data. In short, the rule is: "read up, write down".

446 If a value is read without being written into a variable, for example,
 447 in an `if` statement, the above rule is enforced using the privilege
 448 the code that is executing the read. For example, branching on a
 449 `thread[1]` variable within `block[1]` privilege code is rejected by
 450 our compiler, as it constitutes a "read down".

451 In Figure 6, we can see rule "read up" in action on lines **45 and 46**.
 452 Meanwhile, "write down" is being used to set up variables on line
 453 **18-22**.

454 Intuitively, `@ level[n]` tells Snap the *frequency* at which a value changes in space . For example,
 455 a `grid[1]` value is the same for all threads in a grid, whereas a `block[2]` value changes every
 456 two blocks. The rules described above enforce that the frequency declared for each variable is
 457 respected throughout the program. For instance, the "read up" principle allows variables with
 458 higher replication frequency to read from lower-frequency variables without violating their declared
 459 frequency. Conversely, higher-frequency variables cannot be used to write into lower-frequency
 460 variables, since that would violate the variable's declared frequency.

461 By tracking frequency in this way, Snap can ensure that threads that with a given privilege can
 462 only read variables that change at the same frequency or higher. This guarantees that all threads
 463 with that privilege follow the same control flow, and that our accounting of available privilege class
 464 remains consistent with the program's runtime behavior.

465 Until now, we have discussed up rules for reading and writing to variables that change at different
 466 frequencies, but we have not yet set up a mechanism for actually actually setting up variables that
 467 diverge and change at different frequencies. The main tool for doing this is the `id()` function.

469 3.3 The `id()` Function

470 As opposed to exposing users to special hardware `blockIdx.x`, `threadIdx.x` variables directly, Snap
 471 provides the `id()` function instead. The `id` function returns the relative index, or a "label", for a given
 472 compute privilege. The interpretation of the `id` function depends on both the privilege of the variable
 473 it is being written to, and the privilege class that the `id` function is being invoked in. In our example
 474 in Figure 6, we use a call to `id` in `grid[1]` scope at line **18 and 19** to locate the tile that each block is
 475 ultimately charge of computing. In Figure 3, however, a call to `id` in `thread[32]` scope at line **7** will
 476 not return the global thread ID of the threads launched in the kernel, but rather the *relative* thread ID
 477 within the 32 threads, ranging from 0 to 31.

478 479 3.4 Collective Operations

480 With privilege classes for code and data in place, we can now turn our attention to Snap's reasoning
 481 about collective operations. By statically tracking privilege classes for code, Snap guarantees that
 482 collective operations are semantically valid by ensuring that they are invoked with sufficient privilege.

483 There are two ways to access collective operations in Snap. The first is to use intrinsics provided
 484 by the compiler. These intrinsics have built-in privilege classes attached to them, which the compiler
 485 can directly enforce. In Figure 3, the tensor-core call at line **22** is made using one of the intrinsics
 486 provided by our compiler. The second, more flexible approach allows users to extend these operations
 487 using the `unsafe` construct, explained in Section 3.8. Using unsafe code, users can inline a collective
 488 operation—often a single assembly instruction—into a Snap function and ensure that the function's
 489 top-level requirements, specified with `@require`, faithfully describe the privileges required to execute

```
1 with group(block[1]):
2   id : int @ thread[1] = id()
3   # Illegal read of ID from
4   # block[1] scope.
5   if (id)
6     __syncthreads();
```

Fig. 10. Illegal read of `thread[1]` variable.

491 that instruction. Then, Snap will check the call to the function as any other function call, enforcing
 492 the correct privilege. The rules for checking function calls are described in Section 3.7.

493 A quirk of some collective operations is that, although they need to be invoked with higher compute
 494 privileges, they often produce results with lower data privileges. For example, the tensor-core operation
 495 in Figure 3 is invoked at thread[32] granularity but generates output at the thread[1] level.
 496 That is, even though the call to the tensor call is a convergent operation requiring the thread[32]
 497 compute privilege, it produces divergent output with thread[1] data privilege. To be able to handle
 498 these operations we must be able to allow thread[1] state to coexist with the thread[32] privilege.

499 One potential solution is to artificially coarsen the granularity of the operation by having the
 500 intrinsic internally scatter the thread[1] output into a thread[32] variable. However, this approach
 501 is often unacceptable for low-level CUDA performance tuning. Often, performance engineers need
 502 to manually schedule and interleave memory and arithmetic instructions to hide latency, with
 503 loads or stores occurring well in advance of their dependent operations and interleaved with other
 504 instructions. If we coarsen these operations, we will disallow these optimizations entirely.

505 In Snap, however, this behavior does not pose a problem. Because of the "write down" rule, registers
 506 can be set up in advance of the operation (e.g., line 14-20 in Figure 3) without violating correctness.

508 3.5 Memory

509 In Snap, there are three mechanisms for obtaining a pointer type. The first method is to directly
 510 allocate data. The second and third are to obtain offsets into existing data by using `partition` or
 511 `claim`. These constructs mirror the `group` and `split` constructs, and are used to lower the data
 512 privilege associated with memory.

514 3.5.1 Allocations in Different Memory Spaces.

515 In Section 2.2, we saw that the GPU has distinct memory spaces, like global, local, and shared
 516 memory. Local allocations reside in either the stack or registers, which Snap can distinguish upfront.
 517 Shared memory pointers are explicitly annotated with the `shared` keyword, while global pointers
 518 are unannotated.

519 The primary memory resource Snap needs to manage are shared memory allocations. Local
 520 memory allocations are just stack allocations, and don't need special treatment. We also assume that
 521 sufficient global memory has been allocated before the kernel launch and don't provide a mechanism
 522 for runtime allocations—this is typically the case in real-world programs.

523 In CUDA, shared memory is a limited resource, typically a few hundred kilobytes per block.
 524 Programmers manually manage this memory by bumping pointers to allocate offsets for different
 525 data. Each allocation's pointer must be carefully tracked throughout the program, and exceeding the
 526 available space will lead to a runtime error.

527 Snap simplifies shared memory management by requiring all shared memory allocations have a
 528 static size. Since shared memory is only visible to all threads in a block, a shared memory allocation
 529 is only valid with `block[1]` privilege. An example of such an allocation is shown on lines 28-30
 530 of Figure 6. During compilation, Snap will automatically handle the necessary pointer arithmetic to
 531 assign each allocation an appropriate offset within the shared memory space.

532 The `@requires` annotation specifies the amount of shared memory a function expects to have
 533 available, and Snap uses this information to ensure that a function's allocations do not exceed its
 534 declared limit, and checks whether there is enough shared memory available at its call sites. Figure
 535 6 declares the amount of shared memory it will require on 2. We use standard techniques to
 536 conservatively bound memory usage in Snap's type system.

538 3.5.2 Privileges for Memory.

540 In Section 3.2.2, we focused on tracking for primitive data types and stack allocations, where each
 541 thread maintained its own local copy of the data. In contrast, pointers into memory, such as global or
 542 shared memory, do not represent distinct copies, but instead refer to the same underlying storage
 543 that may be accessed simultaneously by multiple threads. To use them, we must first *lower* their
 544 privileges, allowing individual threads to write different values to different locations, and then *restore*
 545 them back to their original privilege before execution continues.

546 As with variables in Section 3.2.2, lower-privilege data may read from higher-privilege data
 547 without restrictions. For writes, however, Snap enforces that all pointers into memory be written
 548 with `thread[1]` privilege; they therefore must be lowered before they can be safely updated.

550 **Partition.** The `partition` construct plays the same role for memory that `group` plays for compute.
 551 It is used to lower a memory object from a higher privilege to a lower one by assigning each lower
 552 privilege a distinct offset into the original pointer. It takes the original name of a memory object,
 553 `base_name`, an indexing function `f`, and a target privilege `level[n]` to which the memory is being
 554 lowered, and produces a new name, `new_name`, for the partitioned memory. In Figure 6, for example,
 555 we partition the C buffer to `block[1]` privilege on line 25 first, and then lower it further to `thread[1]`
 556 privilege on line 78 so that we can actually write to it.

557 Once a memory object has been partitioned, the original name referring to it goes out of scope,
 558 and the memory can only be accessed through the new name introduced by `partition`. Each use of
 559 this new name applies the indexing function `f` to compute the true offset of the access. Within the
 560 scope of the `partition` statement, the original name cannot be repartitioned. Because `split` denotes
 561 unordered composition, Snap prohibits sibling `split` branches from partitioning the same memory
 562 object. When the scope of the `partition` statement ends, Snap automatically inserts the necessary
 563 synchronization to restore the memory to its original privilege. The mechanism for inserting and
 564 optimizing these synchronization points is described in Section 3.9.

565 At this point, it is necessary to consider how different indexing functions affect the possibility
 566 of data races. If the indexing function is injective, each lower-privilege privilege receives a distinct
 567 offset into the underlying array, and the resulting partition is free of data races. When the indexing
 568 function is not injective, multiple privileges may write to the same location, introducing a potential
 569 data race. Out-of-bounds accesses are undefined behavior.

570 As we discussed in Section 2.2, achieving data-race freedom is *not* one of Snap’s goals, and
 571 our guarantees continue to hold even when indexing functions are not injective. Because Snap
 572 automatically inserts the necessary synchronization before the original memory name is restored to
 573 its higher privilege, Snap ensures that all writes have completed before accessing the orginal name.
 574 The subsequent uses of that pointer will observe the value written by the last-writer.

575 We believe that extending Snap to reason about data races and out-of-bounds accesses is a promising
 576 direction for future work.

578 **Claim.** The `claim` construct is the analogue of `split` for memory . Using `claim`, the entire
 579 memory region can be moved to a lower privilege.

580 Unlike `partition`, `claim` does not take an indexing function. The `claim` construct takes the
 581 original memory’s name, the target privilege to lower to, and a new name to assign to the lowered
 582 memory. The `new_name` is accessible only within a single, explicit `split` branch; sibling branches are
 583 not permitted to read or write from this memory. In Figure 6, we use the `claim` operation on line 63
 584 to lower the memory to `thread[32]` privilege.

585 As with `partition`, a `split` construct introduces its own scope, and the original name is inacces-
 586 sible within this scope.

589 **3.6 Asynchrony**

590 **3.7 Function Composition**

591 To check whether a function call is valid, Snap needs to ensure that the call site can provide the set of
 592 privileges that a function will attempt to use over its execution, and that the arguments passed to the
 593 function also respect the privileges declared for its input and output types. Both requirements are
 594 represented explicitly via the `@require` construct.

595 As mentioned in Section 3.2, each function carries a privilege class. We call this top-level privilege
 596 class the function’s *privilege signature*. The signature specifies the top-level privilege class the
 597 function *assumes* will be provided to it. This signature represents the *minimum amount* of compute
 598 and memory resources the function must be called with. In the example above in Figure 6, line 2
 599 defines the privilege signature. Up to this point, we have described programs written *inside* functions
 600 under the assumption that the resource signature can be satisfied. In Section 3.2, for example, we
 601 used this signature to set up the top-level privilege class for the function.

602 **Compute privileges.** At a function’s call site, then, we need to ensure that the function’s privilege
 603 signature can be satisfied. Snap compares the current compute privilege with the one required by the
 604 function, and verifies that only functions whose privilege signatures can be satisfied are called.

605 **Per-argument checks.** In Snap, functions have pass-by-value semantics. For arguments, we
 606 distinguish between primitive data types and pointers to memory. For primitive data types, like `int`,
 607 `bool`, and stack allocations, users can pass either arguments that live at or above the data privilege
 608 specified by the function’s signature. For pointers, we allow passing pointers that live at higher data
 609 privilege if the pointer is marked as `const`, indicating that it will only be used for reading (line 4 and
 610 5 in Figure 3). Otherwise, we require that pointers match the exact privilege of the function signature.

611 **3.8 Unsafe**

612 **3.9 Compiler Internal: Synchronization**

613 **4 Formalism**

614 Having introduced the full Snap language, we now describe Bundl, a core calculus that formalizes its
 615 most fundamental aspects by statically tracking compute and data privileges. We use Bundl to prove
 616 that well-typed Snap programs are not only type-safe, but will also never get stuck trying to execute
 617 operations for which they do not have the required compute resources.

618 In this section, we describe Bundl’s type system and operational semantics—in particular how it
 619 manages compute and data privileges—and build up to a formal proof of type-and-privilege safety.
 620 By instrumenting Bundl’s operational semantics with privileges, our safety theorem can guarantee
 621 that dynamically-realized privileges match the ones inferred by the type system.

622 **4.1 Bundl Type System**

623 The core idea in Bundl is to track, at the type-level, the program’s current compute and data privilege.
 624 To achieve this, we borrow techniques from the literature on *coeffects* [12] and *dependency tracking*
 625 [1]. In particular, the current compute privilege is tracked on the typing judgment, which has the
 626 form $\Gamma \vdash^\pi e : \tau$ for expressions and $\Gamma \vdash^\pi s$ for statements. The π over the \vdash is the compute privilege of e
 627 and s .

628 The π s themselves are pairs of levels h of the GPU hierarchy at which each privilege lives (either
 629 Thread, Block or Grid), and numbers tracking the size of the logical group denoted by each privilege
 630 . So, for example, a `(Thread,4)` in Bundl would directly correspond to a `thread[4]` in Snap.

631 Each variable in the typing context has a data privilege attached to it. Only when these two
 632 privileges match can data be read or written from a variable. This requirement is made manifest in

638	$\boxed{\Gamma \vdash^\pi e : \tau}$	<i>(Expression typing)</i>
639		
640	$\frac{x :^\pi \tau \in \Gamma}{\Gamma \vdash^\pi x : \tau}$ T-VAR	
641		
642	$\boxed{\Gamma \vdash^\pi s}$	<i>(Statement typing)</i>
643		
644		
645	$\frac{\Gamma \vdash^{(h,n_1)} s_1 \quad \Gamma \vdash^{(h,n_2)} s_2 \quad n_1, n_2 \text{ align to } n}{\Gamma \vdash^{(h,n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\}}$ T-SPLIT	
646		
647	where $n_1, n_2 \text{ align to } n ::= (n_1 + n_2 \leq n) \text{ and } (n_1 n) \text{ and } (n_2 n) \text{ and } (n_2 n_1 + n)$	
648		
649	$\frac{\Gamma \vdash^\pi s}{\Gamma \vdash^{q \cdot \pi} \text{group } q \ s}$ T-GROUP	$\frac{\Gamma, y :^{\downarrow \pi} \tau[]^l \vdash^\pi s \quad l \neq \text{Local}}{\Gamma, x :^\pi \tau[]^l \vdash^\pi \text{lower } x \text{ into } y \text{ in } s}$ T-LOWER
650		
651		
652		
653	$\frac{\Gamma, y :^{(h,n/c)} \tau[]^l \vdash^{(h,n)} s \quad c n \quad l \neq \text{Local}}{\Gamma, x :^{(h,n)} \tau[]^l \vdash^{(h,n)} \text{partition } x \text{ into } y \text{ by } c \text{ in } s}$ T-PARTITION	$\frac{\Gamma \vdash^\pi s}{\Gamma \vdash^\pi \text{destruct in } s}$ T-DESTRUCT
654		
655		
656	$\frac{\Gamma, y :^{(h,n')} \tau[]^l \vdash^{(h,n')} s \quad n' \leq n \quad l \neq \text{Local}}{\Gamma, x :^{(h,n)} \tau[]^l \vdash^{(h,n)} \text{claim } x \text{ into } y \text{ at } n' \text{ in } s}$ T-CLAIM	
657		
658		

Fig. 11. Core typing rules of Bundl. The typing rules presented here are a simplified selection of the full rules, which can be found in Appendix A.2.

the T-VAR rule, which can be found in Figure 11. The premise of this rule also reveals that the context Γ maps variables to both types τ and also the privileges π with which they may be read and written.

Figure 11 also shows other key rules, and these rules fall into two main categories: rules for managing the current compute privilege and rules for managing data privileges.

4.1.1 Compute Privilege Management. In Snap, there are two operations which manage compute privileges: group and split. In Bundl, to better model the details of privilege attenuation , we introduce a third construct, destruct.

Bundl’s group statement directly corresponds to Snap’s, and is checked by the T-GROUP rule. Given some privilege that can be divided into q equally sized smaller privileges π , the statement $\text{group } q \ s$ will check with privilege $q \cdot \pi$ provided that s itself checks with privilege π . The starting privilege is of the form $q \cdot \pi$, and encodes Snap’s alignment requirement.

The split statement, meanwhile, is checked by the T-SPLIT rule, and functions like a binary version of the n-ary splitting construct in Snap. It enforces the same divisibility requirements to ensure that the privileges of code and data remain properly aligned, and then checks the two sub-statements s_1 and s_2 with the divided, smaller privileges.

The final rule for managing compute privileges is T-DESTRUCT, which handles downward movement from one level of the GPU hierarchy to another and makes explicit in Bundl programs exactly where such movement occurs. The \downarrow operation on π s “destructs” the privilege into a many smaller privileges at a lower level; $\downarrow(\text{Grid}, 1) = (\text{Block}, B)$ and $\downarrow(\text{Block}, 1) = (\text{Thread}, T)$, where B and T are parameters to a particular instantiation of Bundl to describe the number of blocks per grid and

$$\begin{array}{l}
 \text{687} \quad \frac{L(t), S(b), \Sigma, t, b, 0 \vdash^{(\text{Grid},1)} s \rightsquigarrow s' \dashv \eta', \sigma', \Sigma' \quad P(t, b) = s}{L, S, \Sigma, P \rightsquigarrow L[t \mapsto \eta'], S[b \mapsto \sigma'], \Sigma', P[(t, b) \mapsto s']} \text{ S-PROGRAM} \\
 \text{688} \\
 \text{689} \\
 \text{690} \\
 \text{691} \quad \frac{p < n_1 \quad n_1, n_2 \text{ align to } n \quad \eta, \sigma, \Sigma, t, b, p \vdash^{(h, n_1)} s_1 \rightsquigarrow s'_1 \dashv \eta', \sigma', \Sigma'}{\eta, \sigma, \Sigma, t, b, p \vdash^{(h, n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\} \rightsquigarrow \text{split}(n_1, n_2)\{s'_1\}\{s_2\} \dashv \eta', \sigma', \Sigma'} \text{ S-SPLIT-LEFT} \\
 \text{692} \\
 \text{693} \\
 \text{694} \quad \frac{p \geq n_1 \quad p < n_1 + n_2 \quad n_1, n_2 \text{ align to } n \quad \eta, \sigma, \Sigma, t, b, p - n_1, \vdash^{(h, n_2)} s_2 \rightsquigarrow s'_2 \dashv \eta', \sigma', \Sigma'}{\eta, \sigma, \Sigma, t, b, p \vdash^{(h, n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\} \rightsquigarrow \text{split}(n_1, n_2)\{s_1\}\{s'_2\} \dashv \eta', \sigma', \Sigma'} \text{ S-SPLIT-RIGHT} \\
 \text{695} \\
 \text{696} \\
 \text{697} \\
 \text{698} \quad \frac{\eta, \sigma, \Sigma, t, b, p \text{ mod } n \vdash^{(h, n)} s \rightsquigarrow s' \dashv \eta', \sigma', \Sigma'}{\eta, \sigma, \Sigma, t, b, p \vdash^{(h, q \cdot n)} \text{group } q \ s \rightsquigarrow \text{group } q \ s' \dashv \eta', \sigma', \Sigma'} \text{ S-GROUP} \\
 \text{699} \\
 \text{700} \\
 \text{701} \quad \frac{\eta, \sigma, \Sigma, t, b, t \text{ mod } T \vdash^{(\text{Thread}, T)} s \rightsquigarrow s' \dashv \eta', \sigma', \Sigma'}{\eta, \sigma, \Sigma, t, b, 0 \vdash^{(\text{Block}, 1)} \text{destruct in } s \rightsquigarrow \text{destruct in } s' \dashv \eta', \sigma', \Sigma'} \text{ S-DESTRUCT-BLOCK} \\
 \text{702} \\
 \text{703} \\
 \text{704} \\
 \text{705} \\
 \text{706}
 \end{array}$$

Fig. 12. Core semantic rules of Bundl. As with the typing rules, we present only a simplified selection of the full rules, which can be found in Appendix A.3.

712 threads per block.¹ Because \downarrow is only defined on (Grid,1) and (Block,1), the rule enforces that one
 713 can only destruct their compute privilege when they have privileges for exactly one grid or block.

714
 715 4.1.2 *Data Privilege Management*. The mechanism for managing data privileges mirrors that of
 716 compute privileges, with each operation for data corresponding to an operation for code.

717 The partition operation is analogous to group-ing a compute privilege. The typing rule for this
 718 operation, **T-PARTITION**, requires that the data privilege of the variable to be partitioned, x , is the
 719 same as the current compute privilege. After partitioning, a fresh variable y is introduced with a new
 720 privilege π/c . We use π/c as a shorthand to denote a division of a privilege π into c equally-sized,
 721 smaller privileges. Within the body of the partition statement, we disallow any references to the
 722 original variable x . Then, we continue checking the body of the partition statement against the
 723 original compute privilege π because the partition has no affect on the current compute privilege.

724 Unlike partition, which divides up a piece of data equally into lower privileges, the claim
 725 operation, mirroring the split construct, gives the entire variable to exactly one, smaller privilege.
 726 Accordingly, Bundl needs to ensure that only one branch of a split operation, with the appropriate
 727 π , can refer to the claimed variable. To ensure that this is the case, the **T-CLAIM** rule links the data
 728 privilege of the variable to the compute privilege of the code claiming it by changing *both* at the same
 729 time. This represents a minor difference from Snap, which implements additional static analysis to
 730 ensure that a claimed variable is only accessed in a single split branch.

731 Lastly, the **T-LOWER** rule mirrors the **T-DESTRUCT** rule; it uses the \downarrow operator to move a variable
 732 from one level of the hierarchy to another, distributing it equally among all the child privileges in the
 733 same manner as **T-PARTITION**.

736 4.2 Bundl Semantics

737 To reflect the fact that GPU programs execute in parallel across numerous threads, we model the
 738 semantics of Bundl in the style of Turon et al. [16], using a two-level small step judgment. We present
 739 the key rules of this semantics in Figure 12.

740 The top level (i.e., device-level) judgment has just one rule: **S-PROGRAM**. This rule acts as a “frame”
 741 for the lower level (i.e., thread-level) judgment, and steps a collection of thread-id-indexed local
 742 memories (L), a collection of block-id-indexed shared memories (S), a global memory (Σ), and a *thread*
 743 *pool* to an updated collection of memories and thread pool. The thread pool maps thread and block ids
 744 to code, intuitively representing the program being executed by each thread at the current moment.
 745 The **S-PROGRAM** non-deterministically chooses a thread id and block id and steps it according to
 746 the thread-level judgment. This allows the semantics to model the full range of non-deterministic
 747 behavior arising from the GPU’s thread scheduler.²

748 The thread-level judgment has the shape $\eta, \sigma, \Sigma, t, b, p \vdash^\pi s \rightsquigarrow s' \dashv \eta', \sigma', \Sigma'$. The η s in this judgment
 749 represent the thread’s local memory, the σ s represent the shared memory, the Σ s represent global
 750 memory, and the s ’s represent the statement being executed by the thread. The primed versions
 751 of each of these are the output of the step judgment. Critically, notice that a π also appears on the
 752 thread-level judgment just as it does on the typing judgment. This is because the thread semantics
 753 *dynamically tracks and verifies privileges*. The same way a program can get stuck if a value does not
 754 have the right type, the semantics of Bundl also get stuck if code attempts to access data or invoke
 755 commands with the wrong privilege. This runtime privilege is present in Bundl, but is erased by Snap
 756 during compilation; we will later use it to prove that well-privileged programs will always execute
 757 with the same privilege that the type system checked them against.

758 Other variables in the judgment track the positions of threads within their scopes. The t variable
 759 represents the thread’s id, b is the id of the block in which the thread lives, and p is the relative
 760 position of the thread in the compute privilege with which it is executing. For example, threads
 761 with a (Thread,4) privilege each have a unique p between 0 and 3, while threads with a (Block,16)
 762 privilege will have the same p within in each block, between 0 and 15 depending on the block.

763 The semantic rules for privileges involve manipulating p to track which threads take which code
 764 paths when privileges are split or grouped. Notice that in the **S-PROGRAM** rule, the thread stepping
 765 judgment always begins with privilege (Grid,1): all the privilege management rules are congruences,
 766 handling further evaluation with lower privileges as determined by the particular syntactic construct
 767 used.

768 The first set of these privilege management rules handle the `split` construct. They enforce the
 769 same alignment requirements as the **T-SPLIT** rule. The rules choose which path to take based on the
 770 value of p : if p is less than the size of the left-hand privilege n_1 , execution of s_1 will proceed with
 771 that lower privilege via the **S-SPLIT-LEFT** rule, while if p is between n_1 and n_2 execution of s_2 will
 772 proceed via the **S-SPLIT-RIGHT**. In this latter case, we execute s_2 with a new p value that subtracts off
 773 the value of n_1 . As an example, if we encounter a `split(2,2){s1}{s2}` with a (Thread,4) privilege, the
 774 first two threads will go “left” and execute s_1 with a (Thread,2) privilege, while the third and fourth
 775 threads will go “right” and execute s_2 with a (Thread,2) privilege.

776 The **S-GROUP** rule is much simpler, as every thread with a compute privilege that encounters this
 777 operation will execute the sub-statement s . What changes across threads is how their p value will be
 778 modified by the grouping operation. In each case, the p of the thread is reduced modulo n , where n is
 779

780 ¹Bundl is abstracted over these B and T values, so instead of tracking classes of privileges the way that Snap does, it only
 781 tracks the top-level privilege described in Section 3.2.

782 ²In reality, the GPU’s thread scheduler guarantees that warps execute their threads in lockstep, but this modeling every thread
 783 as completely independent is both simpler and a conservative overestimate of the nondeterministic behavior of the GPU.

785 the size of the privilege with which s is executed. This effectively recolors all the threads privilege
 786 the new, smaller, privilege with which they will see s .

787 The **S-DESTRUCT-BLOCK** rule for `destruct` behaves similarly to **S-GROUP**. However, as p is an
 788 index into the current privilege, it is always 0 when a privilege has size 1. The `destruct` statement
 789 can only be executed with a size 1 privilege, and so p is not particularly useful here. To compute the
 790 new p for the execution of s with a privilege of size T , we reduce the thread id modulo T and continue
 791 executing s with that as the new value of p . A similar rule exists for destructing a `Grid` privilege into
 792 a size B `Block` privilege.

793 These rules take great care to ensure that p always describes the relative position of a compute
 794 resource within its privilege; the payoff is that Bundl’s semantics can later use this p value to model
 795 the way that Snap automatically adjusts indices into data when partitioning a data privilege.

796 Beyond these key rules for compute privilege management, we have modeled many of the other
 797 features of Snap, such as asynchronous operations and thread synchronization, in Bundl. To handle
 798 features like these we equip the operational semantics with additional structure, including sets of
 799 semaphores [6] for thread syncs and a stack of effect handlers for modeling deferred asynchronous
 800 computations in the style of Ahman and Pretnar [2]. We have elided these details here for simplicity,
 801 but interested readers can see further details in Appendix A.3.

802 4.3 Soundness Theorem

803 Together, the type system and operational semantics help us prove the following syntactic soundness
 804 theorem, which says that Bundl programs are type safe and do not get stuck trying to execute
 805 operations for which they lack the required privilege:

806 THEOREM 4.1. (*Type-and-Privilege Safety*). *For any program s such that $\Gamma \vdash^\pi s$, either:*

- 807 (1) s is `skip`, or
- 808 (2) *for any well-typed environments η , σ , and Σ , there is an s' , η' , σ' , and Σ' such that*
 $\eta, \sigma, \Sigma, t, b, p \vdash^\pi \sim s' \dashv \eta', \sigma', \Sigma'$ *where η' , σ' , and Σ' are well-typed by an extension of Γ' .*

809 PROOF. Via the usual progress and preservation lemmas. The full proof of these lemmas can be
 810 found in Appendix A.4. \square

811 It is worth noting that this soundness theorem is a safety theorem, not a liveness theorem. The
 812 theorem does guarantees that if all threads that can reach a point in the program with statically-
 813 determined privilege π do reach that point, then those threads will be logically grouped according to
 814 that same privilege π . It does not, however, guarantee that all threads that *can* reach the point will
 815 eventually do so. Indeed, in the presence of nontermination, liveness does not hold: a subset of the
 816 threads could split off and loop forever. While we believe the liveness version of this theorem holds
 817 for a terminating fragment of Bundl, it is beyond the scope of this paper—such proofs are notoriously
 818 challenging and are often research contributions [3, 7, 16] in and of themselves. We hope to tackle
 819 this proof for Bundl in future work.

820 5 Compilation

821 If a program typechecks, Snap lowers the program into a CUDA file, which does not contain any
 822 run-time checks. The CUDA file is then compiled by nvcc, NVIDIA’s closed-source compiler, to
 823 produce a machine-executable. As Snap is an imperative language that provides the same level of
 824 abstraction as CUDA, there is a one-to-one mapping between most language constructs and their
 825 CUDA counterparts.

826 Snap is implemented as a compiler tool on top of the CLANG/LLVM compiler stack. Textually,
 827 Snap programs are written as the examples described in the paper, but the constructs themselves

lower to traditional C++ statements that our frontend tool can analyze and generate code for. This is accomplished by making liberal use of attributes, macros, and other sugar features of C++.

- Make correspondence to CUDA clear
- Would be awesome to have a super clear figure about how the compilation step actually works (show the one-to-one mapping, though it is slightly more complicated)
- Make optimizations (lazy sync, commit group etc) as a separate optimization section.
- highlight how new function interfaces work out: we add extra arguments

6 Evaluation

We evaluate Snap in the context of the four main questions:

RQ1 Can Snap express a diversity of CUDA programs?

RQ2 Can Snap express programs that use advanced GPU features?

RQ3 Can Snap match the performance of existing, speed-of-light CUDA code?

Although the following question is not central to our paper’s main claims—since library design ultimately reflects choices made by programmers, and our language allows abstraction design much like any other language—we were, nevertheless, curious to explore it based on empirical and anecdotal observation:

RQ4 Can Snap help build compositional libraries that users can use with confidence?

To perform this evaluation, we use two GPUs. The first is the NVIDIA H100 SXM5, a server-grade chip that supports tensor-core operations and a dedicated hardware copy engine, the Tensor Memory Accelerator. Notably, the H100 introduces a new *logical* level called the *warp group* (collection of 4-aligned warps), and we show that our programming model can accommodate this new level. Moreover, because the H100 has historically served as the primary GPU for large-scale AI training, many CUDA kernels on this hardware are already highly optimized and achieve near-speed-of-light performance, providing a rigorous baseline for comparison. To ensure our results generalize beyond the H100, we additionally test programs on a second GPU, the NVIDIA 4070 SuperTi GPU, a consumer-grade chip.

Like mentioned in Section 5, when name typechecks a program, it produces a CUDA file. We compile the CUDA file with nvcc version with flags . To measure performance of our benchmark, we set up the values from an input using a random number generator. We use the average running time sampled across 10 iterations, with a warm-up of 5 iterations.

6.1 RQ1: Can Snap express a diversity of CUDA programs?

To answer this research question, we evaluate Snap on two benchmarks: a sequence of matrix multiplication kernels that grow in complexity and a single-pass parallel prefix scan with decoupled look-back that requires programmers to think carefully about multiple points of convergence.

Matrix Multiplication. We chose matrix multiplication as our first benchmark for two main reasons. First, there exist several implementations that achieve near-peak performance, providing a strong baseline. Second, matrix multiplication allows for a range of increasingly sophisticated implementations that stress different parts of the language, making it ideal for evaluating expressiveness.

To undertake this exploration, we adapt the codebase from to implement a matrix multiplication on the ffloat datatype, also known as sgemm, in Snap. As this is a ffloat benchmark, it *does not need to* use advanced GPU features like asynchrony or tensor cores to achieve speed-of-light performance. We will discuss these advanced features in Section 6.2.

We implement several variants of the sgemm benchmark, reproduced in Appendix . These include a naive baseline, a version that exploits memory coalescing, 2D blocking with shared-memory staging,

883 2D blocking with vectorized loads, and a warp-tiling strategy that effectively introduces an additional
 884 level of tiling from the perspective of a warp. We find that the runtime performance of our Snap
 885 implementations closely matches that of the original versions. This is expected, since the generated
 886 code is nearly identical to the hand-written versions, aside from small differences such as hoisted
 887 expressions and a few index calculations. The downstream compiler (`nvcc`) readily inlines hoisted
 888 expressions that are used only once, so the performance-critical inner loops remain effectively
 889 unchanged.

890
 891 *Single-pass Parallel Prefix Scan with Decoupled Look-Back.* We also implement scan, a widely used
 892 parallel primitive, in Snap. We focus on the prefix-sum scan, which computes, for each position in an
 893 array, the sum of all elements up to that position. Prefix-sum sits in a different corner of the GPU
 894 design space from matrix multiplication: it is memory intensive, requires careful attention to the
 895 convergence behavior of different threads, and traditionally requires multiple passes over data.

896 We implement the single-pass parallel prefix scan with decoupled look-back, introduced by Merrill
 897 and Garland , an elegant algorithm that does not require multiple passes over the input data. In
 898 this algorithm, each block computes a local prefix sum over its assigned region of the array. Once
 899 finished, it writes the final element of its region into a global array. Each block then *looks back* to
 900 accumulate the contributions of prior blocks, allowing them to independently determine the global
 901 prefix without a full sweep over memory. This decoupled look-back mechanism lets blocks progress
 902 at different speeds while still producing correct global results.

903 The algorithm involves several distinct points of convergence. Within each block, work is de-
 904 composed into fine-grained thread-level and warp-level scans. After producing the local result,
 905 blocks publish their partial prefix to global memory. Finally, each block waits until enough prefix
 906 information from earlier blocks becomes available, at which point it accumulates the value and
 907 completes its section of the global scan.

908 We implement this full strategy in Snap, available in Appendix . Our implementation uses the
 909 unsafe feature to implement a global-memory spinlock that lets blocks check when the preceeding
 910 block's data is ready.

911 6.2 RQ2: Can Snap express programs that use advanced GPU features?

912 To test whether Snap can express programs that use advanced features of modern GPUs, we write a
 913 matrix multiplication for the `bf16` datatype for the H100.

914 This benchmark is an acid test of our language, as the H100 `bf16` matrix multiplication pushes
 915 several language features to the extreme. To write a matrix multiplication that can hit peak throughput
 916 on an H100, we need to write a warp-specialized kernel that uses the Tensor Memory Accelerator
 917 (TMA) , an asynchronous hardware copy engine that can move tiles of data at a time, and uses
 918 the warp-group-level tensor core instructions, or `wmma`, specifically introduced for the Hopper
 919 architecture.

920 The implementation works as follows. First, we assign each block on our machine a logical tile of
 921 the output to compute. The block is then divided into a producer warp-group and a set of consumer
 922 warp-groups, where the producer loads data for multiple consumers, and both must signal to each
 923 other when one is done loading and the other is ready to compute. In this way, the benchmark
 924 overlaps computation with data movement by pipelining loads.

925 The implementation in Snap looks different from normal CUDA code, particularly in how pipelining
 926 is expressed. Since Snap uses *names* and subsequent `partition` or `claim` constructs to determine the
 927 synchronization each region requires, when pipelining, we cannot dynamically change the pipeline
 928 slot simply by maintaining an index variable that wraps around based on the pipeline length. Instead,
 929 each pipeline slot must be given a separate name so that Snap can track them independently and
 930

actually overlap compute with data-movement. This leads to pipeline slots that must be individually named and forces the load logic to be effectively “unrolled,” since we can no longer iterate over a pipeline-slot index. In turn, this forces that all pipelines in Snap be statically sized. In practice, these pipelines are statically sized in CUDA anyway because they occupy shared memory, which is a small, finite resource that must be explicitly managed.

Notably, for this benchmark, to get the wmma instruction to work we did not need to add a special privilege in our language; we could simply use `thread[128]`. We did, however, need to add a special TMA asynchronous data-movement construct, since Snap will eventually need to synchronize these transfers.

6.3 RQ3: Can Snap match the performance of existing, speed-of-light CUDA code?

In Section 6.1 and Section 6.2, we examined programs that expressed the same computation in multiple ways, expressed operations that rely on multiple points of convergence, and used advanced GPU features. We now discuss the performance of these programs.

The performance of the matrix multiplication benchmarks is shown in , where we demonstrate that Snap is competitive with cuBLAS.

For the prefix scan, we compare our performance to , the library introduced earlier in Section 1, and show that we can reach approximately % of CUB’s peak bandwidth.

Finally, we evaluate our H100 implementation and show that it is competitive with cuBLAS on square sizes, falling between %. We emphasize that this is an exacting benchmark, and achieving this level of performance requires writing large kernels with precise control over low-level features.

6.4 Case Study: Can Snap help with library design?

Over the course of our evaluation, we found ourselves developing a small library of functions—similar in spirit to the CUB library—that we could call to execute our operations. Based on our experience, we would like to study, qualitatively, if Snap can help programmers design libraries that they can compose with confidence.

As mentioned in Section 1, CUB occupies a unique design point in the GPU library ecosystem. Unlike many other libraries such as cuBLAS, cuDNN, and cuSparse, which provide global interfaces that users can call and configure, CUB provides a device-side library organized into different levels. It makes these levels apparent by prefixing each of its functions with Device, Block, and Thread. The single-scan prefix sum in Section 6.1 used these functions in Snap already.

Now, we turn our attention to a particular function in CUB—the store function—and examine how it is equivalently expressed in Snap, and how, in our experience, Snap’s interfaces the assumptions implicit in CUB apparent to the user.

In CUB, the load function is implemented as a class, as shown below.

```
1 template<typename T, int BlockDimX,
2         int ItemsPerThread, BlockLoadAlgorithm Algorithm = BLOCK_LOAD_DIRECT,
3         int BlockDimY = 1, int BlockDimZ = 1>
4 class BlockStore:
```

To use it, users must first instantiate an object of this class, and then call it with shared memory.

```
1 using BlockLoad = cub::BlockLoad<int, 128, 4, BLOCK_LOAD_DIRECT>;
2 // Allocate shared memory for BlockLoad
3 __shared__ typename BlockLoad::TempStorage temp_storage;
4 int thread_data[4]; // Thread local data
5 BlockLoad(temp_storage).Load(d_data, thread_data);
```

981 CUB exposes a leaky abstraction, where information about the number of threads, block sizes, and
 982 other details leaks through. There are several details that leak through:

- 983 (1) The CUB documentation needs to specify the number of threads that the function can assume
 984 to be available, because within the function, each thread must locate itself in the computation
 985 and use its `threadIdx.x` accordingly. If the starting `threadIdx.x` is not 0, the function must
 986 compute its relative ID internally.
 987 (2) The “item per thread” design is interesting and serves two purposes. The first is related to
 988 performance: if loops have constant bounds, they can be unrolled, enabling downstream
 989 optimizations. The second is related to correctness: the function relies on the assumption
 990 that all threads call the function with an equal number of values to load.
 991 (3) The CUB documentation also needs to specify that `thread_data` is local; if the same pointer
 992 to shared memory is passed, threads may overwrite each other’s values.
 993 (4) Finally, the CUB documentation specifies that if shared memory is being overwritten, a
 994 `__syncthreads()` call must be made to ensure that all reads have completed.
 995

996 On the other hand, the same function has a completely different interface in Snap:
 997

```
998 1 #include <cuda.h>
999 2 #require(block[1], thread[1])
1000 3 def load(src: ptr(const(shared(int))) @ block[1],
1001 4     item_per_thread : int @ block[1],
1002 5     total_size: int @ block[1],
1003 6     thread_data : ptr(int) @ thread[1]:
```

1003 In Snap, we are able to encapsulate code effectively, reducing the need to communicate numerous
 1004 implementation details through documentation:

- 1005 (1) We do not need to pass in the number of threads at all. Whenever Snap calls a function, it
 1006 threads the relative ID through, so each function can be written locally as if it were running
 1007 alone, rather than having to determine where the thread resides in the global array.
 1008 (2) We do not need to make `item_per_thread` a template argument for *correctness*. Its frequency
 1009 is set at the function signature, so Snap will never allow a function to be called with a
 1010 lower-privilege value.
 1011 (3) In our interface, `thread_data` is explicitly set to a thread-local value. Since it is not marked
 1012 as `const`, Snap conservatively assumes it may be written to, and enforces at compile time
 1013 that only `thread[1]` values are passed in.
 1014 (4) Finally, using our synchronization pass outlined in Section 3.9, a `__syncthreads()` call will
 1015 be inserted automatically if `src` is going to be used for writing.
 1016

1017 7 Related Work

1018 *GPU Programming Languages.*

1019 *Theoretic Foundations.* The design of Bundl is heavily inspired by existing work on coeffect systems
 1020 [12, 13] and dependency tracking [1]. Coeffects allow type systems to talk about how programs
 1021 depend on their environments, and have achieved widespread use in Rust [14] in the form of linear
 1022 types. In Snap and Bundl, privileges act like coeffects by describing what compute resources are
 1023 necessary for programs to execute. Dependency tracking calculi, meanwhile, allow type systems to
 1024 track how data and code depend on each other, and have commonly been used to implement secure
 1025 information flow analyses [5]. In Bundl, low privilege data is unable to flow into high privilege contexts,
 1026 and we use dependency tracking to capture this restriction in Bundl’s type system.
 1027

1030 8 Conclusion, Limitations, and Future Work

1031 We have presented Snap, a new low-level GPU language that statically guarantees safe usage of
 1032 compute resources by construction, and have demonstrated that it is possible to achieve this safety
 1033 without sacrificing performance. We hope our design can inform the implementation of existing
 1034 libraries like CUB that attempt to achieve the same goals without static enforcement.

1035 We believe Snap is ripe for extension with additional features and capabilities to further improve its
 1036 expressivity and usefulness to GPU programmers. Currently, Snap does not have support for explicit
 1037 pointers; programmers must use built-in constructs in order to interact with memory. However,
 1038 prior work like Descend [9] provides a promising template for extending Snap with more permissive
 1039 pointer usage. In particular, were we to restrict Snap’s partition construct to a specific set of
 1040 injective functions identified by Descend, we believe we could guarantee data-race freedom in Snap.
 1041

1042 Another promising area for future work is expanding the set of architectures on which Snap and
 1043 its design principles can be used. While Snap currently only supports programming the GPU, we
 1044 believe we can extend the core ideas cleanly to the boundary with the CPU and in general to other
 1045 heterogeneous hardware architectures. Indeed, we do not think that the idea of compute privileges
 1046 is useful only on the GPU, and we would like to investigate how we might adapt Snap to other
 1047 hierarchical programming systems. Conversely, we also believe it is possible for Snap to support
 1048 more GPU architectures featuring even coarser-grained tensor-core operations like Blackwell [4].
 1049

1050 We also would like to improve the ergonomics of programming with Snap, in particular by
 1051 improving the pipelining experience and providing users with more control over synchronization.
 1052 As we saw in Section 6.2, Snap requires pipeline slots to be explicitly named and all pipelines to
 1053 be statically sized. We believe we can improve upon this by adding explicit language constructs
 1054 for organizing pipelines to reduce boilerplate code in this case. Similarly, all syncs are currently
 handled by Snap’s compiler, but we would like to extend the language with user-level synchronization
 1055 operations, leveraging the existing privilege system to avoid deadlocking.

1056 On the theoretical side, we are interested in further exploration of Bundl. In particular, we would
 1057 like to examine a terminating fragment of the calculus and prove the liveness property discussed in
 1058 Section 4.3: that every thread in the logical grouping denoted by privilege π will always reach the
 1059 parts of a program with that code privilege. We believe this is equivalent to showing that all threads
 1060 with the same privilege always execute the same code and observe the same data, and are optimistic
 1061 that this can be proved via logical relation, following in the footsteps of Turon et al. [16].
 1062

References

- 1063 [1] Martín Abadi, Anindya Banerjee, Nevin Heintze, and Jon G. Riecke. 1999. A core calculus of dependency. In *Proceedings of the 26th ACM SIGPLAN-SIGACT symposium on Principles of programming languages* (San Antonio Texas USA, 1999-01). ACM, 147–160. [doi:10.1145/292540.292555](https://doi.org/10.1145/292540.292555)
- 1064 [2] Danel Ahman and Matija Pretnar. 2021. Asynchronous effects. *Proceedings of the ACM on Programming Languages* 5, POPL (Jan. 2021), 1–28. [doi:10.1145/3434305](https://doi.org/10.1145/3434305)
- 1065 [3] Lars Birkedal, Filip Sieczkowski, and Jacob Thamsborg. 2012. A Concurrent Logical Relation. In *Computer Science Logic (CSL’12) - 26th International Workshop/21st Annual Conference of the EACSL (2012)*. Schloss Dagstuhl – Leibniz-Zentrum für Informatik, 107–121. [doi:10.4230/LIPIcs.CSL.2012.107](https://doi.org/10.4230/LIPIcs.CSL.2012.107)
- 1066 [4] NVIDIA Corporation. 2025. NVIDIA Blackwell Architecture Technical Brief. (2025). <https://resources.nvidia.com/en-us-blackwell-architecture>
- 1067 [5] Dorothy E. Denning and Peter J. Denning. 1977. Certification of programs for secure information flow. *Commun. ACM* 20, 7 (July 1977), 504–513. [doi:10.1145/359636.359712](https://doi.org/10.1145/359636.359712)
- 1068 [6] Edsger Wybe Dijkstra. 1963. Over de sequentialiteit van procesbeschrijvingen. (1963). <https://training-ir7.tdl.org/handle/123456789/594>
- 1069 [7] Azadeh Farzan, Zachary Kincaid, and Andreas Podelski. 2016. Proving Liveness of Parameterized Programs. In *Proceedings of the 31st Annual ACM/IEEE Symposium on Logic in Computer Science*. ACM, New York NY USA, 185–196. [doi:10.1145/2933575.2935310](https://doi.org/10.1145/2933575.2935310)

- 1079 [8] Troels Henriksen, Niels G. W. Serup, Martin Elsman, Fritz Henglein, and Cosmin E. Oancea. 2017. Futhark: purely
1080 functional GPU-programming with nested parallelism and in-place array updates. In *Proceedings of the 38th ACM
1081 SIGPLAN Conference on Programming Language Design and Implementation*. ACM, Barcelona Spain, 556–571. [doi:10.1145/3062341.3062354](https://doi.org/10.1145/3062341.3062354)
- 1082 [9] Bastian Köpcke, Sergei Gorlatch, and Michel Steuwer. 2024. Descend: A Safe GPU Systems Programming Language. 8
1083 (2024), 841–864. Issue PLDI. [doi:10.1145/3656411](https://doi.org/10.1145/3656411)
- 1084 [10] NVIDIA Corporation. 2025. *CUB: CUDA Unbound*. <https://docs.nvidia.com/cuda/cub/index.html> CUDA Toolkit
1085 Documentation.
- 1086 [11] NVIDIA Corporation. 2025. *CUDA C++ Programming Guide*. NVIDIA. <https://docs.nvidia.com/cuda/cuda-c-programming-guide/> Accessed: 2025-11-06.
- 1087 [12] Tomas Petricek, Dominic Orchard, and Alan Mycroft. 2013. Coeffects: Unified Static Analysis of Context-Dependence.
1088 In *Automata, Languages, and Programming*, Fedor V. Fomin, Rūsiņš Freivalds, Marta Kwiatkowska, and David Peleg
1089 (Eds.). Springer, Berlin, Heidelberg, 385–397. [doi:10.1007/978-3-642-39212-2_35](https://doi.org/10.1007/978-3-642-39212-2_35)
- 1090 [13] Tomas Petricek, Dominic Orchard, and Alan Mycroft. 2014. Coeffects: a calculus of context-dependent computation.
1091 *SIGPLAN Not.* 49, 9 (Aug. 2014), 123–135. [doi:10.1145/2692915.2628160](https://doi.org/10.1145/2692915.2628160)
- 1092 [14] Rust Team. 2025. Rust Programming Language. <https://rust-lang.org/>
- 1093 [15] Philippe Tillet. 2025. Introducing Triton: Open-source GPU programming for neural networks. <https://openai.com/index/triton/>
- 1094 [16] Aaron J. Turon, Jacob Thamsborg, Amal Ahmed, Lars Birkedal, and Derek Dreyer. 2013. Logical relations for fine-grained
1095 concurrency. In *Proceedings of the 40th annual ACM SIGPLAN-SIGACT symposium on Principles of programming languages*
(Rome Italy, 2013-01-23). ACM, 343–356. [doi:10.1145/2429069.2429111](https://doi.org/10.1145/2429069.2429111)

1097

1098

1099

1100

1101

1102

1103

1104

1105

1106

1107

1108

1109

1110

1111

1112

1113

1114

1115

1116

1117

1118

1119

1120

1121

1122

1123

1124

1125

1126

1127

1128 **A Complete Bundl Type System, Semantics, and Syntactic Soundness Proofs**

1129 **A.1 Basic Definitions**

1130
 1131 Hierarchy Levels $h ::= \text{Grid} \mid \text{Block} \mid \text{Thread}$
 1132 Memory Kinds $l ::= \text{Local} \mid \text{Shared} \mid \text{Global}$
 1133 Privileges $\pi : h \times \mathbb{N}$
 1134 Base Types $b ::= \text{bool} \mid \text{int} \mid \text{float}$
 1135 Types $\tau ::= b \mid b[]^l \mid \text{Fun}(\Gamma, \pi, m) \mid \text{async } \tau$
 1136 Contexts $\Gamma ::= \cdot \mid \Gamma, x :^\pi \tau$
 1137
 1138 Shared Memory Remaining $m : \mathbb{N}$
 1139
 1140

1141 Privileges π are part of an algebra parameterized over some constant values T (the number of threads
 1142 per block) and B (the number of blocks per grid). With these values, we have two isomorphisms:

$$(Block, 1) \cong (\text{Thread}, T)$$

1144 and

$$(\text{Grid}, 1) \cong (\text{Block}, B)$$

1145 The group and split operations of Snap allows us to move along this isomorphism, from left
 1146 to right. For clarity, in Bundl we split these operations into three: a split operation that can split
 1147 privileges into multiple smaller ones, and a destruct operation that directly moves us along the
 1148 isomorphism, and a group that divides our current privilege into equal sized parts.

1149 Privileges π are also lexicographically ordered in the obvious way. hs are ordered such that
 1150 $\text{Thread} \leq \text{Block} \leq \text{Grid}$, and $(h_1, n_1) \leq (h_2, n_2)$ iff $n_1 \leq n_2$ and $h_1 \leq h_2$.

1151 We define scalar multiplication $i \times h$ of natural numbers with hs : $i \times (h, n) = (h, in)$.

1152 We also define division of privileges and hierarchy levels of type $\pi \times \pi \rightarrow \mathbb{N}$. $\text{Grid}/\text{Block} = B$ and
 1153 $\text{Block}/\text{Thread} = T$. We lift this to privileges like so: $(h_1, n_1)/(h_2, n_2) = ((h_1/h_2) \cdot n_1)/n_2$.

1154 Lastly we define a partial \downarrow operator on hs such that $\downarrow \text{Block} = \text{Thread}$ and $\downarrow \text{Grid} = \text{Block}$. Note
 1155 that $\downarrow \text{Thread}$ is undefined. This operator lifts to π s whose second component is 1 and encodes the
 1156 leftward component of the isomorphism above: $\downarrow (Block, 1) = (\text{Thread}, T)$ and $\downarrow (\text{Grid}, 1) = (\text{Block}, B)$.

1157 Note also that the presentation of these rules in the main body of the paper elide the m portion,
 1158 which tracks the maximum amount of memory a given computation is allowed to use. In the full
 1159 system presented here, both the typing rules and the operational semantics carry an additional piece
 1160 of information tracking allocated memory.

1161 **A.2 Typing Rules**

1162 **A.2.1 Expressions.**

$$\frac{x :^\pi \tau \in \Gamma}{\Gamma \vdash^\pi x : \tau} \text{-Var} \quad \frac{}{\Gamma \vdash^\pi n : \text{int}} \text{-Int} \quad \frac{}{\Gamma \vdash^\pi f : \text{float}} \text{-Float}$$

$$\frac{}{\Gamma \vdash^\pi b : \text{bool}} \text{-Bool} \quad \frac{\pi < (\text{Grid}, 1)}{\Gamma \vdash^\pi \text{partition_id} : \text{int}} \text{-Partition-Id}$$

$$\frac{\Gamma \vdash^{\pi'} e_1 : \tau[]^l \quad \Gamma \vdash^\pi e_2 : \text{int} \quad l = \text{Global} \text{ or } l = \text{Local} \quad \pi \leq \pi'}{\Gamma \vdash^\pi e_1[e_2] : \tau} \text{-Arr-Access}$$

1177	$\frac{\Gamma \vdash^{\pi'} e_1 : \tau[]^{\text{Shared}} \quad \Gamma \vdash^{\pi} e_2 : \text{int} \quad \pi \leq (\text{Block}, 1) \quad \pi \leq \pi'}{\Gamma \vdash^{\pi} e_1[e_2] : \tau}$	T-Arr-Access-Shared
1178		
1179		
1180	$\frac{\Gamma \vdash^{\pi} e_1 : \text{int} \quad \Gamma \vdash^{\pi} e_2 : \text{int}}{\Gamma \vdash^{\pi} e_1 \text{ bop } e_2 : \text{int}}$	T-Bop
1181		
1182	$\frac{\Gamma \vdash^{\pi} e_1 : \text{int} \quad \Gamma \vdash^{\pi} e_2 : \text{int}}{\Gamma \vdash^{\pi} e_1 \text{ cmp } e_2 : \text{bool}}$	T-Cmp
1183		
A.2.2	<i>Statements.</i>	
1184	$\frac{f : \pi' \text{Fun}(x_i : \tau_i, \pi, m') \in \Gamma \quad \Gamma \vdash^{\pi} e_i : \tau_i \quad m' \leq m}{\Gamma \vdash_m^{\pi} f(e_1, \dots, e_n)}$	T-Function-Call
1185		
1186		
1187	$\frac{\Gamma \vdash_m^{(h,n_1)} s_1 \quad \Gamma \vdash_m^{(h,n_2)} s_2 \quad n_1, n_2 \text{ align to } n}{\Gamma \vdash_m^{(h,n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\}}$	T-Split
1188		
1189		
1190	$\frac{\Gamma \vdash_m^{(h,n)} s}{\Gamma \vdash_m^{(h,q \cdot n)} \text{group } q \ s}$	T-Group
1191		
1192	$\frac{x : \pi' \tau \in \Gamma \quad \Gamma \vdash^{\pi} e : \tau \quad \pi' \leq \pi \quad \pi' \pi}{\Gamma \vdash_m^{\pi} x = e}$	T-Assn
1193		
1194	$\frac{\Gamma \vdash_m^{\pi} \text{init}_{\psi}}{\Gamma \vdash_m^{\pi} \text{wait}_{\psi}}$	T-Sync-Init
1195		
1196	$\frac{}{\Gamma \vdash_m^{\pi} \text{wait}_{\psi}}$	T-Sync-Wait
1197		
1198	$\frac{\Gamma \vdash^{\pi} e : \tau \quad \Gamma, x : \pi' \tau \vdash_m^{\pi} s \quad \pi' \leq \pi \quad \pi' \pi \quad \tau \text{ not an array type}}{\Gamma \vdash_m^{\pi} x : \tau @ \pi' = e \text{ in } s}$	T-Decl
1199		
1200		
1201	$\frac{\Gamma \vdash^{\pi'} e_1 : \tau[]^l \quad \Gamma \vdash^{\pi} e_2 : \text{int} \quad \Gamma \vdash^{\pi} e_3 : \tau \quad l = \text{Global or } l = \text{Local} \quad \pi' \leq \pi \quad \pi' \pi}{\Gamma \vdash_m^{\pi} e_1[e_2] = e_3}$	T-Arr-Assn
1202		
1203		
1204	$\frac{\Gamma \vdash^{\pi'} e_1 : \tau[]^{\text{Shared}} \quad \Gamma \vdash^{\pi} e_2 : \text{int} \quad \Gamma \vdash^{\pi} e_3 : \tau \quad \pi \leq (\text{Block}, 1) \quad \pi' \leq \pi \quad \pi' \pi}{\Gamma \vdash_m^{\pi} e_1[e_2] = e_3}$	T-Arr-Assn-Shared
1205		
1206		
1207	$\frac{\Gamma \vdash^{\pi} e : \text{bool} \quad \Gamma \vdash_{m_1}^{\pi} s_1 \quad \Gamma \vdash_{m_2}^{\pi} s_2}{\Gamma \vdash_{\max(m_1, m_2)}^{\pi} \text{if } e \text{ then } s_1 \text{ else } s_2}$	T-If
1208		
1209		
1210	$\frac{\Gamma \vdash^{\pi} e : \text{bool} \quad \Gamma \vdash_m^{\pi} s}{\Gamma \vdash_m^{\pi} \text{while } e \text{ do } s}$	T-While
1211		
1212	$\frac{\Gamma \vdash_{\max(m_1, m_2)}^{\pi} s_1 ; s_2}{\Gamma \vdash_{\max(m_1, m_2)}^{\pi} s_1 ; s_2}$	T-Seq
1213		
1214	$\frac{\Gamma, x : \pi' \tau[]^l \vdash_m^{\pi} s \quad l = \text{Global or } l = \text{Local}}{\Gamma \vdash_{m+n \cdot \text{size}(\tau)}^{\pi} x := \text{alloc } l \ \tau \ n \text{ in } s}$	T-Alloc
1215		
1216	$\frac{\Gamma, x : (\text{Block}, 1) \tau[]^{\text{Shared}} \vdash_m^{\pi} s}{\Gamma \vdash_{m+n \cdot \text{size}(\tau)}^{\pi} x := \text{alloc Shared } \tau \ n \text{ in } s}$	T-Alloc-Shared
1217		
1218		
1219	$\frac{\Gamma, y : (h, n/c) \tau[]^l \vdash_m^{(h,n)} s \quad c n \quad l \neq \text{Local}}{\Gamma, x : (h, n) \tau[]^l \vdash_m^{(h,n)} \text{partition}_{\psi} x \text{ into } y \text{ by } c \text{ in } s}$	T-Partition
1220		
1221		
1222	$\frac{\Gamma, y : (h, n') \tau[]^l \vdash_m^{(h,n')} s \quad n' \leq n \quad l \neq \text{Local}}{\Gamma, x : (h, n) \tau[]^l \vdash_m^{(h,n)} \text{claim}_{\psi} x \text{ into } y \text{ at } n' \text{ in } s}$	T-Claim
1223		
1224		
1225		

$$\begin{array}{c}
 \frac{\Gamma, y : \downarrow^\pi \tau[]^l \vdash_m^\pi s \quad l \neq \text{Local}}{\Gamma, x : \pi \tau[]^l \vdash_m^\pi \text{lower}_\psi x \text{ into } y \text{ in } s} \textbf{T-Lower} \\
 \\
 \frac{\Gamma, y : (\text{Thread}, 1) \text{ async } \tau[]^l \vdash_m^{(\text{Thread}, 1)} s}{\Gamma, x : (\text{Thread}, 1) \tau[]^l \vdash_m^{(\text{Thread}, 1)} \text{async_partition}_\phi x \text{ into } y \text{ in } s} \textbf{T-Async-Partition} \\
 \\
 \frac{}{\Gamma, x : (\text{Thread}, 1) \text{ async } \tau[]^l, y : (\text{Thread}, 1) \tau[]^{l'} \vdash_m^{(\text{Thread}, 1)} \text{async_memcpy}(x, y)} \textbf{T-Async-Memcpy} \\
 \\
 \frac{}{\Gamma, x : \pi \tau[]^l, y : \pi \tau[]^{l'} \vdash_m^\pi \text{memcpy}(x, y)} \textbf{T-Memcpy}
 \end{array}$$

A.3 Complete Bundl Semantics

A.3.1 Definitions.

Global Memory $\Sigma ::= \cdot | \Sigma, n \mapsto^\pi v$

Shared Memory $\sigma ::= \cdot | \sigma, n \mapsto^\pi v$

Local Memory $\eta ::= \cdot | \eta, n \mapsto^\pi v$

Block Memory Map $S ::= \forall n \in B, n \mapsto \sigma$

Thread Memory Map $L ::= \forall n \in T, n \mapsto \eta$

Synchronization Map $\Psi : \psi \rightarrow \mathbb{N} \rightarrow \mathbb{N}$

Deferred Computations Map $\Phi : \phi \rightarrow \{s\}$

In real GPUs, thread ids are only unique within their block. However, in this calculus for simplicity we assume thread ids are global. One can convert back and forth between this abstracted notion of a thread id and a block-unique id via addition modulo T . That is, $t_{\text{real}} = t_{\text{simplified}} \bmod T$ and $t_{\text{simplified}} = t_{\text{real}} + b \cdot T$.

By convention the names for local and shared and global memory do not conflict, as on the GPU they will be separate pointer spaces. Additionally, we freely interchange between using names for variables and integer locations.

In the main body of the paper, for simplicity we elide the synchronization map and the deferred computation map from the operational semantics, as our theorems do not make any guarantees about non-interference. However, as they are part of the full semantics, we include them here for completeness. By convention the synchronization and deferred computation maps are a total functions, initialized to map to $\lambda_. 0$ for ψ s and $\lambda_. \{s\}$ for ϕ s not explicitly initialize.

The shape of the judgment for a single thread is $\eta, \sigma, \Sigma, t, b, p, \Psi \vdash_m^\pi s \rightsquigarrow s' \dashv_m \eta', \sigma', \Sigma', \Psi'$. The t here is the thread id, the b is the block id, and the p is the partition id. The last of these three is modified and managed by the rules for split, group and destruct, and tracks the relative position of the thread within a group. This semantics is in a small step style.

The shape of the judgment for expressions is $\eta, \sigma, \Sigma \vdash_\pi^\pi e \Downarrow v$. The two π s represent the ambient compute context (i.e., the context in which resources are being read), while π' , represents the target compute context (i.e., the compute context of the variable into which the result of the expression is going to be written). This is relevant for computing the value of `partition_id`, which divides the two contexts. As a shorthand, we can divide a privilege by a scalar value like so: $(h, n)/c = (h, n)/(h, c)$.

1275 The overall evaluation of a program is expressed as

$$L, S, \Sigma, P, \Psi, \Phi \rightsquigarrow L', S', \Sigma', P', \Psi', \Phi'.$$

1276 In this judgment P serves as a *thread pool*, mapping pairs of thread and block ids (which don't change)
 1277 to statements and memory (which can be updated by stepping). One can think of P as tracking which
 1280 program is running on each thread. This steps according to the following rule:

$$\frac{L(t), S(b), \Sigma, t, b, 0, \Psi, \Phi \vdash_m^{(\text{Grid}, 1)} s \rightsquigarrow s' \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi' \quad P(t, b) = (s, m)}{L, S, \Sigma, P, \Psi, \Phi \rightsquigarrow L[t \mapsto \eta'], S[b \mapsto \sigma'], \Sigma', P[(t, b) \mapsto (s', m')], \Psi', \Phi'} \text{ S-Program}$$

1281 For simplicity of notation, we define an **update** operation that searches the three environments
 1282 for the one that contains the variable being used (by convention, there is no conflict between the
 1283 environments, as in reality they exist in three separate address spaces). We also define a similar **get**
 1284 operation that retrieves a variable from memory, and a **rename** operation that remaps a variable
 1285 with the same value but under a different name.

$$\text{update}(\eta, \sigma, \Sigma, x, v) = (\eta[x \mapsto^\pi v], \sigma, \Sigma) \text{ when } x \in^\pi \eta$$

$$\text{update}(\eta, \sigma, \Sigma, x, v) = (\eta, \sigma[x \mapsto^\pi v], \Sigma) \text{ when } x \in^\pi \sigma$$

$$\text{update}(\eta, \sigma, \Sigma, x, v) = (\eta, \sigma, \Sigma[x \mapsto^\pi v]) \text{ when } x \in^\pi \Sigma$$

$$\text{get}(\eta, \sigma, \Sigma, x) = \eta(x) \text{ when } x \in^\pi \eta$$

$$\text{get}(\eta, \sigma, \Sigma, x) = \sigma(x) \text{ when } x \in^\pi \sigma$$

$$\text{get}(\eta, \sigma, \Sigma, x) = \Sigma(x) \text{ when } x \in^\pi \Sigma$$

$$\text{rename}(\eta, \sigma, \Sigma, x, y, \pi') = (\eta[y \mapsto^\pi \eta(x)], \sigma, \Sigma) \text{ when } x \in^\pi \eta$$

$$\text{rename}(\eta, \sigma, \Sigma, x, y, \pi') = (\eta, \sigma[y \mapsto^\pi \sigma(x)], \Sigma) \text{ when } x \in^\pi \sigma$$

$$\text{rename}(\eta, \sigma, \Sigma, x, y, \pi') = (\eta, \sigma, \Sigma[y \mapsto^\pi \Sigma(x)]) \text{ when } x \in^\pi \Sigma$$

A.3.2 Privilege Management Rules.

$$\frac{p < n_1 \quad n_1, n_2 \text{ align to } n \quad \eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n_1)} s_1 \rightsquigarrow s'_1 \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi'}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\} \rightsquigarrow \text{split}(n_1, n_2)\{s'_1\}\{s_2\} \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi'} \text{ S-Split-Left}$$

$$\frac{p < n_1 \quad n_1, n_2 \text{ align to } n}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n)} \text{split}(n_1, n_2)\{\text{skip}\}\{s_2\} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-Split-Left-Done}$$

$$\frac{p \geq n_1 \quad p < n_1 + n_2 \quad n_1, n_2 \text{ align to } n \quad \eta, \sigma, \Sigma, t, b, p - n_1, \Psi, \Phi \vdash_m^{(h, n_2)} s_2 \rightsquigarrow s'_2 \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi'}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\} \rightsquigarrow \text{split}(n_1, n_2)\{s_1\}\{s'_2\} \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi'} \text{ S-Split-Right}$$

$$\frac{p \geq n_1 \quad p < n_1 + n_2 \quad n_1, n_2 \text{ align to } n}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n)} \text{split}(n_1, n_2)\{s_1\}\{\text{skip}\} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-Split-Right-Done}$$

$$\frac{p \geq n_1 + n_2 \quad n_1, n_2 \text{ align to } n}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-Split-None}$$

1324	
1325	$\eta, \sigma, \Sigma, t, b, t \text{ mod } T, \Psi, \Phi \vdash_m^{(\text{Thread}, T)} s \rightsquigarrow s' \dashv_m \eta', \sigma', \Sigma', \Psi', \Phi'$
1326	$\eta, \sigma, \Sigma, t, b, 0, \Psi, \Phi \vdash_m^{(\text{Block}, 1)} \text{destruct in } s \rightsquigarrow \text{destruct in } s' \dashv_m \eta', \sigma', \Sigma', \Psi', \Phi'$ S-Destruct-Block
1327	
1328	
1329	$\eta, \sigma, \Sigma, t, b, b \text{ mod } B, \Psi, \Phi \vdash_m^{(\text{Block}, B)} s \rightsquigarrow s' \dashv_m \eta', \sigma', \Sigma', \Psi', \Phi'$
1330	$\eta, \sigma, \Sigma, t, b, 0, \Psi, \Phi \vdash_m^{(\text{Grid}, 1)} \text{destruct in } s \rightsquigarrow \text{destruct in } s' \dashv_m \eta', \sigma', \Sigma', \Psi', \Phi'$ S-Destruct-Grid
1331	
1332	
1333	$\eta, \sigma, \Sigma, t, b, 0, \Psi, \Phi \vdash_m^\pi \text{destruct in skip} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi$ S-Destruct-Done
1334	
1335	$\eta, \sigma, \Sigma, t, b, p \text{ mod } n, \Psi, \Phi \vdash_m^{(h, n)} s \rightsquigarrow s' \dashv_m \eta', \sigma', \Sigma', \Psi', \Phi'$
1336	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, q \cdot n)} \text{group } q \ s \rightsquigarrow \text{group } q \ s'; \dashv_m \eta', \sigma', \Sigma', \Psi', \Phi'$ S-Group
1337	
1338	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{group } q \ \text{skip} \rightsquigarrow \text{skip}; \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi$ S-Group-Done
1339	

A.3.3 *Thread Synchronization*. We define a **size** operation on privileges to compute the size of a partition (the number of individual threads contained within it). The operation is defined as follows:

1340	$\text{size}(\text{Thread}, n) = n$
1341	$\text{size}(\text{Block}, n) = n \cdot T$
1342	$\text{size}(\text{Grid}, n) = n \cdot B \cdot T$
1343	
1344	
1345	
1346	
1347	
1348	
1349	$\Psi(\psi)(p) = 0$
1350	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{wait}_\psi \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi$ S-Sync-Wait-Done
1351	
1352	$\Psi(\psi)(p) \neq 0$
1353	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{wait}_\psi \rightsquigarrow \text{wait}_\psi \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi$ S-Sync-Wait-Spin
1354	
1355	$\Psi' = \Psi(\psi)[p \mapsto \Psi(\psi)(p) - 1]$
1356	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{dec}_\psi \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi$ S-Sync-Dec
1357	
1358	$\Psi(\psi)(p) = 0 \quad \Psi' = \Psi(\psi)[p \mapsto \text{size}(\pi)]$
1359	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{init}_\psi \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi$ S-Sync-Init-Zero
1360	
1361	$\Psi(\psi)(p) \neq 0$
1362	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{init}_\psi \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi$ S-Sync-Init-Nonzero
1363	
A.3.4	<i>Asynchrony</i> .
1364	$\text{rename}(\eta, \sigma, \Sigma, x, y, (\text{Thread}, 1)), t, b, p, \Psi, \Phi \vdash_m^{(\text{Thread}, 1)} s \rightsquigarrow s' \dashv_m \eta', \sigma', \Sigma', \Psi', \Phi'$
1365	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(\text{Thread}, 1)} \text{async_partition}_\phi x \text{ into } y \text{ in } s \rightsquigarrow$
1366	$\text{async_partition}_\phi x \text{ into } y \text{ in } s' \dashv_m \eta', \sigma', \Sigma', \Psi', \Phi'$
1367	
1368	
1369	$\Phi = \Phi'[\phi \mapsto \Phi'(\phi) \cup \{s\}]$
1370	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(\text{Thread}, 1)} \text{async_partition}_\phi x \text{ into } y \text{ in } \text{skip} \rightsquigarrow$
1371	$(\text{async_partition}_\phi x \text{ into } y \text{ in } s) \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi'$ S-Async-Partition-Unwind
1372	

1373	
1374	$\Phi(\phi) = \emptyset$
1375	$\frac{}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(\text{Thread}, 1)} \text{async_partition}_\phi x \text{ into } y \text{ in } \text{skip} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi}$ S-Async-Partition-Done
1376	
1377	
1378	
1379	$\frac{}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(\text{Thread}, 1)} \text{async_memcpy}(x, y) \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi[\phi \mapsto \Phi(\phi) \cup \{\text{memcpy}(x, y)\}]}$ S-Async-Memcpy
1380	
1381	
1382	$\frac{(\eta', \sigma', \Sigma') = \text{update}(\eta, \sigma, \Sigma, x, \text{get}(\eta, \sigma, \Sigma, y))}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{memcpy}(x, y) \rightsquigarrow \text{skip} \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi}$ S-Memcpy
1383	
1384	
1385	A.3.5 <i>Variables and Memory.</i>
1386	
1387	$\frac{\eta, \sigma, \Sigma \vdash_{\pi'}^\pi e \Downarrow v}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi x : \tau @ \pi' := e \text{ in } s \rightsquigarrow s \dashv_m \eta[x \mapsto^{\pi'} v], \sigma, \Sigma, \Psi, \Phi}$ S-Decl
1388	
1389	
1390	$\frac{}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{free } n \dashv_{m-n} \eta, \sigma, \Sigma, \Psi, \Phi}$ S-Free
1391	
1392	
1393	$\frac{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi x := \text{alloc Local } \tau n \text{ in } s \rightsquigarrow s; \text{ free } (n \cdot \text{size}(\tau)) \dashv_{m+n \cdot \text{size}(\tau)} \eta[x \mapsto^\pi \langle x, n \rangle], \sigma, \Sigma, \Psi, \Phi}{}$ S-Alloc-Local
1394	
1395	
1396	$\frac{\pi = (\text{Block}, 1)}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi x := \text{alloc Shared } \tau n \text{ in } s \rightsquigarrow s; \text{ free } (n \cdot \text{size}(\tau)) \dashv_{m+n \cdot \text{size}(\tau)} \eta[x \mapsto^\pi \langle x, n \rangle], \sigma, \Sigma, \Psi, \Phi}$ S-Alloc-Shared
1397	
1398	
1399	
1400	$\frac{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi x := \text{alloc Global } \tau n \text{ in } s \rightsquigarrow s; \text{ free } (n \cdot \text{size}(\tau)) \dashv_{m+n \cdot \text{size}(\tau)} \eta, \sigma, \Sigma[x \mapsto^\pi \langle x, n \rangle], \Psi, \Phi}{}$ S-Alloc-Global
1401	
1402	
1403	$\frac{x \in^{\pi'} \eta, \sigma, \Sigma \quad \eta, \sigma, \Sigma \vdash_{\pi'}^\pi e \Downarrow v \quad (\eta', \sigma', \Sigma') = \text{update}(\eta, \sigma, \Sigma, x, v) \quad \pi' \mid \pi}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi, \Phi \vdash_m^\pi x = e \rightsquigarrow \text{skip} \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi}$ S-Assn
1404	
1405	
1406	
1407	$\frac{\eta, \sigma, \Sigma, t, b, p \vdash_\pi^\pi e_1 \Downarrow \langle l, n \rangle \qquad i < n \quad \pi' \mid \pi}{\eta, \sigma, \Sigma, t, b, p \vdash_\pi^\pi e_2 \Downarrow i \qquad x \in^{\pi'} \eta, \sigma, \Sigma}$
1408	$\frac{\eta, \sigma, \Sigma \vdash_{\pi'}^\pi e_3 \Downarrow v \qquad (\eta', \sigma', \Sigma') = \text{update}(\eta, \sigma, \Sigma, l+i, v)}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi e_1[e_2] = e_3 \rightsquigarrow \text{skip} \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi}$ S-Arr-Assn
1409	
1410	
1411	
1412	$\frac{s' = s[(y + c \cdot p)/y] \quad (\eta', \sigma', \Sigma') = \text{rename}(\eta, \sigma, \Sigma, x, y, \pi/c)}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{partition}_\psi x y c s' \rightsquigarrow \text{init}_\psi; s'; \text{dec}_\psi; \text{wait}_\psi \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi}$ S-Partition
1413	
1414	
1415	$\frac{(\eta', \sigma', \Sigma') = \text{rename}(\eta, \sigma, \Sigma, x, y, (h, n_1))}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n_1 + n_2)} \text{claim}_\psi x \text{ into } y \text{ at } n_1 \text{ in } s \rightsquigarrow \text{init}_\psi; \text{split}(n_1, n_2) \{s'\} \{\text{skip}\}; \text{dec}_\psi; \text{wait}_\psi \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi}$ S-Claim
1416	
1417	
1418	
1419	$\frac{(\eta', \sigma', \Sigma') = \text{rename}(\eta, \sigma, \Sigma, x, y, \downarrow \pi)}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{lower}_\psi x \text{ into } y \text{ in } s \rightsquigarrow \text{init}_\psi; s; \text{dec}_\psi; \text{wait}_\psi \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi}$ S-Lower
1420	
1421	

A.3.6 Control Flow.

$$\begin{array}{c}
 \frac{\eta, \sigma, \Sigma \vdash_{\pi}^{\pi} e \Downarrow \text{true}}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{if } e \text{ then } s_1 \text{ else } s_2 \rightsquigarrow s_1 \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-If-True} \\
 \\
 \frac{\eta, \sigma, \Sigma \vdash_{\pi}^{\pi} e \Downarrow \text{false}}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{if } e \text{ then } s_1 \text{ else } s_2 \rightsquigarrow s_2 \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-If-False} \\
 \\
 \frac{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{while } e \text{ do } s \rightsquigarrow \text{if } e \text{ then } (s; \text{while } e \text{ do } s) \text{ else skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} s_1 \rightsquigarrow s'_1; s_2 \dashv_{m'}^{\pi'} \eta', \sigma', \Sigma', \Psi', \Phi'} \text{ S-While} \\
 \\
 \frac{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} s_1 \rightsquigarrow s'_1 \dashv_{m'}^{\pi'} \eta', \sigma', \Sigma', \Psi', \Phi'}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} s_1; s_2 \rightsquigarrow s'_1; s_2 \dashv_{m'}^{\pi'} \eta', \sigma', \Sigma', \Psi', \Phi'} \text{ S-Seq-First} \\
 \\
 \frac{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{skip}; s_2 \rightsquigarrow s_2 \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{skip}; s_2 \rightsquigarrow s_2 \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-Seq-Done} \\
 \\
 \frac{\Sigma(f) = \{[x_1 : \tau_1, \dots, x_n : \tau_n], s\} \quad \sigma, \Sigma \vdash_{\pi}^{\pi} e_i \Downarrow v_i \quad m' \leq m}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} f(e_1, \dots, e_n) \rightsquigarrow s \dashv_m \eta[x_i \mapsto v_i], \sigma, \Sigma, \Psi, \Phi} \text{ S-Function-Call}
 \end{array}$$

A.3.7 Expressions.

$$\frac{\pi < (\text{Grid}, 1)}{\eta, \sigma, \Sigma \vdash_{\pi}^{\pi}, \text{partition_id} \Downarrow \pi / \pi' - 1} \text{ E-Partition-Id}$$

$$\frac{}{\eta, \sigma, \Sigma \vdash_{\pi}^{\pi}, x \Downarrow \text{get}(\eta, \sigma, \Sigma, x)} \text{ E-Var}$$

$$\frac{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, e_1 \Downarrow \langle l, n \rangle \quad \eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, e_2 \Downarrow i \quad i < n \quad \pi' \leq \pi}{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, e_1[e_2] \Downarrow \text{get}(\eta, \sigma, \Sigma, l+i)} \text{ E-Arr-Access}$$

$$\frac{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, \text{int} \Downarrow n}{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, \text{int} \Downarrow n} \text{ E-Int} \quad \frac{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, \text{bool} \Downarrow b}{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, \text{bool} \Downarrow b} \text{ E-Bool}$$

$$\frac{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, e_1 \Downarrow v_1 \quad \eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, e_2 \Downarrow v_2 \quad v = v_1 \text{ bop } v_2}{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, e_1 \text{ bop } e_2 \Downarrow v} \text{ E-Bop}$$

$$\frac{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, e_1 \Downarrow v_1 \quad \eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, e_2 \Downarrow v_2 \quad v = v_1 \text{ cmp } v_2}{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, e_1 \text{ cmp } e_2 \Downarrow v} \text{ E-Cmp}$$

A.4 Theorems and Proofs

Note that in this section we assume no out of bounds array accesses. In general Snap (and by extension Bundl) makes no guarantees about array out of bounds.

A.4.1 More definitions. As a premise to our type safety theorems, we need to assume we have a well-typed environment, written $\Gamma \vdash \eta, \sigma, \Sigma$. We define what this means inductively

$$\begin{array}{c}
 \frac{}{\eta, \sigma, \Sigma \vdash n : \text{int}} \text{ V-Int} \quad \frac{}{\eta, \sigma, \Sigma \vdash b : \text{bool}} \text{ V-Bool} \quad \frac{}{\eta, \sigma, \Sigma \vdash f : \text{float}} \text{ V-Float} \\
 \\
 \frac{\forall i < n, \eta, \sigma, \Sigma \vdash \text{get}(\eta, \sigma, \Sigma, x + i) : \tau}{\eta, \sigma, \Sigma \vdash \langle x, n \rangle : \tau[]^l} \text{ V-Array} \quad \frac{\Gamma, x_i : \pi \tau_i \vdash_m^{\pi} s \quad \Gamma \vdash \cdot, \cdot, \text{fns} \Sigma}{\eta, \sigma, \Sigma \vdash \{x_i : \tau_i, s\} : \text{Fun}(x_i : \tau_i, \pi, m)} \text{ V-Function} \\
 \\
 \frac{\cdot \vdash \eta, \sigma, \Sigma}{\cdot \vdash \eta, \sigma, \Sigma} \text{ G-Empty} \quad \frac{\eta(x) =^{\pi} v \quad \eta, \sigma, \Sigma \vdash v : \text{int}}{\Gamma, x : \pi \text{ int} \vdash \eta, \sigma, \Sigma} \text{ G-Int}
 \end{array}$$

$$\begin{array}{c}
 \frac{\eta(x) =^\pi v \quad \eta, \sigma, \Sigma \vdash v : \text{bool}}{\Gamma, x :^\pi \text{bool} \vdash \eta, \sigma, \Sigma} \text{ G-Bool} \quad \frac{\eta(x) =^\pi v \quad \eta, \sigma, \Sigma \vdash v : \text{float}}{\Gamma, x :^\pi \text{float} \vdash \eta, \sigma, \Sigma} \text{ G-Float} \\
 \\
 \frac{\eta(x) =^\pi v \quad \eta, \sigma, \Sigma \vdash v : \tau[]}{\Gamma, x :^\pi \tau[]^{\text{Local}} \vdash \eta, \sigma, \Sigma} \text{ G-Local} \quad \frac{\sigma(x) =^\pi v \quad \eta, \sigma, \Sigma \vdash v : \tau[]}{\Gamma, x :^\pi \tau[]^{\text{Shared}} \vdash \eta, \sigma, \Sigma} \text{ G-Shared} \\
 \\
 \frac{\Sigma(x) =^\pi v \quad \eta, \sigma, \Sigma \vdash v : \tau[]}{\Gamma, x :^\pi \tau[]^{\text{Global}} \vdash \eta, \sigma, \Sigma} \text{ G-Global} \quad \frac{\Sigma(x) =^\pi v \quad \eta, \sigma, \Sigma \vdash v : \text{Fun}(\Gamma', \pi, m)}{\Gamma, x :^\pi \text{Fun}(\Gamma', \pi, m) \vdash \eta, \sigma, \Sigma} \text{ G-Function}
 \end{array}$$

We can prove a couple simple lemmas about well-typed environments under operations like **rename**, **update**, and **get**.

LEMMA A.1. (*Well-typed get*) If $\Gamma \vdash \eta, \sigma, \Sigma$ and $x :^\pi \tau \in \Gamma$ then $\eta, \sigma, \Sigma \vdash \text{get}(\eta, \sigma, \Sigma, x) : \tau$.

LEMMA A.2. (*Well-typed rename*) If $\Gamma, x :^\pi \tau \vdash \eta, \sigma, \Sigma$ then $\Gamma, x :^\pi \tau, y :^\pi \tau' \vdash \text{rename}(\eta, \sigma, \Sigma, x, y, \tau')$.

LEMMA A.3. (*Well-typed update*) If $\Gamma \vdash \eta, \sigma, \Sigma$ and $\eta, \sigma, \Sigma \vdash v : \tau$ then $\Gamma, x :^\pi \tau \vdash \text{update}(\eta, \sigma, \Sigma, x, v)$.

We also define a well-formedness precondition on p with respect to π :

$$(h, n) \vdash p ::= p < n$$

We also define well-formedness for the async stack:

$$\Gamma \vdash \Phi ::= \forall \phi, s \in \Phi(\phi), \Gamma \vdash_m^{(\text{Thread}, 1)} s$$

A.5 Proofs

LEMMA A.4. (*Expression Safety*) If $\Gamma \vdash^\pi e : \tau$ and $\Gamma \vdash \eta, \sigma, \Sigma$ and $\pi \vdash p$, then there is some v such that $\eta, \sigma, \Sigma \vdash_{\pi'}^\pi e \Downarrow v$ and $v : \tau$.

PROOF. This proof proceeds by induction on the typing relation for expressions. Despite the fact that this property implies termination, we do not need a logical relation to prove it because the expression language is very simple.

The **T-Int**, **T-Float**, and **T-Bool** cases are trivial, using the rules **E-Int**, **E-Bool** and **E-Float** to compute values. In the case for **T-Partition-Id**, the π premises of the typing rules match the premises of the evaluation rules, so these rules are simple as well.

The cases for **T-Bop** and **T-Cmp** follow directly from the inductive hypotheses, assuming a valid and correctly implemented set of binary operators and comparators.

The only interesting cases are **T-Arr-Access** and **T-Arr-Access-Shared**.

In both cases our inductive hypotheses and inversion give us that $\pi' \leq \pi$, and e_1 evaluates to a $\langle x, n \rangle$, and that all the values between x and $x + n$ in the appropriate environment are typed at τ . We also know that e_2 evaluates to an integer i . We assume that all array accesses are in bounds, so $i < l$, which is sufficient to use the **E-Arr-Access** rule to complete this case, and the proof. \square

LEMMA A.5. (*Expression Determinism*) If $\eta, \sigma, \Sigma, t, b, p \vdash_{\pi'}^\pi e \Downarrow v_1$ and $\eta, \sigma, \Sigma, t, b, p \vdash_{\pi'}^\pi e \Downarrow v_2$ then $v_1 = v_2$.

PROOF. Straightforward by induction on the semantic derivation. \square

LEMMA A.6. (*Expression Well-Typedness*) If $\Gamma \vdash_{\pi'}^\pi e : \tau$ and $\Gamma \vdash \eta, \sigma, \Sigma$ and $\pi \vdash p$, and $\eta, \sigma, \Sigma, t, b, p \vdash_{\pi'}^\pi e \Downarrow v$, then $v : \tau$.

PROOF. By our expression safety lemma our well-typed expression must evaluate to a well-typed value v' . By our determinism lemma v' must be the same as v , so v is well-typed. \square

LEMMA A.7. (*Statement Progress*) If $\Gamma \vdash \eta, \sigma, \Sigma$ and $\Gamma \vdash_m^\pi s$ and $\pi \vdash p$, then either s is skip or there is some s' such that $\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi s \rightsquigarrow s' \vdash_m^{\pi'} \eta', \sigma', \Sigma, \Psi', \Phi'$

PROOF. This proceeds by induction on the typing derivation.

1524 Privilege Management Rules.

- 1525 • Case **T-Split**:

1526 In this case we have by our assumption that $\pi \vdash p$ that $p < n$. We also have that n_1, n_2 align to n ,
 1527 so $n_1 + n_2 \leq n$. There are three cases to consider, then: when $p < n_1$, when $p \geq n_1$ and $p < n_1 + n_2$,
 1528 and when $p \geq n_1 + n_2$.

1529 In the first case, we have by our inductive hypothesis that s is either skip or that it can step
 1530 in an (h, n_1) context. In the former case we can use the **S-Split-Left-Done** rule and in the
 1531 latter we can use the **S-Split-Left rule**.

1532 The second case is almost symmetric. The only additional work we have to do is to argue
 1533 that $(h, n_2) \vdash p - n_1$, or equivalently that $p - n_1 < n_2$. This, however, is immediate from our
 1534 assumption that $p < n_1 + n_2$.

1535 In the last case, we just use the **S-Split-None** rule to step to skip.

- 1536 • Case **T-Destruct**

1537 By our inductive hypothesis, we know that s can step at $\downarrow \pi$ if $\downarrow \pi \vdash p$. The \downarrow operation is only
 1538 defined at (Block, 1) or (Grid, 1), so we only need to consider the cases where π is one of
 1539 those.

1540 In the former case p becomes $t \mathbf{mod} T$ while $\downarrow \pi$ is (Thread, T). $t \mathbf{mod} T < T$ for any t so this
 1541 satisfies the requirement that $\pi \vdash p$, which lets us use our inductive hypothesis: s is either
 1542 skip or can step. If it can step, we can use this to satisfy the premise of **S-Destruct-Block** to
 1543 step in this case. If it is skip, then we use the rule **S-Destruct-Done** to step instead.

1544 The latter case is the same, except using the fact that $b \mathbf{mod} B < B$ and the **S-Destruct-Grid**
 1545 rule.

- 1546 • Case **T-Group**

1547 In this case we have by assumption that $(h, q \cdot n) \vdash p$, i.e., that $p < q \cdot n$.

1548 In this case we have our IH that if $(h, n) \vdash p'$ for some p' , then s is either skip or steps with
 1549 (h, n) privilege with p' as our partition id.

1550 We choose p' to be $p \mathbf{mod} n$. This is always $< n$, so $(h, n) \vdash p'$. This lets us use our IH to get
 1551 that s is either skip (in which case we can use the **S-Group-Done** rule to step) or itself steps,
 1552 which lets us use the **S-Group** rule to step.

1553 Thread Synchronization Rules.

- 1555 • Case **T-Sync-Wait**

1556 $\Psi(\psi)(p)$ is either zero or it is not. In the former case we use the **S-Sync-Wait-Done** rule and
 1557 in the latter we use **S-Sync-Wait-Spin**.

- 1558 • Case **T-Sync-Dec**

1559 We use the **S-Sync-Dec** rule to step.

- 1560 • Case **T-Sync-Init**

1561 We use the **S-Sync-Init-Zero** or **S-Sync-Init-Nonzero** rules depending on whether $\Psi(\psi)(p)$
 1562 is zero or not.

1563 Asynchrony Rules.

- 1565 • Case **T-Async-Partition**

1566 In this case, we have via our IH that if $\Gamma, y : (\text{Thread}, 1) \text{ async } \tau[]^l \vdash \eta', \sigma', \Sigma'$, then we can either
 1567 step s with (Thread, 1) privilege under environments η' , σ' , and Σ' , or s is skip.

1569 We have by assumption that $\Gamma, x:(^{\text{Thread},1}) \text{async } \tau[]^l \vdash \eta, \sigma, \Sigma$. By our environment renaming
 1570 lemma, this gives us what we need to use our IH with $(\eta', \sigma', \Sigma')$ as **rename** $(\eta, \sigma, \Sigma, x, y, (\text{Thread}, 1))$.
 1571 Thus s either steps or is skip. In the former case we can step with **S-Async-Partition-Congr**,
 1572 and in the latter we can use either **S-Async-Partition-Unwind** or **S-Async-Partition-Done**
 1573 depending on whether $\Phi(\phi)$ is empty or not.

- Case **T-Async-Memcpy**

1574 Immediate via use of the **S-Async-Memcpy** rule.

- Case **T-Memcpy**

1575 Immediate via use of the **S-Memcpy** rule.

1576 *Memory Rules.*

- Case **T-Decl**

1580 Via our lemma about expression type safety and our hypothesis that e is well-typed, we obtain
 1581 the premises necessary to use the **S-Decl** rule to step.

- Case **T-Arr-Assn**

1584 Each of e_1 , e_2 and e_3 must evaluate to a well-typed value by the expression type safety lemma.
 1585 In particular, both e_1 evaluates to some $\langle l, n \rangle$ and e_2 evaluates to some i . We assume all array
 1586 accesses are in bounds, so this is sufficient to use the **S-Arr-Assn** rule to step.

- Case **T-Arr-Assn-Shared**

1588 Same as previous case.

- Case **T-Free**

1590 Trivial via the **S-Free** rule.

- Case **T-Partition**

1591 Trivial via the **S-Partition** rule.

- Case **T-Claim**

1594 Trivial via the **S-Claim** rule.

- Case **T-Lower**

1595 Trivial via the **S-Lower** rule.

- Case **T-Alloc**

1598 We assume that l is not Shared, so we can use the **S-Alloc-Local** or **S-Alloc-Global** rule,
 1599 depending on whether l is Local or Global.

- Case **T-Alloc-Shared**

1601 Trivial via the **S-Alloc-Shared** rule.

1602 *Control Rules.*

- Case **T-Skip**

1605 Trivial

- Case **T-While**

1607 Trivial, all while loops step via the **S-While** rule

- Case **T-If**

1609 By our proof of expression type safety, the expression e steps to either the boolean value true
 1610 or false. We can thus use either the **S-If-True** or **S-If-False** rules to step.

- Case **T-Seq**

1612 In this case we know by our IH that s_1 is either skip or can step. In the former case we use
 1613 the **S-Seq-Done** rule and in the latter we use the **S-Seq-First** rule.

- Case **T-Function-Call**

1615 In this case we know by our expression safety lemma that each of the arguments will evaluate to
 1616 a well-typed value. We also have by assumption that f has a function type, which by inversion

on the **V-Function** rule tells us that it is a closure type. Additionally our assumption that $\Gamma \vdash \eta, \sigma, \Sigma$ tell us that Σ contains f at the same type that Γ does. These premises are sufficient to use the **S-Function-Call** rule.

□

LEMMA A.8. (Statement Preservation) If $\Gamma \vdash \eta, \sigma, \Sigma$ and $\Gamma \vdash_m^\pi s$ and $\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_{m'}^\pi s \rightsquigarrow s' \dashv_{m''} \eta', \sigma', \Sigma, \Psi', \Phi'$ and $\pi \vdash p$ and $m \geq m'$ and $\Gamma \vdash \Phi$, then there is some Γ' such that $\Gamma \subseteq \Gamma'$ and $\Gamma' \vdash_{m''}^\pi s'$ and $\Gamma' \vdash \eta', \sigma', \Sigma'$ and $m \geq m''$ and $\Gamma' \vdash \Phi'$.

PROOF. We proceed by induction on the derivation of $\Gamma \vdash_m^\pi s$.

Privilege Management Rules.

- Case **T-Split**

In this case we have by assumption that $(h, n) \vdash p$, $\Gamma \vdash \eta, \sigma, \Sigma$, and $\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_{m'}^\pi \text{split}(n_1, n_2)\{s_1\}\{s_2\} \rightsquigarrow s' \dashv_{m''} \eta', \sigma', \Sigma, \Psi', \Phi'$. Our inductive hypotheses give us that if for any p , if $(h, n_1) \vdash p$ and s_1 steps with partition id p , or if $(h, n_2) \vdash p$ and s_2 steps with partition id p then their results are well typed.

By inversion on our semantic derivation, we are in one of 5 cases.

In the **S-Split-Left** case $p < n_1$ and s_1 steps to s'_1 . This is sufficient to tell us that s'_1 is well-typed and the output environments of that relation η', σ' , and Σ' are all well-typed by $\Gamma' \supseteq \Gamma$, and that the memory is properly bounded by the typing rules.

We can thus use the **T-Split-Left** rule to conclude that the result of this case is well-typed. The **S-Split-Left-Done** rule is trivial via the **T-Skip** rule.

The **Right** cases are symmetric, with the observation that when $p \geq n_1$ and $p < n_1 + n_2$ then $p - n_1 < n_2$.

The last **S-Split-None** rule is trivial via the **T-Skip** rule.

- Case **T-Destruct**

In this case we have by assumption that $\downarrow \pi$ is defined, so π is either **(Block,1)** or **(Grid,1)**. We also assume that $\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_{m'}^\pi \text{destruct}$ in $s \rightsquigarrow s' \dashv_{m''} \eta', \sigma', \Sigma, \Psi', \Phi'$. We also have by our inductive hypothesis that for any p such that $\downarrow \pi \vdash p$ and s'' such that s steps to s'' at p , then that step preserved well-typedness.

By inversion on the step relation, we are in one of three cases.

If the rule used **S-Destruct-Block**, then we know that s steps to s'' and p is $t \bmod T$. $(\text{Thread}, T) \vdash t \bmod T$ for any t , so we can use our inductive hypothesis to conclude that the s'' stepped to by s is well-typed, as are its environments and memory usage. The **T-Destruct** rule then gives us our desired goal.

The **S-Destruct-Grid** case proceeds similarly, while the **S-Destruct-Done** case is trivial.

- Case **T-Group**

In this case we have by assumption that

$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_{m'}^{(h, q \cdot n)} \text{group } q \ s \rightsquigarrow s' \dashv_{m''} \eta', \sigma', \Sigma, \Psi', \Phi'$. We have by our inductive hypothesis that s steps to s'' at some partition id p' and $(h, n) \vdash p'$, then s'' is well typed, as are the other outputs of that step.

By inversion on the step relation, we are in one of two cases. The **S-Group-Done** case is trivial, so we shall focus on the **S-Group** case. In this case we have that s steps to s'' at partition id $p \bmod n$. It is always the case that $(h, n) \vdash p \bmod n$ for any p , so we can use our inductive hypothesis to conclude that s'' is well-typed, as are its output environments and memory usage. From there, it is a simple application of the **T-Group** rule to conclude that $\text{group } q \ s'$ is well-typed, and to finish the case.

1667 **Thread Synchronization Rules.** These rules are all trivial: with one exception all thread synchronization
 1668 primitives step to skip without changing environment or memory, and are thus obviously well-typed.
 1669 **S-Sync-Wait-Spin** does not produce skip, but it steps to the same statement as we already assumed
 1670 typechecks in the premise of the lemma, so is straightforward nonetheless.

1671 If we wanted to say something about deadlock freedom we'd have more work here, but we aren't
 1672 doing that, so these rules are easy.

1673 **Asynchrony Rules.**

- 1674 • Case **T-Async-Partition** In this case we have that s is well-typed in a context where x has
 1675 been renamed into y , with $(\text{Thread},1)$ privilege. We also have that $\Gamma, x:(^{\text{Thread},1}) \tau[]^l \vdash \eta, \sigma, \Sigma$
 1676 and $\Gamma, x:(^{\text{Thread},1}) \tau[]^l \Phi$.

1677 By inversion, we are in one of three cases.

1678 In the **S-Async-Partition-Done** case, we are done.

1679 In the **S-Async-Partition-Congr** case, our inductive hypothesis gives us that there is some
 1680 Γ' such that $\Gamma, y:(^{\text{Thread},1}) \text{async } \tau[]^l \subseteq \Gamma'$ and $\Gamma' \vdash \Phi'$ and $\Gamma' \vdash \text{rename}(\eta, \sigma, \Sigma, x, y, (\text{Thread},1))$
 1681 via our well-typed renaming lemma. This lets us use the **T-Async-Partition** rule to check
 1682 this case, with a choice of Γ' as $\Gamma, x:(^{\text{Thread},1}) \tau[]^l$.

1683 In the **S-Async-Partition-Unwind** case, our assumption that Φ is well-typed tells us that
 1684 $\Gamma, y:(^{\text{Thread},1}) \vdash \{ (\text{Thread},1)_m s$. Thus, we can use the **T-Async-Partition** rule to type this case.

- 1685 • Case **T-Async-Memcpy**

1686 In this case the statement and environment typing are trivial, we need only to show that the
 1687 async stack remains well typed.

1688 In this case we have that x and y have the same type at $(\text{Thread},1)$. This is sufficient for us
 1689 to check $x = y$ at $(\text{Thread},1)$, meaning that adding that instruction to the stack maintains its
 1690 well-typedness.

- 1691 • Case **T-Memcpy**

1692 Immediate via use of the **S-Memcpy** rule. We just need to show that the environment remains
 1693 well typed, which we know via our lemmas about **update** and **get**.

1694 **Memory Rules.**

- 1695 • Case **T-Decl**

1696 By inversion, we are using **S-Decl** rule for evaluation. Our well-typed expression lemma gives
 1697 us that v is well-typed, so it follows from our assumptions and our lemmas about extending
 1698 environments that the extended η and s are well-typed by $\Gamma, x:\pi \tau$.

- 1699 • Case **T-Free**

1700 Trivial.

- 1701 • Case **T-Alloc**

1702 By inversion we are either in the **S-Alloc-Local** or **S-Alloc-Global** rules. In either case, we
 1703 assume that $\Gamma, x:\pi \tau[]^l$ checks s , meaning we can use our extended environment lemmas and
 1704 the **T-Seq** and **T-Free** rules to check these cases.

- 1705 • Case **T-Alloc-Shared**

1706 Same as previous case.

- 1707 • Case **T-Partition**

1708 In this case we have by assumption that $\Gamma, y:(^{(h,n/c)} \tau[]^l s$ and l is not local and c divides n . By
 1709 inversion on our step relation we must be in the **S-Partition** case, so we have $\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi$
 1710 $\text{partition}/xycs \sim \text{init}_\psi; s'; \text{dec}_\psi; \text{wait}_\psi \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi$ where $s' = s[(y + c \cdot p)/y]$ and
 1711 $(\eta', \sigma', \Sigma') = \text{rename}(\eta, \sigma, \Sigma, x, y, \pi/c)$. Via our well-typed renaming lemma contexts we know
 1712 that we can check the renamed environments in the extended environment $\Gamma, y:\pi/c \tau[]^l, x:\pi$

$\tau[]^l$, and this context is also sufficient to check s' (via a substitution-preserves-typing lemma that is obvious). Using the **T-Skip** rule this is exactly what we need to show to complete this case, as the thread sync primitives check trivially via their typing rules.

- Case **T-Claim**

Essentially the same as **T-Partition**.

- Case **T-Lower**

Essentially the same as **T-Partition**.

Control Rules.

- Case **T-If**

We have by assumption that $\text{if } e \text{ then } s_1 \text{ else } s_2$ steps to some s' , and by inversion we know that either e evaluates to true and s' is s_1 , or e evaluates to false and s' is s_2 .

In either case, our inductive hypotheses is sufficient to tell us that these are well-typed. In particular, in both cases our IHs tell us that the amount of memory used by stepping each branch of the **if** is less than the amount of memory computed by the type system for each branch. Because the whole **if** expression checks using the greater of the memory usage of m_1 or m_2 (i.e., the memory usage on each branch), the resulting usage for the whole conditional is also bounded by the type system.

- Case **T-Skip**

Trivial: **skip** does not step

- Case **T-Seq**

In this case we have that s_1 and s_2 are both well-typed (with m_1 memory and m_2 memory respectively), and $m = \mathbf{max}(m_1, m_2)$. We also have by inversion that s_1 either steps to **skip** or s'_1 , and our inductive hypothesis tells us that s'_1 is well-typed.

In the former case we can use the **S-Seq-Done** rule to trivially finish the case. In the latter, our IH allows us to finish the case, since m_1 is always $\leq \mathbf{max}(m_1, m_2)$

- Case **T-While**

By inversion, we have that

$$\begin{aligned} & \eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{while } e \text{ do } s \\ & \rightsquigarrow \text{if } e \text{ then } (s; \text{while } e \text{ do } s) \text{ else skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi. \end{aligned}$$

We also have by assumption that e and s are well-typed. With this information, through use of the **T-If**, **T-Seq**, **T-While**, and **T-Skip** rules, we can conclude that the result of this rule is also well-typed.

- Case **T-Function-Call**

By inversion we have that

$$\begin{aligned} & \eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi f(e_1, \dots, e_n) \\ & \rightsquigarrow \text{call } s \text{ with } (x_i : \pi_i \tau_i \mapsto v_i) @ \{\eta, \sigma, \Sigma, m'\} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi, \end{aligned}$$

and also that $\Sigma(f) = \{[x_1 : \tau_1, \dots, x_n : \tau_n], s\}$, and $\eta, \sigma, \Sigma, t, b, p \vdash^\pi e_i \Downarrow v_i$, and $m' \leq m$.

We also have from the premises of our case that $f : \pi \mathbf{Fun}(x_1 : \pi_1 \tau_1, \dots, x_n : \pi_n \tau_n, \pi, m') \in \Gamma$, and $\Gamma \vdash^\pi e_i : \tau_i$, and $m' \leq m$.

Our lemma for expression well-typedness tells us that that each v_i is a well-typed value, and our assumption that $\Gamma \vdash \eta, \sigma, \Sigma$ tells us that $\Sigma(f)$ is a well-typed function and thus that $\mathbf{fns} \Gamma, x_i : \pi_i \tau_i \vdash_m^\pi s$. We can take the union of this with Γ to produce $\Gamma, x_i : \pi_i \tau_i \vdash_m^\pi s$, which clearly checks s and the output environments.

□