

1 Modular GPU Programming with Typed Privileges

2 ANONYMOUS AUTHOR(S)

3 Abstract

4 1 Introduction

5 CUDA [10] is a low-level, imperative programming language for NVIDIA GPUs. These GPUs are organized into a hierarchy of compute resources: threads, blocks, and a grid. *Threads* are the basic units of sequential execution; *blocks* are groups of threads that can cooperate through shared scratchpad memory; and the *grid* is the full collection of blocks launched for a GPU kernel. CUDA programs are written from the perspective of a single thread, and replicated in parallel across the entire grid.

6 Although programs are written per-thread, some operations supported by the hardware are only semantically valid when executed collectively by a group of threads. For example, the `__syncthreads()` primitive, which synchronizes all threads within a block, causes the machine to deadlock if executed by a subset of threads within that block. Similarly, tensor core instructions—hardware-accelerated matrix-multiply operations—are only well-defined when invoked by a 32-thread group (known as a *warp*) whose starting index is 32-aligned. In essence, these operations require the programmer to carefully marshal compute resources; coordinating which threads execute which blocks of code. As a result, collective operations break the illusion of CUDA programs executing as independent threads.

7 The conceptual clash between the per-thread programming paradigm and these collective operations hinders compositional reasoning. In CUDA, function calls — like all other statements — are independently executed on all threads. However, functions may contain collective operations that are only valid when executed together by a group of threads. This creates a semantic gap: the call site appears to be a per-thread operation, but the function body may require collective execution by a specific configuration of the device’s compute resources, such as an entire block or warp.

8 This gap between the syntax and semantics of a CUDA program additionally creates tension between abstraction and correctness. This tension is apparent even in widely used CUDA libraries that package common functionality via function interfaces . Consider the snippet of documentation presented in Figure 1 taken directly from CUB [9], a library from NVIDIA that provides wrappers around parallel primitives like scan, sort, and histogram, among others. The documentation describes a function called `BlockReduce` , in which all threads in a block collaboratively apply a reduction operation, such as a maximum or prefix sum, over an array.

9 The CUB documentation attempts to specify several preconditions that must be satisfied by callers of `BlockReduce`. First, the function is only well-defined when invoked by an entire block’s worth of threads. Second, the function accesses shared memory, a scratchpad space visible to all threads in a block. Users wishing to reuse this storage must first ensure all the block’s threads have finished using it by inserting a `__syncthreads()` barrier before the reuse. Whether this synchronization is required cannot be determined from the function alone; it depends on how the storage is used afterward.

10 In effect, CUB attempts to informally retrofit CUDA with information about the compute and memory requirements of collective operations.

11 By prefixing functions with identifiers such as `Block`, CUB effectively creates “namespaces” for different kinds of these operations. Even then, this organization is purely conventional: it encodes hierarchy through naming and cannot be enforced by the compiler. Ultimately, correctness depends not only on the user carefully reading and interpreting the documentation, but also on the library implementer upholding the invariants of each namespace.

12 Previous attempts to resolve the mismatch between the syntax and semantics of CUDA programs have generally done so by raising the level of abstraction. Triton [14], a tile-based, multi-dimensional

```

50
51 //include <cub/cub.cuh> // or equivalently <cub/block/block_reduce.cuh>
52
53 __global__ void ExampleKernel(...)

54     // Specialize BlockReduce for a 1D block of 128 threads of type int
55     using BlockReduce = cub::BlockReduce<int, 128>;
56     // Allocate shared memory for BlockReduce
57     __shared__ typename BlockReduce::TempStorage temp_storage;
58     // Obtain a segment of consecutive items that are blocked across threads
59     int thread_data[4];
60     ...
61     // Compute the block-wide max for thread0
62     int aggregate = BlockReduce(temp_storage).Reduce(thread_data, cuda::maximum<>());
63 }

64 
```

Computes a block-wide reduction for thread₀ using the specified binary reduction functor. The first num_valid threads each contribute one input element.

- The return value is undefined in threads other than thread₀.
- For multi-dimensional blocks, threads are linearly ranked in row-major order.
- A subsequent __syncthreads() threadblock barrier should be invoked after calling this method if the collective's temporary storage (e.g., temp_storage) is to be reused or repurposed.

Fig. 1. CUB documentation for the BlockReduce function

array language, circumvents the challenges of compositional reasoning by restricting users to a single level of the hierarchy, namely the block level. While this makes for a simplified programming model, the lack of fine-grained control prevents GPU programmers from writing the highest-performance kernels.

Meanwhile, Pallas is a recent entrant similar to Triton that restricts users to programming at the level of “warpgroups” (four warps).

A variety of functional languages, such as Futhark [7], provide compile-time guarantees through their type systems but do not expose low-level control over the hardware. Some other efforts, like Descend [8], aim to provide a low-level, safe GPU systems programming language, but lack support for collective operations entirely. As a result, CUDA remains the de-facto standard for writing high-performance kernels on modern GPUs.

Rather than sacrificing low-level control for safety, we aim to make fine-grained GPU programming both safe and compositional by design. Our core insight is that GPU programmers must carefully map programs to available compute resources, and unlike CUDA, which obscures this mapping, we can expose it explicitly in syntax. By making this mapping explicit, we can provide low-level access to collective operations while ensuring compile-time safety: the compiler can check that requirements about what hardware resources are required to execute code are satisfied by the context in which that code is executed.

We introduce Snap, a new low-level GPU programming language that guarantees correct usage of hardware resources by construction. Inspired by coeffect systems [11] and dependency calculi [1], Snap statically tracks the configurations of compute and memory resources at the type level.

We reify the notion of compute resources via *privileges*, so named because they dictate which operations are permitted in which contexts. Snap uses these privileges to verify that every collective operation is executed with the necessary resources. To enable modular reasoning, function interfaces include *privilege signatures* that declare the privilege each function requires for correct execution.

To validate the design of Snap, we build Bundl, a core calculus model of Snap that tracks privilege in both its type system and operational semantics. We provide formal rules for manipulating the

privilege of both code and data, and prove *type-and-privilege safety*, ensuring that Bundl is sound, and that all threads are statically privileged to run the operations they attempt at run time.

A parallel goal, in addition to correctness, is performance. To that end, we incorporate modern GPU features such as tensor cores and asynchronous data movement into Snap, and demonstrate that Snap can achieve the same performance as hand-written, highly optimized code on an H100 and a 4070 Ti Super.

Concretely, our contributions are:

- (1) We provide background on GPU compute architecture, CUDA’s per-thread programming model, and the convergence requirements of collective operations. We motivate the need for compile-time checking of resource usage by example (Section 2).
- (2) We present Snap, a low-level GPU programming language with type-level privileges that track compute and memory resources (Section 3).
- (3) We formalize the core concepts of Snap in Bundl, a calculus that tracks privileges in both its type system and operational semantics. We prove type-and-privilege safety, ensuring that threads only execute operations for which they are statically privileged (Section 4).
- (4) We compile Snap programs into CUDA (Section 5), implement various GPU kernels in Snap, and demonstrate that they match the performance of hand-optimized CUDA implementations (Section 6).

Section 7 addresses the limitations of our approach, Section 8 discusses related work, and Section 9 concludes and outlines future work.

2 Background & Motivation

Before diving into the design of Snap, we begin with an overview of the GPU’s compute and memory hierarchies and outline the challenges posed by reasoning about them collectively.

2.1 Compute Hierarchy

In CUDA, programmers launch computations that run on thousands of threads. On the hardware, work is scheduled at the granularity of blocks on a *streaming multiprocessor*, or an *SM*. Within an SM, execution follows the Single Instruction, Multiple Threads (SIMT) model. Instructions are issued to groups of 32 threads, or a *warp*, at a time, and each thread in the warp executes the instruction in lockstep.

Users define the behavior of different levels of the compute hierarchy by defining the behavior of a single thread. CUDA programs can read built-in identifiers like `threadIdx.x`—to determine a thread’s position within a block—and `blockIdx.x`—to determine the block’s position within the grid—at runtime. Combining these two IDs, threads can locate themselves in their share of the full computation.

A natural and tempting way to interpret these built-in identifiers is to think of them as indices of implicit “parallel for-loops” where each iteration is executing simultaneously. While this view is sufficient to understand CUDA’s programming model when threads do independent work, it quickly breaks down in the presence of *collective operations*.

Unlike most instructions, which are executed by a single thread, collective operations must be executed collaboratively by a group of threads. For example, a warp-level tensor-core operation, shown in Figure 2, is only *meaningful* when invoked by a collection of 32

```

1 int tid = threadIdx.x;
2 if (tid >= 0 && tid < 31){
3     float A[4];
4     float B[2];
5     float C[4] = { 0 };
6     # Populate A with unique values
7     for (int i = 0; i < 4; i++)
8         A[i] = tid * 4 + i;
9     # Populate B with unique values
10    for (int i = 0; i < 2; i++)
11        B[i] = tid * 4 + 1;
12    # Issue a warp-level tensor-core
13    # operation. D = A * B + C
14    asm("mma.sync.aligned.m16n8k8... "
15        "%0, %1, %2, %3,    /*A*/ "
16        "%4, %5, %6, %7,    /*B*/ "
17        "%8, %9,           /*B*/ "
18        "%10, %11, %12, %13;" /*C*/
19        : "r"(C[0]), "r"(C[1]), ...
20        : "r"(A[0]), "r"(A[1]), ...
21        : "r"(B[0]), "r"(B[1]),
22        : "r"(C[0]), "r"(C[1]), ...);
```

Fig. 2. Invoking a warp-level tensor-core instruction in CUDA.

148 threads—a “warp”—acting together. If the operation was invoked by fewer than 32 threads, the
 149 instruction would be invalid . To make matters worse, the restriction is not only on the *number* of
 150 threads executing the operation, but also on their *alignment*. In this case, the starting ID of the group
 151 of 32 participating threads must be aligned to a multiple of 32. For example, if the condition on line 2
 152 of Figure 2 were instead `tid > 1 && tid < 33`, the instruction would still be invalid, even though 32
 153 threads would execute it. Once the instruction has been correctly invoked, the tensor-core executes
 154 once collectively for all threads, and its result is distributed across the registers of the participating
 155 threads.

156 Collective operations make reasoning about CUDA programs challenging because they force
 157 programmers to track the convergence behavior of threads. Specifically, programmers must reason
 158 about how many threads reach a particular point in the program, and how those threads are arranged,
 159 accounting for alignment. This difficulty is further amplified due to two reasons. First, programs may
 160 have *multiple points of convergence*, requiring programmers to mentally track the relative ID of a
 161 thread within a logical group as it changes over the course of a program’s execution. Second, threads
 162 may be participating in *different kinds of convergence* within the same program. In our example,
 163 we only considered the warp-level tensor-core operation, but there are other collective operations
 164 that require convergence at different granularities like the *warp-group*-level tensor-core operations ,
 165 which must be issued collectively by four warps, or the `__syncthreads()` synchronization primitive,
 166 which must be executed by all threads within a block.

167 Reasoning about collective operations is already error-prone within the context of a single function,
 168 but becomes even more difficult when reasoning *across* functions . A callee may assume a certain
 169 number of threads and have internally structured its computation based on that assumption, including
 170 reasoning about convergence behavior. However, that assumption is not visible in the callee’s function
 171 interface, which only has information about the input and output types. To invoke that function
 172 correctly, users must read documentation, or worse , read the callee’s implementation to understand
 173 its assumptions, breaking modular reasoning.

174 In this work, we make the programmer’s implicit assumptions
 175 about a program’s convergence behavior explicit in the program’s
 176 source. Consider the example in Figure 3 which shows an equivalent
 177 rewrite of the CUDA program in Figure 2 using Snap. In Figure 3,
 178 the call to the `mma` intrinsic is legal only because it is executed within
 179 a group of 32 threads, as made explicit on line 2. Since Snap’s `group`
 180 construct already enforces both the size and alignment of participating
 181 threads, the legality of the `mma` operation is guaranteed by
 182 construction.

183 Now, if we had attempted to introduce a condition based on `tid`
 184 *within* the group’s scope—for example, `if (tid < 16)` — which
 185 would cause the number of active threads in the group to change
 186 during its execution , Snap would reject the program at compile time
 187 . This is because Snap explicitly tracks the frequency at which values diverge, using the `@` syntax,
 188 and code executing at a lower divergence frequency than a given piece of data does not have the
 189 *privilege* to read or branch on that data. We explain the syntax and rules of legal Snap programs in
 190 greater detail in Section 3.

191 2.2 Memory Hierarchy

192 GPUs also have a *memory hierarchy*, which, like the compute hierarchy, is organized into three levels.
 193 All threads in a grid can read and write from *global memory*, where operands typically reside at the
 194 start of a computation and where results are eventually written. Each block has access to a limited
 195

```

1  tid : int @ thread[1] = id();
2  with group[thread[32]]:
3      A : float[4] @ thread[1]
4      B : float[2] @ thread[1]
5      C : float[4] @ thread[1]
6      # Populate A with unique values
7      for i in range(0, 4, 1):
8          A[i] = tid * 4 + 1
9          C[i] = 0
10     # Populate B with unique values
11     for i in range(0, 2, 1):
12         B[i] = tid * 4 + 1
13     # Issue a warp-level tensor-core
14     intrinsic.mma(
15         A[0], A[1], A[2], A[3],
16         B[0], B[1],
17         C[0], C[1], C[2], C[3],
18         out=[A[0], A[1], A[2], A[3]]
19     )
  
```

Fig. 3. Invoking a warp-level tensor-core instruction in Snap.

amount of *shared memory*, a programmer-managed scratchpad typically used to stage data that will be repeatedly used. Finally, each thread maintains its private state in *registers* and *local memory*, which is used for storing the thread's stack.

This memory hierarchy has its own notion of convergence and divergence, again mirroring that of compute. For example, when launching a kernel on the GPU, a pointer may initially belong to the grid because every thread sees the same pointer value at the start of the computation. To write to that memory, each thread computes an offset from the original base address. In doing so, the pointer *diverges* across threads within the larger memory region so each thread can write independently. After these writes, the memory must *reconverge*, ensuring that all threads have completed their updates before the memory can safely be returned to its original logical owner in the hierarchy.

CUDA does not explicitly model these memory spaces, nor does it track whether allocations in a given memory space, such as shared memory, exceed device limits. Snap, on the other hand, reifies these spaces in its type system by explicitly decorating each variable with its associated memory space and restricting shared memory to static allocations, throwing an error if an allocation exceeds device limits. We will discuss these aspects in detail in Section 3.5. For now, however, we focus on how memory diverges and reconverges in tandem with the compute hierarchy.

Let us reconsider the tensor-core example in Figure 2, this time initializing the operands of the tensor-core operation from pointers into memory. In Figure 4, the registers A, B, and C are populated by first loading from a_mem and b_mem, which point into global memory. After the tensor-core operation completes, the result is written back to c_mem.

In this program, c_mem is accessed from two different levels of the compute heirarchy.

At the start, c_mem is accessible collectively by the group of 32 threads, which determines the base offset into the global array with which the group interacts . This is the ownership assumed when entering this code . Then, c_mem is further partitioned across individual threads so that each thread can write back the portion of the result it computed . To restore c_mem back to its 32 thread-level ownership, all 32 threads must first synchronize to ensure all per-thread writes have completed. In Section 1 , we saw a similar requirement being documented in the CUB library.

Similarly to the compute hierarchy (Section 2.1), CUDA programmers are responsible for tracking the logical ownership of a piece of memory as it evolves over the course of a program, and for ensuring that appropriate synchronization occurs whenever that ownership changes.

Snap, by contrast, makes the evolution of a memory object's logical ownership explicit in the type system and *automatically* inserts the synchronization required

```

1 int x = threadIdx.x;
2 if (x >= 0 && x < 31) {
3     float A[4];
4     float B[2];
5     float C[4] = { 0 };
6
7     // k is A's stride
8     A[0] = a_mem[(x/4)*k+(x%4)];
9     A[1] = a_mem[(x/4)*k+(x%4)+4];
10    A[2] = a_mem[(x/4+8)*k+(x%4)];
11    A[3] = a_mem[(x/4+8)*k+(x%4)+4];
12
13    // n is B's stride
14    B[0] = b_mem[(x%4)*n+(x/4)];
15    B[1] = b_mem[(x%4+4)*n+(x/4)];
16
17    asm("mma.sync...");
18
19    // Write back into c_mem
20    // n is C's stride
21    c_mem[(x/4)*n+2*(x%4)] = C[0];
22    c_mem[(x/4)*n+2*(x%4)+1] = C[1];
23    c_mem[(x/4+8)*n+2*(x%4)] = C[2];
24    c_mem[(x/4+8)*n+2*(x%4)+1] = C[3];

```

Fig. 4. Invoking a warp-level tensor-core instruction in CUDA after loading data from memory.

```

1 #include<"device">
2 #requires thread[32]
3 def simple_mma(
4     a: ptr(const float)) @ thread[32],
5     b: ptr(const float)) @ thread[32],
6     c: ptr(float) @ thread[32]:
7     x: int @ thread[] = id()
8     with group(thread[32]):
9         A: float[4] @ thread[1]
10        B: float[2] @ thread[1]
11        C: float[4] @ thread[1]
12
13        # Reads do not need to be lowered.
14        A[0] = a_mem[(x/4)*k+(x%4)];
15        A[1] = a_mem[(x/4)*k+(x%4)+4];
16        A[2] = a_mem[(x/4+8)*k+(x%4)];
17        A[3] = a_mem[(x/4+8)*k+(x%4)+4];
18
19        B[0] = b_mem[(x%4)*n+(x/4)];
20        B[1] = b_mem[(x%4+4)*n+(x/4)];
21        # Skipping initialize C to 0...
22        intrinsic.mma(
23            A[0], A[1], A[2], A[3],
24            B[0], B[1],
25            C[0], C[1], C[2], C[3],
26            out=[C[0], C[1], C[2], C[3]])
27
28        # Must be at thread[1] to write.
29        idx: int @ thread[1] =
30            lambda ro, co: (x/4+ro)*n+2*(x%4)+co
31        with partition(c_mem, p=thread[32],
32            f=idx) as c_thrd:
33            c_thrd[0, 0] = C[0]
34            c_thrd[0, 1] = C[1]
35            c_thrd[8, 0] = C[2]
36            c_thrd[8, 1]) = C[3]
37
38        # <-- Sync point inferred by name

```

Fig. 5. Invoking a warp-level tensor-core instruction in Snap after loading data from memory.

246 to restore memory to its original ownership . In Figure 5, we show how memory is lowered through
 247 the compute hierarchy for the same tensor-core operation in introduced in Figure 4. The lowering
 248 of `c_mem` is required in this program as Snap only permits writes when data is owned by a single
 249 thread. To lower `c_mem`, we use Snap’s partition operation on line 31. This takes the name of the
 250 memory to partition, `c_mem`, and lowers it to a compute hierarchy below the original one, a single
 251 thread, and assigns it a new name, `c_thrd`; Snap will not allow references to the old name `c_mem`
 252 within the partition’s associated scope. The `partition` function also takes an indexing function,
 253 and each time `c_thrd` is accessed, this indexing function is implicitly applied. Once the `partition`’s
 254 scope ends, Snap will insert a synchronization before the first use of the original name.

255 Finally, let’s examine the function interface on line 3 . The function marks `c_mem` as a non-const
 256 pointer, signaling to Snap that it may be used for writing. In the calling context of this function, we
 257 also synchronize `c_mem`’s logical owner, if one exists, before restoring `c_mem`’s parent in its scope.

258 We would like to emphasize that preventing data races is not an one of Snap’s goals. In Snap, data
 259 races *can* occur; however, since the data is eventually synchronized before it is reused, the last-writer
 260 wins. Out-of-bounds accesses are considered undefined behavior. We believe that prior work, such
 261 as Descend [8], lays the foundation for reasoning about data-race free programs, and that a Descend-
 262 like memory model can be adapted to Snap. Snap is instead focused on the *interaction* between the
 263 compute and memory hierarchies, and on reasoning about them simultaneously to ensure that an
 264 operation is executed only when sufficient compute resources are available, a guarantee that Descend
 265 cannot provide.

266

267

268

269

270

271

272

273

274

275

276

277

278

279

280

281

282

283

284

285

286

287

288

289

290

291

292

293

294

3 The Snap Language

Snap is an imperative, low-level programming language designed at a level of abstraction comparable to CUDA. Snap's surface syntax explicitly represents how compute (grid, block, threads) and data (global, shared and local memory) resources are used and distributed over the course of its execution. Using this explicit representation, Snap can enforce safety properties that CUDA cannot. Namely, Snap rejects programs that attempt to execute operations with insufficient compute and memory resources.

In this section, we introduce Snap's core constructs through a running example, a tensor core matrix multiplication, shown in Figure 6. This is a standard computation a GPU programmer might write while being sufficiently complex to reveal the interesting challenges in GPU programming.

The program in Figure 6 computes a matrix multiplication between two float arrays, A and B to produce an output matrix C. In this kernel, each block computes an independent tile of the output C. To perform the computation, each block first locates its assigned tile index, then loads the corresponding rows and columns from A and B (line 27). Since matrix multiplication offers ample data reuse, to amortize the cost of loading data, we stage A and B in shared memory (lines 34). At this point, each thread within the block diverges and loads different elements of A and B. Finally, the program invokes an Ampere warp-level tensor-core instruction to compute the output, requiring threads in a warp to converge (line 61). We encapsulate this tensor-core instruction in a function, demonstrating how function composition works in Snap. This function's implementation is identical to the one shown in Figure 5.

3.1 Levels

Snap models the machine's compute hierarchy through *levels*. There are three levels in Snap—grid, block and thread—which are organized as expected: a grid consists of multiple units of the block level, which in turn consist of multiple units of the thread level. Levels have a strict

```

1  @cuda("global")
2  # Top-level privilege class and required shared memory usage.
3  @requires(grid[1], block[], thread[32], smem=1280)
4  def mmaTF32NaiveKernel(A: ptr(const(float)) @ grid[1],
5                           B: ptr(const(float)) @ grid[1],
6                           C: ptr(float) @ grid[1],
7                           M : int @ grid[1],
8                           N : int @ grid[1],
9                           K : int @ grid[1]):
10     # Replicate this code across 1 grid.
11     with group(grid[1]):
12         # @ grid[1] inferred from compute privilege
13         MMA_N : constexpr(int) = 8
14         ...
15         num_blocks_n : const(int) = (N + MMA_N - 1) / MMA_N
16         # id() function returns the block id, inferred from
17         # @ block[1].
18         block_row : const(int) @ block[1] = id() / num_blocks_n
19         block_col : const(int) @ block[1] = id() % num_blocks_n
20
21         warp_row: const(int) @ block[1] = block_row * MMA_M
22         warp_col : const(int) @ block[1] = block_col * MMA_N
23         # Give each block an offset into C
24         offset = lambda x: warp_row * N + warp_col + x
25         with partition(C, p=block[1], f=offset) as C_blk:
26             with group(block[1]):
27                 # SMEM declarations only allowed with block[1] privilege.
28                 A_smem : shared(float[16 * 8]) @ block[1]
29                 B_smem : shared(float[8 * 8]) @ block[1]
30                 C_smem : shared(float[16 * 8]) @ block[1]
31                 # Now, id() returns the thread id
32                 idx : int @ thread[1] = id() * 4
33                 # To write to C_smem, must drop to thread[1] privilege
34                 with partition(C_smem, p=thread[1], offset=idx) as C_thrd:
35                     for i in range(0, 4, 1):
36                         with group(thread[1]):
37                             C_thrd[i] = 0
38
39                     for i in range(0, K_tiles, 1):
40                         a_idx : int @ thread[1] = id() * 4
41                         # <-- Name will insert a sync point
42                         # because this partition is a parent partition
43                         # of itself (back edge from for loop)
44                         for j in range(0, 4, 1):
45                             global_row : int @ thread[1] = warp_row + row
46                             global_col: int @ thread[1] = i * MMA_K + col
47                             with partition(A_smem, p=thread[1], f=...) as As_thrd:
48                                 with group(thread[1]):
49                                     As_thrd[0] = A[global_row * K + global_col]
50
51                         b_idx : int @ thread[1] = id() * 2
52                         # <-- Name will insert a sync point
53                         # (back edge from for loop)
54                         for j in range(0, 2, 1):
55                             # Similar to write into C_smem ...
56                             with partition(B_smem, p=thread[1], f=...) as Bs_thrd:
57                                 with group(thread[1]):
58                                     Bs_thrd[0] = B[global_row_b * N + global_col_b]
59
60                         # <-- Name will insert a sync point
61                         # (back edge from for loop)
62                         # Give each warp an offset into C_smem
63                         with claim(C_smem, p=thread[32]) as Cs_warp:
64                             match split(thread):
65                                 case 32:
66                                     # Can pass A_smem and B_smem
67                                     # as arguments to simple_mma for
68                                     # those two are declared const.
69                                     # "read up" is okay.
70                                     simple_mma(A_smem, B_smem, Cs_warp)
71
72                         # Write back final result to C
73                         for j in range(0, 4, 1):
74                             flat_idx_c : int @ thread[1] = id() * 4 + j
75                             row_c : int @ thread[1] = flat_idx_c / MMA_K
76                             col_c : int @ thread[1] = flat_idx_c % MMA_K
77                             offset = lambda x: row_c * N + col_c + x
78                             with partition(C_blk, p=thread[1], f=offset) as C_thrd:
79                                 with group(thread[1]):
80                                     C_thrd[0] = C_smem[row_c * MMA_N + col_c]
81
82     return

```

Fig. 6. Tensor Float 32 matrix Multiplication in Snap.

344 order defined over them. In particular, thread
 345 < block < grid.

346 There are two key differences between Snap's
 347 levels and those in CUDA. First, unlike CUDA,
 348 Snap does not model a a three-dimensional grid
 349 or block structure, say through different levels for grid_x, grid_y, or grid_z. Second, there are
 350 two "levels", warp and warp group, noticeably absent from our hierarchy.

351 These differences are a deliberate design choice. On the hardware, the units of each level are
 352 arranged in a single linear order, and the three-dimensional structure is simply *interpretations* of this
 353 linear ordering, not distinct hardware resources. Similarly, warp and warp group are organizational
 354 constructs defined in terms of existing levels . Namely, a warp is a group of 32 threads whose first
 355 thread ID is aligned to 32. A warp group, which only became a meaningful construct with the release
 356 of the Hopper architecture , consists of 128 threads aligned to 128.

357 Rather than embedding these interpretations into the language as special cases by introducing
 358 new levels, Snap provides constructs that let users group units of a given level . With these constructs,
 359 users can express multi-dimensional structures and define custom groupings with specific quantities
 360 and alignments.

361 The primary mechanism for organizing units at each level is a *privilege*.

362 3.2 Privileges & Privilege Classes

363 Privileges are the central concept of in Snap, allowing it to determine which compute resources are
 364 being requested by the programmer, whether they are available in the program's context, and, once
 365 provided, if those resources are sufficient for a given operation.

366 A *privilege* is a level—grid, block, or thread—parametrized by a static constant n, specifying the
 367 number of units at that level. A privilege is written as level[n]. For example, thread[2] denotes a
 368 privilege of two threads, block[4] denotes a scope of four blocks, and so on.

369 Privileges also carry *alignment information*: a privilege parameterized by n aligned to n. In this way,
 370 a warp is is simply a desugaring of thread[32], and a warp group is a desugaring of thread[128].

371 Finally, privileges have a *partial order* defined on them. We say that level₂[n₂] is a *higher privilege*
 372 than level₁[n₁] iff either level₁ < level₂, or, if level₁ = level₂, then n₂ % n₁ = 0. The condition n₂
 373 % n₁ = 0 may appear somewhat mysterious at first, but we will see why it is required in Section 3.2.2.

374 In Snap, both code and data are bound to privileges . Code is bound to a set of privileges, called a
 375 privilege class meanwhile data is bound to a single privilege.

376 3.2.1 Privileges for Code, or Privilege Classes.

377 Privilege classes capture what compute resources a program has access to. Each function starts
 378 out with a top-level privilege class that defines the total available resources that it *assumes* are
 379 available to it. As a shorthand, when we refer to a code's privilege, we mean the highest privilege
 380 available to it in its class. In Figure 6, the privilege class is declared on line using the notation
 381 @requires(grid=..,block=..) Starting from this top-level declaration, programmers shape the
 382 current privilege class in the function's body through two constructs, group and split.
 383

384 **Group.** The group construct bundles privileges available in the
 385 current privilege class into "groups" of a particular granularity. Intuitively, group narrows the current privilege class and replicates it
 386 across collections of a finer granularity. For example, Figure 6, line
 387 shows a grouping of a single block, while line in Figure 5 shows
 388 a grouping of 32 warps, that will ultimately invoke a tensor-core
 389 operation.

```

1 # Example 1
2 with group(thread[2]):
3     # Illegal because block > thread
4     with group(block[1]):
5         pass
6
7 # Example 2
8 with group(block[1]):
9     # Illegal because 1 % 2 != 0
10    with group([block[2]]):
11        pass

```

Fig. 7. Illegal uses of group.

393 An invocation of `group(level[n])` is allowed if and only if the
 394 privilege class contains privilege higher than `level[n]`. For example,
 395 `group(warp[32])` on in Figure 5 is allowed because the privilege class of the function declares that
 396 the top level on line starts out with at least 32 threads.

397 Once `group(level[n])` is invoked, it modifies the current privilege class in two key ways. First,
 398 it removes all parents of `level[n]` from it. Second, the number of available units for `level` is set to
 399 exactly `n`. Combined, these rules eliminate malformed programs that would result in unsatisfiable
 400 requirements. We shown some examples of illegal group statements in Figure 7.

401 **Split.** Unlike `group` that is used for replication over smaller, equal sized resource sets, `split` is used
 402 for sharding current available units into unequal sized resource sets. For example, line in Figure 6
 403 shows a `split` of a `thread[32]` privilege that is then used to execute a warp-level tensor-core
 404 operation.
 405

406 The `split` operation reflects unordered composition that is not
 407 mere replication, and all branches of the `split` execute in parallel.
 408 Once `split` `level` is invoked, the privilege classes of its branches
 409 diverge. A `split` construct will eventually compile to an `if`-statement
 410 during code generation.

411 Therefore, to satisfy a `split` operation, Snap checks that the
 412 *cumulative sum* of the privileges requested by each branch of the
 413 `split` can be satisfied. For example, if we were to change the program
 414 in Figure 6 slightly (on line), to the one shown in Figure 8, Snap
 415 would throw a compile-time error. Finally, because each privilege
 416 assumes alignment, every branch of the `split` must also be aligned. In
 417 other words, not all splits whose sizes sum to less than or equal to
 418 the available units are necessarily valid. Figure 9 shows an example
 419 of this constraint. This alignment constraint is also why we compare
 420 privileges by checking whether one quantity is a multiple of the
 421 other.

422 Inside each branch that requests `n` units, all parent levels of
 423 `level[n]` are removed from its privilege class, and the number of available units for `level` is
 424 set to exactly `n`, and, within that scope, all resource requirements are checked against that privilege
 425 class.

426 3.2.2 Privileges for Data.

427 Like code, data also has a privilege attached to it that dictates the set of privileges that must be held
 428 for reading and writing to it. Ultimately, it is this interaction between privilege classes for code and
 429 privileges for data that let us make meaningful statements about the set of privileges available at
 430 runtime. The precise details of this interaction are formalized in Section 4.

431 In this section, we focus tracking thread-local stack allocations and local variables. We defer
 432 discussion of pointers to memory to Section 3.5, where we will show how shared views of memory
 433 are lowered and restored through the compute hierarchy.

434 For thread-local data, its privilege is specified in its declaration using the `@ level[n]` syntax.
 435 For example, an integer variable `v` is declared with `thread[1]` privilege with the syntax `v : int @`
 436 `thread[1]`. If not explicitly annotated, Snap infers a variable's privilege to be the privilege of the
 437 code where it was declared. For example, `MMA_N`, `num_blocks_n` are all inferred to be at `grid[1]`
 438 scope on lines in Figure 6.

439 The rules for reading and writing to data are asymmetric: higher-privilege data can flow into lower-
 440 privilege data, but not the other way around. In particular, a value coming from a higher-privilege

```
1 # Example 1
2 with group(thread[32]):
3     match split(thread):
4         case 32:
5             simple_mma(..)
6             # Illegal! 32 + 1 > 33
7         case 1:
8             pass
```

Fig. 8. Illegal split that exceeds available privileges.

```
1 # Example 2
2 with group(thread[33]):
3     match split(thread):
4         case 1:
5             pass
6             # Illegal! 32 + 1 <= 33,
7             # but the second branch
8             # is not aligned by 32.
9         case 32:
10            simple_mma(..)
```

Fig. 9. Illegal split that violates alignment.

variable can be written into lower-privilege data, but a value coming from a lower-privilege variable cannot be written into higher-privilege data. In short, the rule is: "read up, write down".

If a value is read without being written into a variable, for example, in an `if` statement, the above rule is enforced using the privilege the code that is executing the read. For example, branching on a `thread[1]` variable within `block[1]` privilege code is rejected by our compiler, as it constitutes a "read down".

In Figure 6, we can see rule "read up" in action on lines . Meanwhile, "write down" is being used to set up variables on line .

Intuitively, `@ level[n]` tells Snap the *frequency* at which a value changes in space . For example, a `grid[1]` value is the same for all threads in a grid, whereas a `block[2]` value changes every two blocks. The rules described above enforce that the frequency declared for each variable is respected throughout the program. For instance, rule (1) allows variables with higher replication frequency to read from lower-frequency variables without violating their declared frequency. Conversely, higher-frequency variables cannot be used to write into lower-frequency variables, since that would violate the variable's declared frequency.

By tracking frequency in this way, Snap can ensure that threads that with a given privilege can only read variables that change at the same frequency or higher. This guarantees that all threads with that privilege follow the same control flow, and that our accounting of available privilege class remains consistent with the program's runtime behavior.

Until now, we have discussed rules for reading and writing to variables that change at different frequencies, but we have not yet set up a mechanism for actually setting up variables that diverge and change at different frequencies. The main tool for doing this is the `id()` function.

3.3 The `id()` Function

As opposed to exposing users to special hardware `blockIdx.x`, `threadIdx.x` variables directly, Snap provides the `id()` function instead. The `id` function returns the relative index, or a "label", for a given compute privilege. The interpretation of the `id` function depends on both the privilege of the variable it is being written to, and the privilege class that the `id` function is being invoked in. In our example in Figure 6, we use a call to `id` in `grid[1]` scope at line to locate the tile that each block is ultimately charge of computing. In Figure 3, however, a call to `id` in `thread[32]` scope at line will not return the global thread ID of the threads launched in the kernel, but rather the *relative* thread ID within the 32 threads, ranging from 0 to 31.

3.4 Collective Operations

With privilege classes for code and data in place, we can now turn our attention to Snap's reasoning about collective operations. By explicitly tracking privilege classes for code, Snap statically guarantees that collective operations are semantically valid by ensuring that they are invoked with sufficient privileges.

There are two ways to access collective operations in Snap. The first is to use intrinsics provided by the compiler. These intrinsics have built-in privilege classes attached to them, which the compiler can directly enforce. In Figure 3, the tensor-core call at line is made using one of the intrinsics provided by our compiler. The second, more flexible approach allows users to extend these operations using the `unsafe` construct, explained in Section 3.8. Using unsafe code, users can inline a collective operation—often a single assembly instruction—into a Snap function and ensure that the function's top-level requirements, specified with `@require`, faithfully describe the privileges required to execute that instruction. Then, Snap will check the call to the function as any other function call, enforcing the correct privilege. The rules for checking function calls are described in Section 3.7.

```

1 with group(block[1]):
2   id : int @ thread[1] = id()
3   # Illegal read of ID from
4   # block[1] scope.
5   if (id)
6     __syncthreads();

```

Fig. 10. Illegal read of `thread[1]` variable.

A quirk of some collective operations is that, although they need to be invoked with higher compute privileges, they often produce results with lower data privileges. For example, the tensor-core operation in Figure 3 is invoked at `thread[32]` granularity but generates output at the `thread[1]` level. That is, even though the call to the tensor call is a convergent operation requiring the `thread[32]` compute privilege, it produces divergent output with `thread[1]` data privilege. To be able to handle these operations we must be able to allow `thread[1]` state to coexist with the `thread[32]` privilege.

One potential solution is to artificially coarsen the granularity of the operation by having the intrinsic internally scatter the `thread[1]` output into a `thread[32]` variable. However, this approach is often unacceptable for low-level CUDA performance tuning. Often, performance engineers need to manually schedule and interleave memory and arithmetic instructions to hide latency, with loads or stores occurring well in advance of their dependent operations and interleaved with other instructions. If we coarsen these operations, we will disallow these optimizations entirely.

In Snap, however, this behavior does not pose a problem. Because of the "write down" rule, registers can be set up well in advance of the operation (on line in Figure 3, for example) without violating correctness.

3.5 Memory

In Snap, there are three mechanisms for obtaining a pointer type. The first method is to directly allocate data. The second method is to obtain offsets into existing data by using the `partition` construct and the `claim` construct. The `partition` and `claim` constructs mirror the `group` and `split` constructs, and are used to lower the data privilege associated with memory.

3.5.1 Allocations in Different Memory Spaces.

In Section 2.2, we saw that the GPU has distinct memory spaces, like global, local, and shared memory. Local allocations reside in either the stack or registers, which Snap can distinguish upfront. Shared memory pointers are explicitly annotated with the `shared` keyword, while global pointers are unannotated.

The primary memory resource Snap needs to manage are shared memory allocations. Local memory allocations are just stack allocations, and don't need special treatment. We also assume that sufficient global memory has been allocated before the kernel launch and don't provide a mechanism for runtime allocations—this is typically the case in real-world programs.

In CUDA, shared memory is a limited resource, typically a few hundred kilobytes per block. Programmers manually manage this memory by bumping pointers to allocate offsets for different data. Each allocation's pointer must be carefully tracked throughout the program, and exceeding the available space will lead to a runtime error.

Snap simplifies shared memory management by requiring all shared memory allocations have a static size. Since shared memory is only visible to all threads in a block, a shared memory allocation is only valid with `block[1]` privilege. An example of such an allocation is shown on lines of Figure 6. During compilation, Snap will automatically handle the necessary pointer arithmetic to assign each allocation an appropriate offset within the shared memory space.

The `@requires` annotation specifies the amount of shared memory a function expects to have available, and Snap uses this information to ensure that a function's allocations do not exceed its declared limit, and checks whether there is enough shared memory available at its call sites. Figure 6 declares the amount of shared memory it will require on . We use standard techniques to conservatively bound memory usage in Snap's type system.

3.5.2 Privileges for Memory.

In Section 3.2.2, we focused on tracking for primitive data types and stack allocations, where each thread maintained its own local copy of the data. In contrast, pointers into memory, such as global or

540 shared memory, do not represent distinct copies, but instead refer to the same underlying storage
 541 that may be accessed simultaneously by multiple threads. To use them, we must first *lower* their
 542 privileges, allowing individual threads to write different values to different locations, and then *restore*
 543 them back to their original privilege before execution continues.

544 As with variables in Section 3.2.2, lower-privilege data may read from higher-privilege data
 545 without restrictions. For writes, however, Snap enforces that all pointers into memory be written
 546 with `thread[1]` privilege; they therefore must be lowered before they can be safely updated.

547
548

549 **Partition.** The `partition` construct plays the same role for memory that `group` plays for compute.
 550 It is used to *lower* a memory object from a higher privilege to a lower one by assigning each lower
 551 privilege a distinct offset into the original pointer. It takes the original name of a memory object,
 552 `base_name`, an indexing function `f`, and a target privilege `level[n]` to which the memory is being
 553 lowered, and produces a new name, `new_name`, for the partitioned memory. In Figure 6, for example,
 554 we partition the C buffer to `block[1]` privilege on line first, and then lower it further to `thread[1]`
 555 privilege on line so that we can actually write to it.

556 Once a memory object has been partitioned, the original name referring to it goes out of scope,
 557 and the memory can only be accessed through the new name introduced by `partition`. Each use of
 558 this new name applies the indexing function `f` to compute the true offset of the access. Within the
 559 scope of the `partition` statement, the original name cannot be repartitioned. Because `split` denotes
 560 unordered composition, Snap prohibits sibling `split` branches from partitioning the same memory
 561 object. When the scope of the `partition` statement ends, Snap automatically inserts the necessary
 562 synchronization to restore the memory to its original privilege. The mechanism for inserting and
 563 optimizing these synchronization points is described in Section 3.9.

564 At this point, it is necessary to consider how different indexing functions affect the possibility
 565 of data races. If the indexing function is injective, each lower-privilege privilege receives a distinct
 566 offset into the underlying array, and the resulting partition is free of data races. When the indexing
 567 function is not injective, multiple privileges may write to the same location, introducing a potential
 568 data race. Out-of-bounds accesses are undefined behavior.

569 As we discussed in Section 2.2, achieving data-race freedom is *not* one of Snap’s goals, and
 570 our guarantees continue to hold even when indexing functions are not injective. Because Snap
 571 automatically inserts the necessary synchronization before the original memory name is restored to
 572 its higher privilege, Snap ensures that all writes have completed before accessing the orginal name.
 573 The subsequent uses of that pointer will observe the value written by the last-writer.

574 We believe that extending Snap to reason about data races and out-of-bounds accesses is a promising
 575 direction for future work.

576
577

578 **Claim.** The `claim` construct is the analogue of `split` for memory . Using `claim`, the entire
 579 memory region can be moved to a lower privilege.

580 Unlike `partition`, `claim` does not take an indexing function. The `claim` construct takes the
 581 original memory’s name, the target privilege to lower to, and a new name to assign to the lowered
 582 memory. The `new_name` is accessible only within a single, explicit `split` branch; sibling branches are
 583 not permitted to read or write from this memory. In Figure 6, we use the `claim` operation on line to
 584 lower the memory to `thread[32]` privilege.

585 As with `partition`, a `split` construct introduces its own scope, and the original name is inacces-
 586 sible within this scope.

587
588

589 **3.6 Asynchrony**

590 **3.7 Function Composition**

591 To check whether a function is valid, Snap needs to ensure that the call site can provide the set of
 592 privileges that a function will attempt to request over its execution, and that the arguments passed to
 593 the function also respect the privileges declared for its input and output types. Both requirements are
 594 represented explicitly in the function signature in the `@require` construct that we seen in Section 3.2
 595 and Section 3.5.

596 Like mentioned in Section 3.2, each function carries a privilege class. We call this top-level privilege
 597 class the function’s *privilege signature*. The signature specifies the top-level privilege class the function
 598 *assumes* will be provided to it. This signature represents the *minimum amount* of compute and memory
 599 resources the function must be called with. In the example above in Figure 6, line defines the privilege
 600 signature of a function. Up to this point, we have described programs written *inside* the function
 601 under the assumption that the resource signature can be satisfied. In Section 3.2, for example, we
 602 used this signature to set up the top-level privilege class for the function.
 603

604 **Compute privileges.** At a function’s call site, then, we need to ensure that the function’s privilege
 605 signature can be satisfied. Using the same mechanism as Section 3.2 and Section 3.5, Snap can track
 606 the current compute privilege of the program, as well as how much shared memory is available for
 607 use. Through this, Snap can verify that only functions whose privilege signature can be satisfied are
 608 called.
 609

610 **Per-argument checks.** In Snap, functions have pass-by-value semantics. For arguments, we
 611 distinguish between primitive data types and pointers to memory. For primitive data types, like `int`,
 612 `bool`, and stack allocations, users can pass either arguments that live at the exact data privilege, or at
 613 the higher data privilege. For pointers, we allow passing in pointers that live at higher data privilege
 614 if the pointer is marked as `const`, indicating that it will only be used for reading (line in Figure 3).
 615 Otherwise, we require that pointers match the exact scope of the function signature.

616 **3.8 Unsafe**

617 **3.9 Compiler Internal: Synchronization**

618 **4 Formalism**

619 Having introduced the full Snap language, we now describe Bundl, a core calculus that formalizes its
 620 most fundamental aspects. Bundl statically tracks compute and data privileges, and we would like to
 621 use it to prove that well-typed Snap programs successfully evaluate and will never get stuck trying
 622 to execute operations for which they do not have the required compute resources. To that end, in
 623 this section we describe the Bundl’s type system and operational semantics—in particular how it
 624 manages compute and data privileges—and build up to a formal proof of type-and-privilege safety.
 625

626 **4.1 Bundl Type System**

627 Bundl’s core idea is type-level tracking of the program’s current compute privilege. To achieve this,
 628 we borrow techniques from the literature on *coeffects* [11] and *dependency tracking* [1]. In particular,
 629 the current compute privilege is tracked on the typing judgment, which has the form $\Gamma \vdash^\pi e : \tau$ for
 630 expressions and $\Gamma \vdash^\pi s$ for statements. The π over the \vdash is the compute privilege of e and s .

631 The π s themselves are pairs of levels h of the GPU hierarchy at which each privilege lives (either
 632 Thread, Block or Grid), and numbers tracking how many copies of that level each privilege possesses.
 633 So, for example example, a (Thread,4) in Bundl would directly correspond to a `thread[4]` in Snap.

634 Meanwhile, each variable in the typing context has a data privilege attached to it. Only when these
 635 two privileges match can data be read or written from a variable. This requirement is made manifest
 636

| | |
|---|--|
| 638 $\boxed{\Gamma \vdash^\pi e : \tau}$ | <i>(Expression typing)</i> |
| 639 $\frac{x :^\pi \tau \in \Gamma}{\Gamma \vdash^\pi x : \tau}$ T-VAR | |
| 640 | |
| 641 | |
| 642 $\boxed{\Gamma \vdash^\pi s}$ | <i>(Statement typing)</i> |
| 643 | |
| 644 $\frac{\Gamma \vdash^{(h,n_1)} s_1 \quad \Gamma \vdash^{(h,n_2)} s_2 \quad n_1, n_2 \text{ align to } n}{\Gamma \vdash^{(h,n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\}}$ T-SPLIT | |
| 645 | |
| 646 | |
| 647 where $n_1, n_2 \text{ align to } n ::= (n_1 + n_2 \leq n) \text{ and } (n_1 n) \text{ and } (n_2 n) \text{ and } (n_2 n_1 + n)$ | |
| 648 | |
| 649 $\frac{\Gamma \vdash^\pi s}{\Gamma \vdash^{q \cdot \pi} \text{group } q \ s}$ T-GROUP | $\frac{\Gamma, y :^{\downarrow \pi} \tau[]^l \vdash^\pi s \quad l \neq \text{Local}}{\Gamma, x :^\pi \tau[]^l \vdash^\pi \text{lower } x \text{ into } y \text{ in } s}$ T-LOWER |
| 650 | |
| 651 | |
| 652 $\frac{\Gamma, y :^{(h,n/c)} \tau[]^l \vdash^{(h,n)} s \quad c n \quad l \neq \text{Local}}{\Gamma, x :^{(h,n)} \tau[]^l \vdash^{(h,n)} \text{partition } x \text{ into } y \text{ by } c \text{ in } s}$ T-PARTITION | $\frac{\Gamma \vdash^\pi s}{\Gamma \vdash^\pi \text{destruct in } s}$ T-DESTRUCT |
| 653 | |
| 654 | |
| 655 $\frac{\Gamma, y :^{(h,n')} \tau[]^l \vdash^{(h,n')} s \quad n' \leq n \quad l \neq \text{Local}}{\Gamma, x :^{(h,n)} \tau[]^l \vdash^{(h,n)} \text{claim } x \text{ into } y \text{ at } n' \text{ in } s}$ T-CLAIM | |
| 656 | |
| 657 | |
| 658 | |

Fig. 11. Core typing rules of Bundl. The typing rules presented here are a simplified selection of the full rules, which can be found in Appendix A.2.

in the **T-VAR** rule, which can be found in Figure 11, along with the other key typing rules of Bundl. The premise of this rule also reveals that the context Γ maps variables to both types τ and also the privileges π with which they may be read and written.

The rules we discuss in detail in the section fall into two main categories: rules for managing the current compute privilege and rules for managing data privileges.

4.1.1 Compute Privilege Management. In Snap, there are two operations which manage compute privileges: `codegroup` and `split`. In Bundl, to better tease apart the subtleties of privilege attenuation, we introduce a third construct, `destruct`.

Bundl’s `group` statement directly corresponds to Snap’s, and is checked by the **T-GROUP** rule. Given some privilege that can be divided into q equally sized smaller privileges π , the statement `group` q s will check with privilege $q \cdot \pi$ provided that s itself checks with privilege π . Because the starting privilege is of the form $q \cdot \pi$, we do not need to check for alignment.

The `split` statement, meanwhile, is checked by the **T-SPLIT** rule, and functions like a binary version of the n-ary splitting construct in Snap. It enforces the same divisibility requirements to ensure that the privileges of code and data remain properly aligned, and then checks the two sub-statements s_1 and s_2 with the divided, smaller privileges.

The final rule for managing compute privileges is **T-DESTRUCT**, which handles downward movement from one level of the GPU hierarchy to another and makes explicit in Bundl programs exactly where such movement occurs. The \downarrow operation on π s “destructs” the privilege into a many smaller privileges at a lower level; $\downarrow(\text{Grid}, 1) = (\text{Block}, B)$ and $\downarrow(\text{Block}, 1) = (\text{Thread}, T)$, where B and T are parameters to a particular instantiation of Bundl to describe the number of blocks per grid and threads per block. Because Bundl is abstracted over these B and T values, instead of tracking classes of privileges the way that Snap does, it only tracks the top-level privilege described in Section 3.2.

$$\begin{array}{l}
 \frac{L(t), S(b), \Sigma, t, b, 0 \vdash^{(\text{Grid}, 1)} s \rightsquigarrow s' \dashv \eta', \sigma', \Sigma' \quad P(t, b) = s}{L, S, \Sigma, P \rightsquigarrow L[t \mapsto \eta'], S[b \mapsto \sigma'], \Sigma', P[(t, b) \mapsto s']} \text{ S-PROGRAM} \\
 \\
 \frac{p < n_1 \quad n_1, n_2 \text{ align to } n \quad \eta, \sigma, \Sigma, t, b, p \vdash^{(h, n_1)} s_1 \rightsquigarrow s'_1 \dashv \eta', \sigma', \Sigma'}{\eta, \sigma, \Sigma, t, b, p \vdash^{(h, n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\} \rightsquigarrow \text{split}(n_1, n_2)\{s'_1\}\{s_2\} \dashv \eta', \sigma', \Sigma'} \text{ S-SPLIT-LEFT} \\
 \\
 \frac{p \geq n_1 \quad p < n_1 + n_2 \quad n_1, n_2 \text{ align to } n \quad \eta, \sigma, \Sigma, t, b, p - n_1, \vdash^{(h, n_2)} s_2 \rightsquigarrow s'_2 \dashv \eta', \sigma', \Sigma'}{\eta, \sigma, \Sigma, t, b, p \vdash^{(h, n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\} \rightsquigarrow \text{split}(n_1, n_2)\{s_1\}\{s'_2\} \dashv \eta', \sigma', \Sigma'} \text{ S-SPLIT-RIGHT} \\
 \\
 \frac{\eta, \sigma, \Sigma, t, b, p \text{ mod } n \vdash^{(h, n)} s \rightsquigarrow s' \dashv \eta', \sigma', \Sigma'}{\eta, \sigma, \Sigma, t, b, p \vdash^{(h, q \cdot n)} \text{group } q \ s \rightsquigarrow \text{group } q \ s' \dashv \eta', \sigma', \Sigma'} \text{ S-GROUP} \\
 \\
 \frac{\eta, \sigma, \Sigma, t, b, t \text{ mod } T \vdash^{(\text{Thread}, T)} s \rightsquigarrow s' \dashv \eta', \sigma', \Sigma'}{\eta, \sigma, \Sigma, t, b, 0 \vdash^{(\text{Block}, 1)} \text{destruct in } s \rightsquigarrow \text{destruct in } s' \dashv \eta', \sigma', \Sigma'} \text{ S-DESTRUCT-BLOCK}
 \end{array}$$

Fig. 12. Core semantic rules of Bundl. As with the typing rules, we present only a simplified selection of the full rules, which can be found in Appendix A.3.

Because \downarrow is only defined on two inputs, the rule enforces that one can only destruct their compute privilege when in possession of exactly one grid or block.

4.1.2 Data Privilege Management. The data privilege management rules mirror those handling compute privilege, with one operation for data corresponding to each operation for code.

The `partition` operation corresponds to grouping a compute privilege. The typing rule for this operation, `T-PARTITION`, requires that the partitioned variable x is in the context with the same data privilege as the current compute privilege. However, it then renames the variable into a fresh variable y with a new privilege π/c , a division of all the computational resources in π among c equally-sized smaller privileges. Checking of s then proceeds with the original privilege π : partitioning data with a new data privilege does not change the current compute privilege. Note, however, that the original name of the variable is not visible in s ; only the new name may be accessed from within the partition.

In the `T-CLAIM` rule, on the other hand, we do change our compute privilege and our data privilege at the same time. Unlike `partition`, which divides up a piece of data equally among all the privileges that divide its current one, the `claim` operation gives the entire variable one smaller privilege, and then checking of s proceeds with that attenuated privilege. This represents a minor difference from Snap: whereas Snap can use additional static analysis to ensure that the claimed variables are not accessed multiple times across sibling splits, Bundl instead links the data privilege of the variable to the compute privilege of the code claiming it, ensuring that they remain linked.

Lastly, the `T-LOWER` rule mirrors the `T-DESTRUCT` rule; it uses the \downarrow operator to move a variable from one level of the hierarchy to another, distributing it equally among all the child privileges in the same manner as `T-PARTITION`.

4.2 Bundl Semantics

To reflect the fact that GPU programs execute in parallel across numerous threads, we model the semantics of Bundl in the style of Turon et al. [15], using a two-level small step judgment. We present the key rules of this semantics in Figure 12.

The top level (i.e., device-level) judgment has just one rule: **S-PROGRAM**. This rule acts as a “frame” for the lower level (i.e., thread-level) judgment, and steps a collection of thread-id-indexed local memories (L), a collection of block-id-indexed shared memories (S), a global memory (Σ), and a *thread pool* to an updated collection of memories and thread pool. The thread pool maps thread and block ids to code, intuitively representing the program being executed by each thread at the current moment. The **S-PROGRAM** non-deterministically chooses a thread id and block id and steps it according to the thread-level judgment. This allows the semantics to model the full range of non-deterministic behavior arising from the GPU’s thread scheduler.

The thread-level judgment has the shape $\eta, \sigma, \Sigma, t, b, p \vdash^\pi s \rightsquigarrow s' \dashv \eta', \sigma', \Sigma'$. The η s in this judgment represent the thread’s local memory, the σ s represent the shared memory, the Σ s represent global memory, and the s ’s represent the statement being executed by the thread. The primed versions of each of these are the output of the step judgment. Critically, notice that a π also appears on the thread-level judgment just as it does on the typing judgment. This is because the thread semantics *dynamically tracks and verifies privileges*. The same way a program can get stuck if a value does not have the right type, the semantics of Bundl also get stuck if code attempts to access data or invoke commands with the wrong privilege. This is an artifact of Bundl and is not present in Snap itself, but will allow us to prove later in Section 4.3 that well-privileged programs will always execute with the same privilege that the type system checked them against.

The other variables in the judgment track the positions of threads within their scopes. The t is the thread’s id, the b is the id of the block in which the thread lives, and the p is the relative position of the thread in the compute privilege with which it is executing. For example, threads with a (Thread,4) privilege each have a unique p between 0 and 3, while threads with a (Block,16) privilege will have the same p within in each block, between 0 and 15 depending on the block.

The semantic rules for privilege management involve manipulating this p value to track which threads take which code paths when privileges are split or grouped. Notice that in the **S-PROGRAM** rule, the thread stepping judgment always begins with privilege (Grid,1): all the privilege management rules are congruences, handling further evaluation with lower privileges as determined by the particular syntactic construct used.

The first set of these privilege management rules handle the `split` construct. They enforce the same alignment requirements as the **T-SPLIT** rule. The rules choose which path to take based on the value of p : if p is less than the size of the left-hand privilege n_1 , execution of s_1 will proceed with that lower privilege via the **S-SPLIT-LEFT** rule, while if p is between n_1 and n_2 execution of s_2 will proceed via the **S-SPLIT-RIGHT**. In this latter case, we execute s_2 with a new p value that subtracts off the value of n_1 . As an example, if we encounter a `split(2,2){s1}{s2}` with a (Thread,4) privilege, the first two threads will go “left” and execute s_1 with a (Thread,2) privilege, while the third and fourth threads will go “right” and execute s_2 with a (Thread,2) privilege.

The **S-GROUP** rule is much simpler, as every thread with a compute privilege that encounters this operation will execute the sub-statement s . What changes across threads is how their p value will be modified by the grouping operation. In each case, the p of the thread is reduced modulo n , where n is the size of the privilege with which s is executed. This effectively recolors all the threads privilege the new, smaller, privilege with which they will see s .

The **S-DESTRUCT-BLOCK** rule for `destruct` behaves similarly to **S-GROUP**. However, as p is an index into the current privilege, it is always 0 when a privilege has size 1. The `destruct` statement can only be executed with a size 1 privilege, and so p is not particularly useful here. To compute the new p for the execution of s with a privilege of size T , we reduce the thread id modulo T and continue executing s with that as the new value of p . A similar rule exists for destructing a Grid privilege into a size B Block privilege.

These rules take great care to ensure that p always describes the relative position of a compute resource within its privilege; the payoff is that Bundl’s semantics can later use this p value to model the way that Snap automatically adjusts indices into data when partitioning a data privilege.

Beyond these key rules for compute privilege management, we have modeled many of the other features of Snap, such as asynchronous operations and thread synchronization, in Bundl. To handle features like these we equip the operational semantics with additional structure, including sets of semaphores [5] for thread syncs and a stack of effect handlers for handling deferred asynchronous computations in the style of Ahman and Pretnar [2]. We have elided these details here for simplicity, but interested readers can see further details in Appendix A.3.

4.3 Bundl Theorems

THEOREM 4.1. (*Type-and-Privilege Safety*). *For any program s such that $\Gamma \vdash^\pi s$, either:*

- (1) s is skip, or
- (2) *for any well-typed environments η, σ , and Σ , there is an s', η', σ' , and Σ' such that $\eta, \sigma, \Sigma, t, b, p \vdash^\pi s' \rightarrow \eta', \sigma', \Sigma'$ where η', σ' , and Σ' are well-typed by an extension of Γ .*

PROOF. Via the usual progress and preservation lemmas. The full proof of these lemmas can be found in Appendix A.4. \square

In addition to the usual guarantees about type safety, this theorem also guarantees that programs will never get stuck trying to execute operations for which they do not have the required compute resources. By instrumenting Bundl’s operational semantics with privileges, we allow our safety theorem to guarantee that the dynamically-realized privilege matches the one inferred by the type system. It is, however, worth noting that this soundness theorem is a safety theorem, not a liveness theorem: it guarantees that if all threads that can reach a point in the program with statically-determined privilege π do reach that point, then those threads will be logically grouped according to that same privilege π . It does not, however, guarantee that all threads that can reach the point will eventually do so. Indeed, in the presence of nontermination, liveness does not hold: a subset of the threads could split off and loop forever. While we believe the liveness version of this theorem holds for a terminating fragment of Bundl, it is beyond the scope of this paper—such proofs are notoriously challenging and are often research contributions [3, 6, 15] in and of themselves. We hope to tackle this proof for Bundl in future work.

5 Compilation

If a program typechecks, Snap lowers the program into a CUDA file, which does not contain any run-time checks. The CUDA file is then compiled by nvcc, NVIDIA’s closed-source compiler, to produce a machine-executable. As Snap is an imperative language that provides the same level of abstraction as CUDA, there is a one-to-one mapping between most language constructs and their CUDA counterparts.

Snap is implemented as a compiler tool on top of the CLANG/LLVM compiler stack. Textually, Snap programs are written as the examples described in the paper, but the constructs themselves lower to traditional C++ statements that our frontend tool can analyze and generate code for. This is accomplished by making liberal use of attributes, macros, and other sugaring features of C++.

- Make correspondence to CUDA clear
- Would be awesome to have a super clear figure about how the compilation step actually works (show the one-to-one mapping, though it is slightly more complicated)
- Make optimizations (lazy sync, commit group etc) as a separate optimization section.
- highlight how new function interfaces work out: we add extra arguments

834 6 Evaluation

835 We evaluate Snap in the context of the four main questions:

836 **RQ1** Can Snap express a diversity of CUDA programs?

837 **RQ2** Can Snap express programs that use advanced GPU features?

838 **RQ3** Can Snap achieve performance competitive with existing, speed-of-light CUDA code?

839 **RQ4** Can Snap catch bugs that CUDA would otherwise allow?

840 Although the following question is not central to our paper’s main claims—since library design
 841 ultimately reflects choices made by programmers, and our language allows abstraction design much
 842 like any other language—we were, nevertheless, curious to explore it based on empirical and anecdotal
 843 observation:

844 **RQ5** Can Snap help build compositional libraries that users can use with confidence?

845 To perform this evaluation, we use two GPUs. The first is the NVIDIA H100 SXM5, a server-grade
 846 chip that supports tensor-core operations and a dedicated hardware copy engine, the Tensor Memory
 847 Accelerator. Notably, the H100 introduces a new *logical* level called the *warpgroup* (collection of
 848 4-aligned warps), and we show that our programming model can accommodate this new level.
 849 Moreover, because the H100 has historically served as the primary GPU for large-scale AI training,
 850 many CUDA kernels on this hardware are already highly optimized and achieve near-speed-of-
 851 light performance, providing a rigorous baseline for comparison. To ensure our results generalize
 852 beyond the H100, we additionally test programs on a second GPU, the NVIDIA 4070 SuperTi GPU, a
 853 consumer-grade chip.

854 Machine Specification

855 The NVIDIA H100 SXM5

856 The NVIDIA 4070 SuperTi GPU

857 Experimental Setup

858 When Snap produces a CUDA program, we compile the program with nvcc version with flags .

859 To measure performance, we run

860 6.1 RQ1: Can Snap express a diversity of CUDA programs?

861 Single-pass Parallel Prefix Scan with Decoupled Look-back here!

862 6.2 RQ2: Can Snap express programs that use advanced GPU features?

863 To test whether Snap can express programs that use advanced features of modern GPUs, we write a
 864 matrix multiplication for the bf16 datatype for the **H100**.

865 The H100 bf16 matrix multiplication pushes several language features to the extreme.

866 6.3 RQ3: Can Snap achieve performance competitive with existing, speed-of-light CUDA 867 code?

868 6.4 RQ4: Can Snap catch bugs that CUDA would otherwise allow?

869 To

870 6.5 Can Snap help build compositional libraries that users can use with confidence?

871 We use Snap to build an equivalent of NVIDIA’s CUB library []. CUB

883 CUB occupies a unique design point in the GPU library ecosystem. Unlike other many libraries
 884 like cuBLAS, cuDNN, cuSparse which provide global interfaces that users can call and configure,
 885 CUB provides a device-side library, organized into different tiers.

886 CUB also serves as an interesting comparison due it's focus on *collective reduction operations* that
 887 call into hardware primitives (like warpshuffles).

888 7 Limitations

889 Make sure intro highlights all of this/at least signals to them.

- 891 • Can make sub-optimal decision regarding commit group
- 892 • Grid sync: we don't handle right now:
- 893 • Static pipeline only
- 894 • Not about CPU & GPU: only about GPU

895 8 Related Work

896 *GPU Programming Languages.*

897 *Theoretic Foundations.* The design of Bundl is heavily inspired by existing work on coeffect systems
 898 [11, 12] and dependency tracking [1]. Coeffects allow type systems to talk about how programs
 899 depend on their environments, and have achieved widespread use in Rust [13] in the form of linear
 900 types. In Snap and Bundl, privileges act like coeffects by describing what compute resources are
 901 necessary for programs to execute. Dependency tracking calculi, meanwhile, allow type systems to
 902 track how data and code depend on each other, and have commonly been used to implement secure
 903 information flow analyses [4]. In Bundl, low privilege data is unable flow into high privilege contexts,
 904 and we use dependency tracking to capture this restriction in Bundl's type system.

905 9 Conclusion

906 References

- 907 [1] Martín Abadi, Anindya Banerjee, Nevin Heintze, and Jon G. Riecke. 1999. A core calculus of dependency. In *Proceedings*
 908 *of the 26th ACM SIGPLAN-SIGACT symposium on Principles of programming languages* (San Antonio Texas USA, 1999-01).
 909 ACM, 147–160. [doi:10.1145/292540.292555](https://doi.org/10.1145/292540.292555)
- 910 [2] Danel Ahman and Matija Pretnar. 2021. Asynchronous effects. *Proceedings of the ACM on Programming Languages* 5,
 911 POPL (Jan. 2021), 1–28. [doi:10.1145/3434305](https://doi.org/10.1145/3434305)
- 912 [3] Lars Birkedal, Filip Sieczkowski, and Jacob Thamsborg. 2012. A Concurrent Logical Relation. In *Computer Science Logic
 913 (CSL'12) - 26th International Workshop/21st Annual Conference of the EACSL (2012)*. Schloss Dagstuhl – Leibniz-Zentrum
 914 für Informatik, 107–121. [doi:10.4230/LIPIcs.CSL.2012.107](https://doi.org/10.4230/LIPIcs.CSL.2012.107)
- 915 [4] Dorothy E. Denning and Peter J. Denning. 1977. Certification of programs for secure information flow. *Commun. ACM*
 916 20, 7 (July 1977), 504–513. [doi:10.1145/359636.359712](https://doi.org/10.1145/359636.359712)
- 917 [5] Edsger Wybe Dijkstra. 1963. Over de sequentiaaliteit van procesbeschrijvingen. (1963). <https://training-ir7.tdl.org/handle/123456789/594>
- 918 [6] Azadeh Farzan, Zachary Kincaid, and Andreas Podelski. 2016. Proving Liveness of Parameterized Programs. In
 919 *Proceedings of the 31st Annual ACM/IEEE Symposium on Logic in Computer Science*. ACM, New York NY USA, 185–196.
 920 [doi:10.1145/2933575.2935310](https://doi.org/10.1145/2933575.2935310)
- 921 [7] Troels Henriksen, Niels G. W. Serup, Martin Elsman, Fritz Henglein, and Cosmin E. Oancea. 2017. Futhark: purely
 922 functional GPU-programming with nested parallelism and in-place array updates. In *Proceedings of the 38th ACM
 923 SIGPLAN Conference on Programming Language Design and Implementation*. ACM, Barcelona Spain, 556–571. [doi:10.1145/3062341.3062354](https://doi.org/10.1145/3062341.3062354)
- 924 [8] Bastian Köpcke, Sergei Gorlatch, and Michel Steuwer. 2024. Descend: A Safe GPU Systems Programming Language. 8
 925 (2024), 841–864. Issue PLDI. [doi:10.1145/3656411](https://doi.org/10.1145/3656411)
- 926 [9] NVIDIA Corporation. 2025. *CUB: CUDA Unbound*. <https://docs.nvidia.com/cuda/cub/index.html> CUDA Toolkit
 927 Documentation.
- 928 [10] NVIDIA Corporation. 2025. *CUDA C++ Programming Guide*. NVIDIA. <https://docs.nvidia.com/cuda/cuda-c-programming-guide/> Accessed: 2025-11-06.

- 932 [11] Tomas Petricek, Dominic Orchard, and Alan Mycroft. 2013. Coeffects: Unified Static Analysis of Context-Dependence.
933 In *Automata, Languages, and Programming*, Fedor V. Fomin, Rūsiņš Freivalds, Marta Kwiatkowska, and David Peleg
934 (Eds.). Springer, Berlin, Heidelberg, 385–397. [doi:10.1007/978-3-642-39212-2_35](https://doi.org/10.1007/978-3-642-39212-2_35)
- 935 [12] Tomas Petricek, Dominic Orchard, and Alan Mycroft. 2014. Coeffects: a calculus of context-dependent computation.
936 *SIGPLAN Not.* 49, 9 (Aug. 2014), 123–135. [doi:10.1145/2692915.2628160](https://doi.org/10.1145/2692915.2628160)
- 937 [13] Rust Team. 2025. Rust Programming Language. <https://rust-lang.org/>
- 938 [14] Philippe Tillet. 2025. Introducing Triton: Open-source GPU programming for neural networks. <https://openai.com/index/triton/>
- 939 [15] Aaron J. Turon, Jacob Thamsborg, Amal Ahmed, Lars Birkedal, and Derek Dreyer. 2013. Logical relations for fine-grained
940 concurrency. In *Proceedings of the 40th annual ACM SIGPLAN-SIGACT symposium on Principles of programming languages*
941 (Rome Italy, 2013-01-23). ACM, 343–356. [doi:10.1145/2429069.2429111](https://doi.org/10.1145/2429069.2429111)

942

943

944

945

946

947

948

949

950

951

952

953

954

955

956

957

958

959

960

961

962

963

964

965

966

967

968

969

970

971

972

973

974

975

976

977

978

979

980

981 A Complete Bundl Type System, Semantics, and Syntactic Soundness Proofs

982 A.1 Basic Definitions

984 Hierarchy Levels $h ::= \text{Grid} \mid \text{Block} \mid \text{Thread}$
 985 Memory Kinds $l ::= \text{Local} \mid \text{Shared} \mid \text{Global}$
 986 Privileges $\pi : h \times \mathbb{N}$
 987 Base Types $b ::= \text{bool} \mid \text{int} \mid \text{float}$
 988 Types $\tau ::= b \mid b[]^l \mid \text{Fun}(\Gamma, \pi, m) \mid \text{async } \tau$
 989 Contexts $\Gamma ::= \cdot \mid \Gamma, x : \pi \tau$
 990 Shared Memory Remaining $m : \mathbb{N}$

994 Privileges π are part of an algebra parameterized over some constant values T (the number of threads
 995 per block) and B (the number of blocks per grid). With these values, we have two isomorphisms:

$$997 \quad (\text{Block}, 1) \cong (\text{Thread}, T)$$

998 and

$$1000 \quad (\text{Grid}, 1) \cong (\text{Block}, B)$$

1001 The group and split operations of Snap allows us to move along this isomorphism, from left
 1002 to right. For clarity, in Bundl we split these operations into three: a split operation that can split
 1003 privileges into multiple smaller ones, and a destruct operation that directly moves us along the
 1004 isomorphism, and a group that divides our current privilege into equal sized parts.

1005 Privileges π are also lexicographically ordered in the obvious way. hs are ordered such that
 1006 $\text{Thread} \leq \text{Block} \leq \text{Grid}$, and $(h_1, n_1) \leq (h_2, n_2)$ iff $n_1 \leq n_2$ and $h_1 \leq h_2$.

1007 We define scalar multiplication $i \times h$ of natural numbers with hs : $i \times (h, n) = (h, in)$.

1008 We also define division of privileges and hierarchy levels of type $\pi \times \pi \rightarrow \mathbb{N}$. $\text{Grid}/\text{Block} = B$ and
 1009 $\text{Block}/\text{Thread} = T$. We lift this to privileges like so: $(h_1, n_1)/(h_2, n_2) = ((h_1/h_2) \cdot n_1)/n_2$.

1010 Lastly we define a partial \downarrow operator on hs such that $\downarrow \text{Block} = \text{Thread}$ and $\downarrow \text{Grid} = \text{Block}$. Note
 1011 that $\downarrow \text{Thread}$ is undefined. This operator lifts to π s whose second component is 1 and encodes the
 1012 leftward component of the isomorphism above: $\downarrow (\text{Block}, 1) = (\text{Thread}, T)$ and $\downarrow (\text{Grid}, 1) = (\text{Block}, B)$.

1013 Note also that the presentation of these rules in the main body of the paper elide the m portion,
 1014 which tracks the maximum amount of memory a given computation is allowed to use. In the full
 1015 system presented here, both the typing rules and the operational semantics carry an additional piece
 1016 of information tracking allocated memory.

1018 A.2 Typing Rules

1019 A.2.1 Expressions.

$$1020 \quad \frac{x : \pi \tau \in \Gamma}{\Gamma \vdash^\pi x : \tau} \text{-Var} \quad \frac{}{\Gamma \vdash^\pi n : \text{int}} \text{-Int} \quad \frac{}{\Gamma \vdash^\pi f : \text{float}} \text{-Float}$$

$$1023 \quad \frac{}{\Gamma \vdash^\pi b : \text{bool}} \text{-Bool} \quad \frac{\pi < (\text{Grid}, 1)}{\Gamma \vdash^\pi \text{partition_id} : \text{int}} \text{-Partition-Id}$$

$$1027 \quad \frac{\Gamma \vdash^{\pi'} e_1 : \tau[]^l \quad \Gamma \vdash^\pi e_2 : \text{int} \quad l = \text{Global or } l = \text{Local} \quad \pi \leq \pi'}{\Gamma \vdash^\pi e_1[e_2] : \tau} \text{-Arr-Access}$$

| | | |
|------|--|----------------------------|
| 1030 | | |
| 1031 | $\frac{\Gamma \vdash^{\pi'} e_1 : \tau[]^{\text{Shared}} \quad \Gamma \vdash^{\pi} e_2 : \text{int} \quad \pi \leq (\text{Block}, 1) \quad \pi \leq \pi'}{\Gamma \vdash^{\pi} e_1[e_2] : \tau}$ | T-Arr-Access-Shared |
| 1032 | | |
| 1033 | | |
| 1034 | $\frac{\Gamma \vdash^{\pi} e_1 : \text{int} \quad \Gamma \vdash^{\pi} e_2 : \text{int}}{\Gamma \vdash^{\pi} e_1 \text{ bop } e_2 : \text{int}}$ | T-Bop |
| 1035 | $\frac{\Gamma \vdash^{\pi} e_1 : \text{int} \quad \Gamma \vdash^{\pi} e_2 : \text{int}}{\Gamma \vdash^{\pi} e_1 \text{ cmp } e_2 : \text{bool}}$ | T-Cmp |
| 1036 | A.2.2 <i>Statements</i> . | |
| 1037 | $\frac{f : \pi' \text{Fun}(x_i : \tau_i, \pi, m') \in \Gamma \quad \Gamma \vdash^{\pi} e_i : \tau_i \quad m' \leq m}{\Gamma \vdash_m^{\pi} f(e_1, \dots, e_n)}$ | T-Function-Call |
| 1038 | | |
| 1039 | | |
| 1040 | $\frac{\Gamma \vdash_m^{(h,n_1)} s_1 \quad \Gamma \vdash_m^{(h,n_2)} s_2 \quad n_1, n_2 \text{ align to } n}{\Gamma \vdash_m^{(h,n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\}}$ | T-Split |
| 1041 | | |
| 1042 | $\frac{\Gamma \vdash_m^{\downarrow\pi} s}{\Gamma \vdash_m^{\pi} \text{destruct in } s}$ | T-Destruct |
| 1043 | | |
| 1044 | $\frac{\Gamma \vdash_m^{(h,n)} s}{\Gamma \vdash_m^{(h,q \cdot n)} \text{group } q \ s}$ | T-Group |
| 1045 | $\frac{x : \pi' \tau \in \Gamma \quad \Gamma \vdash^{\pi} e : \tau \quad \pi' \leq \pi \quad \pi' \pi}{\Gamma \vdash_m^{\pi} x = e}$ | T-Assn |
| 1046 | | |
| 1047 | $\frac{\Gamma \vdash_m^{\pi} \text{init}_{\psi}}{\Gamma \vdash_m^{\pi} \text{wait}_{\psi}}$ | T-Sync-Init |
| 1048 | | |
| 1049 | $\frac{\Gamma \vdash_m^{\pi} \text{skip}}{\Gamma \vdash_m^{\pi} \text{skip}}$ | T-Skip |
| 1050 | $\frac{n \leq m}{\Gamma \vdash_m^{\pi} \text{free } n}$ | T-Free |
| 1051 | | |
| 1052 | $\frac{\Gamma \vdash^{\pi} e : \tau \quad \Gamma, x : \pi' \tau \vdash_m^{\pi} s \quad \pi' \leq \pi \quad \pi' \pi \quad \tau \text{ not an array type}}{\Gamma \vdash_m^{\pi} x : \tau @ \pi' = e \text{ in } s}$ | T-Decl |
| 1053 | | |
| 1054 | $\frac{\Gamma \vdash^{\pi'} e_1 : \tau[]^l \quad \Gamma \vdash^{\pi} e_2 : \text{int} \quad \Gamma \vdash^{\pi} e_3 : \tau \quad l = \text{Global or } l = \text{Local} \quad \pi' \leq \pi \quad \pi' \pi}{\Gamma \vdash_m^{\pi} e_1[e_2] = e_3}$ | T-Arr-Assn |
| 1055 | | |
| 1056 | | |
| 1057 | | |
| 1058 | $\frac{\Gamma \vdash^{\pi'} e_1 : \tau[]^{\text{Shared}} \quad \Gamma \vdash^{\pi} e_2 : \text{int} \quad \Gamma \vdash^{\pi} e_3 : \tau \quad \pi \leq (\text{Block}, 1) \quad \pi' \leq \pi \quad \pi' \pi}{\Gamma \vdash_m^{\pi} e_1[e_2] = e_3}$ | T-Arr-Assn-Shared |
| 1059 | | |
| 1060 | | |
| 1061 | $\frac{\Gamma \vdash^{\pi} e : \text{bool} \quad \Gamma \vdash_m^{\pi} s_1 \quad \Gamma \vdash_m^{\pi} s_2}{\Gamma \vdash_{\max(m_1, m_2)}^{\pi} \text{if } e \text{ then } s_1 \text{ else } s_2}$ | T-If |
| 1062 | | |
| 1063 | | |
| 1064 | $\frac{\Gamma \vdash^{\pi} e : \text{bool} \quad \Gamma \vdash_m^{\pi} s}{\Gamma \vdash_m^{\pi} \text{while } e \text{ do } s}$ | T-While |
| 1065 | $\frac{\Gamma \vdash_{\max(m_1, m_2)}^{\pi} s_1 ; s_2}{\Gamma \vdash_m^{\pi} s_1 ; s_2}$ | T-Seq |
| 1066 | | |
| 1067 | $\frac{\Gamma, x : \pi' \tau[]^l \vdash_m^{\pi} s \quad l = \text{Global or } l = \text{Local}}{\Gamma \vdash_{m+n \cdot \text{size}(\tau)}^{\pi} x := \text{alloc } l \ \tau \ n \text{ in } s}$ | T-Alloc |
| 1068 | | |
| 1069 | | |
| 1070 | $\frac{\Gamma, x : (\text{Block}, 1) \tau[]^{\text{Shared}} \vdash_m^{\pi} s}{\Gamma \vdash_{m+n \cdot \text{size}(\tau)}^{\pi} x := \text{alloc Shared } \tau \ n \text{ in } s}$ | T-Alloc-Shared |
| 1071 | | |
| 1072 | | |
| 1073 | $\frac{\Gamma, y : (h, n/c) \tau[]^l \vdash_m^{(h, n)} s \quad c n \quad l \neq \text{Local}}{\Gamma, x : (h, n) \tau[]^l \vdash_m^{(h, n)} \text{partition}_{\psi} x \text{ into } y \text{ by } c \text{ in } s}$ | T-Partition |
| 1074 | | |
| 1075 | | |
| 1076 | $\frac{\Gamma, y : (h, n') \tau[]^l \vdash_m^{(h, n')} s \quad n' \leq n \quad l \neq \text{Local}}{\Gamma, x : (h, n) \tau[]^l \vdash_m^{(h, n)} \text{claim}_{\psi} x \text{ into } y \text{ at } n' \text{ in } s}$ | T-Claim |
| 1077 | | |
| 1078 | | |

$$\begin{array}{c}
 \frac{\Gamma, y : \downarrow^\pi \tau[]^l \vdash_m^\pi s \quad l \neq \text{Local}}{\Gamma, x : \pi \tau[]^l \vdash_m^\pi \text{lower}_\psi x \text{ into } y \text{ in } s} \text{ T-Lower} \\
 \\
 \frac{\Gamma, y : (\text{Thread}, 1) \text{ async } \tau[]^l \vdash_m^{(\text{Thread}, 1)} s}{\Gamma, x : (\text{Thread}, 1) \tau[]^l \vdash_m^{(\text{Thread}, 1)} \text{async_partition}_\phi x \text{ into } y \text{ in } s} \text{ T-Async-Partition} \\
 \\
 \frac{}{\Gamma, x : (\text{Thread}, 1) \text{ async } \tau[]^l, y : (\text{Thread}, 1) \tau[]^{l'} \vdash_m^{(\text{Thread}, 1)} \text{async_memcpy}(x, y)} \text{ T-Async-Memcpy} \\
 \\
 \frac{}{\Gamma, x : \pi \tau[]^l, y : \pi \tau[]^{l'} \vdash_m^\pi \text{memcpy}(x, y)} \text{ T-Memcpy}
 \end{array}$$

A.3 Complete Bundl Semantics

A.3.1 Definitions.

Global Memory $\Sigma ::= \cdot | \Sigma, n \mapsto^\pi v$

Shared Memory $\sigma ::= \cdot | \sigma, n \mapsto^\pi v$

Local Memory $\eta ::= \cdot | \eta, n \mapsto^\pi v$

Block Memory Map $S ::= \forall n \in B, n \mapsto \sigma$

Thread Memory Map $L ::= \forall n \in T, n \mapsto \eta$

Synchronization Map $\Psi : \psi \rightarrow \mathbb{N} \rightarrow \mathbb{N}$

Deferred Computations Map $\Phi : \phi \rightarrow \{s\}$

In real GPUs, thread ids are only unique within their block. However, in this calculus for simplicity we assume thread ids are global. One can convert back and forth between this abstracted notion of a thread id and a block-unique id via addition modulo T . That is, $t_{\text{real}} = t_{\text{simplified}} \bmod T$ and $t_{\text{simplified}} = t_{\text{real}} + b \cdot T$.

By convention the names for local and shared and global memory do not conflict, as on the GPU they will be separate pointer spaces. Additionally, we freely interchange between using names for variables and integer locations.

In the main body of the paper, for simplicity we elide the synchronization map and the deferred computation map from the operational semantics, as our theorems do not make any guarantees about non-interference. However, as they are part of the full semantics, we include them here for completeness. By convention the synchronization and deferred computation maps are a total functions, initialized to map to $\lambda_. 0$ for ψ s and $\lambda_. \{s\}$ for ϕ s not explicitly initialize.

The shape of the judgment for a single thread is $\eta, \sigma, \Sigma, t, b, p, \Psi \vdash_m^\pi s \rightsquigarrow s' \dashv_m \eta', \sigma', \Sigma', \Psi'$. The t here is the thread id, the b is the block id, and the p is the partition id. The last of these three is modified and managed by the rules for split, group and destruct, and tracks the relative position of the thread within a group. This semantics is in a small step style.

The shape of the judgment for expressions is $\eta, \sigma, \Sigma \vdash_\pi^\pi e \Downarrow v$. The two π s represent the ambient compute context (i.e., the context in which resources are being read), while π' , represents the target compute context (i.e., the compute context of the variable into which the result of the expression is going to be written). This is relevant for computing the value of `partition_id`, which divides the two contexts. As a shorthand, we can divide a privilege by a scalar value like so: $(h, n)/c = (h, n)/(h, c)$.

1128 The overall evaluation of a program is expressed as

$$L, S, \Sigma, P, \Psi, \Phi \rightsquigarrow L', S', \Sigma', P', \Psi', \Phi'.$$

1131 In this judgment P serves as a *thread pool*, mapping pairs of thread and block ids (which don't change)
 1132 to statements and memory (which can be updated by stepping). One can think of P as tracking which
 1133 program is running on each thread. This steps according to the following rule:

$$\frac{L(t), S(b), \Sigma, t, b, 0, \Psi, \Phi \vdash_m^{(\text{Grid}, 1)} s \rightsquigarrow s' \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi' \quad P(t, b) = (s, m)}{L, S, \Sigma, P, \Psi, \Phi \rightsquigarrow L[t \mapsto \eta'], S[b \mapsto \sigma'], \Sigma', P[(t, b) \mapsto (s', m')], \Psi', \Phi'} \text{ S-Program}$$

1137 For simplicity of notation, we define an **update** operation that searches the three environments
 1138 for the one that contains the variable being used (by convention, there is no conflict between the
 1139 environments, as in reality they exist in three separate address spaces). We also define a similar **get**
 1140 operation that retrieves a variable from memory, and a **rename** operation that remaps a variable
 1141 with the same value but under a different name.

$$\begin{aligned} \text{update}(\eta, \sigma, \Sigma, x, v) &= (\eta[x \mapsto^\pi v], \sigma, \Sigma) \text{ when } x \in^\pi \eta \\ \text{update}(\eta, \sigma, \Sigma, x, v) &= (\eta, \sigma[x \mapsto^\pi v], \Sigma) \text{ when } x \in^\pi \sigma \\ \text{update}(\eta, \sigma, \Sigma, x, v) &= (\eta, \sigma, \Sigma[x \mapsto^\pi v]) \text{ when } x \in^\pi \Sigma \end{aligned}$$

$$\begin{aligned} \text{get}(\eta, \sigma, \Sigma, x) &= \eta(x) \text{ when } x \in^\pi \eta \\ \text{get}(\eta, \sigma, \Sigma, x) &= \sigma(x) \text{ when } x \in^\pi \sigma \\ \text{get}(\eta, \sigma, \Sigma, x) &= \Sigma(x) \text{ when } x \in^\pi \Sigma \end{aligned}$$

$$\begin{aligned} \text{rename}(\eta, \sigma, \Sigma, x, y, \pi') &= (\eta[y \mapsto^\pi \eta(x)], \sigma, \Sigma) \text{ when } x \in^\pi \eta \\ \text{rename}(\eta, \sigma, \Sigma, x, y, \pi') &= (\eta, \sigma[y \mapsto^\pi \sigma(x)], \Sigma) \text{ when } x \in^\pi \sigma \\ \text{rename}(\eta, \sigma, \Sigma, x, y, \pi') &= (\eta, \sigma, \Sigma[y \mapsto^\pi \Sigma(x)]) \text{ when } x \in^\pi \Sigma \end{aligned}$$

1158 A.3.2 Privilege Management Rules.

$$\frac{p < n_1 \quad n_1, n_2 \text{ align to } n \quad \eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n_1)} s_1 \rightsquigarrow s'_1 \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi'}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\} \rightsquigarrow \text{split}(n_1, n_2)\{s'_1\}\{s_2\} \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi'} \text{ S-Split-Left}$$

$$\frac{p < n_1 \quad n_1, n_2 \text{ align to } n}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n)} \text{split}(n_1, n_2)\{\text{skip}\}\{s_2\} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-Split-Left-Done}$$

$$\frac{p \geq n_1 \quad p < n_1 + n_2 \quad n_1, n_2 \text{ align to } n \quad \eta, \sigma, \Sigma, t, b, p - n_1, \Psi, \Phi \vdash_m^{(h, n_2)} s_2 \rightsquigarrow s'_2 \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi'}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\} \rightsquigarrow \text{split}(n_1, n_2)\{s_1\}\{s'_2\} \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi'} \text{ S-Split-Right}$$

$$\frac{p \geq n_1 \quad p < n_1 + n_2 \quad n_1, n_2 \text{ align to } n}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n)} \text{split}(n_1, n_2)\{s_1\}\{\text{skip}\} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-Split-Right-Done}$$

$$\frac{p \geq n_1 + n_2 \quad n_1, n_2 \text{ align to } n}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-Split-None}$$

$$\begin{array}{c}
1177 \\
1178 \quad \frac{\eta, \sigma, \Sigma, t, b, t \text{ mod } T, \Psi, \Phi \vdash_m^{(\text{Thread}, T)} s \rightsquigarrow s' \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi'}{\eta, \sigma, \Sigma, t, b, 0, \Psi, \Phi \vdash_m^{(\text{Block}, 1)} \text{destruct in } s \rightsquigarrow \text{destruct in } s' \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi'} \text{ S-Destruct-Block} \\
1179 \\
1180 \\
1181 \\
1182 \quad \frac{\eta, \sigma, \Sigma, t, b, b \text{ mod } B, \Psi, \Phi \vdash_m^{(\text{Block}, B)} s \rightsquigarrow s' \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi'}{\eta, \sigma, \Sigma, t, b, 0, \Psi, \Phi \vdash_m^{(\text{Grid}, 1)} \text{destruct in } s \rightsquigarrow \text{destruct in } s' \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi'} \text{ S-Destruct-Grid} \\
1183 \\
1184 \\
1185 \quad \frac{}{\eta, \sigma, \Sigma, t, b, 0, \Psi, \Phi \vdash_m^{\pi} \text{destruct in skip} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-Destruct-Done} \\
1186 \\
1187 \\
1188 \quad \frac{\eta, \sigma, \Sigma, t, b, p \text{ mod } n, \Psi, \Phi \vdash_m^{(h, n)} s \rightsquigarrow s' \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi'}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, q \cdot n)} \text{group } q \ s \rightsquigarrow \text{group } q \ s'; \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi'} \text{ S-Group} \\
1189 \\
1190 \\
1191 \quad \frac{}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{group } q \ \text{skip} \rightsquigarrow \text{skip}; \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-Group-Done} \\
1192
\end{array}$$

A.3.3 *Thread Synchronization.* We define a **size** operation on privileges to compute the size of a partition (the number of individual threads contained within it). The operation is defined as follows:

$$\begin{array}{c}
1193 \\
1194 \\
1195 \\
1196 \quad \mathbf{size}(\text{Thread}, n) = n \\
1197 \\
1198 \quad \mathbf{size}(\text{Block}, n) = n \cdot T \\
1199 \\
1200 \\
1201 \\
1202 \quad \frac{\Psi(\psi)(p) = 0}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{wait}_{\psi} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-Sync-Wait-Done} \\
1203 \\
1204 \\
1205 \quad \frac{\Psi(\psi)(p) \neq 0}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{wait}_{\psi} \rightsquigarrow \text{wait}_{\psi} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-Sync-Wait-Spin} \\
1206 \\
1207 \\
1208 \quad \frac{\Psi' = \Psi(\psi)[p \mapsto \Psi(\psi)(p) - 1]}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{dec}_{\psi} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi', \Phi} \text{ S-Sync-Dec} \\
1209 \\
1210 \\
1211 \quad \frac{\Psi(\psi)(p) = 0 \quad \Psi' = \Psi(\psi)[p \mapsto \mathbf{size}(\pi)]}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{init}_{\psi} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi', \Phi} \text{ S-Sync-Init-Zero} \\
1212 \\
1213 \\
1214 \quad \frac{\Psi(\psi)(p) \neq 0}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{init}_{\psi} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-Sync-Init-Nonzero} \\
1215
\end{array}$$

A.3.4 *Asynchrony.*

$$\begin{array}{c}
1216 \\
1217 \quad \frac{\mathbf{rename}(\eta, \sigma, \Sigma, x, y, (\text{Thread}, 1)), t, b, p, \Psi, \Phi \vdash_m^{(\text{Thread}, 1)} s \rightsquigarrow s' \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi'}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(\text{Thread}, 1)} \text{async_partition}_{\phi} x \text{ into } y \text{ in } s \rightsquigarrow \text{async_partition}_{\phi} x \text{ into } y \text{ in } s' \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi'} \text{ S-Async-Partition-Congr} \\
1218 \\
1219 \\
1220 \\
1221 \\
1222 \quad \frac{\Phi = \Phi'[\phi \mapsto \Phi'(\phi) \cup \{s\}]}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(\text{Thread}, 1)} \text{async_partition}_{\phi} x \text{ into } y \text{ in skip} \rightsquigarrow (\text{async_partition}_{\phi} x \text{ into } y \text{ in } s) \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi'} \text{ S-Async-Partition-Unwind} \\
1223 \\
1224 \\
1225
\end{array}$$

| | | |
|------|--|-------------------------------|
| 1226 | | |
| 1227 | $\Phi(\phi) = \emptyset$ | |
| 1228 | $\frac{}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(\text{Thread}, 1)} \text{async_partition}_\phi x \text{ into } y \text{ in } \text{skip} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi}$ | S-Async-Partition-Done |
| 1229 | | |
| 1230 | | |
| 1231 | | |
| 1232 | $\frac{}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(\text{Thread}, 1)} \text{async_memcpy}(x, y) \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi[\phi \mapsto \Phi(\phi) \cup \{\text{memcpy}(x, y)\}]}$ | S-Async-Memcpy |
| 1233 | | |
| 1234 | | |
| 1235 | $\frac{(\eta', \sigma', \Sigma') = \text{update}(\eta, \sigma, \Sigma, \text{get}(\eta, \sigma, \Sigma, y))}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{memcpy}(x, y) \rightsquigarrow \text{skip} \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi}$ | S-Memcpy |
| 1236 | | |
| 1237 | | |
| 1238 | A.3.5 <i>Variables and Memory.</i> | |
| 1239 | | |
| 1240 | $\frac{\eta, \sigma, \Sigma \vdash_{\pi'}^\pi e \Downarrow v}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi x : \tau @ \pi' := e \text{ in } s \rightsquigarrow s \dashv_m \eta[x \mapsto^{\pi'} v], \sigma, \Sigma, \Psi, \Phi}$ | S-Decl |
| 1241 | | |
| 1242 | | |
| 1243 | $\frac{}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{free } n \dashv_{m-n} \eta, \sigma, \Sigma, \Psi, \Phi}$ | S-Free |
| 1244 | | |
| 1245 | | |
| 1246 | $\frac{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi x := \text{alloc Local } \tau n \text{ in } s \rightsquigarrow s; \text{ free } (n \cdot \text{size}(\tau)) \dashv_{m+n \cdot \text{size}(\tau)} \eta[x \mapsto^\pi \langle x, n \rangle], \sigma, \Sigma, \Psi, \Phi}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi x := \text{alloc Shared } \tau n \text{ in } s \rightsquigarrow s; \text{ free } (n \cdot \text{size}(\tau)) \dashv_{m+n \cdot \text{size}(\tau)} \eta[x \mapsto^\pi \langle x, n \rangle], \sigma, \Sigma, \Psi, \Phi}$ | S-Alloc-Local |
| 1247 | | |
| 1248 | | |
| 1249 | $\frac{\pi = (\text{Block}, 1)}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi x := \text{alloc Global } \tau n \text{ in } s \rightsquigarrow s; \text{ free } (n \cdot \text{size}(\tau)) \dashv_{m+n \cdot \text{size}(\tau)} \eta[x \mapsto^\pi \langle x, n \rangle], \sigma, \Sigma, \Psi, \Phi}$ | S-Alloc-Shared |
| 1250 | | |
| 1251 | | |
| 1252 | | |
| 1253 | $\frac{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi x := \text{alloc Global } \tau n \text{ in } s \rightsquigarrow s; \text{ free } (n \cdot \text{size}(\tau)) \dashv_{m+n \cdot \text{size}(\tau)} \eta, \sigma, \Sigma[x \mapsto^\pi \langle x, n \rangle], \Psi, \Phi}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi, \Phi \vdash_m^\pi x = e \rightsquigarrow \text{skip} \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi}$ | S-Alloc-Global |
| 1254 | | |
| 1255 | | |
| 1256 | $\frac{x \in^{\pi'} \eta, \sigma, \Sigma \quad \eta, \sigma, \Sigma \vdash_{\pi'}^\pi e \Downarrow v \quad (\eta', \sigma', \Sigma') = \text{update}(\eta, \sigma, \Sigma, v) \quad \pi' \mid \pi}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi, \Phi \vdash_m^\pi x = e \rightsquigarrow \text{skip} \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi}$ | S-Assn |
| 1257 | | |
| 1258 | | |
| 1259 | | |
| 1260 | $\frac{\eta, \sigma, \Sigma, t, b, p \vdash_\pi^\pi e_1 \Downarrow \langle l, n \rangle \quad i < n \quad \pi' \mid \pi}{\eta, \sigma, \Sigma, t, b, p \vdash_\pi^\pi e_2 \Downarrow i \quad x \in^{\pi'} \eta, \sigma, \Sigma}$ | |
| 1261 | | |
| 1262 | $\frac{\eta, \sigma, \Sigma \vdash_{\pi'}^\pi e_3 \Downarrow v \quad (\eta', \sigma', \Sigma') = \text{update}(\eta, \sigma, \Sigma, l+i, v)}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi e_1[e_2] = e_3 \rightsquigarrow \text{skip} \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi}$ | S-Arr-Assn |
| 1263 | | |
| 1264 | | |
| 1265 | $\frac{s' = s[(y + c \cdot p)/y] \quad (\eta', \sigma', \Sigma') = \text{rename}(\eta, \sigma, \Sigma, x, y, \pi/c)}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{partition}_\psi xy cs \rightsquigarrow \text{init}_\psi; s'; \text{dec}_\psi; \text{wait}_\psi \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi}$ | S-Partition |
| 1266 | | |
| 1267 | | |
| 1268 | $\frac{(\eta', \sigma', \Sigma') = \text{rename}(\eta, \sigma, \Sigma, x, y, (h, n_1))}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n_1 + n_2)} \text{claim}_\psi x \text{ into } y \text{ at } n_1 \text{ in } s \rightsquigarrow \text{init}_\psi; \text{split}(n_1, n_2)\{s'\}\{\text{skip}\}; \text{dec}_\psi; \text{wait}_\psi \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi}$ | S-Claim |
| 1269 | | |
| 1270 | | |
| 1271 | | |
| 1272 | $\frac{(\eta', \sigma', \Sigma') = \text{rename}(\eta, \sigma, \Sigma, x, y, \downarrow \pi)}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{lower}_\psi x \text{ into } y \text{ in } s \rightsquigarrow \text{init}_\psi; s; \text{dec}_\psi; \text{wait}_\psi \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi}$ | S-Lower |
| 1273 | | |
| 1274 | | |

1275 A.3.6 Control Flow.

$$\begin{array}{c}
 \frac{\eta, \sigma, \Sigma \vdash_{\pi}^{\pi} e \Downarrow \text{true}}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{if } e \text{ then } s_1 \text{ else } s_2 \rightsquigarrow s_1 \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-If-True} \\
 \\
 \frac{\eta, \sigma, \Sigma \vdash_{\pi}^{\pi} e \Downarrow \text{false}}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{if } e \text{ then } s_1 \text{ else } s_2 \rightsquigarrow s_2 \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-If-False} \\
 \\
 \frac{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{while } e \text{ do } s \rightsquigarrow \text{if } e \text{ then } (s; \text{while } e \text{ do } s) \text{ else skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} s_1 \rightsquigarrow s'_1 \dashv_{m'}^{\pi'} \eta', \sigma', \Sigma', \Psi', \Phi'} \text{ S-While} \\
 \\
 \frac{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} s_1 \rightsquigarrow s'_1; s_2 \rightsquigarrow s'_2 \dashv_{m'}^{\pi'} \eta', \sigma', \Sigma', \Psi', \Phi'}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{skip}; s_2 \rightsquigarrow s_2 \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-Seq-Done} \\
 \\
 \frac{\Sigma(f) = \{[x_1 : \tau_1, \dots, x_n : \tau_n], s\} \quad \sigma, \Sigma \vdash_{\pi}^{\pi} e_i \Downarrow v_i \quad m' \leq m}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} f(e_1, \dots, e_n) \rightsquigarrow s \dashv_m \eta[x_i \mapsto v_i], \sigma, \Sigma, \Psi, \Phi} \text{ S-Function-Call}
 \end{array}$$

1290 A.3.7 Expressions.

$$\begin{array}{c}
 \frac{\pi < (\text{Grid}, 1)}{\eta, \sigma, \Sigma \vdash_{\pi}^{\pi}, \text{partition_id} \Downarrow \pi / \pi' - 1} \text{ E-Partition-Id} \\
 \\
 \frac{}{\eta, \sigma, \Sigma \vdash_{\pi}^{\pi}, x \Downarrow \text{get}(\eta, \sigma, \Sigma, x)} \text{ E-Var} \\
 \\
 \frac{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, e_1 \Downarrow \langle l, n \rangle \quad \eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, e_2 \Downarrow i \quad i < n \quad \pi' \leq \pi}{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, e_1[e_2] \Downarrow \text{get}(\eta, \sigma, \Sigma, l + i)} \text{ E-Arr-Access} \\
 \\
 \frac{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, \text{Int} \Downarrow n}{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, \text{Int} \Downarrow n} \text{ E-Int} \quad \frac{}{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, b \Downarrow b} \text{ E-Bool} \\
 \\
 \frac{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, e_1 \Downarrow v_1 \quad \eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, e_2 \Downarrow v_2 \quad v = v_1 \text{ bop } v_2}{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, e_1 \text{ bop } e_2 \Downarrow v} \text{ E-Bop} \\
 \\
 \frac{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, e_1 \Downarrow v_1 \quad \eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, e_2 \Downarrow v_2 \quad v = v_1 \text{ cmp } v_2}{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, e_1 \text{ cmp } e_2 \Downarrow v} \text{ E-Cmp}
 \end{array}$$

1309 A.4 Theorems and Proofs

1310 Note that in this section we assume no out of bounds array accesses. In general Snap (and by extension
 1311 Bundl) makes no guarantees about array out of bounds.
 1312

1313 A.4.1 More definitions. As a premise to our type safety theorems, we need to assume we have a
 1314 well-typed environment, written $\Gamma \vdash \eta, \sigma, \Sigma$. We define what this means inductively
 1315

$$\begin{array}{c}
 \frac{}{\eta, \sigma, \Sigma \vdash n : \text{int}} \text{ V-Int} \quad \frac{}{\eta, \sigma, \Sigma \vdash b : \text{bool}} \text{ V-Bool} \quad \frac{}{\eta, \sigma, \Sigma \vdash f : \text{float}} \text{ V-Float} \\
 \\
 \frac{\forall i < n, \eta, \sigma, \Sigma \vdash \text{get}(\eta, \sigma, \Sigma, x + i) : \tau}{\eta, \sigma, \Sigma \vdash \langle x, n \rangle : \tau[]^l} \text{ V-Array} \quad \frac{\Gamma, x_i : \pi \tau_i \vdash_m^{\pi} s \quad \Gamma \vdash \cdot, \cdot, \text{fns } \Sigma}{\eta, \sigma, \Sigma \vdash \{x_i : \tau_i, s\} : \text{Fun}(x_i : \tau_i, \pi, m)} \text{ V-Function} \\
 \\
 \frac{\cdot \vdash \eta, \sigma, \Sigma}{\cdot \vdash \eta, \sigma, \Sigma} \text{ G-Empty} \quad \frac{\eta(x) =^{\pi} v \quad \eta, \sigma, \Sigma \vdash v : \text{int}}{\Gamma, x : \pi \text{ int} \vdash \eta, \sigma, \Sigma} \text{ G-Int}
 \end{array}$$

$$\begin{array}{c}
 \frac{\eta(x) =^\pi v \quad \eta, \sigma, \Sigma \vdash v : \text{bool}}{\Gamma, x :^\pi \text{bool} \vdash \eta, \sigma, \Sigma} \text{ G-Bool} \quad \frac{\eta(x) =^\pi v \quad \eta, \sigma, \Sigma \vdash v : \text{float}}{\Gamma, x :^\pi \text{float} \vdash \eta, \sigma, \Sigma} \text{ G-Float} \\
 \\
 \frac{\eta(x) =^\pi v \quad \eta, \sigma, \Sigma \vdash v : \tau[]}{\Gamma, x :^\pi \tau[]^{\text{Local}} \vdash \eta, \sigma, \Sigma} \text{ G-Local} \quad \frac{\sigma(x) =^\pi v \quad \eta, \sigma, \Sigma \vdash v : \tau[]}{\Gamma, x :^\pi \tau[]^{\text{Shared}} \vdash \eta, \sigma, \Sigma} \text{ G-Shared} \\
 \\
 \frac{\Sigma(x) =^\pi v \quad \eta, \sigma, \Sigma \vdash v : \tau[]}{\Gamma, x :^\pi \tau[]^{\text{Global}} \vdash \eta, \sigma, \Sigma} \text{ G-Global} \quad \frac{\Sigma(x) =^\pi v \quad \eta, \sigma, \Sigma \vdash v : \mathbf{Fun}(\Gamma', \pi, m)}{\Gamma, x :^\pi \mathbf{Fun}(\Gamma', \pi, m) \vdash \eta, \sigma, \Sigma} \text{ G-Function}
 \end{array}$$

We can prove a couple simple lemmas about well-typed environments under operations like **rename**, **update**, and **get**.

LEMMA A.1. (*Well-typed get*) If $\Gamma \vdash \eta, \sigma, \Sigma$ and $x :^\pi \tau \in \Gamma$ then $\eta, \sigma, \Sigma \vdash \mathbf{get}(\eta, \sigma, \Sigma, x) : \tau$.

LEMMA A.2. (*Well-typed rename*) If $\Gamma, x :^\pi \tau \vdash \eta, \sigma, \Sigma$ then $\Gamma, x :^\pi \tau, y :^\pi \tau' \vdash \mathbf{rename}(\eta, \sigma, \Sigma, x, y, \tau')$.

LEMMA A.3. (*Well-typed update*) If $\Gamma \vdash \eta, \sigma, \Sigma$ and $\eta, \sigma, \Sigma \vdash v : \tau$ then $\Gamma, x :^\pi \tau \vdash \mathbf{update}(\eta, \sigma, \Sigma, x, v)$.

We also define a well-formedness precondition on p with respect to π :

$$(h, n) \vdash p ::= p < n$$

We also define well-formedness for the async stack:

$$\Gamma \vdash \Phi ::= \forall \phi, s \in \Phi(\phi), \Gamma \vdash_m^{(\text{Thread}, 1)} s$$

A.5 Proofs

LEMMA A.4. (*Expression Safety*) If $\Gamma \vdash^\pi e : \tau$ and $\Gamma \vdash \eta, \sigma, \Sigma$ and $\pi \vdash p$, then there is some v such that $\eta, \sigma, \Sigma \vdash_{\pi'}^\pi e \Downarrow v$ and $v : \tau$.

PROOF. This proof proceeds by induction on the typing relation for expressions. Despite the fact that this property implies termination, we do not need a logical relation to prove it because the expression language is very simple.

The **T-Int**, **T-Float**, and **T-Bool** cases are trivial, using the rules **E-Int**, **E-Bool** and **E-Float** to compute values. In the case for **T-Partition-Id**, the π premises of the typing rules match the premises of the evaluation rules, so these rules are simple as well.

The cases for **T-Bop** and **T-Cmp** follow directly from the inductive hypotheses, assuming a valid and correctly implemented set of binary operators and comparators.

The only interesting cases are **T-Arr-Access** and **T-Arr-Access-Shared**.

In both cases our inductive hypotheses and inversion give us that $\pi' \leq \pi$, and e_1 evaluates to a $\langle x, n \rangle$, and that all the values between x and $x + n$ in the appropriate environment are typed at τ . We also know that e_2 evaluates to an integer i . We assume that all array accesses are in bounds, so $i < l$, which is sufficient to use the **E-Arr-Access** rule to complete this case, and the proof. \square

LEMMA A.5. (*Expression Determinism*) If $\eta, \sigma, \Sigma, t, b, p \vdash_{\pi'}^\pi e \Downarrow v_1$ and $\eta, \sigma, \Sigma, t, b, p \vdash_{\pi'}^\pi e \Downarrow v_2$ then $v_1 = v_2$.

PROOF. Straightforward by induction on the semantic derivation. \square

LEMMA A.6. (*Expression Well-Typedness*) If $\Gamma \vdash_{\pi'}^\pi e : \tau$ and $\Gamma \vdash \eta, \sigma, \Sigma$ and $\pi \vdash p$, and $\eta, \sigma, \Sigma, t, b, p \vdash_{\pi'}^\pi e \Downarrow v$, then $v : \tau$.

PROOF. By our expression safety lemma our well-typed expression must evaluate to a well-typed value v' . By our determinism lemma v' must be the same as v , so v is well-typed. \square

1373 LEMMA A.7. (*Statement Progress*) If $\Gamma \vdash \eta, \sigma, \Sigma$ and $\Gamma \vdash_m^\pi s$ and $\pi \vdash p$, then either s is skip or there is
 1374 some s' such that $\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi s \rightsquigarrow s' \vdash_m^{\pi'} \eta', \sigma', \Sigma, \Psi', \Phi'$

1375 PROOF. This proceeds by induction on the typing derivation.
 1376

1377 *Privilege Management Rules.*

- 1378 • Case **T-Split**:

1379 In this case we have by our assumption that $\pi \vdash p$ that $p < n$. We also have that n_1, n_2 align to n ,
 1380 so $n_1 + n_2 \leq n$. There are three cases to consider, then: when $p < n_1$, when $p \geq n_1$ and $p < n_1 + n_2$,
 1381 and when $p \geq n_1 + n_2$.

1382 In the first case, we have by our inductive hypothesis that s is either skip or that it can step
 1383 in an (h, n_1) context. In the former case we can use the **S-Split-Left-Done** rule and in the
 1384 latter we can use the **S-Split-Left rule**.

1385 The second case is almost symmetric. The only additional work we have to do is to argue
 1386 that $(h, n_2) \vdash p - n_1$, or equivalently that $p - n_1 < n_2$. This, however, is immediate from our
 1387 assumption that $p < n_1 + n_2$.

1388 In the last case, we just use the **S-Split-None** rule to step to skip.

- 1389 • Case **T-Destruct**

1390 By our inductive hypothesis, we know that s can step at $\downarrow \pi$ if $\downarrow \pi \vdash p$. The \downarrow operation is only
 1391 defined at (Block, 1) or (Grid, 1), so we only need to consider the cases where π is one of
 1392 those.

1393 In the former case p becomes $t \mathbf{mod} T$ while $\downarrow \pi$ is (Thread, T). $t \mathbf{mod} T < T$ for any t so this
 1394 satisfies the requirement that $\pi \vdash p$, which lets us use our inductive hypothesis: s is either
 1395 skip or can step. If it can step, we can use this to satisfy the premise of **S-Destruct-Block** to
 1396 step in this case. If it is skip, then we use the rule **S-Destruct-Done** to step instead.

1397 The latter case is the same, except using the fact that $b \mathbf{mod} B < B$ and the **S-Destruct-Grid**
 1398 rule.

- 1399 • Case **T-Group**

1400 In this case we have by assumption that $(h, q \cdot n) \vdash p$, i.e., that $p < q \cdot n$.

1401 In this case we have our IH that if $(h, n) \vdash p'$ for some p' , then s is either **skip** or steps with
 1402 (h, n) privilege with p' as our partition id.

1403 We choose p' to be $p \mathbf{mod} n$. This is always $< n$, so $(h, n) \vdash p'$. This lets us use our IH to get
 1404 that s is either skip (in which case we can use the **S-Group-Done** rule to step) or itself steps,
 1405 which lets us use the **S-Group** rule to step.

1406 *Thread Synchronization Rules.*

- 1407 • Case **T-Sync-Wait**

1408 $\Psi(\psi)(p)$ is either zero or it is not. In the former case we use the **S-Sync-Wait-Done** rule and
 1409 in the latter we use **S-Sync-Wait-Spin**.

- 1410 • Case **T-Sync-Dec**

1411 We use the **S-Sync-Dec** rule to step.

- 1412 • Case **T-Sync-Init**

1413 We use the **S-Sync-Init-Zero** or **S-Sync-Init-Nonzero** rules depending on whether $\Psi(\psi)(p)$
 1414 is zero or not.

1415 *Asynchrony Rules.*

- 1416 • Case **T-Async-Partition**

1417 In this case, we have via our IH that if $\Gamma, y : (\text{Thread}, 1) \text{ async } \tau[]^l \vdash \eta', \sigma', \Sigma'$, then we can either
 1418 step s with (Thread, 1) privilege under environments η' , σ' , and Σ' , or s is skip.

1422 We have by assumption that $\Gamma, x:(\text{Thread}, 1) \text{ async } \tau[] \vdash \eta, \sigma, \Sigma$. By our environment renaming
 1423 lemma, this gives us what we need to use our IH with $(\eta', \sigma', \Sigma')$ as **rename** $(\eta, \sigma, \Sigma, x, y, (\text{Thread}, 1))$.
 1424 Thus s either steps or is skip. In the former case we can step with **S-Async-Partition-Congr**,
 1425 and in the latter we can use either **S-Async-Partition-Unwind** or **S-Async-Partition-Done**
 1426 depending on whether $\Phi(\phi)$ is empty or not.

- Case **T-Async-Memcpy**

1427 Immediate via use of the **S-Async-Memcpy** rule.

- Case **T-Memcpy**

1428 Immediate via use of the **S-Memcpy** rule.

1429 *Memory Rules.*

- Case **T-Decl**

1430 Via our lemma about expression type safety and our hypothesis that e is well-typed, we obtain
 1431 the premises necessary to use the **S-Decl** rule to step.

- Case **T-Arr-Assn**

1432 Each of e_1 , e_2 and e_3 must evaluate to a well-typed value by the expression type safety lemma.
 1433 In particular, both e_1 evaluates to some $\langle l, n \rangle$ and e_2 evaluates to some i . We assume all array
 1434 accesses are in bounds, so this is sufficient to use the **S-Arr-Assn** rule to step.

- Case **T-Arr-Assn-Shared**

1435 Same as previous case.

- Case **T-Free**

1436 Trivial via the **S-Free** rule.

- Case **T-Partition**

1437 Trivial via the **S-Partition** rule.

- Case **T-Claim**

1438 Trivial via the **S-Claim** rule.

- Case **T-Lower**

1439 Trivial via the **S-Lower** rule.

- Case **T-Alloc**

1440 We assume that l is not Shared, so we can use the **S-Alloc-Local** or **S-Alloc-Global** rule,
 1441 depending on whether l is Local or Global.

- Case **T-Alloc-Shared**

1442 Trivial via the **S-Alloc-Shared** rule.

1443 *Control Rules.*

- Case **T-Skip**

1444 Trivial

- Case **T-While**

1445 Trivial, all while loops step via the **S-While** rule

- Case **T-If**

1446 By our proof of expression type safety, the expression e steps to either the boolean value true
 1447 or false. We can thus use either the **S-If-True** or **S-If-False** rules to step.

- Case **T-Seq**

1448 In this case we know by our IH that s_1 is either skip or can step. In the former case we use
 1449 the **S-Seq-Done** rule and in the latter we use the **S-Seq-First** rule.

- Case **T-Function-Call**

1450 In this case we know by our expression safety lemma that each of the arguments will evaluate to
 1451 a well-typed value. We also have by assumption that f has a function type, which by inversion

on the **V-Function** rule tells us that it is a closure type. Additionally our assumption that $\Gamma \vdash \eta, \sigma, \Sigma$ tell us that Σ contains f at the same type that Γ does. These premises are sufficient to use the **S-Function-Call** rule.

□

1475

LEMMA A.8. (*Statement Preservation*) If $\Gamma \vdash \eta, \sigma, \Sigma$ and $\Gamma \vdash_m^\pi s$ and $\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_{m'}^\pi s \rightsquigarrow s' \dashv_{m''} \eta', \sigma', \Sigma, \Psi', \Phi'$ and $\pi \vdash p$ and $m \geq m'$ and $\Gamma \vdash \Phi$, then there is some Γ' such that $\Gamma \subseteq \Gamma'$ and $\Gamma' \vdash_m^\pi s'$ and $\Gamma' \vdash \eta', \sigma', \Sigma'$ and $m \geq m''$ and $\Gamma' \vdash \Phi'$.

1479

PROOF. We proceed by induction on the derivation of $\Gamma \vdash_m^\pi s$.

1481

Privilege Management Rules.

1482

- Case **T-Split**

In this case we have by assumption that $(h, n) \vdash p$, $\Gamma \vdash \eta, \sigma, \Sigma$, and $\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_{m'}^\pi \text{split}(n_1, n_2)\{s_1\}\{s_2\} \rightsquigarrow s' \dashv_{m''} \eta', \sigma', \Sigma, \Psi', \Phi'$. Our inductive hypotheses give us that if for any p , if $(h, n_1) \vdash p$ and s_1 steps with partition id p , or if $(h, n_2) \vdash p$ and s_2 steps with partition id p then their results are well typed.

By inversion on our semantic derivation, we are in one of 5 cases.

In the **S-Split-Left** case $p < n_1$ and s_1 steps to s'_1 . This is sufficient to tell us that s'_1 is well-typed and the output environments of that relation η', σ' , and Σ' are all well-typed by $\Gamma' \supseteq \Gamma$, and that the memory is properly bounded by the typing rules.

We can thus use the **T-Split-Left** rule to conclude that the result of this case is well-typed.

The **S-Split-Left-Done** rule is trivial via the **T-Skip** rule.

The **Right** cases are symmetric, with the observation that when $p \geq n_1$ and $p < n_1 + n_2$ then $p - n_1 < n_2$.

The last **S-Split-None** rule is trivial via the **T-Skip** rule.

- Case **T-Destruct**

In this case we have by assumption that $\downarrow \pi$ is defined, so π is either **(Block,1)** or **(Grid,1)**. We also assume that $\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_{m'}^\pi \text{destruct}$ in $s \rightsquigarrow s' \dashv_{m''} \eta', \sigma', \Sigma, \Psi', \Phi'$. We also have by our inductive hypothesis that for any p such that $\downarrow \pi \vdash p$ and s'' such that s steps to s'' at p , then that step preserved well-typedness.

By inversion on the step relation, we are in one of three cases.

If the rule used **S-Destruct-Block**, then we know that s steps to s'' and p is $t \bmod T$. $(\text{Thread}, T) \vdash t \bmod T$ for any t , so we can use our inductive hypothesis to conclude that the s'' stepped to by s is well-typed, as are its environments and memory usage. The **T-Destruct** rule then gives us our desired goal.

The **S-Destruct-Grid** case proceeds similarly, while the **S-Destruct-Done** case is trivial.

- Case **T-Group**

In this case we have by assumption that

$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_{m'}^{(h, q \cdot n)} \text{group } q \ s \rightsquigarrow s' \dashv_{m''} \eta', \sigma', \Sigma, \Psi', \Phi'$. We have by our inductive hypothesis that s steps to s'' at some partition id p' and $(h, n) \vdash p'$, then s'' is well typed, as are the other outputs of that step.

By inversion on the step relation, we are in one of two cases. The **S-Group-Done** case is trivial, so we shall focus on the **S-Group** case. In this case we have that s steps to s'' at partition id $p \bmod n$. It is always the case that $(h, n) \vdash p \bmod n$ for any p , so we can use our inductive hypothesis to conclude that s'' is well-typed, as are its output environments and memory usage. From there, it is a simple application of the **T-Group** rule to conclude that $\text{group } q \ s'$ is well-typed, and to finish the case.

1519

Thread Synchronization Rules. These rules are all trivial: with one exception all thread synchronization primitives step to skip without changing environment or memory, and are thus obviously well-typed. **S-Sync-Wait-Spin** does not produce skip, but it steps to the same statement as we already assumed typechecks in the premise of the lemma, so is straightforward nonetheless.

If we wanted to say something about deadlock freedom we'd have more work here, but we aren't doing that, so these rules are easy.

Asynchrony Rules.

- Case **T-Async-Partition** In this case we have that s is well-typed in a context where x has been renamed into y , with $(\text{Thread},1)$ privilege. We also have that $\Gamma, x:(^{\text{Thread},1}) \tau[]^l \vdash \eta, \sigma, \Sigma$ and $\Gamma, x:(^{\text{Thread},1}) \tau[]^l \Phi$.

By inversion, we are in one of three cases.

In the **S-Async-Partition-Done** case, we are done.

In the **S-Async-Partition-Congr** case, our inductive hypothesis gives us that there is some Γ' such that $\Gamma, y:(^{\text{Thread},1}) \text{async } \tau[]^l \subseteq \Gamma'$ and $\Gamma' \vdash \Phi'$ and $\Gamma' \vdash \text{rename}(\eta, \sigma, \Sigma, x, y, (\text{Thread},1))$ via our well-typed renaming lemma. This lets us use the **T-Async-Partition** rule to check this case, with a choice of Γ' as $\Gamma, x:(^{\text{Thread},1}) \tau[]^l$.

In the **S-Async-Partition-Unwind** case, our assumption that Φ is well-typed tells us that $\Gamma, y:(^{\text{Thread},1}) \vdash \{ (\text{Thread},1)_m s$. Thus, we can use the **T-Async-Partition** rule to type this case.

- Case **T-Async-Memcpy**

In this case the statement and environment typing are trivial, we need only to show that the `async` stack remains well typed.

In this case we have that x and y have the same type at $(\text{Thread},1)$. This is sufficient for us to check $x = y$ at $(\text{Thread},1)$, meaning that adding that instruction to the stack maintains its well-typedness.

- Case **T-Memcpy**

Immediate via use of the **S-Memcpy** rule. We just need to show that the environment remains well typed, which we know via our lemmas about **update** and **get**.

Memory Rules.

- Case **T-Decl**

By inversion, we are using **S-Decl** rule for evaluation. Our well-typed expression lemma gives us that v is well-typed, so it follows from our assumptions and our lemmas about extending environments that the extended η and s are well-typed by $\Gamma, x:\pi \tau$.

- Case **T-Free**

Trivial.

- Case **T-Alloc**

By inversion we are either in the **S-Alloc-Local** or **S-Alloc-Global** rules. In either case, we assume that $\Gamma, x:\pi \tau[]^l$ checks s , meaning we can use our extended environment lemmas and the **T-Seq** and **T-Free** rules to check these cases.

- Case **T-Alloc-Shared**

Same as previous case.

- Case **T-Partition**

In this case we have by assumption that $\Gamma, y:(^{(h,n/c)} \tau[]^l s$ and l is not local and c divides n . By inversion on our step relation we must be in the **S-Partition** case, so we have $\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{partition} \psi \text{ycs} \rightsquigarrow \text{init}_\psi; s'; \text{dec}_\psi; \text{wait}_\psi \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi$ where $s' = s[(y + c \cdot p)/y]$ and $(\eta', \sigma', \Sigma') = \text{rename}(\eta, \sigma, \Sigma, x, y, \pi/c)$. Via our well-typed renaming lemma contexts we know that we can check the renamed environments in the extended environment $\Gamma, y:\pi/c \tau[]^l, x:\pi$

1569 $\tau[]^l$, and this context is also sufficient to check s' (via a substitution-preserves-typing lemma
 1570 that is obvious). Using the **T-Skip** rule this is exactly what we need to show to complete this
 1571 case, as the thread sync primitives check trivially via their typing rules.

- 1572 • Case **T-Claim**

1573 Essentially the same as **T-Partition**.

- 1574 • Case **T-Lower**

1575 Essentially the same as **T-Partition**.

1576 *Control Rules.*

- 1577 • Case **T-If**

1578 We have by assumption that **if** e then s_1 else s_2 steps to some s' , and by inversion we
 1579 know that either e evaluates to **true** and s' is s_1 , or e evaluates to **false** and s' is s_2 .

1580 In either case, our inductive hypotheses is sufficient to tell us that these are well-typed. In
 1581 particular, in both cases our IHs tell us that the amount of memory used by stepping each
 1582 branch of the **if** is less than the amount of memory computed by the type system for each
 1583 branch. Because the whole **if** expression checks using the greater of the memory usage of m_1
 1584 or m_2 (i.e., the memory usage on each branch), the resulting usage for the whole conditional
 1585 is also bounded by the type system.

- 1586 • Case **T-Skip**

1587 Trivial: **skip** does not step

- 1588 • Case **T-Seq**

1589 In this case we have that s_1 and s_2 are both well-typed (with m_1 memory and m_2 memory
 1590 respectively), and $m = \mathbf{max}(m_1, m_2)$. We also have by inversion that s_1 either steps to **skip** or
 1591 s'_1 , and our inductive hypothesis tells us that s'_1 is well-typed.

1592 In the former case we can use the **S-Seq-Done** rule to trivially finish the case. In the latter,
 1593 our IH allows us to finish the case, since m_1 is always $\leq \mathbf{max}(m_1, m_2)$

- 1594 • Case **T-While**

1595 By inversion, we have that

$$\begin{aligned} 1597 \quad & \eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{while } e \text{ do } s \\ 1598 \quad & \rightsquigarrow \text{if } e \text{ then } (s; \text{while } e \text{ do } s) \text{ else skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi. \end{aligned}$$

1600 We also have by assumption that e and s are well-typed. With this information, through use
 1601 of the **T-If**, **T-Seq**, **T-While**, and **T-Skip** rules, we can conclude that the result of this rule is
 1602 also well-typed.

- 1603 • Case **T-Function-Call**

1604 By inversion we have that

$$\begin{aligned} 1605 \quad & \eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi f(e_1, \dots, e_n) \\ 1606 \quad & \rightsquigarrow \text{call } s \text{ with } (x_i : \pi_i \tau_i \mapsto v_i) @ \{\eta, \sigma, \Sigma, m'\} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi, \end{aligned}$$

1607 and also that $\Sigma(f) = \{[x_1 : \tau_1, \dots, x_n : \tau_n], s\}$, and $\eta, \sigma, \Sigma, t, b, p \vdash^\pi e_i \Downarrow v_i$, and $m' \leq m$.

1608 We also have from the premises of our case that $f : \pi \mathbf{Fun}(x_1 : \pi_1 \tau_1, \dots, x_n : \pi_n \tau_n, \pi, m') \in \Gamma$, and
 1609 $\Gamma \vdash^\pi e_i : \tau_i$, and $m' \leq m$.

1610 Our lemma for expression well-typedness tells us that that each v_i is a well-typed value,
 1611 and our assumption that $\Gamma \vdash \eta, \sigma, \Sigma$ tells us that $\Sigma(f)$ is a well-typed function and thus that
 1612 $\mathbf{fns} \Gamma, x_i : \pi_i \tau_i \vdash_m^\pi s$. We can take the union of this with Γ to produce $\Gamma, x_i : \pi_i \tau_i \vdash_m^\pi s$, which
 1613 clearly checks s and the output environments.

1614 □