

# 1 Modular GPU Programming with Typed Perspectives

2 ANONYMOUS AUTHOR(S)

3 Abstract

## 4 1 Introduction

5 CUDA [22] is a low-level, imperative programming language for NVIDIA GPUs. These GPUs are  
6 organized into a hierarchy of compute resources. *Threads* are the basic unit of sequential execution;  
7 *blocks* are groups of threads that can cooperate through shared scratchpad memory; and the *grid* is  
8 the full collection of blocks launched for a GPU kernel. A CUDA program executes in parallel across  
9 this hierarchy, but is written from the *perspective of a single thread*.

10 While operations are specified per-thread, some operations are only semantically valid when ex-  
11 ecuted *collectively* by a group of threads. The `__syncthreads()` primitive, for instance, synchronizes  
12 all threads within a block, and will cause a deadlock if executed by only some of the threads within that  
13 block. There are many such collective operations, including tensor core instructions, warp shuffle  
14 operations [10], and other kinds of synchronization primitives. These operations, therefore, require  
15 programmers to carefully marshal compute resources to coordinate which threads execute which lines  
16 of code. As a result, such collective operations *break the illusion* of threads executing independently.

17 The conceptual clash between regular statements (executed by a single thread), and collective  
18 operations (executed by a group of threads) hinders compositional reasoning. In CUDA, function  
19 calls, like all other statements, are per-thread operations. However, a function may contain code that  
20 executes a collective operation, creating a gap between the per-thread syntax for its invocation and  
21 the cooperative semantics of its execution.

22 This gap creates tension between abstraction and correctness, a tension apparent even in widely  
23 used CUDA libraries that package common functionality via function interfaces. Consider the  
24 following snippet of documentation taken directly from CUB [21], a library of parallel primitives. The  
25 documentation describes the `BlockReduce` function [6], in which all threads in a block collaboratively  
26 apply a reduction operation, such as a maximum or prefix sum, over an array:

27 Computes a block-wide reduction for `thread0` using the specified binary reduction functor. The first `num_valid` threads  
28 each contribute one input element.

- 29 • The return value is undefined in threads other than `thread0`.
- 30 • A subsequent `__syncthreads()` threadblock barrier should be invoked after calling this method if the collective's  
31 temporary storage (e.g., `temp_storage`) is to be reused or repurposed.

32 Through informal documentation, CUB attempts to convey several assumptions implicit in  
33 `BlockReduce`'s implementation. First, the function is only well-defined when invoked by all threads  
34 in a block. Second, because it accesses memory shared among threads, any subsequent reuse of that  
35 storage requires synchronization to ensure all threads have completed their accesses.

36 In effect, CUB attempts to retrofit CUDA with information about the compute and memory re-  
37 quirements of collective operations. In fact, by prefixing functions with identifiers such as `Block`,  
38 CUB creates ad-hoc "namespaces" for different functions that assume similar invariants. However,  
39 without a type system to statically enforce these requirements in CUDA, correct usage of this function  
40 depends on the user carefully reading and interpreting the documentation.

41 In this work, we ask the following question: can we provide low-level access to collective operations  
42 while statically guaranteeing that they execute with the necessary configuration of compute re-  
43 sources? By reifying these configurations as type-level *perspectives*—so named because they describe  
44 the view of GPU resources against which a program is written—we find that we can. Our key insight  
45

is that GPU programmers naturally map computations onto different compute resources, and unlike CUDA, which obscures this mapping, we can expose it in a type system that tracks perspectives. This tracking enforces that collective operations are executed with the necessary resources, while still allowing users to access low-level interfaces. Further, it allows users to define and compose their own collective operations by specifying the perspective with which their code must be run.

Our approach is a departure from previous work, which attempts to resolve the mismatch between per-thread syntax and collective execution by restricting access to low-level operations. Triton [32], a tile-based, multi-dimensional array language, limits the user’s perspective to the block level, which prevents programmers from writing the highest-performance kernels. A variety of functional languages, such as Futhark [18], provide compile-time guarantees through their type systems but are machine-agnostic: they cannot expose low-level control over hardware. Some other efforts, like Descend [20], aim to provide a low-level, safe GPU systems programming language, but lack support for collective operations like tensor cores. As a result, CUDA remains the de-facto standard for writing high-performance kernels on modern GPUs.

We introduce Parallax, a new low-level GPU programming language that guarantees correct usage of hardware resources by construction. Inspired by dependency calculi [1], Parallax statically tracks the configurations of compute and memory resources with type-level perspectives. We also develop Prism, a core calculus underpinning Parallax, provide formal rules for manipulating the perspectives of both code and data, and prove *type-and-perspective safety*. This ensures that Prism is sound, and that code always has the right perspective to execute operations at run time.

A parallel goal, alongside to correctness, is performance. We incorporate modern GPU features such as tensor cores and asynchronous data movement into Parallax, and demonstrate that Parallax can achieve the same performance as hand-written, highly optimized code on an H100 and a 4070 Ti Super. Our contributions are:

- (1) Parallax, a low-level GPU programming language with type-level perspectives that track the logical grouping of compute and memory resources (Section 3);
- (2) Prism, a formal model of Parallax that tracks perspectives in both its type system and operational semantics, along with a soundness theorem guaranteeing that programs execute operations only when they have been statically proven to possess the necessary perspective.
- (3) An implementation of Parallax (Section 5) that demonstrates that it can support modern GPU features like tensor cores and asynchronous data movement, achieving performance comparable to hand-optimized CUDA implementations (Section 6).

Section 7 discusses related work, and Section 8 discusses the limitations of our approach and outlines future work.

## 2 Background & Motivation

Before diving into the design of Parallax, we begin with an overview of the GPU’s compute and memory hierarchies and outline the challenges posed by reasoning about them collectively.

### 2.1 Compute Hierarchy

In CUDA, programmers launch computations that run on thousands of threads. These threads are organized into a *compute hierarchy* that defines how work is distributed and scheduled on the GPU. At the top of this hierarchy is the *grid*, representing all threads launched as part of a single computation. The grid is divided into *blocks*, each containing a user-specified number of threads. A *thread*, in turn, is the finest unit of execution, and it is the machine’s basic unit of sequential control.

Users define the behavior of their computation by writing a single program which is replicated across all threads. This program controls the behavior of individual threads by reading built-in

```

99   1 int tid = threadIdx.x;
100  2 if (tid >= 0 && tid < 31){
101  3     float A[4];
102  4     float B[2];
103  5     float C[4] = { 0 };
104  6     # Populate A with unique values
105  7     for (int i = 0; i < 4; i++)
106  8         A[i] = tid * 4 + i;
107  9     # Populate B with unique values
108 10    for (int i = 0; i < 2; i++)
109 11        B[i] = tid * 4 + i;
110 12    # Issue a warp-level tensor-core
111 13    # operation: D = A * B + C
112 14
113 15    ← asm("mma.sync.aligned.m16n8k8.."
114 16        " {X0, X1, %2, %3},    /*D*/ "
115 17        " {X4, X5, %6, %7},    /*A*/ "
116 18        " {X8, X9},           /*B*/ "
117 19        " {X10, X11, %12, %13}; /*C*/ "
118 20        : "=r"(C[0]), "r"(C[1]), ...
119 21        : "r"(A[0]), "r"(A[1]), ...
120 22        : "r"(B[0]), "r"(B[1]),
121 23        : "r"(C[0]), "r"(C[1]), ...);
122

```

Fig. 1. Invoking a warp-level tensor-core instruction in CUDA. We elide some typecasts necessary to invoke the assembly operation.

```

1 tid : int @thread[1] = id();
2 with group(thread[32]):
3     A : float[4] @thread[1]
4     B : float[2] @thread[1]
5     C : float[4] @thread[1]
6
7     # Populate A with unique values
8     for i in range(0, 4, 1):
9         A[i] = tid * 4 + 1
10    C[i] = 0
11
12     # Populate B with unique values
13     for i in range(0, 2, 1):
14         B[i] = tid * 4 + 1
15
16     # Issue tensor-core op.
17     intrinsic.mma(
18         A[0], A[1], A[2], A[3],
19         B[0], B[1],
20         C[0], C[1], C[2], C[3],
21         out=[A[0], A[1], A[2],
22              → A[3]])

```

Fig. 2. Invoking a warp-level tensor-core instruction in Parallax.

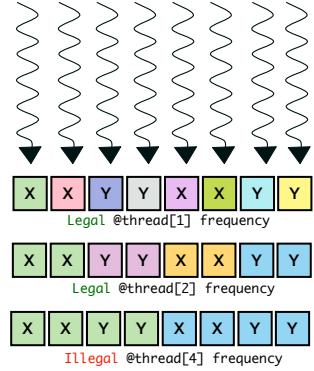


Fig. 3. The @ syntax represents the frequency at which a variable varies across threads T1–T32. Each @ denotes a coloring of the variable; threads that “view” the variable with the same color must observe the same value.

identifiers like `threadIdx.x`—to determine a thread’s position within a block—and `blockIdx.x`—to determine the block’s position within the grid—at runtime. By making control flow decisions based on these two variables, users can assign each thread to its share of the full computation.

A natural and tempting way to interpret these built-in variables is to think of them as the indices of implicit “parallel `for`-loops” surrounding the program, where each iteration is executing simultaneously. While this view is sufficient to understand CUDA’s programming model when threads do independent work, it quickly breaks down in the presence of collective operations.

Unlike most instructions, which are executed by a single thread, collective operations must be executed collaboratively by a group of threads. For example, on line 14 of Figure 1 we perform a “warp-level” tensor core operation [?], which is only *meaningful* when invoked by a collection of 32 threads—a *warp*—acting together. For the call, each thread sets up its portion of the operands, and the operation is performed *once* for the entire warp, with the results scattered across participating threads. The first 32 threads in a block are tasked with this operation; programmers mentally group these 32 threads into a collective unit, and then reflect that grouping in the program via the `if`-statement on line 2. If the condition on line 2 were instead `tid >= 0 && tid < 30`, making fewer than 32 threads reach line 14, the instruction would be undefined. To make matters worse, the restriction is not only on the number of threads executing the operation, but also on their alignment. In this case, the starting ID of the group of 32 participating threads must be aligned to a multiple of 32. So, if the condition on line 2 were instead `tid >= 1 && tid < 32`, the instruction would still be invalid, even though 32 threads would execute it.

Collective operations make reasoning about CUDA programs challenging because they force programmers to track the *convergence behavior* of threads. Specifically, programmers must reason both about how many threads reach a particular point in the program and how those threads are arranged, accounting for alignment. This difficulty is further amplified due to two reasons. First, programs may have multiple *points* of convergence, requiring programmers to mentally track the relative ID of a thread within a logical group as that group evolves over the course of a program’s execution. Second, threads may be participating in multiple *levels* of convergence within the same program. In

our example, we only considered the warp-level tensor core operation, but there are other collective operations that require convergence at different granularities like the *warp-group*-level tensor core operations, which must be issued collectively by four warps, or the block-level `__syncthreads()` synchronization primitive, which must be executed by all threads within a block.

Reasoning about collective operations is already error-prone within the context of a single function, but becomes even more difficult when reasoning interprocedurally *across* functions. A callee may assume a certain number of threads or blocks and may structure its computation, including convergence behavior, around that assumption. However, such assumptions are not visible in the callee’s function interface, which only exposes the input and output types. To invoke that function correctly, users must read documentation, or worse, read the callee’s implementation to understand its assumptions, breaking modular reasoning.

In Parallax, we materialize the programmer’s mental grouping of the compute hierarchy explicitly in the program’s source. Consider the example in Figure 2 which shows an equivalent rewrite of the CUDA program in Figure 1 using Parallax. In Figure 2, the call to the `mma` intrinsic is valid only because it is executed within a `group` of 32 threads, as made explicit on line 2. Since Parallax’s `group` construct already enforces both the size and alignment of participating threads, the validity of the `mma` operation is guaranteed by construction.

## 2.2 Data and the Memory Hierarchy

Data on the GPU is organized into a *memory hierarchy*, mirroring the compute hierarchy. All threads in a grid can read and write from *global memory*, where operands typically reside at the start of a computation and where results are eventually written. Each block has access to a limited amount of *shared memory*, a programmer-managed scratchpad typically used to stage repeatedly-used data. Finally, each thread maintains its own state in *registers* and *local memory*, which is used for storing the thread’s stack. Crucially, registers and local memory are per-thread data, while shared and global memory are accessible by multiple threads.

*Thread-Local Data.* Because registers and local memory are owned by individual threads, a variable with the same name in a program can have different values on different threads at run time. In CUDA, threads can use this to induce divergent behavior. We’ve already seen an example of this: the `tid` variable in Figure 1. Branching on this variable is an instance of data-dependent control flow; in the general case, the possibility of divergence lurks anywhere that execution branches on data.

A key challenge for Parallax is ensuring that all threads organized by other language constructs like `group` remain logically unified, even in the face of data-dependent control flow, which is necessary to compute anything interesting. To do this, Parallax must track the frequency with which values vary in space. The `@` syntax attached to each variable declaration denotes the rate at which a value changes across compute resources on the machine. For example, a `thread[1]` value can vary for every thread, while a `block[1]` can vary for every block, but not for threads within that block. Intuitively, the `@` construct “colors” a variable across threads, and any threads that share a color are required to agree on the variable’s value. Figure 3 illustrates valid and invalid colorings of a variable.

Using this information, Parallax enforces rules for reading from and writing to data. So, for example, had we attempted to introduce a condition based on `tid` *within* the `group`’s scope in Figure 2, Parallax would reject that program at compile time; code that diverges with low frequency cannot read from

```

1 int x = threadIdx.x;
2 if (x >= 0 && x < 31){
3     float A[4];
4     float B[2];
5     float C[4] = { 0 };
6
7     // k is A's stride
8     A[0] = a_mem[(x/4)*k+(x%4)];
9     A[1] = a_mem[(x/4)*k+(x%4)+4];
10    A[2] = a_mem[(x/4+8)*k+(x%4)];
11    A[3] = a_mem[(x/4+8)*k+(x%4)+4];
12
13    // n is B's stride
14    B[0] = b_mem[(x/4)*n+(x/4)];
15    B[1] = b_mem[(x/4+4)*n+(x/4)];
16
17    asm("mma.sync...");
18
19    // Write back into c_mem
20    // n is C's stride
21    c_mem[(x/4)*n+2*(x%4)] = C[0];
22    c_mem[(x/4)*n+2*(x%4)+1] = C[1];
23    c_mem[(x/4+8)*n+2*(x%4)] = C[2];
24
25    → c_mem[(x/4+8)*n+2*(x%4)+1]=C[3];

```

Fig. 4. Invoking a warp-level tensor core instruction in CUDA after loading data from memory.

variables that change at a higher frequency. We will explain the syntax and rules of legal Parallax programs in greater detail in Section 3.

*Thread-Shared Data.* Shared and global memory, on the other hand, is not replicated by every thread; instead, all threads in a block see the same shared memory, while all threads on the grid see the same global memory.

CUDA does not explicitly model shared and global memory spaces, nor does it track whether allocations in a given memory space exceed device limits. Parallax, on the other hand, distinguishes these spaces and restricts shared memory to static allocations, throwing an error if an allocation exceeds device limits. We will discuss these aspects in detail in ???. For now, however, we focus on how memory converges and diverges in tandem with the compute hierarchy.

Because threads share their view of global and shared memory, these spaces are subject to convergent and divergent behavior the same way that threads are. For example, when launching a kernel on the GPU, a global memory pointer may initially belong to the grid because every thread sees the same pointer value at the start of the computation. To write to that memory, each thread computes an offset from the original base address. In doing so, the pointer effectively *diverges* across threads within the larger memory region so each thread can write independently. After these writes, the memory must *reconverge*, ensuring that all threads have completed their updates before the memory can safely be returned to its original logical owner.

Let us reconsider the tensor core example in Figure 1, this time initializing the operands of the tensor core operation from pointers into memory. In Figure 4, A and B are now populated from data in global memory, **a\_mem** and **b\_mem**. After the tensor core operation completes, the result is written back to **c\_mem**, also in global memory. This variable **c\_mem** is logically accessed from two different levels of the compute hierarchy. When entering this code, each thread in a group of 32 threads sees the same value for **c\_mem**. Then, each thread within these 32 locates an offset within **c\_mem** that it can write to (lines 21–24). To restore **c\_mem** back to its 32 thread-level ownership, all 32 threads must first synchronize to ensure all per-thread writes have completed. This is similar to the requirement we saw documented in CUB in Section 1.

Similarly to the compute hierarchy (Section 2.1), CUDA programmers are responsible for tracking the logical owner of a piece of memory as it evolves over the course of a program, and for ensuring that appropriate synchronization occurs whenever that ownership changes.

Parallax, by contrast, makes the evolution of a memory object’s logical ownership explicit in the type system and *automatically* inserts the synchronization required to restore memory to its original ownership. In Figure 5, we show how memory is lowered through the compute hierarchy for the same tensor core operation introduced in Figure 4. The lowering of **c\_mem** is required in this program as Parallax only permits writes when data is owned by a single thread. To lower **c\_mem**, we use Parallax’s **partition** operation on line 31. The operation takes the name of the memory to partition, **c\_mem**, and lowers it to a single thread, and assigns it a new name, **c\_thrd**; Parallax will not allow references to the old name **c\_mem** within the partition’s scope. The **partition** function also takes an indexing function, and each time **c\_thrd** is

```

1  @cuda("device")
2  @requires(thread[32])
3  def simple_mma(
4      a: ptr(const float) @ thread[32],
5      b: ptr(const float) @ thread[32],
6      c: ptr(float) @ thread[32]):
7      x : int @ thread[1] = id()
8      with group(thread[32]):
9          A : float[4] @ thread[1]
10         B : float[2] @ thread[1]
11         C : float[4] @ thread[1]
12
13         # Reads do not need to be lowered.
14         A[0] = a_mem[(x/4)*k+(x%4)];
15         A[1] = a_mem[(x/4)*k+(x%4)+1];
16         A[2] = a_mem[(x/4+8)*k+(x%4)];
17         A[3] = a_mem[(x/4+8)*k+(x%4)+1];
18         B[0] = b_mem[(x%4)*n+(x/4)];
19         B[1] = b_mem[(x%4+4)*n+(x/4)];
20
21         # Skipping initialize C to 0...
22         intrinsic.mma(
23             A[0], A[1], A[2], A[3],
24             B[0], B[1],
25             C[0], C[1], C[2], C[3],
26             out=[C[0], C[1], C[2], C[3]])
27
28         # Must be at thread[1] to write.
29         idx : int @ thread[1] =
30             lambda ro, co: (x/4*ro)*n+2*(x%4)+co
31             with part-
32                 tition(c_mem, p=thread[32], f=idx) as
33                 c_thrd:
34                     c_thrd[0, 0] = C[0]
35                     c_thrd[0, 1] = C[1]
36                     c_thrd[0, 2] = C[2]
37                     c_thrd[0, 3] = C[3]
38
39             # <-- Sync point inferred by name

```

Fig. 5. Invoking a warp-level tensor core instruction in Parallax after loading data from memory.

```

246 1 @ccuda("global")
247 2 # Top-level perspective bounds and shared memory usage.
248 3 @requires(grid[1], block[1], thread[32], smem=1280)
249 4 def mmaTF32NaiveKernel(A: ptr(const(float)) @ grid[1],
250 5     B: ptr(const(float)) @ grid[1],
251 6     C: ptr(float) @ grid[1],
252 7     M : int @ grid[1],
253 8     N : int @ grid[1],
254 9     K : int @ grid[1]):
255 10
256 11     # Starts out with grid[1] perspective.
257 12     with group(grid[1]):
258 13         # @grid[1] inferred from current perspective
259 14         # Each block computes an 16x8 tile
260 15         num_blocks_n : const(int) = (N + 8 - 1) / 8
261 16
262 17         # id() function returns the block id
263 18         # inferred from @block[1].
264 19         blk_row : const(int) @ block[1] = (id() / num_blocks_n) * 16
265 20             ↪ |label[line]{blkRow}|
266 21         blk_col : const(int) @ block[1] = (id() % num_blocks_n) * 8
267 22
268 23         # Give each block an offset into C
269 24         offset = lambda x: blk_row * N + blk_col + x
270 25         with partition(C, p=block[1], f=offset) as C_blk:
271 26             with group(block[1]):
272 27                 # SMEM declarations are only allowed
273 28                 # with a block[1] perspective.
274 29                 A_smem : shared(float[16 * 8]) @ block[1]
275 30                 B_smem : shared(float[8 * 8]) @ block[1]
276 31                 C_smem : shared(float[16 * 8]) @ block[1]
277 32             # Now, id() returns the thread id
278 33             idx : int @ thread[1] = id() * 4
279 34             # To write to C_smem, drop to thread[1] perspective
280 35             with partition(C_smem, p=thread[1], offset=idx) as C_th:
281 36                 for i in range(0, 4, 1):
282 37                     with group(thread[1]):
283 38                         C_th[i] = 0
284 39
285 40
286 41     for i in range(0, K_tiles, 1):
287 42
288 43         # SYNC POINT (back edge from for loop)
289 44         for j in range(0, 4, 1):
290 45             global_row : int @ thread[1] = warp_row + row
291 46             global_col: int @ thread[1] = i * 8 + col
292 47             with partition(A_smem, p=thread[1], f=...) as A_th:
293 48                 with group(thread[1]):
294 49                     A_thrd[0] = A[global_row * K + global_col]
295 50
296 51         # SYNC POINT (back edge from for loop)
297 52         for j in range(0, 2, 1):
298 53             # Similar to write into A_smem ...
299 54             with partition(B_smem, p=thread[1], f=...) as B_th:
300 55                 with group(thread[1]):
301 56                     B_th[0] = B[global_row_b * N + global_col_b]
302 57
303 58         # Give each warp an offset into C_smem
304 59         # SYNC POINT (back edge from for loop)
305 60         with claim(C_smem, p=thread[32]) as Cs_warp:
306 61             match split(thread):
307 62                 case 32:
308 63                     # Call tensor core function inside
309 64                     # a function interface.
310 65                     simple_mma(A_smem, B_smem, Cs_warp)
311 66
312 67         # Write back final result to C
313 68         for j in range(0, 4, 1):
314 69             flat_idx_c : int @ thread[1] = id() * 4 + j
315 70             row_c : int @ thread[1] = flat_idx_c / MMA_K
316 71             col_c : int @ thread[1] = flat_idx_c % MMA_K
317 72             idx = lambda x: row_c * N + col_c + x
318 73             with partition(C_smem, p=thread[1], f=idx) as C_th:
319 74                 with group(thread[1]):
320 75                     C_th[0] = C_smem[row_c * MMA_N + col_c]
321 76
322 77     return

```

Fig. 6. Tensor Float 32 matrix Multiplication in Parallax.

accessed, this indexing function is implicitly applied. Once the `partition`'s scope ends, Parallax will insert a synchronization before the first use of the original name, so that all per-thread writes have completed.

We would like to emphasize that preventing data races is not one of Parallax's goals. In Parallax, data races *can* occur; however, since the data is eventually synchronized before it is reused, the last-writer wins. Out-of-bounds accesses are considered undefined behavior. We believe that prior work, such as Descend [20], lays the foundation for reasoning about data-race free programs, and that a Descend-like memory model can be adapted to Parallax. Parallax is instead focused on the *interaction* between the compute and memory hierarchies, and on reasoning about them simultaneously to ensure that an operation is executed only when sufficient compute resources are available, a guarantee that prior work like Descend cannot provide.

### 3 The Parallax Language

Parallax is an imperative, low-level programming language designed at a level of abstraction comparable to CUDA. Unlike CUDA, Parallax's syntax materializes the mapping of a computation onto the compute and memory hierarchy explicitly in the program source. Using this information, Parallax enforces that programs only execute collective operations with sufficient resources.

To guide our discussion about Parallax's design, we use the program in Figure 6 as a running example. It computes a matrix multiplication between two float arrays, **A** and **B**, to produce an output matrix **C**. In this program, each block computes a  $16 \times 8$  independent tile of the output. To do so, the it first locates the tile index assigned to it (line 19-20). Next, threads in each block load corresponding rows and columns from **A** and **B** (line 41-53) into shared memory. Finally, the program invokes a warp-level, tensor-core instruction to compute the output on (line 62), requiring threads in a warp to

295     converge at the point. We encapsulate this tensor-core instruction in a function, demonstrating how  
 296     function composition works in Parallax. The function is the same as Figure 5.

297

### 298     3.1 Levels

299     Parallax models the machine’s compute hierarchy through *levels*. There are three levels in Parallax—  
 300     **grid**, **block** and **thread**—which are organized as expected: a **grid** consists of multiple units of the  
 301     **block** level, which in turn consist of multiple units of the **thread** level. Levels have an order defined  
 302     over them, with **thread** < **block** < **grid**.

303     There are two key differences between Parallax’s levels and those in CUDA. First, Parallax does  
 304     not model CUDA’s three-dimensional grid or block structure. Second, there are two commonly-used  
 305     “levels”, *warp* and *warpgroup*, noticeably absent from our hierarchy. On the hardware, the units of each  
 306     level are arranged in a single linear order, and the three-dimensional structures of CUDA are simply  
 307     *interpretations* of this ordering, not distinct hardware resources. Similarly, warps and warpgroups are  
 308     organizational constructs defined in terms of existing levels. Namely, a warp is a group of 32 threads  
 309     whose first thread ID is aligned to 32. A warpgroup, which only became a meaningful construct with  
 310     the release of the Hopper architecture, consists of 128 threads aligned to 128.

311     Rather than baking these interpretations into Parallax by adding new dimensions and levels, we  
 312     let users express multi-dimensional structures and define groupings of custom sizes within Parallax.

313

### 314     3.2 Perspectives

315     Perspectives are the central concept of Parallax. A perspective represents the layer of the hierarchy  
 316     from which the programmer is defining the machine’s behavior. They are what allow Parallax to  
 317     determine which compute resources the programmer is controlling, whether they are available in  
 318     the program’s context, and, once provided, if those resources are sufficient for a given operation.

319     A *perspective* is a level—**grid**, **block**, or **thread**—paired with a static constant **n**, specifying the  
 320     number of units at that level. A perspective is written as **level** [**n**]. For example, **thread** [2] denotes  
 321     a perspective of two threads, **block** [4] denotes a perspective of four blocks, and so on. Perspectives  
 322     also carry *alignment information*: a perspective of size **n** is aligned to **n**. In this way, a warp is simply  
 323     a desugaring of **thread** [32], and a warpgroup is a desugaring of **thread** [128].

324     Finally, perspectives have a partial order defined on them. We say that **level** <sub>2</sub> [**n**<sub>2</sub>] is *wider* than  
 325     **level** <sub>1</sub> [**n**<sub>1</sub>], or that **level** <sub>1</sub> [**n**<sub>1</sub>] is *narrower* than **level** <sub>2</sub> [**n**<sub>2</sub>], if either:

326

- (1) **level** <sub>1</sub> < **level** <sub>2</sub>; or,
- (2) **level** <sub>1</sub> = **level** <sub>2</sub> and **n**<sub>2</sub> % **n**<sub>1</sub> = 0.

327     The condition **n**<sub>2</sub> % **n**<sub>1</sub> = 0 enforces alignment, and we will see why it is required in Section 3.4.

328

### 329     3.3 Perspectives on Code

330     Code is associated with a set of perspectives, called a *perspective bound*. Perspective bounds directly  
 331     correspond to the compute resources a program expects to exercise control over the course of its  
 332     execution. At every point in the program, the perspective bound for that point indicates which  
 333     layer of the hierarchy is being programmed, and how that layer can be destructed into narrower  
 334     perspectives. For example, a line of code with a {**block** [1], **thread** [4]} perspective bound tells  
 335     Parallax that the current line of code is being programmed at the **block** [1] perspective and that  
 336     the **block** [1] perspective can be destructed into a some number of **thread** [4] perspectives. As a  
 337     shorthand, when we refer to a code’s perspective, we mean the broadest perspective available in its  
 338     perspective bounds.

339

344     Each function starts out with a top-level perspective bound defined using the notation `@requires(grid[n1], b`  
 345    In Figure 6, the perspective bound is defined on line 2. Starting from this top-level declaration, pro-  
 346    grammers can shape the program's current perspective using two constructs `group` and `split`.

347     **Group.** The `group` construct lets programmers shift from a broader perspective to some number  
 348    of narrower perspectives contained in it. Operationally, the `group` construct does this by replicating  
 349    code written from the perspective of the narrower perspective across the broader one. That is, `group`  
 350    forks many copies of a narrower perspective.  
 351

352     Let's consider this in context of our example. In Figure 6, execution begins at `grid[1]` on line 12.  
 353    At that point, a programmer logically controls the whole grid's behavior. To produce different output  
 354    tiles, the programmer needs to be able to control the behavior of individual blocks, and they shift  
 355    their perspective to `block[1]` on `todo`. Inside the `group(block[1])`'s scope, the behavior of the  
 356    machine can be defined with respect to a single block, and code is replicated across all blocks in the  
 357    grid.

358     As the reader may have already guessed, not all uses of `group` are  
 359    valid. The examples in Figure 7 have no meaningful interpretation on  
 360    the hardware. On line 4, the program tries to broaden its perspective  
 361    to `block[1]` from a `thread[1]` perspective, which is clearly illegal.  
 362    On the other hand, in the second example, the program tries to  
 363    narrow its perspective to a `block[5]` from a `block[6]` perspective.  
 364    While  $5 < 6$ , Parallax cannot evenly replicate `block[5]` perspective  
 365    inside a `block[6]` perspective, and so rejects this program. Recall,  
 366    from Section 3.2, that the condition on whether a perspective is  
 367    narrower than another is actually a condition on divisibility, not just size.  
 368

```

1  # Example 1
2  with group(thread[2]):
3      # Illegal because block > thread
4      with group(block[1]):
5          pass
6
7  # Example 2
8  with group(block[6]):
9      # Illegal because 6 % 5 != 0
10     with group([block[5]]):
11         pass

```

Fig. 7. Illegal uses of `group`.

368     To eliminate such malformed cases, Parallax only allows an invocation of `group(level[n])` if the  
 369    current perspective bounds contains a perspective wider than `level[n]`. Once `group(level[n])` is  
 370    invoked, it modifies the current perspective bounds in two key ways. First, it removes all perspective  
 371    wider than `level[n]` from it. Second, it sets the widest perspective within that `group`'s scope to be  
 372    `level[n]`.  
 373

373     **Split.** Unlike `group`, which is used for replication into multiple narrower perspectives (all of equal  
 374    size), `split` is used for sharding the current perspective unequally. For example, Figure 8 shows  
 375    a `split` from `thread[4]` perspective to one branch with `thread[2]` perspective and two with  
 376    `thread[1]` perspective. The three arms of the `split` execute independently in parallel<sup>1</sup>, as a form  
 377    of unordered composition. When `split(level)` is invoked, the perspectives of its branches diverge.  
 378    At the end of the `split`, they reconverge, and continue execution with the original perspective.  
 379

379     Use of `split` is necessary to write warp-specialized code, a  
 380    programming pattern used in high-performance kernels. Another  
 381    important use of `split`—masking off threads—can be found in our  
 382    running example. Line 63 in Figure 6 shows a `split` that requests  
 383    the first warp in the block, narrowing from `block[1]` into a single  
 384    `thread[32]`. Later in Figure 8, this warp is used to execute a tensor-  
 385    core operation.

```

1  with group(thread[4]):
2      with split(thread):
3          case 2:
4              ...
5          case 1:
6              ...
7          case 1:
8              ...

```

Fig. 8. Simple example of `split`.

386     Because `split` corresponds to unordered composition, it must  
 387    provide each of its branches their requested perspectives simultaneously. Parallax thus checks that  
 388    the cumulative sum of the perspectives requested by each branch of the `split` can be satisfied.  
 389    For example, the program in Figure 9 does not type check. Finally, because perspectives enforce  
 390

<sup>1</sup>Despite the use of the pattern-matching `case` syntax, *all* branches of the `split` execute.

393 alignment, every branch of the **split** must also be aligned; not all **splits** whose sizes sum to less  
 394 than or equal to the available units are valid. Figure 10 shows an example of this constraint being  
 395 violated.

396 Once **split** is invoked, for each branch that requests  $n$  units, all  
 397 perspectives wider than **level** [ $n$ ] are removed from its perspective  
 398 bound, and the available units for **level** are set to exactly  $n$ .

### 400 3.4 Perspectives on Data

401 Section 3.2 described how a programmer can control different layers  
 402 of the hierarchy by changing the perspective of code through **group**  
 403 and **split**. To actually respect the structure of these constructs,  
 404 Parallax must ensure that all units inside a perspective remain log-  
 405 ically bundled together within code written from that perspective,  
 406 and do not diverge. That is, Parallax must be able to reason about  
 407 how control flow introduced through **for**, **while** and **if** statements  
 408 effects the flow of perspectives through a program. To see why, con-  
 409 sider a group of threads seen from **block** [1] perspective, branching  
 410 on a variable which has different values across those threads. In this  
 411 case, different threads will take different paths, and the structure of  
 412 the **block** [1] perspective will be lost.

413 To reason about data-dependent control flow, Parallax must tracks  
 414 which perspectives a piece of data lives at. Parallax distinguishes between two kinds of data: thread-  
 415 local data (or, register and stack data), and thread-shared data (or, shared memory and global memory).

416 3.4.1 *Thread-local data.* Thread-local data is data stored in registers and stack-variables that is  
 417 private to each thread. In the program source in CUDA, a thread-local that has the same name may  
 418 contain different values at runtime. Indeed, to compute anything useful, threads must operate on  
 419 different input values, and being able to reason about this divergent state is essential.

420 Instead of imposing severe restrictions on data-dependent control flow, Parallax solves this problem  
 421 by understanding *how* values diverge across threads. That is, users tell Parallax which perspective  
 422 a thread-local variable lives at. Intuitively, the perspective that a variable lives at tells Parallax the  
 423 *frequency* at which its values change in space. For example, a **grid** [1] value is the same for all  
 424 threads in a grid, meanwhile a **block** [2] value changes every two blocks. Crucially, this frequency  
 425 tells us that a **block** [2] variable is *indistinguishable* by threads within 2 blocks where the starting  
 426 block ID is aligned to 2. Figure 3 shows different frequencies as different colorings of a variable.

427 Programmers specify the perspective that each variables live at in its declaration. A variable **v** of  
 428 type **int** is declared at **thread** [1] perspective using the syntax **v : int @ thread** [1].<sup>2</sup> Variables  
 429 are only available for reading and writing from certain perspective. The rule can be summarized  
 430 as follows: Parallax allows programs to “write down” to narrower perspectives and “read up” from  
 431 wider ones.

432 **Write-down.** Only values that live at wider perspectives can be written to values that live at  
 433 narrower perspectives. For example, a **block** [1] value can be used to write into a **thread** [1] value,  
 434 but a **thread** [1] value cannot be written into a **block** [1] value

436 **Read-up.** Reads from variables are dually constrained. Wider  
 437 perspective values can only be read from values that live narrower  
 438 perspectives. Figure 11 gives an example of an illegal read. Because

```
1 # Example 1
2 with group(thread[4]):
3   match split(thread):
4     case 4:
5       ...
6     # Illegal! 4 + 1 > 4
7     case 1:
8       ...
```

Fig. 9. Illegal split exceeds the available perspectives.

```
1 # Example 2
2 with group(thread[3]):
3   match split(thread):
4     case 2:
5       ...
6     case 1:
7       ...
8     # Illegal! 2 + 1 <= 3,
9     # but the second branch
10    # is not aligned by 2.
```

Fig. 10. Illegal split violates alignment.

```
1 flag : bool @ thread[1] = ...
2 with group(block[1]):
3   if (flag)
4     __syncthreads();
```

Fig. 11. Illegal read of **flag** variable.

<sup>2</sup>If not explicitly annotated, Parallax infers a variable’s perspective to be the perspective of the code where it was declared of **thread** [1] variable.

442 the variable **flag** varies across threads, branching on it causes a sub-  
 443 set of threads in the block to reach the `__syncthreads()`, causing  
 444 a deadlock, despite the fact that `__syncthreads()` should always  
 445 be safe in `block[1]`.

446 Together, the “read up” and “write down” rules ensure that information only flows from high  
 447 perspective variables to low perspective ones . In Figure 6, we can see rule “read up” in action on  
 448 lines 45 and 46. Meanwhile, “write down” is being used to set up variables on line 18-22.

449 3.4.2 *Thread-Shared data*. Unlike thread-local data, which is literally replicated across threads and  
 450 is backed by distinct physical storage, thread-shared data consists of pointers into shared memory  
 451 and global that are visible to some collection of threads. As a result, the perspective such data inhabits  
 452 can evolve as the program executes.  
 453

454 In Parallax, there are three mechanisms for obtaining a thread-shared data. The first method is to  
 455 directly allocate data residing in shared memory (Parallax assumes that all global memory allocations  
 456 have been made before launching the CUDA kernel, as is typically the case with CUDA programs).  
 457 The second and third are to obtain offsets into existing data by using `partition` or `claim`. These  
 458 constructs mirror the `group` and `split` constructs, and are used to widen the perspective associated  
 459 with a pointer.  
 460

461 Parallax simplifies shared memory allocation by requiring all shared memory allocations have a  
 462 static size. Since shared memory is only visible to all threads in a block, a shared memory allocation  
 463 is only valid with if code has a `block[1]` perspective. An example of such an allocation is shown on  
 464 lines 28-30 of Figure 6. During compilation, Parallax will automatically handle the necessary pointer  
 465 arithmetic to assign each allocation an appropriate offset within the shared memory space. The  
 466 `@requires` annotation specifies the amount of shared memory a function expects to have available,  
 467 and Parallax uses this information to ensure that a function’s allocations do not exceed its declared  
 468 limit, and checks whether there is enough shared memory available at its call sites. Figure 6 declares  
 469 the amount of shared memory it will require on 2. We use standard techniques to statically bound  
 470 memory usage with Parallax’s type system.  
 471

472 Now, given an allocation into global or shared memory, Parallax introduces two constructs  
 473 `partition` or `claim` that help programmers change the perspective that a pointer resides at.  
 474

475 **Partition.** The `partition` construct plays the same role for memory that `group` plays for compute.  
 476 It is used to narrow the perspective of a pointer by assigning each narrow perspective a distinct offset  
 477 into the original pointer  
 478

479 The `partition` construct takes the original name of the pointer, `base_name`, an indexing function  
 480 `f`, and a target perspective `level[n]` to which the memory is being narrowed, and produces a new  
 481 variable, `new_name`. Once a memory object has been partitioned, the original name referring to it  
 482 goes out of scope, and the memory can only be accessed through the new name. Each use of this  
 483 new name applies the indexing function `f` to compute the true offset of the access. In Figure 6, for  
 484 example, we partition the C buffer to `block[1]` perspective on line 25 first, and then narrow it  
 485 further to `thread[1]` perspective on line 78 so that we can actually write to it. When the scope  
 486 of the `partition` statement ends, Parallax automatically inserts the necessary synchronization to  
 487 restore the memory to its original perspective. The mechanism for inserting and optimizing these  
 488 synchronization points is described in ??  
 489

490 At this point, it is necessary to consider how different indexing functions affect the possibility of  
 491 data races. If the indexing function is injective, each narrower perspective receives a distinct offset  
 492 into the underlying array, the resulting partition is free of data races. When the indexing function  
 493 is not injective, multiple threads may write to the same location, introducing a potential data race.  
 494 Out-of-bounds accesses are undefined behavior.  
 495

As we discussed in Section 2.2, achieving data-race freedom is *not* one of Parallax’s goals, and our guarantees continue to hold even when indexing functions are not injective. Because Parallax automatically inserts the necessary synchronization before the original memory name is restored to its higher perspective, Parallax ensures that all writes have completed before accessing the original name. The subsequent uses of that pointer will observe the value written by the last writer.

**Claim.** The `claim` construct is the analogue of `split` for memory. Using `claim`, the entire memory region can be moved to a narrower perspective.

Unlike `partition`, `claim` does not take an indexing function. The `claim` construct takes the original memory’s name, the target perspective to narrow to, and a new name to assign to the narrowed memory. The `new_name` is accessible only within a single, explicit `split` branch; sibling branches are not permitted to read or write from this memory. In Figure 6, we use the `claim` operation on line 63 to narrow the memory to `thread[32]` perspective.

As with `partition`, a `split` construct introduces its own scope, and the original name is inaccessible within this scope.

### 3.5 The `id()` Function

As opposed to exposing users to special hardware `blockIdx.x`, `threadIdx.x` variables directly, Parallax provides the `id()` function instead. The `id` function returns the relative index from a given perspective. The interpretation of the `id` function depends on both the perspective of the variable it is being written to, and the perspective that the `id` function is being invoked in. In our example in Figure 6, we use a call to `id` in `grid[1]` scope at line 18 and 19 to locate the tile that each block is ultimately charge of computing. In Figure 2, however, a call to `id` in `thread[32]` scope at line 7 will not return the global thread ID of the threads launched in the kernel, but rather the *relative* thread ID within the 32 threads, ranging from 0 to 31.

### 3.6 Collective Operations

With perspectives on code and data in place, we can now turn our attention to Parallax’s reasoning about collective operations.

There are two ways to access collective operations in Parallax. The first is to use intrinsics provided by the compiler. These intrinsics have built-in perspective classes attached to them, and Parallax enforces that the operation is only ever executed if code is being executed from the correct perspective. In Figure 2, the tensor-core call at line 22 is calling into intrinsics provided by our compiler. The second, more flexible approach allows users to extend these operations using the `unsafe` feature. Using unsafe code, users can inline a collective operation—often a single assembly instruction—into a Parallax function and ensure that the function’s top-level requirements, specified with `@require`, faithfully describe the perspectives required to execute that instruction. Then, Parallax will check the call to the function as any other function call, and will check whether it is being called with a compatible perspective. The rules for checking function calls are described in Section 3.8.

### 3.7 Asynchrony

Asynchronous data-movement works similarly to other data operations.

Users can mark storage as asynchronous with the `async` construct. As with `partition` and `claim`, invoking `async` removes the old name and exposes a fresh name that is only accessible within the scope of the `async` statement. Inside this scope, Parallax restricts the new name to be used solely with the asynchronous data-movement intrinsics that the language

```
1 with async(old_name) as new_name:
2     match thread:
3         case 1:
4             cp_async(smem, new_name, 16);
```

Fig. 12. Example of an `cp.async` instruction in Parallax.

provides. These intrinsics operate on a per-thread level, so Parallax will require users to program from the `thread[1]` perspective.

Parallax models two primary forms of asynchronous data movement: `async mempy` memory copies and tensor-memory accelerator instructions . We show an example of the former in Figure 12. As with the data operations in ??, Parallax is responsible for inserting the necessary synchronization before the program’s first use of the original name, ensuring that the asynchronous transfer has completed.

### 3.8 Function Composition

To check whether a function call is valid, Parallax needs to ensure that the call site can provide the set of perspectives that a function will attempt to use over its execution, and that the arguments passed to the function also respect the perspectives declared for its input and output types. Both requirements are represented explicitly via the `@require` construct.

As mentioned in Section 3.2, each function carries a perspective class. We call this top-level perspective class the function’s *perspective signature*. The signature specifies the top-level perspective class the function *assumes* will be provided to it. This signature represents the *minimum amount* of compute and memory resources the function must be called with. In the example above in Figure 6, line 2 defines the perspective signature. Up to this point, we have described programs written *inside* functions under the assumption that the resource signature can be satisfied. In Section 3.2, for example, we used this signature to set up the top-level perspective class for the function.

**Compute perspectives.** At a function’s call site, then, we need to ensure that the function’s perspective signature can be satisfied. Snap compares the current compute perspective with the one required by the function, and verifies that only functions whose perspective signatures can be satisfied are called.

**Per-argument checks.** In Parallax, functions have pass-by-value semantics. For arguments, we distinguish between primitive data types and pointers to memory. For primitive data types, like `int`, `bool`, and stack allocations, users can pass either arguments that live at or above the data perspective specified by the function’s signature. For pointers, we allow passing pointers that live at higher data perspective if the pointer is marked as `const`, indicating that it will only be used for reading (line 4 and 5 in Figure 2). Otherwise, we require that pointers match the exact perspective of the function signature.

## 4 Formalization in Prism

Having introduced the full Parallax language, we now describe Prism, a core calculus that formalizes its most fundamental aspects by statically tracking perspectives on code and data. We use Prism to argue that well-typed Parallax programs are not only type-safe, but will also never get stuck trying to execute operations for which they lack the correct perspective.

In this section, we describe Prism’s type system and operational semantics—in particular how it manages compute and data perspectives—and build up to a formal proof of type-and-perspective safety. We instrument Prism’s operational semantics with runtime perspective enforcement: its rules will get stuck if they encounter an operation for which they have the wrong perspective. This runtime enforcement means that our safety theorem can guarantee that dynamically-realized perspectives match the ones inferred by the type system.

### 4.1 Prism Type System

The core idea in Prism is to track, at the type-level, the program’s perspective on code and data. To achieve this, we borrow techniques from the literature on *dependency tracking* [1]; in particular, the

589 $\boxed{\Gamma \vdash^\pi e : \tau}$	<i>(Expression typing)</i>
590 $x :^\pi \tau \in \Gamma \quad \text{T-VAR} \quad \frac{\Gamma \vdash^{\pi'} e_1 : \tau[] \quad \Gamma \vdash^\pi e_2 : \text{int} \quad \pi \leq \pi'}{\Gamma \vdash^\pi e_1[e_2] : \tau} \quad \text{T-ARR-ACCESS}$	
593 $\boxed{\Gamma \vdash^\pi s}$	<i>(Statement typing)</i>
595 $\frac{\Gamma \vdash^{(h,n_1)} s_1 \quad \Gamma \vdash^{(h,n_2)} s_2 \quad n_1, n_2 \text{ align to } n}{\Gamma \vdash^{(h,n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\}} \quad \text{T-SPLIT}$	
598      where $n_1, n_2 \text{ align to } n ::= (n_1 + n_2 \leq n) \text{ and } (n_1   n) \text{ and } (n_2   n) \text{ and } (n_2   n_1 + n)$	
600 $\frac{\Gamma \vdash^\pi s}{\Gamma \vdash^{q \cdot \pi} \text{group } q \ s} \quad \text{T-GROUP} \quad \frac{\Gamma, y : \downarrow^\pi \tau[] \vdash^\pi s \quad l \neq \text{Local}}{\Gamma, x :^\pi \tau[] \vdash^\pi \text{lower } x \text{ into } y \text{ in } s} \quad \text{T-LOWER}$	
604 $\frac{\Gamma, y :^{(h,n/c)} \tau[] \vdash^{(h,n)} s \quad c   n \quad l \neq \text{Local}}{\Gamma, x :^{(h,n)} \tau[] \vdash^{(h,n)} \text{partition } x \text{ into } y \text{ by } c \text{ in } s} \quad \text{T-PARTITION} \quad \frac{\Gamma \vdash^\pi s}{\Gamma \vdash^\pi \text{destruct in } s} \quad \text{T-DESTRUCT}$	
607 $\frac{\Gamma, y :^{(h,n')} \tau[] \vdash^{(h,n')} s \quad n' \leq n \quad l \neq \text{Local}}{\Gamma, x :^{(h,n)} \tau[] \vdash^{(h,n)} \text{claim } x \text{ into } y \text{ at } n' \text{ in } s} \quad \text{T-CLAIM}$	

601 Fig. 13. Core typing rules of Prism. The typing rules presented here are a simplified selection of the full rules,  
602 which can be found in Appendix A.2.

613 code perspective is tracked on the typing judgment, which has the form  $\Gamma \vdash^\pi e : \tau$  for expressions and  
614  $\Gamma \vdash^\pi s$  for statements. The  $\pi$  over the  $\vdash$  is the code perspective on  $e$  and  $s$ , and is comprised of a level  
615 and a size, the same structure as a perspective in Parallax.

616 The typing context also tracks the perspective at which each variable lives; data can only be read  
617 or written from a variable when its perspective is compatible with that of the code interacting with it.  
618 This requirement is made manifest in the T-VAR rule, which can be found in Figure 13. Observe that  
619 the  $\pi$  in the variable rule must match exactly between data and code; principles like “read up” and  
620 “write down” are instead encoded directly in the rules for reading and writing, like T-ARR-ACCESS,  
621 which views the array being read with broader perspective than the current code perspective.

622 Figure 13 also shows other key rules, and these rules fall into two main categories: rules for  
623 managing the code perspective and rules for managing data perspective.

626 **4.1.1 Managing Perspectives on Code.** In Parallax, there are two operations which manage code  
627 perspectives: **group** and **split**. In Prism, to better model the details of how perspectives shift  
628 throughout programs, we introduce a third construct, **destruct**.

629 Prism’s **group** statement directly corresponds to Parallax’s, and is checked by the T-GROUP rule.  
630 Given some perspective that can be divided into  $q$  equally sized narrower  $\pi$ s, the statement **group**  $q \ s$   
631 will check with perspective  $q \cdot \pi$  provided that  $s$  itself checks with perspective  $\pi$ . The starting  
632 perspective is of the form  $q \cdot \pi$ , and encodes Parallax’s alignment requirement.

633 The **split** statement, meanwhile, is checked by the T-SPLIT rule, and functions like a binary  
634 version of the n-ary **split** construct in Parallax. It enforces the same divisibility requirements to  
635 ensure that the perspectives on code and data remain properly aligned, and then checks the two  
636 sub-statements  $s_1$  and  $s_2$  with the divided, narrower perspectives.

The final rule for managing code perspective is **T-DESTRUCT**, which shifts perspective down the GPU hierarchy and makes explicit in Prism programs exactly where such shifts occur. The  $\downarrow$  operation on  $\pi s$  “destructs” the perspective into many narrower perspectives at a lower level;  $\downarrow(\text{Grid},1) = (\text{Block},B)$  and  $\downarrow(\text{Block},1) = (\text{Thread},T)$ , where  $B$  and  $T$  are parameters to a particular instantiation of Prism to describe the number of blocks per grid and threads per block.<sup>3</sup> Because  $\downarrow$  is only defined on  $(\text{Grid},1)$  and  $(\text{Block},1)$ , the rule enforces that one can only **destruct** their code perspective with exactly one grid or block.

**4.1.2 Managing Perspectives on Data.** The mechanism for managing data perspective mirrors that of code perspective, with each operation for data corresponding to an operation for code.

The **partition** operation is analogous to **group**-ing a code perspective. The typing rule for this operation, **T-PARTITION**, requires that the data perspective on the variable to be partitioned,  $x$ , is the same as the current perspective on code. After partitioning, a fresh variable  $y$  is introduced with a new perspective  $\pi/c$ , which we use as a shorthand to denote a division of a perspective  $\pi$  into  $c$  equally-sized narrower perspectives. Within the body of the **partition**, we disallow any references to the original variable  $x$  and continue checking with the original code perspective  $\pi$ ; the **partition** has no effect on the current code perspective.

Unlike **partition**, which divides up a piece of data equally among narrower perspectives, the **claim** operation views the claimed data with exactly one narrower perspective. Accordingly, Prism needs to ensure that only one branch of a **split** operation, with the appropriate  $\pi$ , can refer to the claimed variable. To ensure that this is the case, the **T-CLAIM** rule links the data perspective of the variable to the compute perspective of the code claiming it by changing *both* at the same time. This represents a minor difference from Parallax, which uses additional static analysis to ensure that a claimed variable is only accessed in a single **split** branch.

Lastly, the **T-LOWER** rule mirrors the **T-DESTRUCT** rule; it uses the  $\downarrow$  operator to move a variable from one level of the hierarchy to another, distributing it equally among all the narrower perspectives at that level in the same manner as **T-PARTITION**.

## 4.2 Prism Semantics

Having explained the key rules of the type system, we can move on to discuss Prism’s operational semantics. To reflect the fact that GPU programs execute in parallel across numerous threads, we model the semantics of Prism in the style of Turon et al. [35], using a two-level small step judgment. We present the key rules of this semantics in Figure 14.

The top level (i.e., device-level) judgment has just one rule: **S-PROGRAM**. This rule acts as a “frame” for the lower level (i.e., thread-level) judgment, and steps a collection of thread-ID-indexed local memories ( $L$ ), a collection of block-ID-indexed shared memories ( $S$ ), a global memory ( $\Sigma$ ), and a *thread pool* to an updated collection of memories and thread pool. The thread pool maps thread and block IDs to code, intuitively representing the program being executed by each thread at the current moment. The **S-PROGRAM** non-deterministically chooses a thread ID and block ID and steps it according to the thread-level judgment. This allows the semantics to model the full range of non-deterministic behavior arising from the GPU’s thread scheduler.<sup>4</sup>

The thread-level judgment has the shape  $\eta, \sigma, \Sigma, t, b, p \vdash^\pi s \rightsquigarrow s' + \eta', \sigma', \Sigma'$ , where

- $\eta$  denotes thread-local memory,
- $\sigma$  denotes shared memory,

<sup>3</sup>Prism is abstracted over these  $B$  and  $T$  values, so instead of tracking perspective bounds the way that Parallax does, it only tracks the top-level perspective described in Section 3.2.

<sup>4</sup>In reality, the GPU’s thread scheduler guarantees that the threads of a warp execute in lockstep, but modeling every thread as completely independent is both simpler and a conservative overestimate of the nondeterministic behavior of the GPU.

$$\begin{array}{l}
687 \quad \frac{L(t), S(b), \Sigma, t, b, 0 \vdash^{(\text{Grid},1)} s \rightsquigarrow s' + \eta', \sigma', \Sigma' \quad P(t, b) = s}{L, S, \Sigma, P \rightsquigarrow L[t \mapsto \eta'], S[b \mapsto \sigma'], \Sigma', P[(t, b) \mapsto s']} \text{ S-PROGRAM} \\
688 \\
689 \\
690 \\
691 \quad \frac{p < n_1 \quad n_1, n_2 \text{ align to } n \quad \eta, \sigma, \Sigma, t, b, p \vdash^{(h, n_1)} s_1 \rightsquigarrow s'_1 + \eta', \sigma', \Sigma'}{\eta, \sigma, \Sigma, t, b, p \vdash^{(h, n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\} \rightsquigarrow \text{split}(n_1, n_2)\{s'_1\}\{s_2\} + \eta', \sigma', \Sigma'} \text{ S-SPLIT-LEFT} \\
692 \\
693 \\
694 \quad \frac{p \geq n_1 \quad p < n_1 + n_2 \quad n_1, n_2 \text{ align to } n \quad \eta, \sigma, \Sigma, t, b, p - n_1, \vdash^{(h, n_2)} s_2 \rightsquigarrow s'_2 + \eta', \sigma', \Sigma'}{\eta, \sigma, \Sigma, t, b, p \vdash^{(h, n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\} \rightsquigarrow \text{split}(n_1, n_2)\{s_1\}\{s'_2\} + \eta', \sigma', \Sigma'} \text{ S-SPLIT-RIGHT} \\
695 \\
696 \\
697 \\
698 \quad \frac{\eta, \sigma, \Sigma, t, b, p \text{ mod } n \vdash^{(h, n)} s \rightsquigarrow s' + \eta', \sigma', \Sigma'}{\eta, \sigma, \Sigma, t, b, p \vdash^{(h, q \cdot n)} \text{group } q \ s \rightsquigarrow \text{group } q \ s'; + \eta', \sigma', \Sigma'} \text{ S-GROUP} \\
699 \\
700 \\
701 \quad \frac{\eta, \sigma, \Sigma, t, b, t \text{ mod } T \vdash^{(\text{Thread}, T)} s \rightsquigarrow s' + \eta', \sigma', \Sigma'}{\eta, \sigma, \Sigma, t, b, 0 \vdash^{(\text{Block}, 1)} \text{destruct in } s \rightsquigarrow \text{destruct in } s' + \eta', \sigma', \Sigma'} \text{ S-DESTRUCT-BLOCK} \\
702 \\
703 \\
704 \\
705 \quad \frac{}{\eta, \sigma, \Sigma, t, b, p \vdash^{(\text{Block}, 1)} x := \text{alloc Shared } \tau \ n \ \text{in } s \rightsquigarrow s + \eta, \sigma[x \mapsto \pi(x, n)], \Sigma} \text{ S-ALLOC-SHARED} \\
706
\end{array}$$

Fig. 14. Core semantic rules of Prism. As with the typing rules, we present only a simplified selection of the full rules, which can be found in Appendix A.3.

- 711 •  $\Sigma$  denotes global memory,
- 712 •  $t$  denotes the thread's ID,
- 713 •  $b$  denotes the ID of the block in which the thread lives, and
- 714 •  $p$  denotes the relative position of the thread within  $\pi$  (the perspective ID).

715 Critically, notice that a  $\pi$  also appears on the thread-level judgment just as it does on the typing  
716 judgment. This is because the thread semantics *dynamically tracks and enforces* perspectives. The  
717 same way evaluation of a program “gets stuck” if a value does not have the right type, the semantics of  
718 Prism also get stuck if code attempts to access data or invoke commands with the wrong perspective.  
719 As an example, observe the **S-ALLOC-SHARED** rule in Figure 14, which requires a (**Block**,1) perspective  
720 and will fail to step if encountered with a different one. This runtime perspective is present in Prism,  
721 but is erased by Parallax during compilation; we will later use it to prove that well-typed programs  
722 will always execute with the same perspective that the type system viewed them with.

723 The semantic rules for perspectives involve manipulating  $p$  to track which threads take which code  
724 paths when perspectives are **split** or **grouped**. Notice that in **S-PROGRAM**, the thread-level judgment  
725 always begins with perspective (**Grid**,1): all the perspective management rules are congruences,  
726 narrowing the perspective of further evaluation as determined by the particular rule used.

727 These rules take great care to ensure that  $p$  always describes the relative position of a compute  
728 resource within its perspective; the payoff is that Prism’s semantics can later use this  $p$  value to model  
729 the way that Parallax automatically adjusts indices into data when partitioning a data perspective.  
730 Beyond these key rules for managing perspective on code, we have modeled many of the other  
731 features of Parallax, such as asynchronous operations and thread synchronization, in Prism. To  
732 handle features like these we equip the operational semantics with additional structure, including  
733 sets of semaphores [12] for synchronization and a stack of effect handlers for modeling deferred  
734

asynchronous computations inspired by Ahman and Pretnar [3]. We have elided these details here for simplicity, but interested readers can see further details in Appendix A.3.

### 4.3 Soundness Theorem

Together, the type system and operational semantics allow us to prove the following syntactic soundness theorem, which says that Prism programs are type safe and do not get stuck trying to execute operations for which they lack the required perspective:

**THEOREM 4.1. (Type-and-Perspective Safety).** *For any program  $s$  such that  $\Gamma \vdash^\pi s$ , either:*

- (1)  $s$  is `skip`, or
- (2) for any well-typed environments  $\eta, \sigma$ , and  $\Sigma$ , there is an  $s'$ ,  $\eta'$ ,  $\sigma'$ , and  $\Sigma'$  such that  $\eta, \sigma, \Sigma, t, b, p \vdash^\pi s \rightsquigarrow^* s' + \eta', \sigma', \Sigma'$  such that  $\Gamma' \vdash^\pi s'$ , where  $\Gamma'$  is an extension of  $\Gamma$ , and  $\eta', \sigma'$ , and  $\Sigma'$  are well-typed with respect to  $\Gamma'$ .

**PROOF.** Via the usual progress and preservation lemmas, available in Appendix A.4.  $\square$

It is worth noting that this soundness theorem guarantees a syntactic safety property, not a liveness property: it does not guarantee that all threads sharing perspective  $\pi$  that *can* reach a program point typed with  $\pi$  *will* eventually do so. Indeed, in the presence of nontermination, liveness does not hold—some of the threads could `split` off and loop forever. While we believe the liveness version of this theorem holds for a terminating fragment of Prism, it is not provable with syntactic methods; the proof would require semantic techniques that are notoriously challenging and would be a research contribution [4, 15, 35] in and of itself. We hope to tackle this proof for Prism in future work.

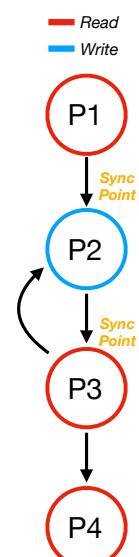
## 5 Implementation

Parallax is implemented as an embedded language in Python. Once a program type checks, Parallax lowers it to a CUDA file. All perspective information is erased during this step, and the generated CUDA contains no run-time checks. The file can then be compiled by `nvcc`, NVIDIA’s closed-source compiler, to produce an executable. Because Parallax operates at roughly the same abstraction level as CUDA, there is a one-to-one mapping between most language constructs and their CUDA counterparts. A notable change is the addition of three parameters to each device function: the thread’s relative ID, the block’s ID, and an offset into shared memory that it can use for allocations.

*Inserting Synchronization.* Most of Parallax’s implementation is straightforward, but inserting synchronization points is more involved. As described in ??, once data has been partitioned, Parallax is responsible for inserting the appropriate synchronization points when the partition’s scope ends.

To determine where these points need to be inserted, Parallax constructs a *data-control-flow graph* from the program. Each node corresponds to a partition, and each edge captures program-order precedence: the parent partition must complete before the child begins. The graph can have backedges introduced through loops. In this graph, each partition is categorized as a *read* or a *write* partition. Synchronization points are inserted according to the following scheme:

- (1) If the parent partition is a write partition, a synchronization point is inserted to ensure that the current partition observes the most recent data; or
- (2) If the current partition is a write partition, a synchronization point is inserted to ensure that all preceding reads on the same memory have completed.



785     Figure 15 shows an example graph with the synchronization points derived from  
 786     these two conditions. The inferred synchronization points in Figure 6 have also  
 787     been marked on lines **TODO**.

788     Using this information, Parallax emits *wait* operations before a partition begins  
 789     (to wait for its parent) and *arrive* operations when a partition ends (to signal com-  
 790     pletion). CUDA exposes a general *split-barrier* primitive that we use to implement  
 791     these waits and arrives. For special cases, such as synchronizing an entire block or  
 792     warp, Parallax instead uses primitives like `__syncthreads()` or `__syncwarp()`.

793     Synchronization for asynchronous data movement is handled in exactly the  
 794     same way and uses the same underlying graph. CUDA allows asynchronous load  
 795     operations to be associated with a split barrier, so Parallax binds each asynchronous  
 796     transfer to the appropriate wait–arrive pair inferred from the graph. Certain features,  
 797     such as commit-group-style synchronization , require additional reasoning, and  
 798     Parallax performs further static analysis to insert the necessary synchronization  
 799     around these operations.

800     It is worth noting that naively inserting synchronization immediately after each partition would be  
 801     correct but prohibitively slow. To avoid this, Parallax applies two optimization passes. A wait-motion  
 802     pass pushes waits downward toward the first use of the partitioned name, and an arrive-motion pass  
 803     pulls arrives upward toward its last use, reducing unnecessary stalls while preserving correctness.

## 804     6 Evaluation

805     Having explained the design of Parallax, we now evaluate it in the context of four main questions:

806     **RQ1** Can Parallax express a diversity of CUDA programs?

807     **RQ2** Can Parallax express programs that use advanced GPU features?

808     **RQ3** Can Parallax match the performance of existing, speed-of-light CUDA code?

809     Although the following question is not central to our paper’s main claims—since library design  
 810     ultimately reflects choices made by programmers, and our language allows abstraction design much  
 811     like any other language—we were, nevertheless, curious to explore it based on empirical and anecdotal  
 812     observation:

813     **RQ4** Can Parallax help build compositional libraries that users can use with confidence?

814     To perform this evaluation, we use two GPUs. The first is the NVIDIA H100 SXM5, a server-grade  
 815     chip that supports tensor-core operations and a dedicated hardware copy engine, the Tensor Memory  
 816     Accelerator. Notably, the H100 introduces a new *logical* level called the *warpgroup* (collection of  
 817     4-aligned warps), and we show that our programming model can accommodate this new level.  
 818     Moreover, because the H100 has historically served as the primary GPU for large-scale AI training,  
 819     many CUDA kernels on this hardware are already highly optimized and achieve near-speed-of-  
 820     light performance, providing a rigorous baseline for comparison. To ensure our results generalize  
 821     beyond the H100, we additionally test programs on a second GPU, the NVIDIA 4070 SuperTi GPU, a  
 822     consumer-grade chip.

823     Like mentioned in Section 5, when name typechecks a program, it produces a CUDA file. We  
 824     compile the CUDA file with `nvc` version with flags . To measure performance of our benchmark,  
 825     we set up the values from an input using a random number generator. We use the average running  
 826     time sampled across 10 iterations, with a warm-up of 5 iterations.

## 834 6.1 RQ1: Can Parallax express a diversity of CUDA programs?

835 To answer this research question, we evaluate Parallax on two benchmarks: a sequence of matrix  
 836 multiplication kernels that grow in complexity and a single-pass parallel prefix scan with decoupled  
 837 look-back that requires programmers to think carefully about multiple points of convergence.  
 838

839 *Matrix Multiplication.* We chose matrix multiplication as our first benchmark for two main reasons.  
 840 First, there exist several implementations that achieve near-peak performance, providing a strong  
 841 baseline. Second, matrix multiplication allows for a range of increasingly sophisticated implementa-  
 842 tions that stress different parts of the language, making it ideal for evaluating expressiveness.  
 843

844 To undertake this exploration, we adapt the codebase from to implement a matrix multiplication  
 845 on the `f16` datatype, also known as `sgemm`, in Parallax. As this is a `f16` benchmark, it *does*  
 846 *not need to* use advanced GPU features like asynchrony or tensor cores to achieve speed-of-light  
 847 performance. We will discuss these advanced features in Section 6.2.

848 We implement several variants of the `sgemm` benchmark, reproduced in Appendix . These include a  
 849 naive baseline, a version that exploits memory coalescing, 2D blocking with shared-memory staging,  
 850 2D blocking with vectorized loads, and a warp-tiling strategy that effectively introduces an additional  
 851 level of tiling from the perspective of a warp. We find that the runtime performance of our Parallax  
 852 implementations closely matches that of the original versions. This is expected, since the generated  
 853 code is nearly identical to the hand-written versions, aside from small differences such as hoisted  
 854 expressions and a few index calculations. The downstream compiler (`nvcc`) readily inlines hoisted  
 855 expressions that are used only once, so the performance-critical inner loops remain effectively  
 856 unchanged.

857 *Single-pass Parallel Prefix Scan with Decoupled Look-Back.* We also implement scan, a widely used  
 858 parallel primitive, in Parallax. We focus on the prefix-sum scan, which computes, for each position in  
 859 an array, the sum of all elements up to that position. Prefix-sum sits in a different corner of the GPU  
 860 design space from matrix multiplication: it is memory intensive, requires careful attention to the  
 861 convergence behavior of different threads, and traditionally requires multiple passes over data.  
 862

863 We implement the single-pass parallel prefix scan with decoupled look-back, introduced by Merrill  
 864 and Garland , an elegant algorithm that does not require multiple passes over the input data. In  
 865 this algorithm, each block computes a local prefix sum over its assigned region of the array. Once  
 866 finished, it writes the final element of its region into a global array. Each block then *looks back* to  
 867 accumulate the contributions of prior blocks, allowing them to independently determine the global  
 868 prefix without a full sweep over memory. This decoupled look-back mechanism lets blocks progress  
 869 at different speeds while still producing correct global results.

870 The algorithm involves several distinct points of convergence. Within each block, work is de-  
 871 composed into fine-grained thread-level and warp-level scans. After producing the local result,  
 872 blocks publish their partial prefix to global memory. Finally, each block waits until enough prefix  
 873 information from earlier blocks becomes available, at which point it accumulates the value and  
 874 completes its section of the global scan.

875 We implement this full strategy in Parallax, available in Appendix . Our implementation uses the  
 876 unsafe feature to implement a global-memory spinlock that lets blocks check when the preceeding  
 877 block's data is ready.

## 878 6.2 RQ2: Can Parallax express programs that use advanced GPU features?

879 To test whether Parallax can express programs that use advanced features of modern GPUs, we write  
 880 a matrix multiplication for the `bfloat16` datatype for the H100.  
 881

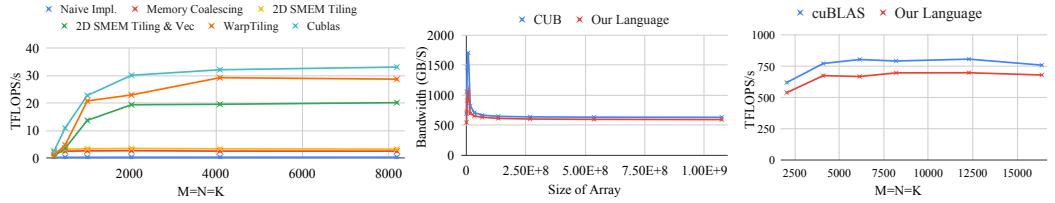


Fig. 16. Performance of `sgemm` on square matrices as matrix dimension  $M=N=K$  increases.  
 Fig. 17. Performance for prefix sum as input array size increases.  
 Fig. 18. Performance of `hgemm` on square matrices as matrix dimension  $M=N=K$  increases.

This benchmark is an acid test of our language, as the H100 bf16 matrix multiplication pushes several language features to the extreme. To write a matrix multiplication that can hit peak throughput on an H100, we need to write a warp-specialized kernel that uses the Tensor Memory Accelerator (TMA), an asynchronous hardware copy engine that can move tiles of data at a time, and uses the warp-group-level tensor core instructions, or `wmma`, specifically introduced for the Hopper architecture.

The implementation works as follows. First, we assign each block on our machine a logical tile of the output to compute. The block is then divided into a producer warp-group and a set of consumer warp-groups, where the producer loads data for multiple consumers, and both must signal to each other when one is done loading and the other is ready to compute. In this way, the benchmark overlaps computation with data movement by pipelining loads.

The implementation in Parallax looks different from normal CUDA code, particularly in how pipelining is expressed. Since Parallax uses *names* and subsequent `partition` or `claim` constructs to determine the synchronization each region requires, when pipelining, we cannot dynamically change the pipeline slot simply by maintaining an index variable that wraps around based on the pipeline length. Instead, each pipeline slot must be given a separate name so that Parallax can track them independently and actually overlap compute with data-movement. This leads to pipeline slots that must be individually named and forces the load logic to be effectively “unrolled,” since we can no longer iterate over a pipeline-slot index. In turn, this forces that all pipelines in Parallax be statically sized. In practice, these pipelines are statically sized in CUDA anyway because they occupy shared memory, which is a small, finite resource that must be explicitly managed.

Notably, for this benchmark, to get the `wmma` instruction to work we did not need to add a new perspective to our language; we could simply use `thread[128]`. We did, however, need to add a special TMA asynchronous data-movement construct, since Parallax will eventually need to synchronize these transfers.

### 6.3 RQ3: Can Parallax match the performance of existing, speed-of-light CUDA code?

In Section 6.1 and Section 6.2, we examined programs that expressed the same computation in multiple ways, expressed operations that rely on multiple points of convergence, and used advanced GPU features. We now discuss the performance of these programs.

The performance of the matrix multiplication benchmarks is shown in , where we demonstrate that Parallax is competitive with cuBLAS [7].

For the prefix scan, we compare our performance to , the library introduced earlier in Section 1, and show that we can reach approximately % of CUB’s peak bandwidth.

Finally, we evaluate our H100 implementation and show that it is competitive with cuBLAS on square sizes, falling between %. We emphasize that this is an exacting benchmark, and achieving this level of performance requires writing large kernels with precise control over low-level features.

## 932 6.4 Case Study: Can Parallax help with library design?

933 Over the course of our evaluation, we found ourselves developing a small library of functions—similar  
 934 in spirit to the CUB library—that we could call to execute our operations. Based on our experience,  
 935 we would like to study, qualitatively, if Parallax can help programmers design libraries that they can  
 936 compose with confidence.

937 As mentioned in Section 1, CUB occupies a unique design point in the GPU library ecosystem.  
 938 Unlike many other libraries such as cuBLAS [7], cuDNN [8], and cuSPARSE [9], which provide  
 939 global interfaces that users can call and configure, CUB provides a device-side library organized into  
 940 different levels. It makes these levels apparent by prefixing each of its functions with **Device**, **Block**,  
 941 and **Thread**. The single-scan prefix sum in Section 6.1 used these functions in Parallax already.

942 Let’s turn our attention to a particular CUB function—the store function—and examine how it is  
 943 equivalently expressed in Parallax, and how Parallax’s type system reifies the assumptions implicit  
 944 in CUB.

945 In CUB, the store function is implemented as a class, as shown below:

```
947 1 template<typename T, int BlockDimX,
948   2     int ItemsPerThread, BlockLoadAlgorithm Algorithm = BLOCK_LOAD_DIRECT,
949   3     int BlockDimY = 1, int BlockDimZ = 1>
4 class BlockStore:
```

950 To use it, users must first instantiate an object of this class, and then call it with shared memory.

```
951 1 using BlockLoad = cub::BlockLoad<int, 128, 4, BLOCK_LOAD_DIRECT>;
952 2 // Allocate shared memory for BlockLoad
953 3 __shared__ typename BlockLoad::TempStorage temp_storage;
954 4 int thread_data[4]; // Thread local data
955 5 BlockLoad(temp_storage).Load(d_data, thread_data);
```

956 CUB exposes a leaky abstraction, where information about the number of threads, block sizes, and  
 957 other details seeps through:

- 958 (1) The CUB documentation needs to specify the number of threads that the function can assume  
 959 to be available, because within the function, each thread must locate itself in the computation  
 960 and use its **threadIdx.x** accordingly. If the starting **threadIdx.x** is not 0, the function must  
 961 compute its relative ID internally.
- 962 (2) The “item per thread” design is interesting and serves two purposes. The first is related to  
 963 performance: if loops have constant bounds, they can be unrolled, enabling downstream  
 964 optimizations. The second is related to correctness: the function relies on the assumption  
 965 that all threads call the function with an equal number of values to load.
- 966 (3) The CUB documentation also needs to specify that **thread\_data** can be data local to each  
 967 thread.
- 968 (4) Finally, the CUB documentation specifies that if shared memory is being overwritten, a  
 969 **\_syncthreads()** call must be made to ensure that all reads have completed.

971 On the other hand, the same function has a completely different interface in Parallax:

```
972 1 #include "device"
973 2 #include "require.h"
974 3 def load(src: ptr(const(shared(int))) block[1],
975 4     item_per_thread: int block[1],
976 5     total_size: int block[1],
977 6     thread_data : ptr(int) thread[1]):
```

978 In Parallax, we are able to encapsulate code effectively, reducing the need to communicate numerous  
 979 implementation details through documentation:

- (1) We do not need to pass in the number of threads at all. Whenever Parallax calls a function, it threads the relative ID through, so each function can be written locally as if it were running alone, rather than having to determine where the thread resides in the global array.
- (2) We do not need to make `item_per_thread` a template argument for *correctness*. Its frequency is set at the function signature, so Parallax will never allow a function to be called with a value living at a narrower perspective than its signature requires.
- (3) In our interface, `thread_data` is explicitly set to a thread-local value. Since it is not marked as `const`, Parallax conservatively assumes it may be written to, and enforces at compile time that only `thread[1]` values are passed in.
- (4) Finally, using our synchronization pass outlined in Section 5, a `__syncthreads()` call will be inserted automatically if `src` is going to be used for writing.

## 7 Related Work

Parallax builds on a rich tradition for systems languages for GPU programming and theoretical foundations for reasoning about parallel programs. However, Parallax differs from these systems in a crucial way: it treats collective operations, and therefore composability, as first-class concerns. Parallax occupies a distinct niche by pairing low-level control with safety, rather than prioritizing one at the expense of the other.

**Imperative Languages for GPUs.** CUDA [22], ROCm [2], and OpenCL [31] are imperative, C-style programming languages that expose low-level access to GPU hardware. These languages do not model perspectives for code or data, statically or otherwise; instead, programmers must orchestrate computation on the machine by describing how individual threads execute code. The subtle and often silent failures permitted by this model are the motivation for Parallax.

Descend [20], a new, Rust-inspired language, uses a type system to track aspects of the compute and memory hierarchy, and, is thus the closest to our work. However, Descend’s main concern is preventing data races, while ours is ensuring that collective operations execute in valid contexts. Moreover, Descend lacks support for many collective operations like tensor cores; in fact, because Descend allows threads to read their own IDs and thus induce data-dependent control flow, adding such support would require major changes to the language. It is also worth noting that while Descend does formalize a type system, it does not attempt to prove any properties about it.

**Functional Languages for GPUs.** Futhark [19], Accelerate [5] and Vertigo [14] are functional array languages with compilers targeting CUDA. These languages expose a high level interface, abstracting away details of the hardware entirely in exchange for stronger safety guarantees. Parallax, however, makes an intentional decision to expose low-level details of the hardware, and thus does not trade performance for safety.

**Tile-Based Kernel DSLs.** More recently, tile-based GPU languages—Triton [33], Pallas [30], Tilus [13], Helion [23]—offer a middle ground between high-level abstraction and low-level control. However, these languages sidestep the question of composability entirely because they restrict programmers to a single vantage point in the GPU compute hierarchy: a block (Triton, Tilus, Helion), or a warpgroup (Pallas). Gluon [34] and is a new low-level, tile-based language. However, it does not check and cannot enforce that collective operations are executed at the correct layer of the hierarchy.

**Task-Based and Scheduling Languages.** Languages like Cypress [36], Halide [24], Fireiron [16] and RISE/ELEVATE [17] expose a mapping specification or scheduling language that lets users modify an existing reference program through external commands. Because they operate by transforming a fixed source-of-truth rather than generating the program directly, they expose a fundamentally different programming model than Parallax.

**Libraries Built on CUDA.** Many GPU libraries, such as CUTLASS [29], CUB [21], and ThunderKittens [25], expose device functions that operate at various levels of the compute hierarchy. These functions are typically organized via C++ namespaces to reflect their intended usage (e.g., warp-level, block-level). However, this organization is purely conventional: it encodes hierarchy through naming rather than enforcing it with a compiler. As a result, correct use requires careful discipline from both the library implementer (to uphold naming and usage invariants) and the user (to correctly interpret them). Any mismatch or subtle misunderstanding between the intended and actual use of these functions goes unchecked by the compiler, leaving room for fragile errors.

**Theoretical Foundations.** The design of Prism is heavily inspired by existing work on dependency tracking [1]. Dependency tracking calculi allow type systems to track how data and code depend on each other, and have commonly been used to implement secure information flow analyses [11]. In Prism, data that lives at a narrow perspective unable flow into data living at a broader perspective, and we use dependency tracking to capture this restriction in Prism’s type system.

## 8 Conclusion, Limitations, and Future Work

We have presented Parallax, a new, low-level GPU language that statically guarantees safe usage of compute resources by construction, without sacrificing low-level control. Parallax introduces a new mental model for writing GPU code, which we are excited to make more expressive and ergonomic.

Currently, Parallax does not natively support explicit references to pointers; programmers must rely on built-in constructs to interact with memory. Prior work such as Descend [20], which builds on Rust’s ownership model [28], provides a promising direction for extending Parallax with more flexible pointer semantics.

We also aim to improve the ergonomics of programming with Parallax, particularly by enhancing the pipelining experience. As discussed in Section 6.2, Parallax currently requires pipeline slots to be explicitly name; we can improve this by adding language support for generating pipeline-style code.

Parallax can be extended to more architectures; in particular, it can accommodate newer GPUs—such as Blackwell—that support coarser-grained tensor-core operations. More broadly, we hope to generalize the model presented here to hierarchical compute environments, including distributed systems.

On the theoretical side, we plan to explore a terminating fragment of Prism and prove the liveness property discussed in Section 4.3: that all threads sharing perspective  $\pi$  eventually reach the parts of a program viewed at that  $\pi$ . This amounts to showing that threads sharing the same perspective execute the same code and observe the same data, which we hope to prove using logical relations, following Turon et al. [35] and Spies et al. [26, 27].

We believe that Parallax is a promising low-level substrate that enables confident composition and can serve as a foundation for building higher-level libraries and abstractions.

## References

- [1] Martín Abadi, Anindya Banerjee, Nevin Heintze, and Jon G. Riecke. 1999. A core calculus of dependency. In *Proceedings of the 26th ACM SIGPLAN-SIGACT symposium on Principles of programming languages* (San Antonio Texas USA, 1999-01). ACM, 147–160. [doi:10.1145/292540.292555](https://doi.org/10.1145/292540.292555)
- [2] Advanced Micro Devices, Inc. 2024. AMD ROCm™ Software. <https://www.amd.com/en/products/software/rocm.html>. Accessed: 2025-11-10.
- [3] Danel Ahman and Matija Pretnar. 2021. Asynchronous effects. *Proceedings of the ACM on Programming Languages* 5, POPL (Jan. 2021), 1–28. [doi:10.1145/3434305](https://doi.org/10.1145/3434305)
- [4] Lars Birkedal, Filip Sieczkowski, and Jacob Thamsborg. 2012. A Concurrent Logical Relation. In *Computer Science Logic (CSL’12) - 26th International Workshop/21st Annual Conference of the EACSL (2012)*. Schloss Dagstuhl – Leibniz-Zentrum für Informatik, 107–121. [doi:10.4230/LIPIcs.CSL.2012.107](https://doi.org/10.4230/LIPIcs.CSL.2012.107)
- [5] Manuel M.T. Chakravarty, Gabriele Keller, Sean Lee, Trevor L. McDonell, and Vinod Grover. 2011. Accelerating Haskell array codes with multicore GPUs. In *Proceedings of the Sixth Workshop on Declarative Aspects of Multicore*

- 1079        *Programming* (Austin, Texas, USA) (*DAMP '11*). Association for Computing Machinery, New York, NY, USA, 3–14.  
 1080        doi:[10.1145/1926354.1926358](https://doi.org/10.1145/1926354.1926358)
- 1081 [6] NVIDIA Corporation. 2025. cub::BlockReduce. [https://nvidia.github.io/ccl/cub/api/classcub\\_1\\_1BlockReduce.html](https://nvidia.github.io/ccl/cub/api/classcub_1_1BlockReduce.html)
- 1082 [7] NVIDIA Corporation. 2025. cuBLAS. <https://docs.nvidia.com/cuda/cublas/index.html>
- 1083 [8] NVIDIA Corporation. 2025. cuDNN. <https://docs.nvidia.com/deeplearning/cudnn/latest/>
- 1084 [9] NVIDIA Corporation. 2025. cuSPARSE. <https://docs.nvidia.com/cuda/cusparse/index.html>
- 1085 [10] NVIDIA Corporation. 2025. Warp Shuffle Functions. <https://docs.nvidia.com/cuda/cuda-c-programming-guide/#warp-shuffle-functions>
- 1086 [11] Dorothy E. Denning and Peter J. Denning. 1977. Certification of programs for secure information flow. *Commun. ACM* 20, 7 (July 1977), 504–513. doi:[10.1145/359636.359712](https://doi.org/10.1145/359636.359712)
- 1087 [12] Edsger Wybe Dijkstra. 1963. Over de sequentialiteit van procesbeschrijvingen. (1963). <https://training-ir7.tdl.org/handle/123456789/594>
- 1088 [13] Yaoyao Ding, Bohan Hou, Xiao Zhang, Allan Lin, Tianqi Chen, Cody Yu Hao, Yida Wang, and Gennady Pekhimenko. 2025. Tilos: A Tile-Level GPGPU Programming Language for Low-Precision Computation. arXiv:2504.12984 [cs.LG] <https://arxiv.org/abs/2504.12984>
- 1089 [14] C. Elliott. 2005. Vertigo – GPU Compiler & Embedded Language for Graphics Processors. <http://conal.net/Vertigo/>. Accessed: 2025-11-10.
- 1090 [15] Azadeh Farzan, Zachary Kincaid, and Andreas Podelski. 2016. Proving Liveness of Parameterized Programs. In *Proceedings of the 31st Annual ACM/IEEE Symposium on Logic in Computer Science*. ACM, New York NY USA, 185–196. doi:[10.1145/2933575.2935310](https://doi.org/10.1145/2933575.2935310)
- 1091 [16] Bastian Hagedorn, Archibald Samuel Elliott, Henrik Barthels, Rastislav Bodik, and Vinod Grover. 2020. Fireiron: A Data-Movement-Aware Scheduling Language for GPUs. In *Proceedings of the ACM International Conference on Parallel Architectures and Compilation Techniques* (Virtual Event, GA, USA) (*PACT '20*). Association for Computing Machinery, New York, NY, USA, 71–82. doi:[10.1145/3410463.3414632](https://doi.org/10.1145/3410463.3414632)
- 1092 [17] Bastian Hagedorn, Johannes Lenfers, Thomas Kundefinedhler, Xueying Qin, Sergei Gorlatch, and Michel Steuwer. 2020. Achieving high-performance the functional way: a functional pearl on expressing high-performance optimizations as rewrite strategies. *Proc. ACM Program. Lang.* 4, ICFP, Article 92 (Aug. 2020), 29 pages. doi:[10.1145/3408974](https://doi.org/10.1145/3408974)
- 1093 [18] Troels Henriksen, Niels G. W. Serup, Martin Elsman, Fritz Henglein, and Cosmin E. Oancea. 2017. Futhark: purely functional GPU-programming with nested parallelism and in-place array updates. In *Proceedings of the 38th ACM SIGPLAN Conference on Programming Language Design and Implementation*. ACM, Barcelona Spain, 556–571. doi:[10.1145/3062341.3062354](https://doi.org/10.1145/3062341.3062354)
- 1094 [19] Troels Henriksen, Niels G. W. Serup, Martin Elsman, Fritz Henglein, and Cosmin E. Oancea. 2017. Futhark: purely functional GPU-programming with nested parallelism and in-place array updates. *SIGPLAN Not.* 52, 6 (June 2017), 556–571. doi:[10.1145/3140587.3062354](https://doi.org/10.1145/3140587.3062354)
- 1095 [20] Bastian Köpcke, Sergei Gorlatch, and Michel Steuwer. 2024. Descend: A Safe GPU Systems Programming Language. 8 (2024), 841–864. Issue PLDI. doi:[10.1145/3656411](https://doi.org/10.1145/3656411)
- 1096 [21] NVIDIA Corporation. 2025. *CUB: CUDA Unbound*. <https://docs.nvidia.com/cuda/cub/index.html> CUDA Toolkit Documentation.
- 1097 [22] NVIDIA Corporation. 2025. *CUDA C++ Programming Guide*. NVIDIA. <https://docs.nvidia.com/cuda/cuda-c-programming-guide/> Accessed: 2025-11-06.
- 1098 [23] PyTorch Team at Meta. 2025. Helion: A High-Level DSL for Performant and Portable ML Kernels. <https://pytorch.org/blog/helion/>. Accessed: 2025-11-10.
- 1099 [24] Jonathan Ragan-Kelley, Connnelly Barnes, Andrew Adams, Sylvain Paris, Frédo Durand, and Saman Amarasinghe. 2013. Halide: a language and compiler for optimizing parallelism, locality, and recomputation in image processing pipelines. *SIGPLAN Not.* 48, 6 (June 2013), 519–530. doi:[10.1145/2499370.2462176](https://doi.org/10.1145/2499370.2462176)
- 1100 [25] Benjamin Frederick Spector, Simran Arora, Aaryan Singhal, Arjun Parthasarathy, Daniel Y Fu, and Christopher Re. 2025. ThunderKittens: Simple, Fast, and \$textrit{Adorable}\\$ Kernels. In *The Thirteenth International Conference on Learning Representations*. <https://openreview.net/forum?id=0fjfVOSUra>
- 1101 [26] Simon Spies, Lennard Gähler, Daniel Gratzer, Joseph Tassarotti, Robbert Krebbers, Derek Dreyer, and Lars Birkedal. 2021. Transfinite Iris: resolving an existential dilemma of step-indexed separation logic. In *Proceedings of the 42nd ACM SIGPLAN International Conference on Programming Language Design and Implementation*. ACM, Virtual Canada, 80–95. doi:[10.1145/3453483.3454031](https://doi.org/10.1145/3453483.3454031)
- 1102 [27] Simon Spies, Neel Krishnaswami, and Derek Dreyer. 2021. Transfinite step-indexing for termination. *Proceedings of the ACM on Programming Languages* 5, POPL (Jan. 2021), 1–29. doi:[10.1145/3434294](https://doi.org/10.1145/3434294)
- 1103 [28] Rust Team. 2025. Rust Programming Language. <https://rust-lang.org/>
- 1104 [29] Vijay Thakkar, Pradeep Ramani, Cris Cecka, Aniket Shivam, Honghao Lu, Ethan Yan, Jack Kosaian, Mark Hoemmen, Haicheng Wu, Andrew Kerr, Matt Nicely, Duane Merrill, Dustyn Blasig, Aditya Atluri, Fengqi Qiao, Piotr Majcher, Paul

- 1128 Springer, Markus Hohnerbach, Jin Wang, and Manish Gupta. 2023. *CUTLASS*. <https://github.com/NVIDIA/cutlass>
- 1129 [30] The JAX Authors. 2024. Pallas:Mosaic GPU – A JAX Kernel Language for GPUs. <https://docs.jax.dev/en/latest/pallas/gpu/index.html>. Accessed: 2025-11-10.
- 1130 [31] The Khronos Group Inc. 2025. OpenCL™ – The Open Standard for Parallel Programming of Heterogeneous Systems. <https://www.khronos.org/opencl/>. Accessed: 2025-11-10.
- 1131 [32] Philippe Tillet. 2025. Introducing Triton: Open-source GPU programming for neural networks. <https://openai.com/index/triton/>
- 1132 [33] Philippe Tillet, H. T. Kung, and David Cox. 2019. Triton: an intermediate language and compiler for tiled neural network computations. In *Proceedings of the 3rd ACM SIGPLAN International Workshop on Machine Learning and Programming Languages* (Phoenix, AZ, USA) (MAPL 2019). Association for Computing Machinery, New York, NY, USA, 10–19. doi:[10.1145/3315508.3329973](https://doi.org/10.1145/3315508.3329973)
- 1133 [34] triton-lang. 2025. “01-intro.py” – Introduction to Gluon tutorial in Triton. <https://github.com/triton-lang/triton/blob/main/python/tutorials/gluon/01-intro.py>. Accessed: 2025-11-10.
- 1134 [35] Aaron J. Turon, Jacob Thamsborg, Amal Ahmed, Lars Birkedal, and Derek Dreyer. 2013. Logical relations for fine-grained concurrency. In *Proceedings of the 40th annual ACM SIGPLAN-SIGACT symposium on Principles of programming languages* (Rome Italy, 2013-01-23). ACM, 343–356. doi:[10.1145/2429069.2429111](https://doi.org/10.1145/2429069.2429111)
- 1135 [36] Rohan Yadav, Michael Garland, Alex Aiken, and Michael Bauer. 2025. Task-Based Tensor Computations on Modern GPUs. *Proc. ACM Program. Lang.* 9, PLDI, Article 163 (June 2025), 25 pages. doi:[10.1145/3729262](https://doi.org/10.1145/3729262)
- 1136
- 1137
- 1138
- 1139
- 1140
- 1141
- 1142
- 1143
- 1144
- 1145
- 1146
- 1147
- 1148
- 1149
- 1150
- 1151
- 1152
- 1153
- 1154
- 1155
- 1156
- 1157
- 1158
- 1159
- 1160
- 1161
- 1162
- 1163
- 1164
- 1165
- 1166
- 1167
- 1168
- 1169
- 1170
- 1171
- 1172
- 1173
- 1174
- 1175
- 1176

1177 **A Complete Prism Type System, Semantics, and Syntactic Soundness Proofs**

1178 **A.1 Basic Definitions**

1180                  Hierarchy Levels  $h ::= \text{Grid} \mid \text{Block} \mid \text{Thread}$   
 1181                  Memory Kinds  $l ::= \text{Local} \mid \text{Shared} \mid \text{Global}$   
 1182                  Perspectives  $\pi : h \times \mathbb{N}$   
 1183                  Base Types  $b ::= \text{bool} \mid \text{int} \mid \text{float}$   
 1184                  Types  $\tau ::= b \mid b[]^l \mid \text{Fun}(\Gamma, \pi, m) \mid \text{async } \tau$   
 1185                  Contexts  $\Gamma ::= \cdot \mid \Gamma, x : \pi \tau$   
 1186                  Shared Memory Remaining  $m : \mathbb{N}$   
 1187  
 1188

1190                  Perspectives  $\pi$  are part of an algebra parameterized over some constant values  $T$  (the number of  
 1191 threads per block) and  $B$  (the number of blocks per grid). With these values, we have two isomorphisms:  
 1192

$$( \text{Block}, 1 ) \cong ( \text{Thread}, T )$$

1193 and

$$( \text{Grid}, 1 ) \cong ( \text{Block}, B )$$

1197                  The **group** and **split** operations of Parallax allows us to move along this isomorphism, from left  
 1198 to right. For clarity, in Prism we split these operations into three: a **split** operation that can split  
 1199 perspectives into multiple narrower ones, and a **destruct** operation that directly moves us along  
 1200 the isomorphism, and a **group** that divides our current perspective into equal sized parts.

1201                  Perspectives  $\pi$  are also lexicographically ordered in the obvious way.  $hs$  are ordered such that  
 1202  $\text{Thread} \leq \text{Block} \leq \text{Grid}$ , and  $(h_1, n_1) \leq (h_2, n_2)$  iff  $n_1 | n_2$  and  $h_1 \leq h_2$ .

1203                  We define scalar multiplication  $i \times h$  of natural numbers with  $hs$ :  $i \times (h, n) = (h, in)$ .

1204                  We also define division of perspectives and hierarchy levels of type  $\pi \times \pi \rightarrow \mathbb{N}$ .  $\text{Grid}/\text{Block} = B$   
 1205 and  $\text{Block}/\text{Thread} = T$ . We lift this to perspectives like so:  $(h_1, n_1) / (h_2, n_2) = ((h_1/h_2) \cdot n_1) / n_2$ .

1206                  Lastly we define a partial  $\downarrow$  operator on  $hs$  such that  $\downarrow \text{Block} = \text{Thread}$  and  $\downarrow \text{Grid} = \text{Block}$ . Note that  
 1207  $\downarrow \text{Thread}$  is undefined. This operator lifts to  $\pi s$  whose second component is 1 and encodes the leftward  
 1208 component of the isomorphism above:  $\downarrow (\text{Block}, 1) = (\text{Thread}, T)$  and  $\downarrow (\text{Grid}, 1) = (\text{Block}, B)$ .

1209                  Note also that the presentation of these rules in the main body of the paper elide the  $m$  portion,  
 1210 which tracks the maximum amount of memory a given computation is allowed to use. In the full  
 1211 system presented here, both the typing rules and the operational semantics carry an additional piece  
 1212 of information tracking allocated memory.

1214 **A.2 Complete Prism Typing Rules**

1215 **A.2.1 Expressions.**

$$\frac{x : \pi \tau \in \Gamma}{\Gamma \vdash^\pi x : \tau} \text{-Var} \quad \frac{}{\Gamma \vdash^\pi n : \text{int}} \text{-Int} \quad \frac{}{\Gamma \vdash^\pi f : \text{float}} \text{-Float}$$

$$\frac{}{\Gamma \vdash^\pi b : \text{bool}} \text{-Bool} \quad \frac{\pi < (\text{Grid}, 1)}{\Gamma \vdash^\pi \text{partition\_id} : \text{int}} \text{-Partition-Id}$$

$$\frac{\Gamma \vdash^\pi' e_1 : \tau[]^l \quad \Gamma \vdash^\pi e_2 : \text{int} \quad l = \text{Global} \text{ or } l = \text{Local} \quad \pi \leq \pi'}{\Gamma \vdash^\pi e_1[e_2] : \tau} \text{-Arr-Access}$$

1226	$\frac{\Gamma \vdash^{\pi'} e_1 : \tau[]^{\text{Shared}} \quad \Gamma \vdash^{\pi} e_2 : \text{int} \quad \pi \leq (\text{Block}, 1) \quad \pi \leq \pi'}{\Gamma \vdash^{\pi} e_1[e_2] : \tau}$	T-Arr-Access-Shared
1227		
1228	$\frac{\Gamma \vdash^{\pi} e_1 : \text{int} \quad \Gamma \vdash^{\pi} e_2 : \text{int}}{\Gamma \vdash^{\pi} e_1 \text{ bop } e_2 : \text{int}}$	T-Bop
1229	$\frac{\Gamma \vdash^{\pi} e_1 : \text{int} \quad \Gamma \vdash^{\pi} e_2 : \text{int}}{\Gamma \vdash^{\pi} e_1 \text{ cmp } e_2 : \text{bool}}$	T-Cmp
1230		
1231		
1232	A.2.2 <i>Statements</i> .	
1233	$\frac{f : \pi' \text{Fun}(x_i : \tau_i, \pi, m') \in \Gamma \quad \Gamma \vdash^{\pi} e_i : \tau_i \quad m' \leq m}{\Gamma \vdash_m^{\pi} f(e_1, \dots, e_n)}$	T-Function-Call
1234		
1235		
1236	$\frac{\Gamma \vdash_m^{(h, n_1)} s_1 \quad \Gamma \vdash_m^{(h, n_2)} s_2 \quad n_1, n_2 \text{ align to } n}{\Gamma \vdash_m^{(h, n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\}}$	T-Split
1237		
1238	$\frac{\Gamma \vdash_m^{\downarrow \pi} s}{\Gamma \vdash_m^{\pi} \text{destruct in } s}$	T-Destruct
1239		
1240	$\frac{\Gamma \vdash_m^{(h, n)} s}{\Gamma \vdash_m^{(h, q \cdot n)} \text{group } q \ s}$	T-Group
1241	$\frac{x : \pi' \tau \in \Gamma \quad \Gamma \vdash^{\pi} e : \tau \quad \pi' \leq \pi}{\Gamma \vdash_m^{\pi} x = e}$	T-Assn
1242		
1243	$\frac{\Gamma \vdash_m^{\pi} \text{init}_{\psi}}{\Gamma \vdash_m^{\pi} \text{dec}_{\psi}}$	T-Sync-Init
1244		
1245	$\frac{\Gamma \vdash_m^{\pi} \text{wait}_{\psi}}{\Gamma \vdash_m^{\pi} \text{Sync-Wait}}$	T-Sync-Wait
1246	$\frac{\Gamma \vdash_m^{\pi} \text{skip}}{\Gamma \vdash_m^{\pi} \text{Sync-Skip}}$	T-Skip
1247	$\frac{\Gamma \vdash^{\pi} e : \tau \quad \Gamma, x : \pi' \tau \vdash_m^{\pi} s \quad \pi' \leq \pi \quad \tau \text{ not an array type}}{\Gamma \vdash_m^{\pi} x : \tau @ \pi' = e \text{ in } s}$	T-Decl
1248		
1249		
1250	$\frac{\Gamma \vdash^{\pi'} e_1 : \tau[]^l \quad \Gamma \vdash^{\pi} e_2 : \text{int} \quad \Gamma \vdash^{\pi} e_3 : \tau \quad l = \text{Global or } l = \text{Local} \quad \pi' \leq \pi}{\Gamma \vdash_m^{\pi} e_1[e_2] = e_3}$	T-Arr-Assn
1251		
1252		
1253	$\frac{\Gamma \vdash^{\pi'} e_1 : \tau[]^{\text{Shared}} \quad \Gamma \vdash^{\pi} e_2 : \text{int} \quad \Gamma \vdash^{\pi} e_3 : \tau \quad \pi \leq (\text{Block}, 1) \quad \pi' \leq \pi}{\Gamma \vdash_m^{\pi} e_1[e_2] = e_3}$	T-Arr-Assn-Shared
1254		
1255		
1256	$\frac{\Gamma \vdash^{\pi} e : \text{bool} \quad \Gamma \vdash_{m_1}^{\pi} s_1 \quad \Gamma \vdash_{m_2}^{\pi} s_2}{\Gamma \vdash_{\max(m_1, m_2)}^{\pi} \text{if } e \text{ then } s_1 \text{ else } s_2}$	T-If
1257		
1258		
1259	$\frac{\Gamma \vdash^{\pi} e : \text{bool} \quad \Gamma \vdash_m^{\pi} s}{\Gamma \vdash_m^{\pi} \text{while } e \text{ do } s}$	T-While
1260	$\frac{\Gamma \vdash_m^{\pi} s_1 \quad \Gamma \vdash_m^{\pi} s_2}{\Gamma \vdash_{\max(m_1, m_2)}^{\pi} s_1; s_2}$	T-Seq
1261		
1262	$\frac{\Gamma, x : \pi' \tau[]^l \vdash_m^{\pi} s \quad l = \text{Global or } l = \text{Local}}{\Gamma \vdash_{m+n \cdot \text{size}(\tau)}^{\pi} x := \text{alloc } l \ \tau \ n \text{ in } s}$	T-Alloc
1263		
1264		
1265	$\frac{\Gamma, x : (\text{Block}, 1) \tau[]^{\text{Shared}} \vdash_m^{\pi} s}{\Gamma \vdash_{m+n \cdot \text{size}(\tau)}^{\pi} x := \text{alloc Shared } \tau \ n \text{ in } s}$	T-Alloc-Shared
1266		
1267		
1268	$\frac{\Gamma, y : (h, n/c) \tau[]^l \vdash_m^{(h, n)} s \quad c   n \quad l \neq \text{Local}}{\Gamma, x : (h, n) \tau[]^l \vdash_m^{(h, n)} \text{partition}_{\psi} x \text{ into } y \text{ by } c \text{ in } s}$	T-Partition
1269		
1270		
1271		
1272	$\frac{\Gamma, y : (h, n') \tau[]^l \vdash_m^{(h, n')} s \quad n' \leq n \quad l \neq \text{Local}}{\Gamma, x : (h, n) \tau[]^l \vdash_m^{(h, n)} \text{claim}_{\psi} x \text{ into } y \text{ at } n' \text{ in } s}$	T-Claim
1273		
1274		

$$\begin{array}{c}
 \frac{\Gamma, y : \downarrow^\pi \tau[]^l \vdash_m^\pi s \quad l \neq \text{Local}}{\Gamma, x : {}^\pi \tau[]^l \vdash_m^\pi \text{lower}_\psi x \text{ into } y \text{ in } s} \text{-Lower} \\
 \\ 
 \frac{\Gamma, y : {}^{(\text{Thread},1)} \text{async } \tau[]^l \vdash_m^{(\text{Thread},1)} s}{\Gamma, x : {}^{(\text{Thread},1)} \tau[]^l \vdash_m^{(\text{Thread},1)} \text{async\_partition}_\phi x \text{ into } y \text{ in } s} \text{-Async-Partition} \\
 \\ 
 \frac{\Gamma, x : {}^{(\text{Thread},1)} \text{async } \tau[], y : {}^{(\text{Thread},1)} \tau[]^l' \vdash_m^{(\text{Thread},1)} \text{async\_memcpy}(x,y)}{\Gamma, x : {}^\pi \tau[]^l, y : {}^\pi \tau[]^l' \vdash_m^\pi \text{memcpy}(x,y)} \text{-Async-Memcpy} \\
 \\ 
 \frac{}{\Gamma, x : {}^\pi \tau[]^l, y : {}^\pi \tau[]^l' \vdash_m^\pi \text{memcpy}(x,y)} \text{-Memcpy}
 \end{array}$$

### A.3 Complete Prism Semantics

#### A.3.1 Definitions.

Global Memory  $\Sigma ::= \cdot | \Sigma, n \mapsto^\pi v$

Shared Memory  $\sigma ::= \cdot | \sigma, n \mapsto^\pi v$

Local Memory  $\eta ::= \cdot | \eta, n \mapsto^\pi v$

Block Memory Map  $S ::= \forall n \in B, n \mapsto \sigma$

Thread Memory Map  $L ::= \forall n \in T, n \mapsto \eta$

Synchronization Map  $\Psi : \psi \rightarrow \mathbb{N} \rightarrow \mathbb{N}$

Deferred Computations Map  $\Phi : \phi \rightarrow \{s\}$

In real GPUs, thread IDs are only unique within their block. However, in this calculus for simplicity we assume thread IDs are global. One can convert back and forth between this abstracted notion of a thread ID and a block-unique ID via addition modulo  $T$ . That is,  $t_{\text{real}} = t_{\text{simplified}} \bmod T$  and  $t_{\text{simplified}} = t_{\text{real}} + b \cdot T$ .

By convention the names for local and shared and global memory do not conflict, as on the GPU they will be separate pointer spaces. Additionally, we freely interchange between using names for variables and integer locations.

In the main body of the paper, for simplicity we elide the synchronization map and the deferred computation map from the operational semantics, as our theorems do not make any guarantees about non-interference. However, as they are part of the full semantics, we include them here for completeness. By convention the synchronization and deferred computation maps are a total functions, initialized to map to  $\lambda_. 0$  for  $\psi$ s and  $\lambda_. \{\}$  for  $\phi$ s not explicitly initialize.

The shape of the judgment for a single thread is  $\eta, \sigma, \Sigma, t, b, p, \Psi \vdash_m^\pi s \rightsquigarrow s' \dashv_m^\pi \eta', \sigma', \Sigma', \Psi'$ . The  $t$  here is the thread ID, the  $b$  is the block ID, and the  $p$  is the perspective ID. The last of these three is modified and managed by the rules for **split**, **group** and **destruct**, and tracks the relative position of the thread within a group. This semantics is in a small step style.

The shape of the judgment for expressions is  $\eta, \sigma, \Sigma \vdash_\pi^\pi e \Downarrow v$ . The two  $\pi$ s represent the ambient compute context (i.e., the context in which resources are being read), while  $\pi'$ , represents the target compute context (i.e., the compute context of the variable into which the result of the expression is going to be written. This is relevant for computing the value of **partition\_id**, which divides the two contexts. As a shorthand, we can divide a perspective by a scalar value like so:  $(h,n)/c = (h,n)/(h,c)$ .

1324 The overall evaluation of a program is expressed as

$$L, S, \Sigma, P, \Psi, \Phi \rightsquigarrow L', S', \Sigma', P', \Psi', \Phi'.$$

1325  
1326 In this judgment  $P$  serves as a *thread pool*, mapping pairs of thread and block IDs (which don't change)  
1327 to statements and memory (which can be updated by stepping). One can think of  $P$  as tracking which  
1328 program is running on each thread. This steps according to the following rule:

$$\frac{L(t), S(b), \Sigma, t, b, 0, \Psi, \Phi \vdash_m^{(\text{Grid}, 1)} s \rightsquigarrow s' \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi' \quad P(t, b) = (s, m)}{L, S, \Sigma, P, \Psi, \Phi \rightsquigarrow L[t \mapsto \eta'], S[b \mapsto \sigma'], \Sigma', P[(t, b) \mapsto (s', m')], \Psi', \Phi'} \text{ S-Program}$$

1329  
1330 For simplicity of notation, we define an **update** operation that searches the three environments  
1331 for the one that contains the variable being used (by convention, there is no conflict between the  
1332 environments, as in reality they exist in three separate address spaces). We also define a similar **get**  
1333 operation that retrieves a variable from memory, and a **rename** operation that remaps a variable  
1334 with the same value but under a different name.

$$\begin{aligned} \text{update}(\eta, \sigma, \Sigma, x, v) &= (\eta[x \mapsto^\pi v], \sigma, \Sigma) \text{ when } x \in^\pi \eta \\ \text{update}(\eta, \sigma, \Sigma, x, v) &= (\eta, \sigma[x \mapsto^\pi v], \Sigma) \text{ when } x \in^\pi \sigma \\ \text{update}(\eta, \sigma, \Sigma, x, v) &= (\eta, \sigma, \Sigma[x \mapsto^\pi v]) \text{ when } x \in^\pi \Sigma \end{aligned}$$

$$\begin{aligned} \text{get}(\eta, \sigma, \Sigma, x) &= \eta(x) \text{ when } x \in^\pi \eta \\ \text{get}(\eta, \sigma, \Sigma, x) &= \sigma(x) \text{ when } x \in^\pi \sigma \\ \text{get}(\eta, \sigma, \Sigma, x) &= \Sigma(x) \text{ when } x \in^\pi \Sigma \end{aligned}$$

$$\begin{aligned} \text{rename}(\eta, \sigma, \Sigma, x, y, \pi') &= (\eta[y \mapsto^\pi \eta(x)], \sigma, \Sigma) \text{ when } x \in^\pi \eta \\ \text{rename}(\eta, \sigma, \Sigma, x, y, \pi') &= (\eta, \sigma[y \mapsto^\pi \sigma(x)], \Sigma) \text{ when } x \in^\pi \sigma \\ \text{rename}(\eta, \sigma, \Sigma, x, y, \pi') &= (\eta, \sigma, \Sigma[y \mapsto^\pi \Sigma(x)]) \text{ when } x \in^\pi \Sigma \end{aligned}$$

### A.3.2 Perspective Management Rules.

$$\frac{p < n_1 \quad n_1, n_2 \text{ align to } n \quad \eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n_1)} s_1 \rightsquigarrow s'_1 \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi'}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\} \rightsquigarrow \text{split}(n_1, n_2)\{s'_1\}\{s_2\} \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi'} \text{ S-Split-Left}$$

$$\frac{p < n_1 \quad n_1, n_2 \text{ align to } n}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n)} \text{split}(n_1, n_2)\{\text{skip}\}\{s_2\} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-Split-Left-Done}$$

$$\frac{p \geq n_1 \quad p < n_1 + n_2 \quad n_1, n_2 \text{ align to } n \quad \eta, \sigma, \Sigma, t, b, p - n_1, \Psi, \Phi \vdash_m^{(h, n_2)} s_2 \rightsquigarrow s'_2 \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi'}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\} \rightsquigarrow \text{split}(n_1, n_2)\{s_1\}\{s'_2\} \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi'} \text{ S-Split-Right}$$

$$\frac{p \geq n_1 \quad p < n_1 + n_2 \quad n_1, n_2 \text{ align to } n}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n)} \text{split}(n_1, n_2)\{s_1\}\{\text{skip}\} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-Split-Right-Done}$$

$$\frac{p \geq n_1 + n_2 \quad n_1, n_2 \text{ align to } n}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-Split-None}$$

$$\begin{array}{c}
1373 \\
1374 \quad \frac{\eta, \sigma, \Sigma, t, b, t \text{ mod } T, \Psi, \Phi \vdash_m^{(\text{Thread}, T)} s \rightsquigarrow s' \dashv_m \eta', \sigma', \Sigma', \Psi', \Phi'}{\eta, \sigma, \Sigma, t, b, 0, \Psi, \Phi \vdash_m^{(\text{Block}, 1)} \text{destruct in } s \rightsquigarrow \text{destruct in } s' \dashv_m \eta', \sigma', \Sigma', \Psi', \Phi'} \text{ S-Destruct-Block} \\
1375 \\
1376 \\
1377 \\
1378 \quad \frac{\eta, \sigma, \Sigma, t, b, b \text{ mod } B, \Psi, \Phi \vdash_m^{(\text{Block}, B)} s \rightsquigarrow s' \dashv_m \eta', \sigma', \Sigma', \Psi', \Phi'}{\eta, \sigma, \Sigma, t, b, 0, \Psi, \Phi \vdash_m^{(\text{Grid}, 1)} \text{destruct in } s \rightsquigarrow \text{destruct in } s' \dashv_m \eta', \sigma', \Sigma', \Psi', \Phi'} \text{ S-Destruct-Grid} \\
1379 \\
1380 \\
1381 \quad \frac{\eta, \sigma, \Sigma, t, b, 0, \Psi, \Phi \vdash_m^{\pi} \text{destruct in } \text{skip} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi}{\eta, \sigma, \Sigma, t, b, 0, \Psi, \Phi \vdash_m^{\pi} \text{destruct in skip} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-Destruct-Done} \\
1382 \\
1383 \\
1384 \quad \frac{\eta, \sigma, \Sigma, t, b, p \text{ mod } n, \Psi, \Phi \vdash_m^{(h, n)} s \rightsquigarrow s' \dashv_m \eta', \sigma', \Sigma', \Psi', \Phi'}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, q \cdot n)} \text{group } q \ s \rightsquigarrow \text{group } q \ s' \dashv_m \eta', \sigma', \Sigma', \Psi', \Phi'} \text{ S-Group} \\
1385 \\
1386 \quad \frac{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{group } q \ \text{skip} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{group } q \ \text{skip} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-Group-Done} \\
1387 \\
1388
\end{array}$$

A.3.3 *Thread Synchronization.* We define a **size** operation on perspectives to compute the size of a perspective (the number of individual compute resources sharing it). The operation is defined thusly:

$$\begin{array}{c}
1391 \\
1392 \quad \text{size(Thread, } n) = n \\
1393 \quad \text{size(Block, } n) = n \cdot T \\
1394 \quad \text{size(Grid, } n) = n \cdot B \cdot T \\
1395 \\
1396 \\
1397 \\
1398 \quad \frac{\Psi(\psi)(p) = 0}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{wait}_{\psi} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-Sync-Wait-Done} \\
1399 \\
1400 \\
1401 \quad \frac{\Psi(\psi)(p) \neq 0}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{wait}_{\psi} \rightsquigarrow \text{wait}_{\psi} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-Sync-Wait-Spin} \\
1402 \\
1403 \\
1404 \quad \frac{\Psi' = \Psi(\psi)[p \mapsto \Psi(\psi)(p) - 1]}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{dec}_{\psi} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-Sync-Dec} \\
1405 \\
1406 \\
1407 \quad \frac{\Psi(\psi)(p) = 0 \quad \Psi' = \Psi(\psi)[p \mapsto \text{size}(\pi)]}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{init}_{\psi} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-Sync-Init-Zero} \\
1408 \\
1409 \\
1410 \quad \frac{\Psi(\psi)(p) \neq 0}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{init}_{\psi} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-Sync-Init-Nonzero} \\
1411 \\
1412
\end{array}$$

A.3.4 *Asynchrony.*

$$\begin{array}{c}
1413 \quad \frac{\text{rename}(\eta, \sigma, \Sigma, x, y, (\text{Thread}, 1)), t, b, p, \Psi, \Phi \vdash_m^{(\text{Thread}, 1)} s \rightsquigarrow s' \dashv_m \eta', \sigma', \Sigma', \Psi', \Phi'}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_{m'}^{(\text{Thread}, 1)} \text{async\_partition}_{\phi} x \text{ into } y \text{ in } s \rightsquigarrow \text{async\_partition}_{\phi} x \text{ into } y \text{ in } s' \dashv_m \eta', \sigma', \Sigma', \Psi', \Phi'} \text{ S-Async-Partition-Congr} \\
1414 \\
1415 \\
1416 \\
1417 \\
1418 \quad \frac{\Phi = \Phi'[\phi \mapsto \Phi'(\phi) \cup \{s\}]}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(\text{Thread}, 1)} \text{async\_partition}_{\phi} x \text{ into } y \text{ in skip} \rightsquigarrow (\text{async\_partition}_{\phi} x \text{ into } y \text{ in } s) \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-Async-Partition-Unwind} \\
1419 \\
1420 \\
1421
\end{array}$$

1422		$\Phi(\phi) = \emptyset$		
1423			<b>S-Async-Partition-Done</b>	
1424	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(\text{Thread}, 1)}$	<b>async_partition</b> $_{\phi} x \text{ into } y \text{ in } \text{skip} \rightsquigarrow$		
1425		<b>skip</b> $\dashv_m \eta, \sigma, \Sigma, \Psi, \Phi$		
1426				
1427			<b>S-Async-Memcpy</b>	
1428	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(\text{Thread}, 1)}$	<b>async_memcpy</b> $(x, y) \rightsquigarrow$		
1429		<b>skip</b> $\dashv_m \eta, \sigma, \Sigma, \Psi, \Phi[\phi \mapsto \Phi(\phi) \cup \{\text{memcpy}(x, y)\}]$		
1430				
1431		$(\eta', \sigma', \Sigma') = \text{update}(\eta, \sigma, \Sigma, x, \text{get}(\eta, \sigma, \Sigma, y))$		
1432	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{memcpy}(x, y) \rightsquigarrow \text{skip} \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi$	<b>S-Memcpy</b>		
1433				
1434	<i>A.3.5 Variables and Memory.</i>			
1435		$\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi} e \Downarrow v \quad \pi' \leq \pi$		
1436	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} x : \tau @ \pi' := e \text{ in } s \rightsquigarrow s \dashv_m \eta[x \mapsto^{\pi'} v], \sigma, \Sigma, \Psi, \Phi$	<b>S-Decl</b>		
1437				
1438			<b>S-Free</b>	
1439	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{free } n \dashv_{m-n} \eta, \sigma, \Sigma, \Psi, \Phi$			
1440				
1441			<b>S-Alloc-Local</b>	
1442	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} x := \text{alloc Local } \tau n \text{ in } s \rightsquigarrow$			
1443	$s; \text{free } (n \cdot \text{size}(\tau)) \dashv_{m+n \cdot \text{size}(\tau)} \eta[x \mapsto^{\pi} \langle x, n \rangle], \sigma, \Sigma, \Psi, \Phi$			
1444				
1445		$\pi = (\text{Block}, 1)$		
1446	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} x := \text{alloc Shared } \tau n \text{ in } s \rightsquigarrow$	<b>S-Alloc-Shared</b>		
1447	$s; \text{free } (n \cdot \text{size}(\tau)) \dashv_{m+n \cdot \text{size}(\tau)} \eta, \sigma[x \mapsto^{\pi} \langle x, n \rangle], \Sigma, \Psi, \Phi$			
1448				
1449			<b>S-Alloc-Global</b>	
1450	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} x := \text{alloc Global } \tau n \text{ in } s \rightsquigarrow$			
1451	$s; \text{free } (n \cdot \text{size}(\tau)) \dashv_{m+n \cdot \text{size}(\tau)} \eta, \sigma, \Sigma[x \mapsto^{\pi} \langle x, n \rangle], \Psi, \Phi$			
1452	$x \in \pi' \quad \eta, \sigma, \Sigma \quad \eta, \sigma, \Sigma \vdash_{\pi'}^{\pi} e \Downarrow v \quad (\eta', \sigma', \Sigma') = \text{update}(\eta, \sigma, \Sigma, v) \quad \pi' \leq \pi$			
1453	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi, \Phi \vdash_m^{\pi} x = e \rightsquigarrow \text{skip} \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi$	<b>S-Assn</b>		
1454				
1455				
1456	$\eta, \sigma, \Sigma, t, b, p \vdash_{\pi}^{\pi} e_1 \Downarrow \langle l, n \rangle$	$i < n \quad \pi' \leq \pi$		
1457	$\eta, \sigma, \Sigma, t, b, p \vdash_{\pi}^{\pi} e_2 \Downarrow i$	$x \in \pi' \quad \eta, \sigma, \Sigma$		
1458	$\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi} e_3 \Downarrow v \quad (\eta', \sigma', \Sigma') = \text{update}(\eta, \sigma, \Sigma, l+i, v)$			
1459	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} e_1[e_2] = e_3 \rightsquigarrow \text{skip} \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi$	<b>S-Arr-Assn</b>		
1460				
1461	$s' = s[(y + c \cdot p)/y] \quad (\eta', \sigma', \Sigma') = \text{rename}(\eta, \sigma, \Sigma, x, y, \pi/c)$			
1462	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{partition}_{\psi} x y c s \rightsquigarrow \text{init}_{\psi}; s'; \text{dec}_{\psi}; \text{wait}_{\psi} \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi$	<b>S-Partition</b>		
1463				
1464		$(\eta', \sigma', \Sigma') = \text{rename}(\eta, \sigma, \Sigma, x, y, (h, n_1))$	<b>S-Claim</b>	
1465	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n_1 + n_2)} \text{claim}_{\psi} x \text{ into } y \text{ at } n_1 \text{ in } s \rightsquigarrow$			
1466	$\text{init}_{\psi}; \text{split}(n_1, n_2); s' \{ \text{skip} \}; \text{dec}_{\psi}; \text{wait}_{\psi} \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi$			
1467				
1468		$(\eta', \sigma', \Sigma') = \text{rename}(\eta, \sigma, \Sigma, x, y, \downarrow \pi)$		
1469	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{lower}_{\psi} x \text{ into } y \text{ in } s \rightsquigarrow \text{init}_{\psi}; s; \text{dec}_{\psi}; \text{wait}_{\psi} \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi$	<b>S-Lower</b>		
1470				

## 1471 A.3.6 Control Flow.

$$\begin{array}{c}
 \frac{\eta, \sigma, \Sigma \vdash_{\pi}^{\pi} e \Downarrow \text{true}}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{if } e \text{ then } s_1 \text{ else } s_2 \rightsquigarrow s_1 \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-If-True} \\
 \\
 \frac{\eta, \sigma, \Sigma \vdash_{\pi}^{\pi} e \Downarrow \text{false}}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{if } e \text{ then } s_1 \text{ else } s_2 \rightsquigarrow s_2 \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-If-False} \\
 \\
 \frac{}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{while } e \text{ do } s \rightsquigarrow \text{if } e \text{ then } (s; \text{while } e \text{ do } s) \text{ else skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-While} \\
 \\
 \frac{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} s_1 \rightsquigarrow s'_1 \dashv_{m'}^{\pi'} \eta', \sigma', \Sigma', \Psi', \Phi'}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} s_1; s_2 \rightsquigarrow s'_1; s_2 \dashv_{m'}^{\pi'} \eta', \sigma', \Sigma', \Psi', \Phi'} \text{ S-Seq-First} \\
 \\
 \frac{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{skip}; s_2 \rightsquigarrow s_2 \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{skip}} \text{ S-Seq-Done} \\
 \\
 \frac{\Sigma(f) = \{[x_1 : \tau_1, \dots, x_n : \tau_n], s\} \quad \sigma, \Sigma \vdash_{\pi}^{\pi} e_i \Downarrow v_i \quad m' \leq m}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} f(e_1, \dots, e_n) \rightsquigarrow s \dashv_m \eta[x_i \mapsto v_i], \sigma, \Sigma, \Psi, \Phi} \text{ S-Function-Call}
 \end{array}$$

## 1486 A.3.7 Expressions.

$$\begin{array}{c}
 \frac{\pi < (\text{Grid}, 1) \quad \pi' \leq \pi}{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi} \text{partition\_id} \Downarrow \pi / \pi' - 1} \text{ E-Partition-Id} \\
 \\
 \frac{}{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi} x \Downarrow \text{get}(\eta, \sigma, \Sigma, x)} \text{ E-Var} \\
 \\
 \frac{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi} e_1 \Downarrow \langle l, n \rangle \quad \eta, \sigma, \Sigma \vdash_{\pi'}^{\pi} e_2 \Downarrow i \quad i < n \quad \pi' \leq \pi}{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi} e_1[e_2] \Downarrow \text{get}(\eta, \sigma, \Sigma, l+i)} \text{ E-Arr-Access} \\
 \\
 \frac{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi} \text{int} \Downarrow n}{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi} n \Downarrow n} \text{ E-Int} \quad \frac{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi} b \Downarrow b}{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi} b} \text{ E-Bool} \\
 \\
 \frac{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi} e_1 \Downarrow v_1 \quad \eta, \sigma, \Sigma \vdash_{\pi'}^{\pi} e_2 \Downarrow v_2 \quad v = v_1 \text{ bop } v_2}{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi} e_1 \text{ bop } e_2 \Downarrow v} \text{ E-Bop} \\
 \\
 \frac{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi} e_1 \Downarrow v_1 \quad \eta, \sigma, \Sigma \vdash_{\pi'}^{\pi} e_2 \Downarrow v_2 \quad v = v_1 \text{ cmp } v_2}{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi} e_1 \text{ cmp } e_2 \Downarrow v} \text{ E-Cmp}
 \end{array}$$

## 1505 A.4 Theorems and Proofs

1506 Note that in this section we assume no out of bounds array accesses. In general Parallax (and by  
 1507 extension Prism) makes no guarantees about array out of bounds.

1508 A.4.1 More definitions. As a premise to our type safety theorems, we need to assume we have a  
 1509 well-typed environment, written  $\Gamma \vdash \eta, \sigma, \Sigma$ . We define what this means inductively

$$\begin{array}{c}
 \frac{}{\eta, \sigma, \Sigma \vdash n : \text{int}} \text{ V-Int} \quad \frac{}{\eta, \sigma, \Sigma \vdash b : \text{bool}} \text{ V-Bool} \quad \frac{}{\eta, \sigma, \Sigma \vdash f : \text{float}} \text{ V-Float} \\
 \\
 \frac{\forall i < n, \eta, \sigma, \Sigma \vdash \text{get}(\eta, \sigma, \Sigma, x + i) : \tau}{\eta, \sigma, \Sigma \vdash \langle x, n \rangle : \tau[l]} \text{ V-Array} \quad \frac{\Gamma, x_i : \pi \tau_i \vdash_m^{\pi} s \quad \Gamma \vdash \cdot, \cdot, \text{fns } \Sigma}{\eta, \sigma, \Sigma \vdash \{x_i : \tau_i, s\} : \text{Fun}(x_i : \tau_i, \pi, m)} \text{ V-Function} \\
 \\
 \frac{}{\cdot \vdash \eta, \sigma, \Sigma} \text{ G-Empty} \quad \frac{\eta(x) =^{\pi} v \quad \eta, \sigma, \Sigma \vdash v : \text{int}}{\Gamma, x : \pi \text{ int} \vdash \eta, \sigma, \Sigma} \text{ G-Int}
 \end{array}$$

$$\begin{array}{c}
 \frac{\eta(x) =^\pi v \quad \eta, \sigma, \Sigma \vdash v : \text{bool}}{\Gamma, x :^\pi \text{bool} \vdash \eta, \sigma, \Sigma} \text{ G-Bool} \quad \frac{\eta(x) =^\pi v \quad \eta, \sigma, \Sigma \vdash v : \text{float}}{\Gamma, x :^\pi \text{float} \vdash \eta, \sigma, \Sigma} \text{ G-Float} \\
 \\ 
 \frac{\eta(x) =^\pi v \quad \eta, \sigma, \Sigma \vdash v : \tau[]}{\Gamma, x :^\pi \tau[]^{\text{Local}} \vdash \eta, \sigma, \Sigma} \text{ G-Local} \quad \frac{\sigma(x) =^\pi v \quad \eta, \sigma, \Sigma \vdash v : \tau[]}{\Gamma, x :^\pi \tau[]^{\text{Shared}} \vdash \eta, \sigma, \Sigma} \text{ G-Shared} \\
 \\ 
 \frac{\Sigma(x) =^\pi v \quad \eta, \sigma, \Sigma \vdash v : \tau[]}{\Gamma, x :^\pi \tau[]^{\text{Global}} \vdash \eta, \sigma, \Sigma} \text{ G-Global} \quad \frac{\Sigma(x) =^\pi v \quad \eta, \sigma, \Sigma \vdash v : \text{Fun}(\Gamma', \pi, m)}{\Gamma, x :^\pi \text{Fun}(\Gamma', \pi, m) \vdash \eta, \sigma, \Sigma} \text{ G-Function}
 \end{array}$$

We can prove a couple simple lemmas about well-typed environments under operations like **rename**, **update**, and **get**.

LEMMA A.1. (*Well-typed get*) If  $\Gamma \vdash \eta, \sigma, \Sigma$  and  $x :^\pi \tau \in \Gamma$  then  $\eta, \sigma, \Sigma \vdash \text{get}(\eta, \sigma, \Sigma, x) : \tau$ .

LEMMA A.2. (*Well-typed rename*) If  $\Gamma, x :^\pi \tau \vdash \eta, \sigma, \Sigma$  then  $\Gamma, x :^\pi \tau, y : \pi' \tau \vdash \text{rename}(\eta, \sigma, \Sigma, x, y, \pi')$ .

LEMMA A.3. (*Well-typed update*) If  $\Gamma \vdash \eta, \sigma, \Sigma$  and  $\eta, \sigma, \Sigma \vdash v : \tau$  then  $\Gamma, x :^\pi \tau \vdash \text{update}(\eta, \sigma, \Sigma, x, v)$ .

We also define a well-formedness precondition on  $p$  with respect to  $\pi$ :

$$(h, n) \vdash p ::= p < n$$

We also define well-formedness for the async stack:

$$\Gamma \vdash \Phi ::= \forall \phi, s \in \Phi(\phi), \Gamma \vdash_m^{(\text{Thread}, 1)} s$$

## A.5 Proofs

LEMMA A.4. (*Expression Safety*) If  $\Gamma \vdash^\pi e : \tau$  and  $\Gamma \vdash \eta, \sigma, \Sigma$  and  $\pi \vdash p$  and  $\pi' \leq \pi$ , then there is some  $v$  such that  $\eta, \sigma, \Sigma \vdash_{\pi'} e \Downarrow v$  and  $v : \tau$ .

PROOF. This proof proceeds by induction on the typing relation for expressions. Despite the fact that this property implies termination, we do not need a logical relation to prove it because the expression language is very simple.

The **T-Int**, **T-Float**, and **T-Bool** cases are trivial, using the rules **E-Int**, **E-Bool** and **E-Float** to compute values. In the case for **T-Partition-Id**, the  $\pi$  premises of the typing rules and our assumption that  $\pi' \leq \pi$  match the premises of the evaluation rule, so this rule is simple as well.

The cases for **T-Bop** and **T-Cmp** follow directly from the inductive hypotheses, assuming a valid and correctly implemented set of binary operators and comparators.

The only interesting cases are **T-Arr-Access** and **T-Arr-Access-Shared**.

In both cases our inductive hypotheses and inversion give us that  $\pi' \leq \pi$ , and  $e_1$  evaluates to a  $\langle x, n \rangle$ , and that all the values between  $x$  and  $x + n$  in the appropriate environment are typed at  $\tau$ . We also know that  $e_2$  evaluates to an integer  $i$ . We assume that all array accesses are in bounds, so  $i < l$ , which is sufficient to use the **E-Arr-Access** rule to complete this case, and the proof.  $\square$

LEMMA A.5. (*Expression Determinism*) If  $\eta, \sigma, \Sigma, t, b, p \vdash_\pi^\pi e \Downarrow v_1$  and  $\eta, \sigma, \Sigma, t, b, p \vdash_\pi^\pi e \Downarrow v_2$  then  $v_1 = v_2$ .

PROOF. Straightforward by induction on the semantic derivation.  $\square$

LEMMA A.6. (*Expression Well-Typedness*) If  $\Gamma \vdash_{\pi'}^\pi e : \tau$  and  $\Gamma \vdash \eta, \sigma, \Sigma$  and  $\pi \vdash p$ , and  $\eta, \sigma, \Sigma, t, b, p \vdash_{\pi'}^\pi e \Downarrow v$ , then  $v : \tau$ .

PROOF. By our expression safety lemma our well-typed expression must evaluate to a well-typed value  $v'$ . By our determinism lemma  $v'$  must be the same as  $v$ , so  $v$  is well-typed.  $\square$

LEMMA A.7. (*Statement Progress*) If  $\Gamma \vdash \eta, \sigma, \Sigma$  and  $\Gamma \vdash_m^\pi s$  and  $\pi \vdash p$ , then either  $s$  is **skip** or there is some  $s'$  such that  $\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi s' \rightsquigarrow s' \vdash_m^{\pi'} \eta', \sigma', \Sigma, \Psi', \Phi'$

PROOF. This proceeds by induction on the typing derivation.

### Perspective Management Rules.

- Case **T-Split**:

In this case we have by our assumption that  $\pi \vdash p$  that  $p < n$ . We also have that  $n_1, n_2$  align to  $n$ , so  $n_1 + n_2 \leq n$ . There are three cases to consider, then: when  $p < n_1$ , when  $p \geq n_1$  and  $p < n_1 + n_2$ , and when  $p \geq n_1 + n_2$ .

In the first case, we have by our inductive hypothesis that  $s$  is either **skip** or that it can step in an  $(h, n_1)$  context. In the former case we can use the **S-Split-Left-Done** rule and in the latter we can use the **S-Split-Left rule**.

The second case is almost symmetric. The only additional work we have to do is to argue that  $(h, n_2) \vdash p - n_1$ , or equivalently that  $p - n_1 < n_2$ . This, however, is immediate from our assumption that  $p < n_1 + n_2$ .

In the last case, we just use the **S-Split-None** rule to step to **skip**.

- Case **T-Destruct**

By our inductive hypothesis, we know that  $s$  can step at  $\downarrow \pi$  if  $\downarrow \pi \vdash p$ . The  $\downarrow$  operation is only defined at **(Block, 1)** or **(Grid, 1)**, so we only need to consider the cases where  $\pi$  is one of those.

In the former case  $p$  becomes  $t \mathbf{mod} T$  while  $\downarrow \pi$  is **(Thread, T)**.  $t \mathbf{mod} T < T$  for any  $t$  so this satisfies the requirement that  $\pi \vdash p$ , which lets us use our inductive hypothesis:  $s$  is either **skip** or can step. If it can step, we can use this to satisfy the premise of **S-Destruct-Block** to step in this case. If it is **skip**, then we use the rule **S-Destruct-Done** to step instead.

The latter case is the same, except using the fact that  $b \mathbf{mod} B < B$  and the **S-Destruct-Grid** rule.

- Case **T-Group**

In this case we have by assumption that  $(h, q \cdot n) \vdash p$ , i.e., that  $p < q \cdot n$ .

In this case we have our IH that if  $(h, n) \vdash p'$  for some  $p'$ , then  $s$  is either **skip** or steps with  $(h, n)$  perspective with  $p'$  as our perspective ID.

We choose  $p'$  to be  $p \mathbf{mod} n$ . This is always  $< n$ , so  $(h, n) \vdash p'$ . This lets us use our IH to get that  $s$  is either **skip** (in which case we can use the **S-Group-Done** rule to step) or itself steps, which lets us use the **S-Group** rule to step.

### Thread Synchronization Rules.

- Case **T-Sync-Wait**

$\Psi(\psi)(p)$  is either zero or it is not. In the former case we use the **S-Sync-Wait-Done** rule and in the latter we use **S-Sync-Wait-Spin**.

- Case **T-Sync-Dec**

We use the **S-Sync-Dec** rule to step.

- Case **T-Sync-Init**

We use the **S-Sync-Init-Zero** or **S-Sync-Init-Nonzero** rules depending on whether  $\Psi(\psi)(p)$  is zero or not.

### Asynchrony Rules.

- Case **T-Async-Partition**

In this case, we have via our IH that if  $\Gamma, y : (\text{Thread}, 1) \text{ async } \tau[]^l \vdash \eta', \sigma', \Sigma'$ , then we can either step  $s$  with **(Thread, 1)** perspective under environments  $\eta'$ ,  $\sigma'$ , and  $\Sigma'$ , or  $s$  is **skip**.

1618 We have by assumption that  $\Gamma, x : (\text{Thread}, 1) \text{ async } \tau[]^I \vdash \eta, \sigma, \Sigma$ . By our environment renaming  
 1619 lemma, this gives us what we need to use our IH with  $(\eta', \sigma', \Sigma')$  as **rename** $(\eta, \sigma, \Sigma, x, y, (\text{Thread}, 1))$ .  
 1620 Thus  $s$  either steps or is **skip**. In the former case we can step with **S-Async-Partition-Congr**,  
 1621 and in the latter we can use either **S-Async-Partition-Unwind** or **S-Async-Partition-Done**  
 1622 depending on whether  $\Phi(\phi)$  is empty or not.

- Case **T-Async-Memcpy**

1623 Immediate via use of the **S-Async-Memcpy** rule.

- Case **T-Memcpy**

1624 Immediate via use of the **S-Memcpy** rule.

1625 *Memory Rules.*

- Case **T-Decl**

1626 Via our lemma about expression type safety and our hypothesis that  $e$  is well-typed, we obtain  
 1627 the premises necessary to use the **S-Decl** rule to step.

- Case **T-Arr-Assn**

1628 Each of  $e_1, e_2$  and  $e_3$  must evaluate to a well-typed value by the expression type safety lemma.  
 1629 In particular, both  $e_1$  evaluates to some  $\langle l, n \rangle$  and  $e_2$  evaluates to some  $i$ . We assume all array  
 1630 accesses are in bounds, so this is sufficient to use the **S-Arr-Assn** rule to step.

- Case **T-Arr-Assn-Shared**

1631 Same as previous case.

- Case **T-Free**

1632 Trivial via the **S-Free** rule.

- Case **T-Partition**

1633 Trivial via the **S-Partition** rule.

- Case **T-Claim**

1634 Trivial via the **S-Claim** rule.

- Case **T-Lower**

1635 Trivial via the **S-Lower** rule.

- Case **T-Alloc**

1636 We assume that  $l$  is not **Shared**, so we can use the **S-Alloc-Local** or **S-Alloc-Global** rule,  
 1637 depending on whether  $l$  is **Local** or **Global**.

- Case **T-Alloc-Shared**

1638 Trivial via the **S-Alloc-Shared** rule.

1639 *Control Rules.*

- Case **T-Skip**

1640 Trivial

- Case **T-While**

1641 Trivial, all while loops step via the **S-While** rule

- Case **T-If**

1642 By our proof of expression type safety, the expression  $e$  steps to either the boolean value true  
 1643 or false. We can thus use either the **S-If-True** or **S-If-False** rules to step.

- Case **T-Seq**

1644 In this case we know by our IH that  $s_1$  is either **skip** or can step. In the former case we use  
 1645 the **S-Seq-Done** rule and in the latter we use the **S-Seq-First** rule.

- Case **T-Function-Call**

1646 In this case we know by our expression safety lemma that each of the arguments will evaluate to  
 1647 a well-typed value. We also have by assumption that  $f$  has a function type, which by inversion

on the **V-Function** rule tells us that it is a closure type. Additionally our assumption that  $\Gamma \vdash \eta, \sigma, \Sigma$  tell us that  $\Sigma$  contains  $f$  at the same type that  $\Gamma$  does. These premises are sufficient to use the **S-Function-Call** rule.

□

**LEMMA A.8. (Statement Preservation)** If  $\Gamma \vdash \eta, \sigma, \Sigma$  and  $\Gamma \vdash_m^\pi s$  and  $\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_{m'}^\pi s \rightsquigarrow s' \dashv_{m''} \eta', \sigma', \Sigma, \Psi', \Phi'$  and  $\pi \vdash p$  and  $m \geq m'$  and  $\Gamma \vdash \Phi$ , then there is some  $\Gamma'$  such that  $\Gamma \subseteq \Gamma'$  and  $\Gamma' \vdash_m^\pi s'$  and  $\Gamma' \vdash \eta', \sigma', \Sigma'$  and  $m \geq m''$  and  $\Gamma' \vdash \Phi'$ .

**PROOF.** We proceed by induction on the derivation of  $\Gamma \vdash_m^\pi s$ .

### Perspective Management Rules.

- Case **T-Split**

In this case we have by assumption that  $(h, n) \vdash p$ ,  $\Gamma \vdash \eta, \sigma, \Sigma$ , and  $\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_{m'}^\pi \text{split}(n_1, n_2)\{s_1\}\{s_2\} \rightsquigarrow s' \dashv_{m''} \eta', \sigma', \Sigma, \Psi', \Phi'$ . Our inductive hypotheses give us that if for any  $p$ , if  $(h, n_1) \vdash p$  and  $s_1$  steps with perspective ID  $p$ , or if  $(h, n_2) \vdash p$  and  $s_2$  steps with perspective ID  $p$  then their results are well typed.

By inversion on our semantic derivation, we are in one of 5 cases.

In the **S-Split-Left** case  $p < n_1$  and  $s_1$  steps to  $s'_1$ . This is sufficient to tell us that  $s'_1$  is well-typed and the output environments of that relation  $\eta', \sigma'$ , and  $\Sigma'$  are all well-typed by  $\Gamma' \supseteq \Gamma$ , and that the memory is properly bounded by the typing rules.

We can thus use the **T-Split-Left** rule to conclude that the result of this case is well-typed.

The **S-Split-Left-Done** rule is trivial via the **T-Skip** rule.

The **Right** cases are symmetric, with the observation that when  $p \geq n_1$  and  $p < n_1 + n_2$  then  $p - n_1 < n_2$ .

The last **S-Split-None** rule is trivial via the **T-Skip** rule.

- Case **T-Destruct**

In this case we have by assumption that  $\downarrow \pi$  is defined, so  $\pi$  is either **(Block,1)** or **(Grid,1)**. We also assume that  $\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_{m'}^\pi \text{destruct in } s \rightsquigarrow s' \dashv_{m''} \eta', \sigma', \Sigma, \Psi', \Phi'$ . We also have by our inductive hypothesis that for any  $p$  such that  $\downarrow \pi \vdash p$  and  $s''$  such that  $s$  steps to  $s''$  at  $p$ , then that step preserved well-typedness.

By inversion on the step relation, we are in one of three cases.

If the rule used **S-Destruct-Block**, then we know that  $s$  steps to  $s''$  and  $p$  is  $t \bmod T$ .  $(\text{Thread}, T) \vdash t \bmod T$  for any  $t$ , so we can use our inductive hypothesis to conclude that the  $s''$  stepped to by  $s$  is well-typed, as are its environments and memory usage. The **T-Destruct** rule then gives us our desired goal.

The **S-Destruct-Grid** case proceeds similarly, while the **S-Destruct-Done** case is trivial.

- Case **T-Group**

In this case we have by assumption that

$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_{m'}^{(h,q,n)} \text{group } q \ s \rightsquigarrow s' \dashv_{m''} \eta', \sigma', \Sigma, \Psi', \Phi'$ . We have by our inductive hypothesis that  $s$  steps to  $s''$  at some perspective ID  $p'$  and  $(h, n) \vdash p'$ , then  $s''$  is well typed, as are the other outputs of that step.

By inversion on the step relation, we are in one of two cases. The **S-Group-Done** case is trivial, so we shall focus on the **S-Group** case. In this case we have that  $s$  steps to  $s''$  at perspective ID  $p \bmod n$ . It is always the case that  $(h, n) \vdash p \bmod n$  for any  $p$ , so we can use our inductive hypothesis to conclude that  $s''$  is well-typed, as are its output environments and memory usage. From there, it is a simple application of the **T-Group** rule to conclude that **group**  $q$   $s'$  is well-typed, and to finish the case.

**Thread Synchronization Rules.** These rules are all trivial: with one exception all thread synchronization primitives step to skip without changing environment or memory, and are thus obviously well-typed. **S-Sync-Wait-Spin** does not produce skip, but it steps to the same statement as we already assumed typechecks in the premise of the lemma, so is straightforward nonetheless.

If we wanted to say something about deadlock freedom we'd have more work here, but we aren't doing that, so these rules are easy.

### Asynchrony Rules.

- Case **T-Async-Partition** In this case we have that  $s$  is well-typed in a context where  $x$  has been renamed into  $y$ , with **(Thread,1)** perspective. We also have that  $\Gamma, x : \text{Thread} \vdash \tau[]^l \vdash \eta, \sigma, \Sigma$  and  $\Gamma, x : \text{Thread} \vdash \tau[]^l \Phi$ .

By inversion, we are in one of three cases.

In the **S-Async-Partition-Done** case, we are done.

In the **S-Async-Partition-Congr** case, our inductive hypothesis gives us that there is some  $\Gamma'$  such that  $\Gamma, y : \text{Thread} \vdash \text{async } \tau[]^l \subseteq \Gamma'$  and  $\Gamma' \vdash \Phi'$  and  $\Gamma' \vdash \text{rename}(\eta, \sigma, \Sigma, x, y, \text{Thread}, 1)$  via our well-typed renaming lemma. This lets us use the **T-Async-Partition** rule to check this case, with a choice of  $\Gamma'$  as  $\Gamma, x : \text{Thread} \vdash \tau[]^l$ .

In the **S-Async-Partition-Unwind** case, our assumption that  $\Phi$  is well-typed tells us that  $\Gamma, y : \text{Thread} \vdash \{\text{Thread}, 1\}_m s$ . Thus, we can use the **T-Async-Partition** rule to type this case.

- Case **T-Async-Memcpy**

In this case the statement and environment typing are trivial, we need only to show that the async stack remains well typed.

In this case we have that  $x$  and  $y$  have the same type at **(Thread,1)**. This is sufficient for us to check  $x = y$  at **(Thread,1)**, meaning that adding that instruction to the stack maintains its well-typedness.

- Case **T-Memcpy**

Immediate via use of the **S-Memcpy** rule. We just need to show that the environment remains well typed, which we know via our lemmas about **update** and **get**.

### Memory Rules.

- Case **T-Decl**

By inversion, we are using **S-Decl** rule for evaluation. Our well-typed expression lemma gives us that  $v$  is well-typed, so it follows from our assumptions and our lemmas about extending environments that the extended  $\eta$  and  $s$  are well-typed by  $\Gamma, x : \pi \tau$ .

- Case **T-Free**

Trivial.

- Case **T-Alloc**

By inversion we are either in the **S-Alloc-Local** or **S-Alloc-Global** rules. In either case, we assume that  $\Gamma, x : \pi \tau[]^l$  checks  $s$ , meaning we can use our extended environment lemmas and the **T-Seq** and **T-Free** rules to check these cases.

- Case **T-Alloc-Shared**

Same as previous case.

- Case **T-Partition**

In this case we have by assumption that  $\Gamma, y : (h, n/c) \tau[]^l s$  and  $l$  is not local and  $c$  divides  $n$ . By inversion on our step relation we must be in the **S-Partition** case, so we have  $\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m \text{partition}_\psi y c y s \rightsquigarrow \text{init}_\psi; s'; \text{dec}_\psi; \text{wait}_\psi \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi$  where  $s' = s[(y + c \cdot p)/y]$  and  $(\eta', \sigma', \Sigma') = \text{rename}(\eta, \sigma, \Sigma, x, y, \pi/c)$ . Via our well-typed renaming lemma contexts we know

that we can check the renamed environments in the extended environment  $\Gamma, y:\pi/c \tau[]^l, x:\pi\tau[]^l$ , and this context is also sufficient to check  $s'$  (via a substitution-preserves-typing lemma that is obvious). Using the **T-Skip** rule this is exactly what we need to show to complete this case, as the thread sync primitives check trivially via their typing rules.

- Case **T-Claim**

Essentially the same as **T-Partition**.

- Case **T-Lower**

Essentially the same as **T-Partition**.

### Control Rules.

- Case **T-If**

We have by assumption that **if**  $e$  **then**  $s_1$  **else**  $s_2$  steps to some  $s'$ , and by inversion we know that either  $e$  evaluates to **true** and  $s'$  is  $s_1$ , or  $e$  evaluates to **false** and  $s'$  is  $s_2$ .

In either case, our inductive hypotheses is sufficient to tell us that these are well-typed. In particular, in both cases our IHs tell us that the amount of memory used by stepping each branch of the **if** is less than the amount of memory computed by the type system for each branch. Because the whole **if** expression checks using the greater of the memory usage of  $m_1$  or  $m_2$  (i.e., the memory usage on each branch), the resulting usage for the whole conditional is also bounded by the type system.

- Case **T-Skip**

Trivial: **skip** does not step

- Case **T-Seq**

In this case we have that  $s_1$  and  $s_2$  are both well-typed (with  $m_1$  memory and  $m_2$  memory respectively), and  $m = \mathbf{max}(m_1, m_2)$ . We also have by inversion that  $s_1$  either steps to **skip** or  $s'_1$ , and our inductive hypothesis tells us that  $s'_1$  is well-typed.

In the former case we can use the **S-Seq-Done** rule to trivially finish the case. In the latter, our IH allows us to finish the case, since  $m_1$  is always  $\leq \mathbf{max}(m_1, m_2)$

- Case **T-While**

By inversion, we have that

$$\begin{aligned} & \eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \mathbf{while } e \mathbf{ do } s \\ & \rightsquigarrow \mathbf{if } e \mathbf{ then } (s; \mathbf{while } e \mathbf{ do } s) \mathbf{ else skip } \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi. \end{aligned}$$

We also have by assumption that  $e$  and  $s$  are well-typed. With this information, through use of the **T-If**, **T-Seq**, **T-While**, and **T-Skip** rules, we can conclude that the result of this rule is also well-typed.

- Case **T-Function-Call**

By inversion we have that

$$\begin{aligned} & \eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi f(e_1, \dots, e_n) \\ & \rightsquigarrow \mathbf{call } s \mathbf{ with } (x_i : \pi_i \tau_i \mapsto v_i) @ \{\eta, \sigma, \Sigma, m'\} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi, \end{aligned}$$

and also that  $\Sigma(f) = \{[x_1 : \tau_1, \dots, x_n : \tau_n], s\}$ , and  $\eta, \sigma, \Sigma, t, b, p \vdash^\pi e_i \Downarrow v_i$ , and  $m' \leq m$ .

We also have from the premises of our case that  $f : \pi \mathbf{Fun}(x_1 : \pi_1 \tau_1, \dots, x_n : \pi_n \tau_n, \pi, m') \in \Gamma$ , and  $\Gamma \vdash^\pi e_i : \tau_i$ , and  $m' \leq m$ .

Our lemma for expression well-typedness tells us that that each  $v_i$  is a well-typed value, and our assumption that  $\Gamma \vdash \eta, \sigma, \Sigma$  tells us that  $\Sigma(f)$  is a well-typed function and thus that  $\mathbf{fns } \Gamma, x_i : \pi_i \tau_i \vdash_m^\pi s$ . We can take the union of this with  $\Gamma$  to produce  $\Gamma, x_i : \pi_i \tau_i \vdash_m^\pi s$ , which clearly checks  $s$  and the output environments.

□