

1 Modular GPU Programming with Typed Privileges

2 ANONYMOUS AUTHOR(S)

3 Abstract

4 1 Introduction

5 CUDA [11] is a low-level, imperative programming language for NVIDIA GPUs. These GPUs are organized into a hierarchy of compute resources: threads, blocks, and a grid. *Threads* are the basic units of sequential execution; *blocks* are groups of threads that can cooperate through shared scratchpad memory; and the *grid* is the full collection of blocks launched for a GPU kernel. CUDA programs are written from the perspective of a single thread.

6 Although programs are written from the perspective of a single thread, some operations are only semantically valid when executed *collectively* by a group of threads. For example, the `__syncthreads()` primitive, which synchronizes all threads within a block, causes the machine to deadlock if executed by a subset of threads within that block. Similarly, tensor core instructions, which are hardware-accelerated matrix-multiply operations, are only well-defined when invoked by a group of 32-threads whose starting index is 32-aligned. Therefore, these operations require programmers to carefully marshal compute resources and coordinate which threads execute which lines of code. As a result, collective operations *break the illusion* of threads executing independently.

7 The conceptual clash between regular statements (executed by a single thread), and collective operations (executed by a group of threads) hinders compositional reasoning. In CUDA, function calls, like all other statements, are per-thread operations. The functions, however, may contain code that executes a collective operation, creating a semantic gap between the per-thread syntax for a function's invocation and the underlying semantics of the function's execution.

8 This gap between the syntax and semantics of a CUDA program creates tension between abstraction and correctness. This tension is apparent even in widely used CUDA libraries that package common functionality via function interfaces. Consider the snippet of documentation presented in Figure 1 taken directly from CUB [10], a library from NVIDIA that provides wrappers around parallel primitives like scan, sort, and histogram, among others. The documentation describes a function called `BlockReduce`, in which all threads in a block collaboratively apply a reduction operation, such as a maximum or prefix sum, over an array.

9 The CUB documentation attempts to make several implicit assumptions of `BlockReduce`'s implementation clear to the programmer. First, the function is only well-defined when invoked by an entire block's worth of threads. Second, the function accesses shared memory, a scratchpad space visible to all threads in a block. Users wishing to reuse this storage must first ensure all the block's threads have finished using it by inserting a `__syncthreads()` barrier. Whether this synchronization is required cannot be determined from the function alone; it depends on how the storage is used afterward.

10 In effect, CUB attempts to informally retrofit CUDA with information about the compute and memory requirements of collective operations. By prefixing functions with identifiers such as `Block`, CUB effectively creates "namespaces" for different kinds of these operations. Even then, this organization is purely conventional: it encodes hierarchy through naming and cannot be enforced by the compiler. Ultimately, correctness depends not only on the user carefully reading and interpreting the documentation, but also on the library implementer upholding the invariants of each namespace.

11 Previous attempts to resolve the mismatch between the syntax and semantics of CUDA programs have generally done so by raising the level of abstraction. Triton [15], a tile-based, multi-dimensional array language, circumvents the challenges of compositional reasoning by restricting users to a single level of the hierarchy, namely the block level. While this makes for a simplified programming model,

```

50
51 //include <cub/cub.cuh> // or equivalently <cub/block/block_reduce.cuh>
52
53 _global_ void ExampleKernel(...)

54 {   // Specialize BlockReduce for a 1D block of 128 threads of type int
55     using BlockReduce = cub::BlockReduce<int, 128>;
56     // Allocate shared memory for BlockReduce
57     __shared__ typename BlockReduce::TempStorage temp_storage;
58     // Obtain a segment of consecutive items that are blocked across threads
59     int thread_data[4];
60     ...
61     // Compute the block-wide max for thread0
62     int aggregate = BlockReduce(temp_storage).Reduce(thread_data, cuda::maximum<>());
63 }
64
65 Computes a block-wide reduction for thread0 using the specified binary reduction functor. The first num_valid threads
66 each contribute one input element.
67
68 • The return value is undefined in threads other than thread0.
69 • For multi-dimensional blocks, threads are linearly ranked in row-major order.
70 • A subsequent __syncthreads() threadblock barrier should be invoked after calling this method if the
71 collective's temporary storage (e.g., temp_storage) is to be reused or repurposed.
72
73

```

Fig. 1. CUB documentation for the BlockReduce function

the lack of fine-grained control prevents GPU programmers from writing the highest-performance kernels. A variety of functional languages, such as Futhark [8], provide compile-time guarantees through their type systems but do not expose low-level control over the hardware. Some other efforts, like Descend [9], aim to provide a low-level, safe GPU systems programming language, but lack support for collective operations entirely. As a result, CUDA remains the de-facto standard for writing high-performance kernels on modern GPUs.

Rather than sacrificing low-level control for safety, we aim to make fine-grained GPU programming both safe and compositional by design. Our key insight is that GPU programmers naturally map computations onto different compute resources, and unlike CUDA, which obscures this mapping, we can expose it explicitly using the surface syntax of the language. Using this information, we can check at compile time whether the usage of collective operations contradicts the resources available in that context.

We introduce Snap, a new low-level GPU programming language that guarantees correct usage of hardware resources by construction. Inspired by coeffect systems [12] and dependency calculi [1], Snap statically tracks the configurations of compute and memory resources at the type level.

In this work, we reify the notion of compute resources via *privileges*, so named because they dictate which operations are permitted in which contexts. Snap uses these privileges to verify that every collective operation is executed with the necessary resources. To enable modular reasoning, function interfaces include *privilege signatures* that declare the privilege each function requires for correct execution.

We implement Snap and to validate our design, we build Bundl, a core calculus model of Snap that tracks privileges in both its type system and operational semantics. We provide formal rules for manipulating the privilege of both code and data, and prove *type-and-privilege safety*, ensuring that Bundl is sound, and that all threads are statically privileged to run the operations they attempt at run time.

A parallel goal, in addition to correctness, is performance. To that end, we incorporate modern GPU features such as tensor cores and asynchronous data movement into Snap, and demonstrate that Snap can achieve the same performance as hand-written, highly optimized code on an H100 and a 4070 Ti Super. Our contributions are:

- (1) Snap, a low-level GPU programming language with type-level privileges that track compute and memory resources (Section 3);

- 99 (2) Bundl, a calculus for Snap that tracks privileges in both its type system and operational
 100 semantics, ultimately helping us prove a soundness theorem that states that threads only
 101 execute operations for which they are statically privileged (Section 4); and
 102 (3) An implementation of Snap (Section 5) that demonstrates that Snap can expose modern
 103 GPU features like tensor cores and asynchronous data-movement, and achieve performance
 104 comparable to hand-optimized CUDA implementations (Section 6).

2 Background & Motivation

107 Before diving into Snap’s design, we begin with an overview of the GPU’s compute and memory
 108 hierarchies and outline the challenges posed by reasoning about them collectively.

2.1 Compute Hierarchy

111 In CUDA, programmers launch computations that run on thousands of threads. These threads are
 112 organized into a *compute hierarchy* that defines how work is distributed and scheduled on the GPU. At
 113 the top of this hierarchy is the *grid*, representing all threads launched as part of a single computation.
 114 The grid is divided into *blocks*, each containing a user-specified number of threads. A *thread*, in turn,
 115 is the finest unit of execution, and it is the machine’s basic unit of sequential control.

116 Users define the behavior across different levels of the compute hierarchy by defining the behavior of individual threads that constitute
 117 those levels. To do so, CUDA allows programs to read built-in identifiers like `threadIdx.x`—to determine a thread’s position within
 118 a block—and `blockIdx.x`—to determine the block’s position within
 119 the grid—at runtime. Combining these two IDs, threads can locate
 120 themselves in their share of the full computation.,

121 A natural and tempting way to interpret these built-in identifiers
 122 is to think of them as indices of implicit “parallel for-loops” where
 123 each iteration is executing simultaneously. While this view is suf-
 124 ficient to understand CUDA’s programming model when threads
 125 do independent work, it quickly breaks down in the presence of
 126 *collective operations*.

127 Unlike most instructions, which are executed by a single thread,
 128 collective operations must be executed collaboratively by a group
 129 of threads. For example, a “warp-level” tensor-core operation is only *meaningful* when invoked by
 130 a collection of 32 threads—a *warp*—acting together. Consider Figure 2 which shows a tensor-core
 131 operation being invoked by a warp on line 14. For the call, each thread sets up its portion of the
 132 operands, and the operation is performed *once* for the entire warp, with the results scattered across
 133 participating threads. If the condition on line 2 were instead `tid > 0 && tid < 30`, making fewer than
 134 32 threads reach line 14, the instruction would be undefined. To make matters worse, the restriction
 135 is not only on the *number* of threads executing the operation, but also on their *alignment*. In this case,
 136 the starting ID of the group of 32 participating threads must be aligned to a multiple of 32. So, if the
 137 condition on line 2 were instead `tid > 1 && tid < 33`, the instruction would still be invalid, even
 138 though 32 threads would execute it.

139 Collective operations make reasoning about CUDA programs challenging because they force
 140 programmers to track the convergence behavior of threads. Specifically, programmers must reason
 141 about how many threads reach a particular point in the program, and how those threads are arranged,
 142 accounting for alignment. This difficulty is further amplified due to two reasons. First, programs may
 143 have *multiple points of convergence*, requiring programmers to mentally track the relative ID of a
 144 thread within a logical group as it changes over the course of a program’s execution. Second, threads
 145 146

```

1 int tid = threadIdx.x;
2 if (tid >= 0 && tid < 31){
3     float A[4];
4     float B[2];
5     float C[4] = { 0 };
6     # Populate A with unique values
7     for (int i = 0; i < 4; i++)
8         A[i] = tid * 4 + i;
9     # Populate B with unique values
10    for (int i = 0; i < 2; i++)
11        B[i] = tid * 4 + i;
12    # Issue a warp-level tensor-core
13    # operation. D = A * B + C
14    asm("mma.sync.aligned.m16n8k8...""
15        "%0, %1, %2, %3", "/*D*/"
16        "%4, %5, %6, %7", "/*A*/"
17        "%8, %9", "/*B*/"
18        "%10, %11, %12, %13";"/*C*/"
19        : "=r"(C[0]), "=r"(C[1]), ...
20        : "r"(A[0]), "r"(A[1]), ...
21        "r"(B[0]), "r"(B[1]),
22        "r"(C[0]), "r"(C[1]), ...);}
```

Fig. 2. Invoking a warp-level tensor-core instruction in CUDA.

may be participating in *multiple levels of convergence* within the same program. In our example, we only considered the warp-level tensor-core operation, but there are other collective operations that require convergence at different granularities like the *warp-group*-level tensor-core operations, which must be issued collectively by four warps, or the block-level `__syncthreads()` synchronization primitive, which must be executed by all threads within a block.

Reasoning about collective operations is already error-prone within the context of a single function, but becomes even more difficult when reasoning interprocedurally *across* functions. A callee may assume a certain number of threads or blocks and may structure its computation, including convergence behavior, around that assumption. However, such assumptions are not visible in the callee's function interface, which only exposes the input and output types. To invoke that function correctly, users must read documentation, or worse, read the callee's implementation to understand its assumptions, breaking modular reasoning.

In this work, we make the programmer's implicit assumptions about a program's convergence behavior explicit in the program's source. Consider the example in Figure 3 which shows an equivalent rewrite of the CUDA program in Figure 2 using Snap. In Figure 3, the call to the `mma` intrinsic is legal only because it is executed within a group of 32 threads, as made explicit on line 2. Since Snap's `group` construct already enforces both the size and alignment of participating threads, the legality of the `mma` operation is guaranteed by construction.

Further, Snap also tracks the *frequency* at which values diverge in *space*. The `@` syntax attached to each variable deceleration denotes the rate at which a value changes across compute resources on the machine. For example, a `a@ thread[1]` value can diverge per-thread, while a `a@ block[1]` can diverge for every block, but not for threads within that block. Using this information, Snap enforces rules for reading from and writing to data. So, if we had attempted to introduce a condition based on `tid` *within* the group's scope, Snap would reject that program at compile time. This is because code executing at a lower frequency cannot read variables that change at a higher frequency than it. We will explain the syntax and rules of legal Snap programs in greater detail in Section 3.

2.2 Memory Hierarchy

GPUs also have a *memory hierarchy*, which, like the compute hierarchy, is organized into three levels. All threads in a grid can read and write from *global memory*, where operands typically reside at the start of a computation and where results are eventually written. Each block has access to a limited amount of *shared memory*, a programmer-managed scratchpad typically used to stage data that will be repeatedly used. Finally, each thread maintains its private state in *registers* and *local memory*, which is used for storing the thread's stack.

This memory hierarchy has its own notion of convergence and divergence, mirroring that of compute. For example, when launching a kernel on the GPU, a pointer may initially belong to the grid because every thread sees the same pointer value at the start of the computation. To write to that memory, each thread computes an offset from the original base address. In doing so, the pointer *diverges*

```

1  tid : int @ thread[1] = id();
2  with group(thread[32]):
3    A : float[4] @ thread[1]
4    B : float[2] @ thread[1]
5    C : float[4] @ thread[1]
6    # Populate A with unique values
7    for i in range(0, 4, 1):
8      A[i] = tid * 4 + 1
9      C[i] = 0
10   # Populate B with unique values
11   for i in range(0, 2, 1):
12     B[i] = tid * 4 + 1
13   # Issue a warp-level tensor-core
14   intrinsic.mma(
15     A[0], A[1], A[2], A[3],
16     B[0], B[1],
17     C[0], C[1], C[2], C[3],
18     out=[A[0], A[1], A[2], A[3]])
19

```

Fig. 3. Invoking a warp-level tensor-core instruction in Snap.

```

1  int x = threadIdx.x;
2  if (x >= 0 && x < 31{
3    float A[4];
4    float B[2];
5    float C[4] = { 0 };
6
7    // k is A's stride
8    A[0] = a_mem[(x/4)*k+(x%4)];
9    A[1] = a_mem[(x/4)*k+(x%4)+4];
10   A[2] = a_mem[(x/4*8)*k+(x%4)];
11   A[3] = a_mem[(x/4*8)*k+(x%4)+4];
12
13   // n is B's stride
14   B[0] = b_mem[(x%4)*n+(x/4)];
15   B[1] = b_mem[(x%4+4)*n+(x/4)];
16
17   asm("mma.sync...");
18
19   // Write back into c_mem
20   // n is C 's stride
21   c_mem[(x/4)*n+2*(x%4)] = C[0];
22   c_mem[(x/4)*n+2*(x%4)+1] = C[1];
23   c_mem[(x/4*8)*n+2*(x%4)] = C[2];
24   c_mem[(x/4*8)*n+2*(x%4)+1]=C[3];}

```

Fig. 4. Invoking a warp-level tensor-core instruction in CUDA after loading data from memory.

197 across threads within the larger memory region so each thread can write independently. After these
 198 writes, the memory must *reconverge*, ensuring that all threads have completed their updates before
 199 the memory can safely be returned to its original logical owner in the hierarchy.

200 CUDA does not explicitly model these memory spaces, nor does it track whether allocations in a
 201 given memory space, such as shared memory, exceed device limits. Snap, on the other hand, reifies
 202 these spaces in its type system by explicitly decorating each variable with its associated memory
 203 space and restricting shared memory to static allocations, throwing an error if an allocation exceeds
 204 device limits. We will discuss these aspects in detail in Section 3.5. For now, however, we focus on
 205 how memory diverges and converges in tandem with the compute hierarchy.

206 Let us reconsider the tensor-core example in Figure 2, this time initializing the operands of the
 207 tensor-core operation from pointers into memory. In Figure 4, A and B are now populated from data
 208 in global memory, `a_mem` and `b_mem`. After the tensor-core operation completes, the result is written
 209 back to `c_mem`, also in global memory.

210 In this program, `c_mem` is logically accessed from two different levels of the compute hierarchy.
 211 When entering this code, each thread in a group of 32 threads sees the same value for `c_mem`. Then,
 212 each thread within these 32 locates an offset within `c_mem` that it can write to (lines 21-24). To
 213 restore `c_mem` back to its 32 thread-level ownership, all 32 threads must first synchronize to ensure all
 214 per-thread writes have completed. This is similar to the requirement we saw documented in Figure 1.
 215

216 Similarly to the compute hierarchy (Section 2.1), CUDA programmers are responsible for tracking
 217 the logical owner of a piece of memory as it evolves over the course of a program, and for ensuring
 218 that appropriate synchronization occurs whenever that ownership changes.

218 Snap, by contrast, makes the evolution of a memory object's
 219 logical ownership explicit in the type system and *automatically*
 220 inserts the synchronization required to restore memory to its
 221 original ownership. In Figure 5, we show how memory is lowered
 222 through the compute hierarchy for the same tensor-core
 223 operation introduced in Figure 4. The lowering of `c_mem` is
 224 required in this program as Snap only permits writes when
 225 data is owned by a single thread. To lower `c_mem`, we use Snap's
 226 `partition` operation on line 31. The operation takes the name
 227 of the memory to partition, `c_mem`, and lowers it to a single
 228 thread, and assigns it a new name, `c_thrd`; Snap will not allow
 229 references to the old name `c_mem` within the partition's scope.
 230 The `partition` function also takes an indexing function, and
 231 each time `c_thrd` is accessed, this indexing function is implica-
 232 tly applied. Once the `partition`'s scope ends, Snap will insert
 233 a synchronization before the first use of the original name, so
 234 that all per-thread writes have completed.

235 We would like to emphasize that preventing data races is not
 236 an one of Snap's goals. In Snap, data races *can* occur; however,
 237 since the data is eventually synchronized before it is reused,
 238 the last-writer wins. Out-of-bounds accesses are considered
 239 undefined behavior. We believe that prior work, such as De-
 240 scend [9], lays the foundation for reasoning about data-race
 241 free programs, and that a Descend-like memory model can be adapted to Snap. Snap is instead focused
 242 on the *interaction* between the compute and memory hierarchies, and on reasoning about them
 243 simultaneously to ensure that an operation is executed only when sufficient compute resources are
 244 available, a guarantee that prior work like Descend cannot provide.

```

1  @cuda("device")
2  #requires(thread[32])
3  def simple_mma(
4      a: ptr(const(float)) @ thread[32],
5      b: ptr(const(float)) @ thread[32],
6      c: ptr(float) @ thread[32]):
7      x : int @ thread[] = id()
8      with group(thread[32]):
9          A : float[4] @ thread[]
10         B : float[2] @ thread[1]
11         C : float[4] @ thread[1]
12         # Reads do not need to be lowered.
13         A[0] = a_mem[(x/4)*k+(x%4)];
14         A[1] = a_mem[(x/4)*k+(x%4)+1];
15         A[2] = a_mem[(x/4+8)*k+(x%4)];
16         A[3] = a_mem[(x/4+8)*k+(x%4)+4];
17         B[0] = b_mem[(x%4)*n+(x/4)];
18         B[1] = b_mem[(x%4+4)*n+(x/4)];
19         # Skipping initialize C to ...
20         intrinsic.mma(
21             A[0], A[1], A[2], A[3],
22             B[0], B[1],
23             C[0], C[1], C[2], C[3],
24             out=[C[0], C[1], C[2], C[3]])
25         # Must be at thread[1] to write.
26         idx : int @ thread[1] =
27             lambda ro, co: (x/4+ro)*n+2*(x%4)+co
28         with partition(c_mem, p=thread[32], f=idx)
29             as c_thrd:
30                 c_thrd[0, 0] = C[0]
31                 c_thrd[0, 1] = C[1]
32                 c_thrd[8, 0] = C[2]
33                 c_thrd[8, 1] = C[3]
34         # <-- Sync point inferred by name

```

Fig. 5. Invoking a warp-level tensor-core instruction in Snap after loading data from memory.

3 The Snap Language

Snap is an imperative, low-level programming language designed at a level of abstraction comparable to CUDA. Snap's surface syntax explicitly represents how compute (grid, block, threads) and data (global, shared and local memory) resources are used and distributed over the course of its execution. Using this explicit representation, Snap can enforce safety properties that CUDA cannot. Namely, Snap rejects programs that attempt to execute operations with insufficient compute and memory resources.

In this section, we introduce Snap's core constructs through a running example, a tensor-core matrix multiplication, shown in Figure 6. This is a standard computation a GPU programmer might write while being sufficiently complex to reveal the interesting challenges in GPU programming.

The program in Figure 6 computes a matrix multiplication between two float arrays, A and B to produce an output matrix C. In this kernel, each block computes an independent tile of the output C. To perform the computation, each block first locates its assigned tile index, then loads the corresponding rows and columns from A and B (line 20-21). Since matrix multiplication offers ample data reuse, to amortize the cost of loading data, we stage A and B in shared memory (lines 46-59). At this point, each thread within the block diverges and loads different elements of A and B. Finally, the program invokes an Amere warp-level tensor-core instruction to compute the output, requiring threads in a warp to converge (line 71). We encapsulate this tensor-core instruction in a function, demonstrating how function composition works in Snap. This function's implementation is identical to the one shown in Figure 5.

3.1 Levels

Snap models the machine's compute hierarchy through *levels*. There are three levels in Snap—grid, block and thread—which are organized as expected: a grid consists of multiple units of the block level, which in turn consist of multiple units of the thread level. Levels have a strict

```

1 @ccuda("global")
2 # Top-level privilege class and required shared memory usage.
3 @requires(grid[1], block[1], thread[32], smem=1280)
4 def mmaTF32NaiveKernel(A: ptr(const(float)) @ grid[1],
5                         B: ptr(const(float)) @ grid[1],
6                         C: ptr(float) @ grid[1],
7                         M : int @ grid[1],
8                         N : int @ grid[1],
9                         K : int @ grid[1]):
# Replicate this code across 1 grid.
with group(grid[1]):
    # @ grid[1] inferred from compute privilege
    MMA_N : constexpr(int) = 8
    ...
    num_blocks_n : const(int) = (N + MMA_N - 1) / MMA_N
    # id() function returns the block id, inferred from
    # @ block[1].
    block_row : const(int) @ block[1] = id() / num_blocks_n
    block_col : const(int) @ block[1] = id() % num_blocks_n

    warp_row: const(int) @ block[1] = block_row * MMA_M
    warp_col : const(int) @ block[1] = block_col * MMA_N
    # Give each block an offset into C
    offset = lambda x: warp_row * N + warp_col + x
    with partition(C, p=block[1], f=offset) as C_blk:
        with group(block[1]):
            # SEM declarations only allowed with block[1] privilege.
            A_smem : shared(float[16 * 8]) @ block[1]
            B_smem : shared(float[8 * 8]) @ block[1]
            C_smem : shared(float[16 * 8]) @ block[1]
            # Now, id() returns the thread id
            idx : int @ thread[1] = id() * 4
            # To write to C_smem, must drop to thread[1] privilege
            with partition(C_smem, p[thread[1]], offset=idx) as C_thrd:
                for i in range(0, 4, 1):
                    with group(thread[1]):
                        C_thrd[i] = 0

            for i in range(0, K_tiles, 1):
                a_idx : int @ thread[1] = id() * 4
                # <-- Name will insert a sync point
                # because this partition is a parent partition
                # or itself (back edge from for loop)
                for j in range(0, 4, 1):
                    global_row : int @ thread[1] = warp_row + row
                    global_col: int @ thread[1] = i * MMA_K + col
                    with partition(A_smem, p=thread[1], f=...) as As_thrd:
                        with group(thread[1]):
                            As_thrd[0] = A[global_row * K + global_col]

                b_idx : int @ thread[1] = id() * 2
                # <-- Name will insert a sync point
                # (back edge from for loop)
                for j in range(0, 2, 1):
                    # Similar to write into C_smem ...
                    with partition(B_smem, p=thread[1], f=...) as Bs_thrd:
                        with group(thread[1]):
                            Bs_thrd[0] = B[global_row_b * N + global_col_b]

                # <-- Name will insert a sync point
                # (back edge from for loop)
                # Give each warp an offset into C_smem
                with claim(C_smem, p[thread[32]]) as Cs_warp:
                    match split(thread):
                        case 32:
                            # Can pass A_smem and B_smem
                            # as arguments to simple_mma for
                            # those two are declared const.
                            # "read up" is okay.
                            simple_mma(A_smem, B_smem, Cs_warp)

# Write back final result to C
for j in range(0, 4, 1):
    flat_idx_c : int @ thread[1] = id() * 4 + j
    row_c : int @ thread[1] = flat_idx_c / MMA_K
    col_c : int @ thread[1] = flat_idx_c % MMA_K
    offset = lambda x: row_c * N + col_c + x
    with partition(C_blk, p=thread[1], f=offset) as C_thrd:
        with group(thread[1]):
            C_thrd[0] = C_smem[row_c * MMA_N + col_c]
return

```

Fig. 6. Tensor Float 32 matrix Multiplication in Snap.

295 order defined over them. In particular, thread
 296 < block < grid.

297 There are two key differences between Snap's
 298 levels and those in CUDA. First, unlike CUDA,
 299 Snap does not model a a three-dimensional grid
 300 or block structure, say through different levels for grid_x, grid_y, or grid_z. Second, there are two
 301 "levels", *warp* and *warp group*, noticeably absent from our hierarchy.

302 These differences are a deliberate design choice. On the hardware, the units of each level are
 303 arranged in a single linear order, and the three-dimensional structures are simply *interpretations* of this
 304 linear ordering, not distinct hardware resources. Similarly, warp and warp group are organizational
 305 constructs defined in terms of existing levels . Namely, a warp is a group of 32 threads whose first
 306 thread ID is aligned to 32. A warp group, which only became a meaningful construct with the release
 307 of the Hopper architecture , consists of 128 threads aligned to 128.

308 Rather than embedding these interpretations into the language as special cases by introducing
 309 new levels, Snap provides constructs that let users group units of a given level . With these constructs,
 310 users can express multi-dimensional structures and define custom groupings with specific quantities
 311 and alignments.

312 The primary mechanism for organizing units at each level is a *privilege*.

313

314 3.2 Privileges & Privilege Classes

315 Privileges are the central concept in Snap, allowing it to determine which compute resources are
 316 being requested by the programmer, whether they are available in the program's context, and, once
 317 provided, if those resources are sufficient for a given operation.

318 A *privilege* is a level—grid, block, or thread—parametrized by a static constant n, specifying the
 319 number of units at that level. A privilege is written as level[n]. For example, thread[2] denotes a
 320 privilege of two threads, block[4] denotes a scope of four blocks, and so on.

321 Privileges also carry *alignment information*: a privilege parameterized by n is aligned to n. In this
 322 way, a warp is simply a desugaring of thread[32], and a warp group is a desugaring of thread[128].

323 Finally, privileges have a *partial order* defined on them. We say that level₂[n₂] is a *higher privilege*
 324 than level₁[n₁] iff either level₁ < level₂, or, if level₁ = level₂, then n₂ % n₁ = 0. The condition n₂
 325 % n₁ = 0 may appear somewhat mysterious at first, but we will see why it is required in Section 3.2.2.

326 In Snap, both code and data are bound to privileges . Code is bound to a set of privileges, called a
 327 *privilege class* meanwhile data is bound to a single privilege.

328 3.2.1 Privileges for Code, or Privilege Classes.

329 Privilege classes capture what compute resources a program has access to. Each function starts
 330 out with a top-level privilege class that defines the total available resources that it *assumes* are
 331 available to it. As a shorthand, when we refer to a code's privilege, we mean the highest privilege
 332 available to it in its class. In Figure 6, the privilege class is declared on line 2 using the notation
 333 @requires(grid=..,block=..) Starting from this top-level declaration, programmers shape the
 334 current privilege class in the function's body through two constructs, group and split.
 335

336 **Group.** The group construct bundles privileges available in the current privilege class into "groups"
 337 of a particular granularity. Intuitively, group narrows the current privilege class and replicates it
 338 across collections of a finer granularity. For example, Figure 6, line 26 shows a grouping of a single
 339 block, while line 22 in Figure 5 shows a grouping of 32 warps , that will ultimately invoke a tensor-core
 340 operation.

341

342

344 An invocation of `group(level[n])` is allowed if and only if the
 345 privilege class contains privilege higher than `level[n]`. For example,
 346 `group(warp[32])` on 22 in Figure 5 is allowed because the privilege
 347 class of the function declares that the top level on line 22 starts out
 348 with at least 32 threads.

349 Once `group(level[n])` is invoked, it modifies the current privi-
 350 lege class in two key ways. First, it removes all parents of `level[n]`
 351 from it. Second, the number of available units for `level` is set to
 352 exactly `n`. Combined, these rules eliminate malformed programs that would result in unsatisfiable
 353 requirements. We shown some examples of illegal group statements in Figure 7.

```
1 # Example 1
2 with group(thread[2]):
3     # Illegal because block > thread
4     with group(block[1]):
5         pass
6
7 # Example 2
8 with group(block[1]):
9     # Illegal because 1 % 2 != 0
10    with group([block[2]]):
11        pass
```

Fig. 7. Illegal uses of group.

355 **Split.** Unlike `group`, which is used for replication over smaller, equal sized resource sets, `split` is
 356 used for sharding currently available resources into unequal sized sets. For example, line 63 in Figure 6
 357 shows a `split` of a `thread[32]` privilege that is used to execute a warp-level tensor-core operation.

358 The `split` operation reflects unordered composition that is not
 359 mere replication, and all branches of the `split` execute in parallel.
 360 Once `split(level)` is invoked, the privilege classes of its branches
 361 diverge. A `split` construct will eventually compile to an `if`-statement
 362 during code generation.

363 Therefore, to satisfy a `split` operation, Snap checks that the
 364 *cumulative sum* of the privileges requested by each branch of the
 365 `split` can be satisfied. For example, if we were to change the program
 366 in Figure 6 slightly (on line 63-69), to the one shown in Figure 8,
 367 Snap would throw a compile-time error. Finally, because privileges
 368 assume alignment, every branch of the `split` must also be aligned.
 369 In other words, not all splits whose sizes sum to less than or equal to
 370 the available units are necessarily valid. Figure 9 shows an example
 371 of this constraint. This alignment constraint is also why we compare
 372 privileges by checking whether one quantity is a multiple of the
 373 other's.

```
1 # Example 1
2 with group(thread[32]):
3     match split(thread):
4         case 32:
5             simple_mma(..)
6             # Illegal! 32 + 1 > 33
7         case 1:
8             pass
```

Fig. 8. Illegal split that exceeds available privileges.

```
1 # Example 2
2 with group(thread[33]):
3     match split(thread):
4         case 1:
5             pass
6             # Illegal! 32 + 1 <= 33,
7             # but the second branch
8             # is not aligned by 32.
9         case 32:
10            simple_mma(..)
```

Fig. 9. Illegal split that violates alignment.

374 Inside each branch that requests `n` units, all parent levels of
 375 `level[n]` are removed from its privilege class, and the number of available units for `level` is
 376 set to exactly `n`, and, within that scope, all resource requirements are checked against that privilege
 377 class.

378 3.2.2 Privileges for Data.

380 Like code, data also has a privilege attached to it that dictates the set of privileges that must be held
 381 for reading and writing to it. Ultimately, it is this interaction between privilege classes for code and
 382 privileges for data that let us make meaningful statements about the set of privileges available at
 383 runtime. The precise details of this interaction are formalized in Section 4.

384 In this section, we focus tracking thread-local stack allocations and local variables. We defer
 385 discussion of pointers to memory to Section 3.5, where we will show how shared views of memory
 386 are lowered and restored through the compute hierarchy.

387 For thread-local data, its privilege is specified in its declaration using the `@ level[n]` syntax.
 388 For example, an integer variable `v` is declared with `thread[1]` privilege with the syntax `v : int @`
 389 `thread[1]`. If not explicitly annotated, Snap infers a variable's privilege to be the privilege of the
 390 code where it was declared. For example, `MMA_N`, `num_blocks_n` are all inferred to be at `grid[1]`
 391 scope on lines 13 and 15 in Figure 6.

393 The rules for reading and writing to data are asymmetric: higher-privilege data can flow into lower-
 394 privilege data, but not the other way around. In particular, a value coming from a higher-privilege
 395 variable can be written into lower-privilege data, but a value coming from a lower-privilege variable
 396 cannot be written into higher-privilege data. In short, the rule is: "read up, write down".

397 If a value is read without being written into a variable, for example,
 398 in an `if` statement, the above rule is enforced using the privilege
 399 the code that is executing the read. For example, branching on a
 400 `thread[1]` variable within `block[1]` privilege code is rejected by
 401 our compiler, as it constitutes a "read down".

402 In Figure 6, we can see rule "read up" in action on lines **45 and 46**.
 403 Meanwhile, "write down" is being used to set up variables on line
 404 **18-22**.

405 Intuitively, `@ level[n]` tells Snap the *frequency* at which a value changes in space . For example,
 406 a `grid[1]` value is the same for all threads in a grid, whereas a `block[2]` value changes every
 407 two blocks. The rules described above enforce that the frequency declared for each variable is
 408 respected throughout the program. For instance, the "read up" principle allows variables with
 409 higher replication frequency to read from lower-frequency variables without violating their declared
 410 frequency. Conversely, higher-frequency variables cannot be used to write into lower-frequency
 411 variables, since that would violate the variable's declared frequency.

412 By tracking frequency in this way, Snap can ensure that threads that with a given privilege can
 413 only read variables that change at the same frequency or higher. This guarantees that all threads
 414 with that privilege follow the same control flow, and that our accounting of available privilege class
 415 remains consistent with the program's runtime behavior.

416 Until now, we have discussed up rules for reading and writing to variables that change at different
 417 frequencies, but we have not yet set up a mechanism for actually actually setting up variables that
 418 diverge and change at different frequencies. The main tool for doing this is the `id()` function.

420 3.3 The `id()` Function

421 As opposed to exposing users to special hardware `blockIdx.x`, `threadIdx.x` variables directly, Snap
 422 provides the `id()` function instead. The `id` function returns the relative index, or a "label", for a given
 423 compute privilege. The interpretation of the `id` function depends on both the privilege of the variable
 424 it is being written to, and the privilege class that the `id` function is being invoked in. In our example
 425 in Figure 6, we use a call to `id` in `grid[1]` scope at line **18 and 19** to locate the tile that each block is
 426 ultimately charge of computing. In Figure 3, however, a call to `id` in `thread[32]` scope at line **7** will
 427 not return the global thread ID of the threads launched in the kernel, but rather the *relative* thread ID
 428 within the 32 threads, ranging from 0 to 31.

429 3.4 Collective Operations

430 With privilege classes for code and data in place, we can now turn our attention to Snap's reasoning
 431 about collective operations. By statically tracking privilege classes for code, Snap guarantees that
 432 collective operations are semantically valid by ensuring that they are invoked with sufficient privilege.
 433

434 There are two ways to access collective operations in Snap. The first is to use intrinsics provided
 435 by the compiler. These intrinsics have built-in privilege classes attached to them, which the compiler
 436 can directly enforce. In Figure 3, the tensor-core call at line **22** is made using one of the intrinsics
 437 provided by our compiler. The second, more flexible approach allows users to extend these operations
 438 using the `unsafe` construct, explained in Section [3.8](#). Using unsafe code, users can inline a collective
 439 operation—often a single assembly instruction—into a Snap function and ensure that the function's
 440 top-level requirements, specified with `@require`, faithfully describe the privileges required to execute
 441

```
1 with group(block[1]):
2   id : int @ thread[1] = id()
3   # Illegal read of ID from
4   # block[1] scope.
5   if (id)
6     __syncthreads();
```

Fig. 10. Illegal read of `thread[1]` variable.

that instruction. Then, Snap will check the call to the function as any other function call, enforcing the correct privilege. The rules for checking function calls are described in Section 3.7.

A quirk of some collective operations is that, although they need to be invoked with higher compute privileges, they often produce results with lower data privileges. For example, the tensor-core operation in Figure 3 is invoked at `thread[32]` granularity but generates output at the `thread[1]` level. That is, even though the call to the tensor call is a convergent operation requiring the `thread[32]` compute privilege, it produces divergent output with `thread[1]` data privilege. To be able to handle these operations we must be able to allow `thread[1]` state to coexist with the `thread[32]` privilege.

One potential solution is to artificially coarsen the granularity of the operation by having the intrinsic internally scatter the `thread[1]` output into a `thread[32]` variable. However, this approach is often unacceptable for low-level CUDA performance tuning. Often, performance engineers need to manually schedule and interleave memory and arithmetic instructions to hide latency, with loads or stores occurring well in advance of their dependent operations and interleaved with other instructions. If we coarsen these operations, we will disallow these optimizations entirely.

In Snap, however, this behavior does not pose a problem. Because of the "write down" rule, registers can be set up in advance of the operation (e.g., line 14-20 in Figure 3) without violating correctness.

3.5 Memory

In Snap, there are three mechanisms for obtaining a pointer type. The first method is to directly allocate data. The second and third are to obtain offsets into existing data by using `partition` or `claim`. These constructs mirror the `group` and `split` constructs, and are used to lower the data privilege associated with memory.

3.5.1 Allocations in Different Memory Spaces.

In Section 2.2, we saw that the GPU has distinct memory spaces, like global, local, and shared memory. Local allocations reside in either the stack or registers, which Snap can distinguish upfront. Shared memory pointers are explicitly annotated with the `shared` keyword, while global pointers are unannotated.

The primary memory resource Snap needs to manage are shared memory allocations. Local memory allocations are just stack allocations, and don't need special treatment. We also assume that sufficient global memory has been allocated before the kernel launch and don't provide a mechanism for runtime allocations—this is typically the case in real-world programs.

In CUDA, shared memory is a limited resource, typically a few hundred kilobytes per block. Programmers manually manage this memory by bumping pointers to allocate offsets for different data. Each allocation's pointer must be carefully tracked throughout the program, and exceeding the available space will lead to a runtime error.

Snap simplifies shared memory management by requiring all shared memory allocations have a static size. Since shared memory is only visible to all threads in a block, a shared memory allocation is only valid with `block[1]` privilege. An example of such an allocation is shown on lines 28-30 of Figure 6. During compilation, Snap will automatically handle the necessary pointer arithmetic to assign each allocation an appropriate offset within the shared memory space.

The `@requires` annotation specifies the amount of shared memory a function expects to have available, and Snap uses this information to ensure that a function's allocations do not exceed its declared limit, and checks whether there is enough shared memory available at its call sites. Figure 6 declares the amount of shared memory it will require on 2. We use standard techniques to conservatively bound memory usage in Snap's type system.

3.5.2 Privileges for Memory.

491 In Section 3.2.2, we focused on tracking for primitive data types and stack allocations, where each
 492 thread maintained its own local copy of the data. In contrast, pointers into memory, such as global or
 493 shared memory, do not represent distinct copies, but instead refer to the same underlying storage
 494 that may be accessed simultaneously by multiple threads. To use them, we must first *lower* their
 495 privileges, allowing individual threads to write different values to different locations, and then *restore*
 496 them back to their original privilege before execution continues.

497 As with variables in Section 3.2.2, lower-privilege data may read from higher-privilege data
 498 without restrictions. For writes, however, Snap enforces that all pointers into memory be written
 499 with `thread[1]` privilege; they therefore must be lowered before they can be safely updated.

500
 501 **Partition.** The `partition` construct plays the same role for memory that `group` plays for compute.
 502 It is used to lower a memory object from a higher privilege to a lower one by assigning each lower
 503 privilege a distinct offset into the original pointer. It takes the original name of a memory object,
 504 `base_name`, an indexing function `f`, and a target privilege `level[n]` to which the memory is being
 505 lowered, and produces a new name, `new_name`, for the partitioned memory. In Figure 6, for example,
 506 we partition the C buffer to `block[1]` privilege on line 25 first, and then lower it further to `thread[1]`
 507 privilege on line 78 so that we can actually write to it.

508 Once a memory object has been partitioned, the original name referring to it goes out of scope,
 509 and the memory can only be accessed through the new name introduced by `partition`. Each use of
 510 this new name applies the indexing function `f` to compute the true offset of the access. Within the
 511 scope of the `partition` statement, the original name cannot be repartitioned. Because `split` denotes
 512 unordered composition, Snap prohibits sibling `split` branches from partitioning the same memory
 513 object. When the scope of the `partition` statement ends, Snap automatically inserts the necessary
 514 synchronization to restore the memory to its original privilege. The mechanism for inserting and
 515 optimizing these synchronization points is described in Section 3.9.

516 At this point, it is necessary to consider how different indexing functions affect the possibility
 517 of data races. If the indexing function is injective, each lower-privilege privilege receives a distinct
 518 offset into the underlying array, and the resulting partition is free of data races. When the indexing
 519 function is not injective, multiple privileges may write to the same location, introducing a potential
 520 data race. Out-of-bounds accesses are undefined behavior.

521 As we discussed in Section 2.2, achieving data-race freedom is *not* one of Snap’s goals, and
 522 our guarantees continue to hold even when indexing functions are not injective. Because Snap
 523 automatically inserts the necessary synchronization before the original memory name is restored to
 524 its higher privilege, Snap ensures that all writes have completed before accessing the orginal name.
 525 The subsequent uses of that pointer will observe the value written by the last-writer.

526 We believe that extending Snap to reason about data races and out-of-bounds accesses is a promising
 527 direction for future work.

528
 529 **Claim.** The `claim` construct is the analogue of `split` for memory . Using `claim`, the entire
 530 memory region can be moved to a lower privilege.

531 Unlike `partition`, `claim` does not take an indexing function. The `claim` construct takes the
 532 original memory’s name, the target privilege to lower to, and a new name to assign to the lowered
 533 memory. The `new_name` is accessible only within a single, explicit `split` branch; sibling branches are
 534 not permitted to read or write from this memory. In Figure 6, we use the `claim` operation on line 63
 535 to lower the memory to `thread[32]` privilege.

536 As with `partition`, a `split` construct introduces its own scope, and the original name is inacces-
 537 sible within this scope.

540 **3.6 Asynchrony**

541 **3.7 Function Composition**

542 To check whether a function call is valid, Snap needs to ensure that the call site can provide the set of
 543 privileges that a function will attempt to use over its execution, and that the arguments passed to the
 544 function also respect the privileges declared for its input and output types. Both requirements are
 545 represented explicitly via the `@require` construct.
 546

547 As mentioned in Section 3.2, each function carries a privilege class. We call this top-level privilege
 548 class the function’s *privilege signature*. The signature specifies the top-level privilege class the
 549 function *assumes* will be provided to it. This signature represents the *minimum amount* of compute
 550 and memory resources the function must be called with. In the example above in Figure 6, line 2
 551 defines the privilege signature. Up to this point, we have described programs written *inside* functions
 552 under the assumption that the resource signature can be satisfied. In Section 3.2, for example, we
 553 used this signature to set up the top-level privilege class for the function.

554 **Compute privileges.** At a function’s call site, then, we need to ensure that the function’s privilege
 555 signature can be satisfied. Snap compares the current compute privilege with the one required by the
 556 function, and verifies that only functions whose privilege signatures can be satisfied are called.
 557

558 **Per-argument checks.** In Snap, functions have pass-by-value semantics. For arguments, we
 559 distinguish between primitive data types and pointers to memory. For primitive data types, like `int`,
 560 `bool`, and stack allocations, users can pass either arguments that live at or above the data privilege
 561 specified by the function’s signature. For pointers, we allow passing pointers that live at higher data
 562 privilege if the pointer is marked as `const`, indicating that it will only be used for reading (line 4 and
 563 5 in Figure 3). Otherwise, we require that pointers match the exact privilege of the function signature.
 564

565 **3.8 Unsafe**

566 **3.9 Compiler Internal: Synchronization**

567 **4 Formalism**

568 Having introduced the full Snap language, we now describe Bundl, a core calculus that formalizes its
 569 most fundamental aspects by statically tracking compute and data privileges. We use Bundl to prove
 570 that well-typed Snap programs are not only type-safe, but will also never get stuck trying to execute
 571 operations for which they do not have the required compute resources.
 572

573 In this section, we describe Bundl’s type system and operational semantics—in particular how it
 574 manages compute and data privileges—and build up to a formal proof of type-and-privilege safety.
 575 By instrumenting Bundl’s operational semantics with privileges, our safety theorem can guarantee
 576 that dynamically-realized privileges match the ones inferred by the type system.
 577

578 **4.1 Bundl Type System**

579 The core idea in Bundl is to track, at the type-level, the program’s current compute and data privilege.
 580 To achieve this, we borrow techniques from the literature on *coeffects* [12] and *dependency tracking*
 581 [1]. In particular, the current compute privilege is tracked on the typing judgment, which has the
 582 form $\Gamma \vdash^\pi e : \tau$ for expressions and $\Gamma \vdash^\pi s$ for statements. The π over the \vdash is the compute privilege of e
 583 and s .

584 The π s themselves are pairs of levels h of the GPU hierarchy at which each privilege lives (either
 585 Thread, Block or Grid), and numbers tracking the size of the logical group denoted by each privilege.
 586 So, for example, a `(Thread,4)` in Bundl would directly correspond to a `thread[4]` in Snap.

587 Each variable in the typing context has a data privilege attached to it. Only when these two
 588 privileges match can data be read or written from a variable. This requirement is made manifest in
 589

589	$\boxed{\Gamma \vdash^\pi e : \tau}$	<i>(Expression typing)</i>
590	$\frac{x :^\pi \tau \in \Gamma}{\Gamma \vdash^\pi x : \tau}$ T-VAR	
591		
592		
593	$\boxed{\Gamma \vdash^\pi s}$	<i>(Statement typing)</i>
594		
595	$\frac{\Gamma \vdash^{(h,n_1)} s_1 \quad \Gamma \vdash^{(h,n_2)} s_2 \quad n_1, n_2 \text{ align to } n}{\Gamma \vdash^{(h,n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\}}$ T-SPLIT	
596		
597		
598	where $n_1, n_2 \text{ align to } n ::= (n_1 + n_2 \leq n) \text{ and } (n_1 n) \text{ and } (n_2 n) \text{ and } (n_2 n_1 + n)$	
599		
600	$\frac{\Gamma \vdash^\pi s}{\Gamma \vdash^{q \cdot \pi} \text{group } q \ s}$ T-GROUP	$\frac{\Gamma, y :^{\downarrow \pi} \tau[]^l \vdash^\pi s \quad l \neq \text{Local}}{\Gamma, x :^\pi \tau[]^l \vdash^\pi \text{lower } x \text{ into } y \text{ in } s}$ T-LOWER
601		
602		
603	$\frac{\Gamma, y :^{(h,n/c)} \tau[]^l \vdash^{(h,n)} s \quad c n \quad l \neq \text{Local}}{\Gamma, x :^{(h,n)} \tau[]^l \vdash^{(h,n)} \text{partition } x \text{ into } y \text{ by } c \text{ in } s}$ T-PARTITION	$\frac{\Gamma \vdash^\pi s}{\Gamma \vdash^\pi \text{destruct in } s}$ T-DESTRUCT
604		
605		
606	$\frac{\Gamma, y :^{(h,n')} \tau[]^l \vdash^{(h,n')} s \quad n' \leq n \quad l \neq \text{Local}}{\Gamma, x :^{(h,n)} \tau[]^l \vdash^{(h,n)} \text{claim } x \text{ into } y \text{ at } n' \text{ in } s}$ T-CLAIM	
607		
608		
609		

Fig. 11. Core typing rules of Bundl. The typing rules presented here are a simplified selection of the full rules, which can be found in Appendix A.2.

the T-VAR rule, which can be found in Figure 11. The premise of this rule also reveals that the context Γ maps variables to both types τ and also the privileges π with which they may be read and written.

Figure 11 also shows other key rules, and these rules fall into two main categories: rules for managing the current compute privilege and rules for managing data privileges.

4.1.1 Compute Privilege Management. In Snap, there are two operations which manage compute privileges: group and split. In Bundl, to better model the details of privilege attenuation , we introduce a third construct, destruct.

Bundl’s group statement directly corresponds to Snap’s, and is checked by the T-GROUP rule. Given some privilege that can be divided into q equally sized smaller privileges π , the statement $\text{group } q \ s$ will check with privilege $q \cdot \pi$ provided that s itself checks with privilege π . The starting privilege is of the form $q \cdot \pi$, and encodes Snap’s alignment requirement.

The split statement, meanwhile, is checked by the T-SPLIT rule, and functions like a binary version of the n-ary splitting construct in Snap. It enforces the same divisibility requirements to ensure that the privileges of code and data remain properly aligned, and then checks the two sub-statements s_1 and s_2 with the divided, smaller privileges.

The final rule for managing compute privileges is T-DESTRUCT, which handles downward movement from one level of the GPU hierarchy to another and makes explicit in Bundl programs exactly where such movement occurs. The \downarrow operation on π “destructs” the privilege into a many smaller privileges at a lower level; $\downarrow(\text{Grid}, 1) = (\text{Block}, B)$ and $\downarrow(\text{Block}, 1) = (\text{Thread}, T)$, where B and T are parameters to a particular instantiation of Bundl to describe the number of blocks per grid and

$$\begin{array}{l}
638 \quad \frac{L(t), S(b), \Sigma, t, b, 0 \vdash^{(\text{Grid},1)} s \rightsquigarrow s' \dashv \eta', \sigma', \Sigma' \quad P(t, b) = s}{L, S, \Sigma, P \rightsquigarrow L[t \mapsto \eta'], S[b \mapsto \sigma'], \Sigma', P[(t, b) \mapsto s']} \text{ S-PROGRAM} \\
639 \\
640 \\
641 \\
642 \quad \frac{p < n_1 \quad n_1, n_2 \text{ align to } n \quad \eta, \sigma, \Sigma, t, b, p \vdash^{(h, n_1)} s_1 \rightsquigarrow s'_1 \dashv \eta', \sigma', \Sigma'}{\eta, \sigma, \Sigma, t, b, p \vdash^{(h, n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\} \rightsquigarrow \text{split}(n_1, n_2)\{s'_1\}\{s_2\} \dashv \eta', \sigma', \Sigma'} \text{ S-SPLIT-LEFT} \\
643 \\
644 \\
645 \quad \frac{p \geq n_1 \quad p < n_1 + n_2 \quad n_1, n_2 \text{ align to } n \quad \eta, \sigma, \Sigma, t, b, p - n_1, \vdash^{(h, n_2)} s_2 \rightsquigarrow s'_2 \dashv \eta', \sigma', \Sigma'}{\eta, \sigma, \Sigma, t, b, p \vdash^{(h, n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\} \rightsquigarrow \text{split}(n_1, n_2)\{s_1\}\{s'_2\} \dashv \eta', \sigma', \Sigma'} \text{ S-SPLIT-RIGHT} \\
646 \\
647 \\
648 \\
649 \quad \frac{\eta, \sigma, \Sigma, t, b, p \text{ mod } n \vdash^{(h, n)} s \rightsquigarrow s' \dashv \eta', \sigma', \Sigma'}{\eta, \sigma, \Sigma, t, b, p \vdash^{(h, q \cdot n)} \text{group } q \ s \rightsquigarrow \text{group } q \ s' \dashv \eta', \sigma', \Sigma'} \text{ S-GROUP} \\
650 \\
651 \\
652 \quad \frac{\eta, \sigma, \Sigma, t, b, t \text{ mod } T \vdash^{(\text{Thread}, T)} s \rightsquigarrow s' \dashv \eta', \sigma', \Sigma'}{\eta, \sigma, \Sigma, t, b, 0 \vdash^{(\text{Block}, 1)} \text{destruct in } s \rightsquigarrow \text{destruct in } s' \dashv \eta', \sigma', \Sigma'} \text{ S-DESTRUCT-BLOCK} \\
653 \\
654 \\
655 \\
656 \quad \text{Fig. 12. Core semantic rules of Bundl. As with the typing rules, we present only a simplified selection of the full} \\
657 \quad \text{rules, which can be found in Appendix A.3.} \\
658 \\
659 \\
660 \\
661 \\
662 \\
663 \quad \text{threads per block.}¹ Because } \downarrow \text{ is only defined on (Grid,1) and (Block,1), the rule enforces that one} \\
664 \quad \text{can only destruct their compute privilege when they have privileges for exactly one grid or block.} \\
665 \\
666 \\
667 \quad 4.1.2 \ Data Privilege Management. The mechanism for managing data privileges mirrors that of} \\
668 \quad \text{compute privileges, with each operation for data corresponding to an operation for code.} \\
669 \quad \text{The partition operation is analogous to group-ing a compute privilege. The typing rule for this} \\
670 \quad \text{operation, T-PARTITION, requires that the data privilege of the variable to be partitioned, } x, \text{ is the} \\
671 \quad \text{same as the current compute privilege. After partitioning, a fresh variable } y \text{ is introduced with a new} \\
672 \quad \text{privilege } \pi/c. \text{ We use } \pi/c \text{ as a shorthand to denote a division of a privilege } \pi \text{ into } c \text{ equally-sized,} \\
673 \quad \text{smaller privileges. Within the body of the partition statement, we disallow any references to the} \\
674 \quad \text{original variable } x. \text{ Then, we continue checking the body of the partition statement against the} \\
675 \quad \text{original compute privilege } \pi \text{ because the partition has no affect on the current compute privilege.} \\
676 \quad \text{Unlike partition, which divides up a piece of data equally into lower privileges, the claim} \\
677 \quad \text{operation, mirroring the split construct, gives the entire variable to exactly one, smaller privilege.} \\
678 \quad \text{Accordingly, Bundl needs to ensure that only one branch of a split operation, with the appropriate} \\
679 \quad \text{privilege, can refer to the claimed variable. To ensure that this is the case, the T-CLAIM rule links the data} \\
680 \quad \text{privilege of the variable to the compute privilege of the code claiming it by changing both at the same} \\
681 \quad \text{time. This represents a minor difference from Snap, which implements additional static analysis to} \\
682 \quad \text{ensure that a claimed variable is only accessed in a single split branch.} \\
683 \quad \text{Lastly, the T-LOWER rule mirrors the T-DESTRUCT rule; it uses the } \downarrow \text{ operator to move a variable} \\
684 \quad \text{from one level of the hierarchy to another, distributing it equally among all the child privileges in the} \\
685 \quad \text{same manner as T-PARTITION.} \\
686
\end{array}$$

Fig. 12. Core semantic rules of Bundl. As with the typing rules, we present only a simplified selection of the full rules, which can be found in Appendix A.3.

threads per block.¹ Because \downarrow is only defined on (Grid,1) and (Block,1), the rule enforces that one can only destruct their compute privilege when they have privileges for exactly one grid or block.

4.1.2 Data Privilege Management. The mechanism for managing data privileges mirrors that of compute privileges, with each operation for data corresponding to an operation for code.

The partition operation is analogous to group-ing a compute privilege. The typing rule for this operation, **T-PARTITION**, requires that the data privilege of the variable to be partitioned, x , is the same as the current compute privilege. After partitioning, a fresh variable y is introduced with a new privilege π/c . We use π/c as a shorthand to denote a division of a privilege π into c equally-sized, smaller privileges. Within the body of the partition statement, we disallow any references to the original variable x . Then, we continue checking the body of the partition statement against the original compute privilege π because the partition has no affect on the current compute privilege.

Unlike partition, which divides up a piece of data equally into lower privileges, the claim operation, mirroring the split construct, gives the entire variable to exactly one, smaller privilege. Accordingly, Bundl needs to ensure that only one branch of a split operation, with the appropriate privilege, can refer to the claimed variable. To ensure that this is the case, the **T-CLAIM** rule links the data privilege of the variable to the compute privilege of the code claiming it by changing both at the same time. This represents a minor difference from Snap, which implements additional static analysis to ensure that a claimed variable is only accessed in a single split branch.

Lastly, the **T-LOWER** rule mirrors the **T-DESTRUCT** rule; it uses the \downarrow operator to move a variable from one level of the hierarchy to another, distributing it equally among all the child privileges in the same manner as **T-PARTITION**.

687 4.2 Bundl Semantics

688 To reflect the fact that GPU programs execute in parallel across numerous threads, we model the
 689 semantics of Bundl in the style of Turon et al. [16], using a two-level small step judgment. We present
 690 the key rules of this semantics in Figure 12.

691 The top level (i.e., device-level) judgment has just one rule: **S-PROGRAM**. This rule acts as a “frame”
 692 for the lower level (i.e., thread-level) judgment, and steps a collection of thread-id-indexed local
 693 memories (L), a collection of block-id-indexed shared memories (S), a global memory (Σ), and a *thread*
 694 *pool* to an updated collection of memories and thread pool. The thread pool maps thread and block ids
 695 to code, intuitively representing the program being executed by each thread at the current moment.
 696 The **S-PROGRAM** non-deterministically chooses a thread id and block id and steps it according to
 697 the thread-level judgment. This allows the semantics to model the full range of non-deterministic
 698 behavior arising from the GPU’s thread scheduler.²

699 The thread-level judgment has the shape $\eta, \sigma, \Sigma, t, b, p \vdash^\pi s \rightsquigarrow s' \dashv \eta', \sigma', \Sigma'$. The η s in this judgment
 700 represent the thread’s local memory, the σ s represent the shared memory, the Σ s represent global
 701 memory, and the s ’s represent the statement being executed by the thread. The primed versions
 702 of each of these are the output of the step judgment. Critically, notice that a π also appears on the
 703 thread-level judgment just as it does on the typing judgment. This is because the thread semantics
 704 *dynamically tracks and verifies privileges*. The same way a program can get stuck if a value does not
 705 have the right type, the semantics of Bundl also get stuck if code attempts to access data or invoke
 706 commands with the wrong privilege. This runtime privilege is present in Bundl, but is erased by Snap
 707 during compilation; we will later use it to prove that well-privileged programs will always execute
 708 with the same privilege that the type system checked them against.

709 Other variables in the judgment track the positions of threads within their scopes. The t variable
 710 represents the thread’s id, b is the id of the block in which the thread lives, and p is the relative
 711 position of the thread in the compute privilege with which it is executing. For example, threads
 712 with a (Thread,4) privilege each have a unique p between 0 and 3, while threads with a (Block,16)
 713 privilege will have the same p within in each block, between 0 and 15 depending on the block.

714 The semantic rules for privileges involve manipulating p to track which threads take which code
 715 paths when privileges are split or grouped. Notice that in the **S-PROGRAM** rule, the thread stepping
 716 judgment always begins with privilege (Grid,1): all the privilege management rules are congruences,
 717 handling further evaluation with lower privileges as determined by the particular syntactic construct
 718 used.

719 The first set of these privilege management rules handle the `split` construct. They enforce the
 720 same alignment requirements as the **T-SPLIT** rule. The rules choose which path to take based on the
 721 value of p : if p is less than the size of the left-hand privilege n_1 , execution of s_1 will proceed with
 722 that lower privilege via the **S-SPLIT-LEFT** rule, while if p is between n_1 and n_2 execution of s_2 will
 723 proceed via the **S-SPLIT-RIGHT**. In this latter case, we execute s_2 with a new p value that subtracts off
 724 the value of n_1 . As an example, if we encounter a `split(2,2){s1}{s2}` with a (Thread,4) privilege, the
 725 first two threads will go “left” and execute s_1 with a (Thread,2) privilege, while the third and fourth
 726 threads will go “right” and execute s_2 with a (Thread,2) privilege.

727 The **S-GROUP** rule is much simpler, as every thread with a compute privilege that encounters this
 728 operation will execute the sub-statement s . What changes across threads is how their p value will be
 729 modified by the grouping operation. In each case, the p of the thread is reduced modulo n , where n is
 730

731 ¹Bundl is abstracted over these B and T values, so instead of tracking classes of privileges the way that Snap does, it only
 732 tracks the top-level privilege described in Section 3.2.

733 ²In reality, the GPU’s thread scheduler guarantees that warps execute their threads in lockstep, but this modeling every thread
 734 as completely independent is both simpler and a conservative overestimate of the nondeterministic behavior of the GPU.

736 the size of the privilege with which s is executed. This effectively recolors all the threads privilege
 737 the new, smaller, privilege with which they will see s .

738 The **S-DESTRUCT-BLOCK** rule for `destruct` behaves similarly to **S-GROUP**. However, as p is an
 739 index into the current privilege, it is always 0 when a privilege has size 1. The `destruct` statement
 740 can only be executed with a size 1 privilege, and so p is not particularly useful here. To compute the
 741 new p for the execution of s with a privilege of size T , we reduce the thread id modulo T and continue
 742 executing s with that as the new value of p . A similar rule exists for destructing a `Grid` privilege into
 743 a size B `Block` privilege.

744 These rules take great care to ensure that p always describes the relative position of a compute
 745 resource within its privilege; the payoff is that Bundl’s semantics can later use this p value to model
 746 the way that Snap automatically adjusts indices into data when partitioning a data privilege.

747 Beyond these key rules for compute privilege management, we have modeled many of the other
 748 features of Snap, such as asynchronous operations and thread synchronization, in Bundl. To handle
 749 features like these we equip the operational semantics with additional structure, including sets of
 750 semaphores [6] for thread syncs and a stack of effect handlers for modeling deferred asynchronous
 751 computations in the style of Ahman and Pretnar [2]. We have elided these details here for simplicity,
 752 but interested readers can see further details in Appendix A.3.

753 4.3 Soundness Theorem

754 Together, the type system and operational semantics help us prove the following syntactic soundness
 755 theorem, which says that Bundl programs are type safe and do not get stuck trying to execute
 756 operations for which they lack the required privilege:

757 THEOREM 4.1. (*Type-and-Privilege Safety*). *For any program s such that $\Gamma \vdash^\pi s$, either:*

- 758 (1) s is `skip`, or
- 759 (2) *for any well-typed environments η , σ , and Σ , there is an s' , η' , σ' , and Σ' such that*
 $\eta, \sigma, \Sigma, t, b, p \vdash^\pi \sim s' \dashv \eta', \sigma', \Sigma'$ *where η' , σ' , and Σ' are well-typed by an extension of Γ' .*

760 PROOF. Via the usual progress and preservation lemmas. The full proof of these lemmas can be
 761 found in Appendix A.4. \square

762 It is worth noting that this soundness theorem is a safety theorem, not a liveness theorem. The
 763 theorem does guarantees that if all threads that can reach a point in the program with statically-
 764 determined privilege π do reach that point, then those threads will be logically grouped according to
 765 that same privilege π . It does not, however, guarantee that all threads that *can* reach the point will
 766 eventually do so. Indeed, in the presence of nontermination, liveness does not hold: a subset of the
 767 threads could split off and loop forever. While we believe the liveness version of this theorem holds
 768 for a terminating fragment of Bundl, it is beyond the scope of this paper—such proofs are notoriously
 769 challenging and are often research contributions [3, 7, 16] in and of themselves. We hope to tackle
 770 this proof for Bundl in future work.

771 5 Compilation

772 If a program typechecks, Snap lowers the program into a CUDA file, which does not contain any
 773 run-time checks. The CUDA file is then compiled by nvcc, NVIDIA’s closed-source compiler, to
 774 produce a machine-executable. As Snap is an imperative language that provides the same level of
 775 abstraction as CUDA, there is a one-to-one mapping between most language constructs and their
 776 CUDA counterparts.

777 Snap is implemented as a compiler tool on top of the CLANG/LLVM compiler stack. Textually,
 778 Snap programs are written as the examples described in the paper, but the constructs themselves

lower to traditional C++ statements that our frontend tool can analyze and generate code for. This is accomplished by making liberal use of attributes, macros, and other sugar features of C++.

- Make correspondence to CUDA clear
- Would be awesome to have a super clear figure about how the compilation step actually works (show the one-to-one mapping, though it is slightly more complicated)
- Make optimizations (lazy sync, commit group etc) as a separate optimization section.
- highlight how new function interfaces work out: we add extra arguments

6 Evaluation

We evaluate Snap in the context of the four main questions:

RQ1 Can Snap express a diversity of CUDA programs?

RQ2 Can Snap express programs that use advanced GPU features?

RQ3 Can Snap match the performance of existing, speed-of-light CUDA code?

Although the following question is not central to our paper’s main claims—since library design ultimately reflects choices made by programmers, and our language allows abstraction design much like any other language—we were, nevertheless, curious to explore it based on empirical and anecdotal observation:

RQ4 Can Snap help build compositional libraries that users can use with confidence?

To perform this evaluation, we use two GPUs. The first is the NVIDIA H100 SXM5, a server-grade chip that supports tensor-core operations and a dedicated hardware copy engine, the Tensor Memory Accelerator. Notably, the H100 introduces a new *logical* level called the *warp group* (collection of 4-aligned warps), and we show that our programming model can accommodate this new level. Moreover, because the H100 has historically served as the primary GPU for large-scale AI training, many CUDA kernels on this hardware are already highly optimized and achieve near-speed-of-light performance, providing a rigorous baseline for comparison. To ensure our results generalize beyond the H100, we additionally test programs on a second GPU, the NVIDIA 4070 SuperTi GPU, a consumer-grade chip.

Like mentioned in Section 5, when name typechecks a program, it produces a CUDA file. We compile the CUDA file with nvcc version with flags . To measure performance of our benchmark, we set up the values from an input using a random number generator. We use the average running time sampled across 10 iterations, with a warm-up of 5 iterations.

6.1 RQ1: Can Snap express a diversity of CUDA programs?

To answer this research question, we evaluate Snap on two benchmarks: a sequence of matrix multiplication kernels that grow in complexity and a single-pass parallel prefix scan with decoupled look-back that requires programmers to think carefully about multiple points of convergence.

Matrix Multiplication. We chose matrix multiplication as our first benchmark for two main reasons. First, there exist several implementations that achieve near-peak performance, providing a strong baseline. Second, matrix multiplication allows for a range of increasingly sophisticated implementations that stress different parts of the language, making it ideal for evaluating expressiveness.

To undertake this exploration, we adapt the codebase from to implement a matrix multiplication on the float datatype, also known as sgemm, in Snap. As this is a float benchmark, it *does not need to* use advanced GPU features like asynchrony or tensor cores to achieve speed-of-light performance. We will discuss these advanced features in Section 6.2.

We implement several variants of the sgemm benchmark, reproduced in Appendix . These include a naive baseline, a version that exploits memory coalescing, 2D blocking with shared-memory staging,

834 2D blocking with vectorized loads, and a warp-tiling strategy that effectively introduces an additional
 835 level of tiling from the perspective of a warp. We find that the runtime performance of our Snap
 836 implementations closely matches that of the original versions. This is expected, since the generated
 837 code is nearly identical to the hand-written versions, aside from small differences such as hoisted
 838 expressions and a few index calculations. The downstream compiler (`nvcc`) readily inlines hoisted
 839 expressions that are used only once, so the performance-critical inner loops remain effectively
 840 unchanged.

841 *Single-pass Parallel Prefix Scan with Decoupled Look-Back.* We also implement scan, a widely used
 842 parallel primitive, in Snap. We focus on the prefix-sum scan, which computes, for each position in an
 843 array, the sum of all elements up to that position. Prefix-sum sits in a different corner of the GPU
 844 design space from matrix multiplication: it is memory intensive, requires careful attention to the
 845 convergence behavior of different threads, and traditionally requires multiple passes over data.
 846

847 We implement the single-pass parallel prefix scan with decoupled look-back, introduced by Merrill
 848 and Garland , an elegant algorithm that does not require multiple passes over the input data. In
 849 this algorithm, each block computes a local prefix sum over its assigned region of the array. Once
 850 finished, it writes the final element of its region into a global array. Each block then *looks back* to
 851 accumulate the contributions of prior blocks, allowing them to independently determine the global
 852 prefix without a full sweep over memory. This decoupled look-back mechanism lets blocks progress
 853 at different speeds while still producing correct global results.

854 The algorithm involves several distinct points of convergence. Within each block, work is de-
 855 composed into fine-grained thread-level and warp-level scans. After producing the local result,
 856 blocks publish their partial prefix to global memory. Finally, each block waits until enough prefix
 857 information from earlier blocks becomes available, at which point it accumulates the value and
 858 completes its section of the global scan.

859 We implement this full strategy in Snap, available in Appendix . Our implementation uses the
 860 unsafe feature to implement a global-memory spinlock that lets blocks check when the preceeding
 861 block's data is ready.

862 **6.2 RQ2: Can Snap express programs that use advanced GPU features?**

863 To test whether Snap can express programs that use advanced features of modern GPUs, we write a
 864 matrix multiplication for the `bf16` datatype for the H100.

865 This benchmark is an acid test of our language, as the H100 `bf16` matrix multiplication pushes
 866 several language features to the extreme. To write a matrix multiplication that can hit peak throughput
 867 on an H100, we need to write a warp-specialized kernel that uses the Tensor Memory Accelerator
 868 (TMA) , an asynchronous hardware copy engine that can move tiles of data at a time, and uses
 869 the warp-group-level tensor core instructions, or `wmma`, specifically introduced for the Hopper
 870 architecture.

871 The implementation works as follows. First, we assign each block on our machine a logical tile of
 872 the output to compute. The block is then divided into a producer warp-group and a set of consumer
 873 warp-groups, where the producer loads data for multiple consumers, and both must signal to each
 874 other when one is done loading and the other is ready to compute. In this way, the benchmark
 875 overlaps computation with data movement by pipelining loads.

876 The implementation in Snap looks different from normal CUDA code, particularly in how pipelining
 877 is expressed. Since Snap uses *names* and subsequent `partition` or `claim` constructs to determine the
 878 synchronization each region requires, when pipelining, we cannot dynamically change the pipeline
 879 slot simply by maintaining an index variable that wraps around based on the pipeline length. Instead,
 880 each pipeline slot must be given a separate name so that Snap can track them independently and
 881

883 actually overlap compute with data-movement. This leads to pipeline slots that must be individually
 884 named and forces the load logic to be effectively “unrolled,” since we can no longer iterate over a
 885 pipeline-slot index. In turn, this forces that all pipelines in Snap be statically sized. In practice, these
 886 pipelines are statically sized in CUDA anyway because they occupy shared memory, which is a small,
 887 finite resource that must be explicitly managed.

888 Notably, for this benchmark, to get the wmma instruction to work we did not need to add a special
 889 privilege in our language; we could simply use `thread[128]`. We did, however, need to add a special
 890 TMA asynchronous data-movement construct, since Snap will eventually need to synchronize these
 891 transfers.

893 6.3 RQ3: Can Snap match the performance of existing, speed-of-light CUDA code?

894 In Section 6.1 and Section 6.2, we examined programs that expressed the same computation in
 895 multiple ways, expressed operations that rely on multiple points of convergence, and used advanced
 896 GPU features. We now discuss the performance of these programs.

897 The performance of the matrix multiplication benchmarks is shown in , where we demonstrate
 898 that Snap is competitive with cuBLAS.

899 For the prefix scan, we compare our performance to , the library introduced earlier in Section 1,
 900 and show that we can reach approximately % of CUB’s peak bandwidth.

901 Finally, we evaluate our H100 implementation and show that it is competitive with cuBLAS on
 902 square sizes, falling between %. We emphasize that this is an exacting benchmark, and achieving this
 903 level of performance requires writing large kernels with precise control over low-level features.

905 6.4 Case Study: Can Snap help with library design?

906 Over the course of our evaluation, we found ourselves developing a small library of functions—similar
 907 in spirit to the CUB library—that we could call to execute our operations. Based on our experience,
 908 we would like to study, qualitatively, if Snap can help programmers design libraries that they can
 909 compose with confidence.

910 As mentioned in Section 1, CUB occupies a unique design point in the GPU library ecosystem.
 911 Unlike many other libraries such as cuBLAS, cuDNN, and cuSparse, which provide global interfaces
 912 that users can call and configure, CUB provides a device-side library organized into different levels.
 913 It makes these levels apparent by prefixing each of its functions with Device, Block, and Thread.
 914 The single-scan prefix sum in Section 6.1 used these functions in Snap already.

915 Now, we turn our attention to a particular function in CUB—the store function—and examine
 916 how it is equivalently expressed in Snap, and how, in our experience, Snap’s interfaces the assumptions
 917 implicit in CUB apparent to the user.

918 In CUB, the load function is implemented as a class, as shown below.

```
921 1 template<typename T, int BlockDimX,
922   int ItemsPerThread, BlockLoadAlgorithm Algorithm = BLOCK_LOAD_DIRECT,
923   int BlockDimY = 1, int BlockDimZ = 1>
924 4 class BlockStore:
```

925 To use it, users must first instantiate an object of this class, and then call it with shared memory.

```
926
927 1 using BlockLoad = cub::BlockLoad<int, 128, 4, BLOCK_LOAD_DIRECT>;
928 2 // Allocate shared memory for BlockLoad
929 3 __shared__ typename BlockLoad::TempStorage temp_storage;
930 4 int thread_data[4]; // Thread local data
931 5 BlockLoad(temp_storage).Load(d_data, thread_data);
```

932 CUB exposes a leaky abstraction, where information about the number of threads, block sizes, and
 933 other details leaks through. There are several details that leak through:

- 934 (1) The CUB documentation needs to specify the number of threads that the function can assume
 935 to be available, because within the function, each thread must locate itself in the computation
 936 and use its `threadIdx.x` accordingly. If the starting `threadIdx.x` is not 0, the function must
 937 compute its relative ID internally.
 938 (2) The “item per thread” design is interesting and serves two purposes. The first is related to
 939 performance: if loops have constant bounds, they can be unrolled, enabling downstream
 940 optimizations. The second is related to correctness: the function relies on the assumption
 941 that all threads call the function with an equal number of values to load.
 942 (3) The CUB documentation also needs to specify that `thread_data` is local; if the same pointer
 943 to shared memory is passed, threads may overwrite each other’s values.
 944 (4) Finally, the CUB documentation specifies that if shared memory is being overwritten, a
 945 `__syncthreads()` call must be made to ensure that all reads have completed.
 946

947 On the other hand, the same function has a completely different interface in Snap:
 948

```
949 1 #include <cuda.h>
950 2 #require(block[1], thread[1])
951 3 def load(src: ptr(const(shared(int))) block[1],
952 4     item_per_thread : int block[1],
953 5     total_size: int block[1],
954 6     thread_data : ptr(int) thread[1]):
```

954 In Snap, we are able to encapsulate code effectively, reducing the need to communicate numerous
 955 implementation details through documentation:

- 956 (1) We do not need to pass in the number of threads at all. Whenever Snap calls a function, it
 957 threads the relative ID through, so each function can be written locally as if it were running
 958 alone, rather than having to determine where the thread resides in the global array.
 959 (2) We do not need to make `item_per_thread` a template argument for *correctness*. Its frequency
 960 is set at the function signature, so Snap will never allow a function to be called with a
 961 lower-privilege value.
 962 (3) In our interface, `thread_data` is explicitly set to a thread-local value. Since it is not marked
 963 as `const`, Snap conservatively assumes it may be written to, and enforces at compile time
 964 that only `thread[1]` values are passed in.
 965 (4) Finally, using our synchronization pass outlined in Section 3.9, a `__syncthreads()` call will
 966 be inserted automatically if `src` is going to be used for writing.
 967

968 7 Related Work

969 *GPU Programming Languages.*

970 *Theoretic Foundations.* The design of Bundl is heavily inspired by existing work on coeffect systems
 971 [12, 13] and dependency tracking [1]. Coeffects allow type systems to talk about how programs
 972 depend on their environments, and have achieved widespread use in Rust [14] in the form of linear
 973 types. In Snap and Bundl, privileges act like coeffects by describing what compute resources are
 974 necessary for programs to execute. Dependency tracking calculi, meanwhile, allow type systems to
 975 track how data and code depend on each other, and have commonly been used to implement secure
 976 information flow analyses [5]. In Bundl, low privilege data is unable to flow into high privilege contexts,
 977 and we use dependency tracking to capture this restriction in Bundl’s type system.
 978

981 8 Conclusion, Limitations, and Future Work

982 We have presented Snap, a new low-level GPU language that statically guarantees safe usage of
 983 compute resources by construction, and have demonstrated that it is possible to achieve this safety
 984 without sacrificing performance. We hope our design can inform the implementation of existing
 985 libraries like CUB that attempt to achieve the same goals without static enforcement.

986 We believe Snap is ripe for extension with additional features and capabilities to further improve its
 987 expressivity and usefulness to GPU programmers. Currently, Snap does not have support for explicit
 988 pointers; programmers must use built-in constructs in order to interact with memory. However,
 989 prior work like Descend [9] provides a promising template for extending Snap with more permissive
 990 pointer usage. In particular, were we to restrict Snap’s partition construct to a specific set of
 991 injective functions identified by Descend, we believe we could guarantee data-race freedom in Snap.
 992

993 Another promising area for future work is expanding the set of architectures on which Snap and
 994 its design principles can be used. While Snap currently only supports programming the GPU, we
 995 believe we can extend the core ideas cleanly to the boundary with the CPU and in general to other
 996 heterogeneous hardware architectures. Indeed, we do not think that the idea of compute privileges
 997 is useful only on the GPU, and we would like to investigate how we might adapt Snap to other
 998 hierarchical programming systems. Conversely, we also believe it is possible for Snap to support
 999 more GPU architectures featuring even coarser-grained tensor-core operations like Blackwell [4].
 1000

1001 We also would like to improve the ergonomics of programming with Snap, in particular by
 1002 improving the pipelining experience and providing users with more control over synchronization.
 1003 As we saw in Section 6.2, Snap requires pipeline slots to be explicitly named and all pipelines to
 1004 be statically sized. We believe we can improve upon this by adding explicit language constructs
 1005 for organizing pipelines to reduce boilerplate code in this case. Similarly, all syncs are currently
 1006 handled by Snap’s compiler, but we would like to extend the language with user-level synchronization
 1007 operations, leveraging the existing privilege system to avoid deadlocking.
 1008

1009 On the theoretical side, we are interested in further exploration of Bundl. In particular, we would
 1010 like to examine a terminating fragment of the calculus and prove the liveness property discussed in
 1011 Section 4.3: that every thread in the logical grouping denoted by privilege π will always reach the
 1012 parts of a program with that code privilege. We believe this is equivalent to showing that all threads
 1013 with the same privilege always execute the same code and observe the same data, and are optimistic
 1014 that this can be proved via logical relation, following in the footsteps of Turon et al. [16].
 1015

1013 References

- 1014 [1] Martín Abadi, Anindya Banerjee, Nevin Heintze, and Jon G. Riecke. 1999. A core calculus of dependency. In *Proceedings*
 1015 *of the 26th ACM SIGPLAN-SIGACT symposium on Principles of programming languages* (San Antonio Texas USA, 1999-01).
 1016 ACM, 147–160. [doi:10.1145/292540.292555](https://doi.org/10.1145/292540.292555)
- 1017 [2] Danel Ahman and Matija Pretnar. 2021. Asynchronous effects. *Proceedings of the ACM on Programming Languages* 5,
 1018 POPL (Jan. 2021), 1–28. [doi:10.1145/3434305](https://doi.org/10.1145/3434305)
- 1019 [3] Lars Birkedal, Filip Sieczkowski, and Jacob Thamsborg. 2012. A Concurrent Logical Relation. In *Computer Science Logic
 1020 (CSL’12) - 26th International Workshop/21st Annual Conference of the EACSL (2012)*. Schloss Dagstuhl – Leibniz-Zentrum
 1021 für Informatik, 107–121. [doi:10.4230/LIPIcs.CSL.2012.107](https://doi.org/10.4230/LIPIcs.CSL.2012.107)
- 1022 [4] NVIDIA Corporation. 2025. NVIDIA Blackwell Architecture Technical Brief. (2025). <https://resources.nvidia.com/en-us-blackwell-architecture>
- 1023 [5] Dorothy E. Denning and Peter J. Denning. 1977. Certification of programs for secure information flow. *Commun. ACM*
 1024 20, 7 (July 1977), 504–513. [doi:10.1145/359636.359712](https://doi.org/10.1145/359636.359712)
- 1025 [6] Edsger Wybe Dijkstra. 1963. Over de sequentialiteit van procesbeschrijvingen. (1963). <https://training-ir7.tdl.org/handle/123456789/594>
- 1026 [7] Azadeh Farzan, Zachary Kincaid, and Andreas Podelski. 2016. Proving Liveness of Parameterized Programs. In
 1027 *Proceedings of the 31st Annual ACM/IEEE Symposium on Logic in Computer Science*. ACM, New York NY USA, 185–196.
 1028 [doi:10.1145/2933575.2935310](https://doi.org/10.1145/2933575.2935310)

- 1030 [8] Troels Henriksen, Niels G. W. Serup, Martin Elsman, Fritz Henglein, and Cosmin E. Oancea. 2017. Futhark: purely
1031 functional GPU-programming with nested parallelism and in-place array updates. In *Proceedings of the 38th ACM*
1032 *SIGPLAN Conference on Programming Language Design and Implementation*. ACM, Barcelona Spain, 556–571. [doi:10.1145/3062341.3062354](https://doi.org/10.1145/3062341.3062354)
- 1033 [9] Bastian Köpcke, Sergei Gorlatch, and Michel Steuwer. 2024. Descend: A Safe GPU Systems Programming Language. 8
1034 (2024), 841–864. Issue PLDI. [doi:10.1145/3656411](https://doi.org/10.1145/3656411)
- 1035 [10] NVIDIA Corporation. 2025. *CUB: CUDA Unbound*. <https://docs.nvidia.com/cuda/cub/index.html> CUDA Toolkit
1036 Documentation.
- 1037 [11] NVIDIA Corporation. 2025. *CUDA C++ Programming Guide*. NVIDIA. <https://docs.nvidia.com/cuda/cuda-c-programming-guide/> Accessed: 2025-11-06.
- 1038 [12] Tomas Petricek, Dominic Orchard, and Alan Mycroft. 2013. Coeffects: Unified Static Analysis of Context-Dependence.
1039 In *Automata, Languages, and Programming*, Fedor V. Fomin, Rūsiņš Freivalds, Marta Kwiatkowska, and David Peleg
1040 (Eds.). Springer, Berlin, Heidelberg, 385–397. [doi:10.1007/978-3-642-39212-2_35](https://doi.org/10.1007/978-3-642-39212-2_35)
- 1041 [13] Tomas Petricek, Dominic Orchard, and Alan Mycroft. 2014. Coeffects: a calculus of context-dependent computation.
1042 *SIGPLAN Not.* 49, 9 (Aug. 2014), 123–135. [doi:10.1145/2692915.2628160](https://doi.org/10.1145/2692915.2628160)
- 1043 [14] Rust Team. 2025. Rust Programming Language. <https://rust-lang.org/>
- 1044 [15] Philippe Tillet. 2025. Introducing Triton: Open-source GPU programming for neural networks. <https://openai.com/index/triton/>
- 1045 [16] Aaron J. Turon, Jacob Thamsborg, Amal Ahmed, Lars Birkedal, and Derek Dreyer. 2013. Logical relations for fine-grained
1046 concurrency. In *Proceedings of the 40th annual ACM SIGPLAN-SIGACT symposium on Principles of programming languages*
1047 (Rome Italy, 2013-01-23). ACM, 343–356. [doi:10.1145/2429069.2429111](https://doi.org/10.1145/2429069.2429111)
- 1048
- 1049
- 1050
- 1051
- 1052
- 1053
- 1054
- 1055
- 1056
- 1057
- 1058
- 1059
- 1060
- 1061
- 1062
- 1063
- 1064
- 1065
- 1066
- 1067
- 1068
- 1069
- 1070
- 1071
- 1072
- 1073
- 1074
- 1075
- 1076
- 1077
- 1078

1079 **A Complete Bundl Type System, Semantics, and Syntactic Soundness Proofs**

1080 **A.1 Basic Definitions**

1081
 1082 Hierarchy Levels $h ::= \text{Grid} \mid \text{Block} \mid \text{Thread}$
 1083 Memory Kinds $l ::= \text{Local} \mid \text{Shared} \mid \text{Global}$
 1084 Privileges $\pi : h \times \mathbb{N}$
 1085 Base Types $b ::= \text{bool} \mid \text{int} \mid \text{float}$
 1086 Types $\tau ::= b \mid b[]^l \mid \text{Fun}(\Gamma, \pi, m) \mid \text{async } \tau$
 1087 Contexts $\Gamma ::= \cdot \mid \Gamma, x :^\pi \tau$
 1088
 1089 Shared Memory Remaining $m : \mathbb{N}$

1090
 1091
 1092 Privileges π are part of an algebra parameterized over some constant values T (the number of threads
 1093 per block) and B (the number of blocks per grid). With these values, we have two isomorphisms:
 1094

$$(1095) \quad (\text{Block}, 1) \cong (\text{Thread}, T)$$

1096 and

$$(1098) \quad (\text{Grid}, 1) \cong (\text{Block}, B)$$

1099 The group and split operations of Snap allows us to move along this isomorphism, from left
 1100 to right. For clarity, in Bundl we split these operations into three: a split operation that can split
 1101 privileges into multiple smaller ones, and a destruct operation that directly moves us along the
 1102 isomorphism, and a group that divides our current privilege into equal sized parts.

1103 Privileges π are also lexicographically ordered in the obvious way. hs are ordered such that
 1104 $\text{Thread} \leq \text{Block} \leq \text{Grid}$, and $(h_1, n_1) \leq (h_2, n_2)$ iff $n_1 \leq n_2$ and $h_1 \leq h_2$.

1105 We define scalar multiplication $i \times h$ of natural numbers with hs : $i \times (h, n) = (h, in)$.

1106 We also define division of privileges and hierarchy levels of type $\pi \times \pi \rightarrow \mathbb{N}$. $\text{Grid}/\text{Block} = B$ and
 1107 $\text{Block}/\text{Thread} = T$. We lift this to privileges like so: $(h_1, n_1)/(h_2, n_2) = ((h_1/h_2) \cdot n_1)/n_2$.

1108 Lastly we define a partial \downarrow operator on hs such that $\downarrow \text{Block} = \text{Thread}$ and $\downarrow \text{Grid} = \text{Block}$. Note
 1109 that $\downarrow \text{Thread}$ is undefined. This operator lifts to π s whose second component is 1 and encodes the
 1110 leftward component of the isomorphism above: $\downarrow (\text{Block}, 1) = (\text{Thread}, T)$ and $\downarrow (\text{Grid}, 1) = (\text{Block}, B)$.

1111 Note also that the presentation of these rules in the main body of the paper elide the m portion,
 1112 which tracks the maximum amount of memory a given computation is allowed to use. In the full
 1113 system presented here, both the typing rules and the operational semantics carry an additional piece
 1114 of information tracking allocated memory.

1116 **A.2 Typing Rules**

1117 **A.2.1 Expressions.**

$$1118 \quad \frac{x :^\pi \tau \in \Gamma}{\Gamma \vdash^\pi x : \tau} \text{-Var} \quad \frac{}{\Gamma \vdash^\pi n : \text{int}} \text{-Int} \quad \frac{}{\Gamma \vdash^\pi f : \text{float}} \text{-Float}$$

$$1121 \quad \frac{}{\Gamma \vdash^\pi b : \text{bool}} \text{-Bool} \quad \frac{\pi < (\text{Grid}, 1)}{\Gamma \vdash^\pi \text{partition_id} : \text{int}} \text{-Partition-Id}$$

$$1124 \quad \frac{\Gamma \vdash^{\pi'} e_1 : \tau[]^l \quad \Gamma \vdash^\pi e_2 : \text{int} \quad l = \text{Global or } l = \text{Local} \quad \pi \leq \pi'}{\Gamma \vdash^\pi e_1[e_2] : \tau} \text{-Arr-Access}$$

1128		1129
	$\frac{\Gamma \vdash^{\pi'} e_1 : \tau[\cdot]^{\text{Shared}} \quad \Gamma \vdash^{\pi} e_2 : \text{int} \quad \pi \leq (\text{Block}, 1) \quad \pi \leq \pi'}{\Gamma \vdash^{\pi} e_1[e_2] : \tau}$	T-Arr-Access-Shared
1130		
1131		1132
	$\frac{\Gamma \vdash^{\pi} e_1 : \text{int} \quad \Gamma \vdash^{\pi} e_2 : \text{int}}{\Gamma \vdash^{\pi} e_1 \text{ bop } e_2 : \text{int}}$	T-Bop
1133		
1134		A.2.2 Statements.
1135	$\frac{f : \pi' \mathbf{Fun}(x_i : \tau_i, \pi, m') \in \Gamma \quad \Gamma \vdash^{\pi} e_i : \tau_i \quad m' \leq m}{\Gamma \vdash_m^{\pi} f(e_1, \dots, e_n)}$	T-Function-Call
1136		
1137		
1138		1139
	$\frac{\Gamma \vdash_m^{(h,n_1)} s_1 \quad \Gamma \vdash_m^{(h,n_2)} s_2 \quad n_1, n_2 \text{ align to } n}{\Gamma \vdash_m^{(h,n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\}}$	T-Split
1140		
1141		1142
	$\frac{\Gamma \vdash_m^{(h,n)} s}{\Gamma \vdash_m^{(h,q \cdot n)} \text{group } q \ s}$	T-Group
1143		
1144		1145
	$\frac{\Gamma \vdash_m^{\pi} \text{init}_{\psi}}{\Gamma \vdash_m^{\pi} \text{dec}_{\psi}}$	T-Sync-Init
1146		
1147		1148
	$\frac{}{\Gamma \vdash_m^{\pi} \text{wait}_{\psi}}$	T-Sync-Wait
1149		
1150		1151
	$\frac{\Gamma \vdash^{\pi} e : \tau \quad \Gamma, x : \pi' \vdash_m^{\pi} s \quad \pi' \leq \pi \quad \pi' \pi \quad \tau \text{ not an array type}}{\Gamma \vdash_m^{\pi} x : \tau @ \pi' = e \text{ in } s}$	T-Decl
1152		
1153		1154
	$\frac{\Gamma \vdash^{\pi'} e_1 : \tau[\cdot]^l \quad \Gamma \vdash^{\pi} e_2 : \text{int} \quad \Gamma \vdash^{\pi} e_3 : \tau \quad l = \text{Global or } l = \text{Local} \quad \pi' \leq \pi \quad \pi' \pi}{\Gamma \vdash_m^{\pi} e_1[e_2] = e_3}$	T-Arr-Assn
1155		
1156		1157
	$\frac{\Gamma \vdash^{\pi'} e_1 : \tau[\cdot]^{\text{Shared}} \quad \Gamma \vdash^{\pi} e_2 : \text{int} \quad \Gamma \vdash^{\pi} e_3 : \tau \quad \pi \leq (\text{Block}, 1) \quad \pi' \leq \pi \quad \pi' \pi}{\Gamma \vdash_m^{\pi} e_1[e_2] = e_3}$	T-Arr-Assn-Shared
1158		
1159		1160
	$\frac{\Gamma \vdash^{\pi} e : \text{bool} \quad \Gamma \vdash_{m_1}^{\pi} s_1 \quad \Gamma \vdash_{m_2}^{\pi} s_2}{\Gamma \vdash_{\max(m_1, m_2)}^{\pi} \text{if } e \text{ then } s_1 \text{ else } s_2}$	T-If
1161		
1162		1163
	$\frac{\Gamma \vdash^{\pi} e : \text{bool} \quad \Gamma \vdash_m^{\pi} s}{\Gamma \vdash_m^{\pi} \text{while } e \text{ do } s}$	T-While
1164		
1165		1166
	$\frac{\Gamma, x : \pi' \vdash_m^{\pi} s \quad l = \text{Global or } l = \text{Local}}{\Gamma \vdash_{m+n \cdot \text{size}(\tau)}^{\pi} x := \text{alloc } l \ \tau \ n \text{ in } s}$	T-Alloc
1167		
1168		1169
	$\frac{\Gamma, x : (\text{Block}, 1) \ \tau[\cdot]^{\text{Shared}} \vdash_m^{\pi} s}{\Gamma \vdash_{m+n \cdot \text{size}(\tau)}^{(\text{Block}, 1)} x := \text{alloc Shared } \tau \ n \text{ in } s}$	T-Alloc-Shared
1170		
1171		1172
	$\frac{\Gamma, y : (h, n/c) \ \tau[\cdot]^l \vdash_m^{(h,n)} s \quad c n \quad l \neq \text{Local}}{\Gamma, x : (h, n) \ \tau[\cdot]^l \vdash_m^{(h,n)} \text{partition}_{\psi} x \text{ into } y \text{ by } c \text{ in } s}$	T-Partition
1173		
1174		1175
	$\frac{\Gamma, y : (h, n') \ \tau[\cdot]^l \vdash_m^{(h,n')} s \quad n' \leq n \quad l \neq \text{Local}}{\Gamma, x : (h, n) \ \tau[\cdot]^l \vdash_m^{(h,n)} \text{claim}_{\psi} x \text{ into } y \text{ at } n' \text{ in } s}$	T-Claim
1176		

1177 $\frac{\Gamma, y : \downarrow^\pi \tau[]^l \vdash_m^\pi s \quad l \neq \text{Local}}{\Gamma, x :^\pi \tau[]^l \vdash_m^\pi \text{lower}_\psi x \text{ into } y \text{ in } s} \textbf{T-Lower}$

1178
1179
1180

1181 $\frac{\Gamma, y :^{(\text{Thread},1)} \text{async } \tau[]^l \vdash_m^{(\text{Thread},1)} s}{\Gamma, x :^{(\text{Thread},1)} \tau[]^l \vdash_m^{(\text{Thread},1)} \text{async_partition}_\phi x \text{ into } y \text{ in } s} \textbf{T-Async-Partition}$

1182
1183
1184

1185 $\frac{}{\Gamma, x :^{(\text{Thread},1)} \text{async } \tau[]^l, y :^{(\text{Thread},1)} \tau[]'' \vdash_m^{(\text{Thread},1)} \text{async_memcpy}(x,y)} \textbf{T-Async-Memcpy}$

1186
1187
1188 $\frac{}{\Gamma, x :^\pi \tau[]^l, y :^\pi \tau[]'' \vdash_m^\pi \text{memcpy}(x,y)} \textbf{T-Memcpy}$

A.3 Complete Bundl Semantics

A.3.1 Definitions.

Global Memory $\Sigma ::= \cdot | \Sigma, n \mapsto^\pi v$

Shared Memory $\sigma ::= \cdot \mid \sigma, n \mapsto^\pi v$

Local Memory $\eta ::= \cdot \mid \eta, n \mapsto^\pi v$

Block Memory Map $S ::= \forall n \in B, n \mapsto \sigma$

Thread Memory Map $L ::= \forall n \in T, n \mapsto \eta$

Synchronization Map $\Psi: \psi \rightarrow \mathbb{N} \rightarrow \mathbb{N}$

Deferred Computations Map $\Phi: \phi \rightarrow \{s\}$

In real GPUs, thread ids are only unique within their block. However, in this calculus for simplicity we assume thread ids are global. One can convert back and forth between this abstracted notion of a thread id and a block-unique id via addition modulo T . That is, $t_{\text{real}} = t_{\text{simplified}} \bmod T$ and $t_{\text{simplified}} = t_{\text{real}} + b \cdot T$.

By convention the names for local and shared and global memory do not conflict, as on the GPU they will be separate pointer spaces. Additionally, we freely interchange between using names for variables and integer locations.

In the main body of the paper, for simplicity we elide the synchronization map and the deferred computation map from the operational semantics, as our theorems do not make any guarantees about non-interference. However, as they are part of the full semantics, we include them here for completeness. By convention the synchronization and deferred computation maps are a total functions, initialized to map to $\lambda .0$ for ψ s and $\lambda .\{\}$ for ϕ s not explicitly initialize.

The shape of the judgment for a single thread is $\eta, \sigma, \Sigma, t, b, p, \Psi \vdash_m^\pi s \rightsquigarrow s' \dashv_m \eta', \sigma', \Sigma', \Psi'$. The t here is the thread id, the b is the block id, and the p is the partition id. The last of these three is modified and managed by the rules for split, group and destruct, and tracks the relative position of the thread within a group. This semantics is in a small step style.

The shape of the judgment for expressions is $\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi} e \Downarrow v$. The two π s represent the ambient compute context (i.e., the context in which resources are being read), while π' , represents the target compute context (i.e., the compute context of the variable into which the result of the expression is going to be written. This is relevant for computing the value of `partition_id`, which divides the two contexts. As a shorthand, we can divide a privilege by a scalar value like so: $(h,n)/c = (h,n)/(h,c)$.

1226 The overall evaluation of a program is expressed as

$$L, S, \Sigma, P, \Psi, \Phi \rightsquigarrow L', S', \Sigma', P', \Psi', \Phi'.$$

1229 In this judgment P serves as a *thread pool*, mapping pairs of thread and block ids (which don't change)
 1230 to statements and memory (which can be updated by stepping). One can think of P as tracking which
 1231 program is running on each thread. This steps according to the following rule:

$$\frac{L(t), S(b), \Sigma, t, b, 0, \Psi, \Phi \vdash_m^{(\text{Grid}, 1)} s \rightsquigarrow s' \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi' \quad P(t, b) = (s, m)}{L, S, \Sigma, P, \Psi, \Phi \rightsquigarrow L[t \mapsto \eta'], S[b \mapsto \sigma'], \Sigma', P[(t, b) \mapsto (s', m')], \Psi', \Phi'} \text{ S-Program}$$

1235 For simplicity of notation, we define an **update** operation that searches the three environments
 1236 for the one that contains the variable being used (by convention, there is no conflict between the
 1237 environments, as in reality they exist in three separate address spaces). We also define a similar **get**
 1238 operation that retrieves a variable from memory, and a **rename** operation that remaps a variable
 1239 with the same value but under a different name.

$$1241 \quad \text{update}(\eta, \sigma, \Sigma, x, v) = (\eta[x \mapsto^\pi v], \sigma, \Sigma) \text{ when } x \in^\pi \eta$$

$$1242 \quad \text{update}(\eta, \sigma, \Sigma, x, v) = (\eta, \sigma[x \mapsto^\pi v], \Sigma) \text{ when } x \in^\pi \sigma$$

$$1243 \quad \text{update}(\eta, \sigma, \Sigma, x, v) = (\eta, \sigma, \Sigma[x \mapsto^\pi v]) \text{ when } x \in^\pi \Sigma$$

$$1244 \quad \text{get}(\eta, \sigma, \Sigma, x) = \eta(x) \text{ when } x \in^\pi \eta$$

$$1245 \quad \text{get}(\eta, \sigma, \Sigma, x) = \sigma(x) \text{ when } x \in^\pi \sigma$$

$$1246 \quad \text{get}(\eta, \sigma, \Sigma, x) = \Sigma(x) \text{ when } x \in^\pi \Sigma$$

$$1247 \quad \text{rename}(\eta, \sigma, \Sigma, x, y, \pi') = (\eta[y \mapsto^\pi \eta(x)], \sigma, \Sigma) \text{ when } x \in^\pi \eta$$

$$1248 \quad \text{rename}(\eta, \sigma, \Sigma, x, y, \pi') = (\eta, \sigma[y \mapsto^\pi \sigma(x)], \Sigma) \text{ when } x \in^\pi \sigma$$

$$1249 \quad \text{rename}(\eta, \sigma, \Sigma, x, y, \pi') = (\eta, \sigma, \Sigma[y \mapsto^\pi \Sigma(x)]) \text{ when } x \in^\pi \Sigma$$

1250 A.3.2 Privilege Management Rules.

$$1251 \quad \frac{p < n_1 \quad n_1, n_2 \text{ align to } n \quad \eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n_1)} s_1 \rightsquigarrow s'_1 \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi'}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\} \rightsquigarrow \text{split}(n_1, n_2)\{s'_1\}\{s_2\} \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi'} \text{ S-Split-Left}$$

$$1252 \quad \frac{p < n_1 \quad n_1, n_2 \text{ align to } n}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n)} \text{split}(n_1, n_2)\{\text{skip}\}\{s_2\} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-Split-Left-Done}$$

$$1253 \quad \frac{p \geq n_1 \quad p < n_1 + n_2 \quad n_1, n_2 \text{ align to } n \quad \eta, \sigma, \Sigma, t, b, p - n_1, \Psi, \Phi \vdash_m^{(h, n_2)} s_2 \rightsquigarrow s'_2 \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi'}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\} \rightsquigarrow \text{split}(n_1, n_2)\{s_1\}\{s'_2\} \dashv_{m'} \eta', \sigma', \Sigma', \Psi', \Phi'} \text{ S-Split-Right}$$

$$1254 \quad \frac{p \geq n_1 \quad p < n_1 + n_2 \quad n_1, n_2 \text{ align to } n}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n)} \text{split}(n_1, n_2)\{s_1\}\{\text{skip}\} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-Split-Right-Done}$$

$$1255 \quad \frac{p \geq n_1 + n_2 \quad n_1, n_2 \text{ align to } n}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n)} \text{split}(n_1, n_2)\{s_1\}\{s_2\} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-Split-None}$$

1275			
1276	$\eta, \sigma, \Sigma, t, b, t \text{ mod } T, \Psi, \Phi \vdash_m^{(\text{Thread}, T)} s \rightsquigarrow s' \dashv_m \eta', \sigma', \Sigma', \Psi', \Phi'$		
1277		$\eta, \sigma, \Sigma, t, b, 0, \Psi, \Phi \vdash_m^{(\text{Block}, 1)} \text{destruct in } s \rightsquigarrow \text{destruct in } s' \dashv_m \eta', \sigma', \Sigma', \Psi', \Phi'$	S-Destruct-Block
1278			
1279			
1280	$\eta, \sigma, \Sigma, t, b, b \text{ mod } B, \Psi, \Phi \vdash_m^{(\text{Block}, B)} s \rightsquigarrow s' \dashv_m \eta', \sigma', \Sigma', \Psi', \Phi'$		
1281		$\eta, \sigma, \Sigma, t, b, 0, \Psi, \Phi \vdash_m^{(\text{Grid}, 1)} \text{destruct in } s \rightsquigarrow \text{destruct in } s' \dashv_m \eta', \sigma', \Sigma', \Psi', \Phi'$	S-Destruct-Grid
1282			
1283			
1284	$\eta, \sigma, \Sigma, t, b, 0, \Psi, \Phi \vdash_m^\pi \text{destruct in skip} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi$		S-Destruct-Done
1285			
1286	$\eta, \sigma, \Sigma, t, b, p \text{ mod } n, \Psi, \Phi \vdash_m^{(h, n)} s \rightsquigarrow s' \dashv_m \eta', \sigma', \Sigma', \Psi', \Phi'$		
1287		$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, q \cdot n)} \text{group } q \ s \rightsquigarrow \text{group } q \ s'; \dashv_m \eta', \sigma', \Sigma', \Psi', \Phi'$	S-Group
1288			
1289	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{group } q \ \text{skip} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi$		S-Group-Done
1290			

A.3.3 *Thread Synchronization.* We define a **size** operation on privileges to compute the size of a partition (the number of individual threads contained within it). The operation is defined as follows:

1293			
1294	$\text{size(Thread}, n) = n$		
1295		$\text{size(Block}, n) = n \cdot T$	
1296			$\text{size(Grid}, n) = n \cdot B \cdot T$
1297			
1298			
1299			
1300	$\Psi(\psi)(p) = 0$		
1301	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{wait}_\psi \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi$		S-Sync-Wait-Done
1302			
1303	$\Psi(\psi)(p) \neq 0$		
1304	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{wait}_\psi \rightsquigarrow \text{wait}_\psi \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi$		S-Sync-Wait-Spin
1305			
1306	$\Psi' = \Psi(\psi)[p \mapsto \Psi(\psi)(p) - 1]$		
1307	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{dec}_\psi \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi', \Phi$		S-Sync-Dec
1308			
1309	$\Psi(\psi)(p) = 0 \quad \Psi' = \Psi(\psi)[p \mapsto \text{size}(\pi)]$		
1310	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{init}_\psi \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi', \Phi$		S-Sync-Init-Zero
1311			
1312	$\Psi(\psi)(p) \neq 0$		
1313	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{init}_\psi \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi$		S-Sync-Init-Nonzero
1314			
A.3.4	<i>Asynchrony.</i>		
1315	$\text{rename}(\eta, \sigma, \Sigma, x, y, (\text{Thread}, 1)), t, b, p, \Psi, \Phi \vdash_m^{(\text{Thread}, 1)} s \rightsquigarrow s' \dashv_m \eta', \sigma', \Sigma', \Psi', \Phi'$		
1316		$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(\text{Thread}, 1)} \text{async_partition}_\phi x \text{ into } y \text{ in } s \rightsquigarrow$	
1317		$\text{async_partition}_\phi x \text{ into } y \text{ in } s' \dashv_m \eta', \sigma', \Sigma', \Psi', \Phi'$	
1318			
1319			
1320	$\Phi = \Phi'[\phi \mapsto \Phi'(\phi) \cup \{s\}]$		
1321	$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(\text{Thread}, 1)} \text{async_partition}_\phi x \text{ into } y \text{ in skip} \rightsquigarrow$		S-Async-Partition-Unwind
1322	$(\text{async_partition}_\phi x \text{ into } y \text{ in } s) \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi'$		
1323			

1324		
1325	$\Phi(\phi) = \emptyset$	
1326	$\frac{}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(\text{Thread}, 1)} \text{async_partition}_\phi x \text{ into } y \text{ in } \text{skip} \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi}$	S-Async-Partition-Done
1327		
1328		
1329		
1330	$\frac{}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(\text{Thread}, 1)} \text{async_memcpy}(x, y) \rightsquigarrow \text{skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi[\phi \mapsto \Phi(\phi) \cup \{\text{memcpy}(x, y)\}]}$	S-Async-Memcpy
1331		
1332		
1333	$\frac{(\eta', \sigma', \Sigma') = \text{update}(\eta, \sigma, \Sigma, x, \text{get}(\eta, \sigma, \Sigma, y))}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{memcpy}(x, y) \rightsquigarrow \text{skip} \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi}$	S-Memcpy
1334		
1335		
1336	A.3.5 <i>Variables and Memory.</i>	
1337	$\frac{\eta, \sigma, \Sigma \vdash_{\pi'}^\pi e \Downarrow v}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi x : \tau @ \pi' := e \text{ in } s \rightsquigarrow s \dashv_m \eta[x \mapsto^{\pi'} v], \sigma, \Sigma, \Psi, \Phi}$	S-Decl
1338		
1339		
1340	$\frac{}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{free } n \dashv_{m-n} \eta, \sigma, \Sigma, \Psi, \Phi}$	S-Free
1341		
1342		
1343	$\frac{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi x := \text{alloc Local } \tau n \text{ in } s \rightsquigarrow s; \text{ free } (n \cdot \text{size}(\tau)) \dashv_{m+n \cdot \text{size}(\tau)} \eta[x \mapsto^\pi \langle x, n \rangle], \sigma, \Sigma, \Psi, \Phi}{}$	S-Alloc-Local
1344		
1345		
1346	$\frac{\pi = (\text{Block}, 1)}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi x := \text{alloc Shared } \tau n \text{ in } s \rightsquigarrow s; \text{ free } (n \cdot \text{size}(\tau)) \dashv_{m+n \cdot \text{size}(\tau)} \eta[x \mapsto^\pi \langle x, n \rangle], \sigma, \Sigma, \Psi, \Phi}$	S-Alloc-Shared
1347		
1348		
1349		
1350	$\frac{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi x := \text{alloc Global } \tau n \text{ in } s \rightsquigarrow s; \text{ free } (n \cdot \text{size}(\tau)) \dashv_{m+n \cdot \text{size}(\tau)} \eta, \sigma, \Sigma[x \mapsto^\pi \langle x, n \rangle], \Psi, \Phi}{}$	S-Alloc-Global
1351		
1352		
1353	$\frac{x \in^{\pi'} \eta, \sigma, \Sigma \quad \eta, \sigma, \Sigma \vdash_{\pi'}^\pi e \Downarrow v \quad (\eta', \sigma', \Sigma') = \text{update}(\eta, \sigma, \Sigma, x, v) \quad \pi' \mid \pi}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi, \Phi \vdash_m^\pi x = e \rightsquigarrow \text{skip} \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi}$	S-Assn
1354		
1355		
1356		
1357	$\frac{\eta, \sigma, \Sigma, t, b, p \vdash_\pi^\pi e_1 \Downarrow \langle l, n \rangle \qquad i < n \quad \pi' \mid \pi \\ \eta, \sigma, \Sigma, t, b, p \vdash_\pi^\pi e_2 \Downarrow i \qquad x \in^{\pi'} \eta, \sigma, \Sigma \\ \eta, \sigma, \Sigma \vdash_{\pi'}^\pi e_3 \Downarrow v \qquad (\eta', \sigma', \Sigma') = \text{update}(\eta, \sigma, \Sigma, l+i, v)}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi e_1[e_2] = e_3 \rightsquigarrow \text{skip} \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi}$	S-Arr-Assn
1358		
1359		
1360		
1361		
1362	$\frac{s' = s[(y + c \cdot p)/y] \quad (\eta', \sigma', \Sigma') = \text{rename}(\eta, \sigma, \Sigma, x, y, \pi/c)}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{partition}_\psi x y c s \rightsquigarrow \text{init}_\psi; s'; \text{dec}_\psi; \text{wait}_\psi \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi}$	S-Partition
1363		
1364		
1365	$\frac{(\eta', \sigma', \Sigma') = \text{rename}(\eta, \sigma, \Sigma, x, y, (h, n_1))}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{(h, n_1+n_2)} \text{claim}_\psi x \text{ into } y \text{ at } n_1 \text{ in } s \rightsquigarrow \text{init}_\psi; \text{split}(n_1, n_2) \{s'\} \{\text{skip}\}; \text{dec}_\psi; \text{wait}_\psi \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi}$	S-Claim
1366		
1367		
1368		
1369		
1370	$\frac{(\eta', \sigma', \Sigma') = \text{rename}(\eta, \sigma, \Sigma, x, y, \downarrow \pi)}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{lower}_\psi x \text{ into } y \text{ in } s \rightsquigarrow \text{init}_\psi; s; \text{dec}_\psi; \text{wait}_\psi \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi}$	S-Lower
1371		
1372		

A.3.6 Control Flow.

$$\begin{array}{c}
 \frac{\eta, \sigma, \Sigma \vdash_{\pi}^{\pi} e \Downarrow \text{true}}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{if } e \text{ then } s_1 \text{ else } s_2 \rightsquigarrow s_1 \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-If-True} \\
 \\
 \frac{\eta, \sigma, \Sigma \vdash_{\pi}^{\pi} e \Downarrow \text{false}}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{if } e \text{ then } s_1 \text{ else } s_2 \rightsquigarrow s_2 \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-If-False} \\
 \\
 \frac{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{while } e \text{ do } s \rightsquigarrow \text{if } e \text{ then } (s; \text{while } e \text{ do } s) \text{ else skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} s_1 \rightsquigarrow s'_1; s_2 \dashv_{m'}^{\pi'} \eta', \sigma', \Sigma', \Psi', \Phi'} \text{ S-While} \\
 \\
 \frac{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} s_1 \rightsquigarrow s'_1 \dashv_{m'}^{\pi'} \eta', \sigma', \Sigma', \Psi', \Phi'}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} s_1; s_2 \rightsquigarrow s'_1; s_2 \dashv_{m'}^{\pi'} \eta', \sigma', \Sigma', \Psi', \Phi'} \text{ S-Seq-First} \\
 \\
 \frac{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{skip}; s_2 \rightsquigarrow s_2 \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} \text{skip}; s_2 \rightsquigarrow s_2 \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi} \text{ S-Seq-Done} \\
 \\
 \frac{\Sigma(f) = \{[x_1 : \tau_1, \dots, x_n : \tau_n], s\} \quad \sigma, \Sigma \vdash_{\pi}^{\pi} e_i \Downarrow v_i \quad m' \leq m}{\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^{\pi} f(e_1, \dots, e_n) \rightsquigarrow s \dashv_m \eta[x_i \mapsto v_i], \sigma, \Sigma, \Psi, \Phi} \text{ S-Function-Call}
 \end{array}$$

A.3.7 Expressions.

$$\begin{array}{c}
 \frac{\pi < (\text{Grid}, 1)}{\eta, \sigma, \Sigma \vdash_{\pi}^{\pi}, \text{partition_id} \Downarrow \pi / \pi' - 1} \text{ E-Partition-Id} \\
 \\
 \frac{}{\eta, \sigma, \Sigma \vdash_{\pi}^{\pi}, x \Downarrow \text{get}(\eta, \sigma, \Sigma, x)} \text{ E-Var} \\
 \\
 \frac{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, e_1 \Downarrow \langle l, n \rangle \quad \eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, e_2 \Downarrow i \quad i < n \quad \pi' \leq \pi}{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, e_1[e_2] \Downarrow \text{get}(\eta, \sigma, \Sigma, l + i)} \text{ E-Arr-Access} \\
 \\
 \frac{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, \text{int} \Downarrow n}{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, \text{int} \Downarrow n} \text{ E-Int} \quad \frac{}{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, b \Downarrow b} \text{ E-Bool} \\
 \\
 \frac{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, e_1 \Downarrow v_1 \quad \eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, e_2 \Downarrow v_2 \quad v = v_1 \text{ bop } v_2}{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, e_1 \text{ bop } e_2 \Downarrow v} \text{ E-Bop} \\
 \\
 \frac{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, e_1 \Downarrow v_1 \quad \eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, e_2 \Downarrow v_2 \quad v = v_1 \text{ cmp } v_2}{\eta, \sigma, \Sigma \vdash_{\pi'}^{\pi}, e_1 \text{ cmp } e_2 \Downarrow v} \text{ E-Cmp}
 \end{array}$$

A.4 Theorems and Proofs

Note that in this section we assume no out of bounds array accesses. In general Snap (and by extension Bundl) makes no guarantees about array out of bounds.

A.4.1 *More definitions.* As a premise to our type safety theorems, we need to assume we have a well-typed environment, written $\Gamma \vdash \eta, \sigma, \Sigma$. We define what this means inductively

$$\begin{array}{c}
 \frac{}{\eta, \sigma, \Sigma \vdash n : \text{int}} \text{ V-Int} \quad \frac{}{\eta, \sigma, \Sigma \vdash b : \text{bool}} \text{ V-Bool} \quad \frac{}{\eta, \sigma, \Sigma \vdash f : \text{float}} \text{ V-Float} \\
 \\
 \frac{\forall i < n, \eta, \sigma, \Sigma \vdash \text{get}(\eta, \sigma, \Sigma, x + i) : \tau}{\eta, \sigma, \Sigma \vdash \langle x, n \rangle : \tau[]^l} \text{ V-Array} \quad \frac{\Gamma, x_i : \pi \tau_i \vdash_m^{\pi} s \quad \Gamma \vdash \cdot, \cdot, \text{fns} \Sigma}{\eta, \sigma, \Sigma \vdash \{x_i : \tau_i, s\} : \text{Fun}(x_i : \tau_i, \pi, m)} \text{ V-Function} \\
 \\
 \frac{}{\cdot \vdash \eta, \sigma, \Sigma} \text{ G-Empty} \quad \frac{\eta(x) =^{\pi} v \quad \eta, \sigma, \Sigma \vdash v : \text{int}}{\Gamma, x : \pi \text{ int} \vdash \eta, \sigma, \Sigma} \text{ G-Int}
 \end{array}$$

$$\begin{array}{c}
 \frac{\eta(x) =^\pi v \quad \eta, \sigma, \Sigma \vdash v : \text{bool}}{\Gamma, x :^\pi \text{bool} \vdash \eta, \sigma, \Sigma} \text{ G-Bool} \quad \frac{\eta(x) =^\pi v \quad \eta, \sigma, \Sigma \vdash v : \text{float}}{\Gamma, x :^\pi \text{float} \vdash \eta, \sigma, \Sigma} \text{ G-Float} \\
 \\
 \frac{\eta(x) =^\pi v \quad \eta, \sigma, \Sigma \vdash v : \tau[]}{\Gamma, x :^\pi \tau[]^{\text{Local}} \vdash \eta, \sigma, \Sigma} \text{ G-Local} \quad \frac{\sigma(x) =^\pi v \quad \eta, \sigma, \Sigma \vdash v : \tau[]}{\Gamma, x :^\pi \tau[]^{\text{Shared}} \vdash \eta, \sigma, \Sigma} \text{ G-Shared} \\
 \\
 \frac{\Sigma(x) =^\pi v \quad \eta, \sigma, \Sigma \vdash v : \tau[]}{\Gamma, x :^\pi \tau[]^{\text{Global}} \vdash \eta, \sigma, \Sigma} \text{ G-Global} \quad \frac{\Sigma(x) =^\pi v \quad \eta, \sigma, \Sigma \vdash v : \mathbf{Fun}(\Gamma', \pi, m)}{\Gamma, x :^\pi \mathbf{Fun}(\Gamma', \pi, m) \vdash \eta, \sigma, \Sigma} \text{ G-Function}
 \end{array}$$

We can prove a couple simple lemmas about well-typed environments under operations like **rename**, **update**, and **get**.

LEMMA A.1. (*Well-typed get*) If $\Gamma \vdash \eta, \sigma, \Sigma$ and $x :^\pi \tau \in \Gamma$ then $\eta, \sigma, \Sigma \vdash \mathbf{get}(\eta, \sigma, \Sigma, x) : \tau$.

LEMMA A.2. (*Well-typed rename*) If $\Gamma, x :^\pi \tau \vdash \eta, \sigma, \Sigma$ then $\Gamma, x :^\pi \tau, y :^\pi \tau' \vdash \mathbf{rename}(\eta, \sigma, \Sigma, x, y, \tau')$.

LEMMA A.3. (*Well-typed update*) If $\Gamma \vdash \eta, \sigma, \Sigma$ and $\eta, \sigma, \Sigma \vdash v : \tau$ then $\Gamma, x :^\pi \tau \vdash \mathbf{update}(\eta, \sigma, \Sigma, x, v)$.

We also define a well-formedness precondition on p with respect to π :

$$(h, n) \vdash p ::= p < n$$

We also define well-formedness for the async stack:

$$\Gamma \vdash \Phi ::= \forall \phi, s \in \Phi(\phi), \Gamma \vdash_m^{(\text{Thread}, 1)} s$$

A.5 Proofs

LEMMA A.4. (*Expression Safety*) If $\Gamma \vdash^\pi e : \tau$ and $\Gamma \vdash \eta, \sigma, \Sigma$ and $\pi \vdash p$, then there is some v such that $\eta, \sigma, \Sigma \vdash_{\pi'}^\pi e \Downarrow v$ and $v : \tau$.

PROOF. This proof proceeds by induction on the typing relation for expressions. Despite the fact that this property implies termination, we do not need a logical relation to prove it because the expression language is very simple.

The **T-Int**, **T-Float**, and **T-Bool** cases are trivial, using the rules **E-Int**, **E-Bool** and **E-Float** to compute values. In the case for **T-Partition-Id**, the π premises of the typing rules match the premises of the evaluation rules, so these rules are simple as well.

The cases for **T-Bop** and **T-Cmp** follow directly from the inductive hypotheses, assuming a valid and correctly implemented set of binary operators and comparators.

The only interesting cases are **T-Arr-Access** and **T-Arr-Access-Shared**.

In both cases our inductive hypotheses and inversion give us that $\pi' \leq \pi$, and e_1 evaluates to a $\langle x, n \rangle$, and that all the values between x and $x + n$ in the appropriate environment are typed at τ . We also know that e_2 evaluates to an integer i . We assume that all array accesses are in bounds, so $i < l$, which is sufficient to use the **E-Arr-Access** rule to complete this case, and the proof. \square

LEMMA A.5. (*Expression Determinism*) If $\eta, \sigma, \Sigma, t, b, p \vdash_{\pi'}^\pi e \Downarrow v_1$ and $\eta, \sigma, \Sigma, t, b, p \vdash_{\pi'}^\pi e \Downarrow v_2$ then $v_1 = v_2$.

PROOF. Straightforward by induction on the semantic derivation. \square

LEMMA A.6. (*Expression Well-Typedness*) If $\Gamma \vdash_{\pi'}^\pi e : \tau$ and $\Gamma \vdash \eta, \sigma, \Sigma$ and $\pi \vdash p$, and $\eta, \sigma, \Sigma, t, b, p \vdash_{\pi'}^\pi e \Downarrow v$, then $v : \tau$.

PROOF. By our expression safety lemma our well-typed expression must evaluate to a well-typed value v' . By our determinism lemma v' must be the same as v , so v is well-typed. \square

LEMMA A.7. (*Statement Progress*) If $\Gamma \vdash \eta, \sigma, \Sigma$ and $\Gamma \vdash_m^\pi s$ and $\pi \vdash p$, then either s is skip or there is some s' such that $\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi s \rightsquigarrow s' \vdash_m^{\pi'} \eta', \sigma', \Sigma, \Psi', \Phi'$

PROOF. This proceeds by induction on the typing derivation.

Privilege Management Rules.

- Case **T-Split**:

In this case we have by our assumption that $\pi \vdash p$ that $p < n$. We also have that n_1, n_2 align to n , so $n_1 + n_2 \leq n$. There are three cases to consider, then: when $p < n_1$, when $p \geq n_1$ and $p < n_1 + n_2$, and when $p \geq n_1 + n_2$.

In the first case, we have by our inductive hypothesis that s is either skip or that it can step in an (h, n_1) context. In the former case we can use the **S-Split-Left-Done** rule and in the latter we can use the **S-Split-Left rule**.

The second case is almost symmetric. The only additional work we have to do is to argue that $(h, n_2) \vdash p - n_1$, or equivalently that $p - n_1 < n_2$. This, however, is immediate from our assumption that $p < n_1 + n_2$.

In the last case, we just use the **S-Split-None** rule to step to skip.

- Case **T-Destruct**

By our inductive hypothesis, we know that s can step at $\downarrow \pi$ if $\downarrow \pi \vdash p$. The \downarrow operation is only defined at (Block, 1) or (Grid, 1), so we only need to consider the cases where π is one of those.

In the former case p becomes $t \mathbf{mod} T$ while $\downarrow \pi$ is (Thread, T). $t \mathbf{mod} T < T$ for any t so this satisfies the requirement that $\pi \vdash p$, which lets us use our inductive hypothesis: s is either skip or can step. If it can step, we can use this to satisfy the premise of **S-Destruct-Block** to step in this case. If it is skip, then we use the rule **S-Destruct-Done** to step instead.

The latter case is the same, except using the fact that $b \mathbf{mod} B < B$ and the **S-Destruct-Grid** rule.

- Case **T-Group**

In this case we have by assumption that $(h, q \cdot n) \vdash p$, i.e., that $p < q \cdot n$.

In this case we have our IH that if $(h, n) \vdash p'$ for some p' , then s is either skip or steps with (h, n) privilege with p' as our partition id.

We choose p' to be $p \mathbf{mod} n$. This is always $< n$, so $(h, n) \vdash p'$. This lets us use our IH to get that s is either skip (in which case we can use the **S-Group-Done** rule to step) or itself steps, which lets us use the **S-Group** rule to step.

Thread Synchronization Rules.

- Case **T-Sync-Wait**

$\Psi(\psi)(p)$ is either zero or it is not. In the former case we use the **S-Sync-Wait-Done** rule and in the latter we use **S-Sync-Wait-Spin**.

- Case **T-Sync-Dec**

We use the **S-Sync-Dec** rule to step.

- Case **T-Sync-Init**

We use the **S-Sync-Init-Zero** or **S-Sync-Init-Nonzero** rules depending on whether $\Psi(\psi)(p)$ is zero or not.

Asynchrony Rules.

- Case **T-Async-Partition**

In this case, we have via our IH that if $\Gamma, y : (\text{Thread}, 1) \text{ async } \tau[]^l \vdash \eta', \sigma', \Sigma'$, then we can either step s with $(\text{Thread}, 1)$ privilege under environments η' , σ' , and Σ' , or s is skip.

1520 We have by assumption that $\Gamma, x:(\text{Thread}, 1) \text{async } \tau[]^l \vdash \eta, \sigma, \Sigma$. By our environment renaming
 1521 lemma, this gives us what we need to use our IH with $(\eta', \sigma', \Sigma')$ as **rename** $(\eta, \sigma, \Sigma, x, y, (\text{Thread}, 1))$.
 1522 Thus s either steps or is skip. In the former case we can step with **S-Async-Partition-Congr**,
 1523 and in the latter we can use either **S-Async-Partition-Unwind** or **S-Async-Partition-Done**
 1524 depending on whether $\Phi(\phi)$ is empty or not.

- Case **T-Async-Memcpy**

1526 Immediate via use of the **S-Async-Memcpy** rule.

- Case **T-Memcpy**

1528 Immediate via use of the **S-Memcpy** rule.

1529 *Memory Rules.*

- Case **T-Decl**

1532 Via our lemma about expression type safety and our hypothesis that e is well-typed, we obtain
 1533 the premises necessary to use the **S-Decl** rule to step.

- Case **T-Arr-Assn**

1535 Each of e_1, e_2 and e_3 must evaluate to a well-typed value by the expression type safety lemma.
 1536 In particular, both e_1 evaluates to some $\langle l, n \rangle$ and e_2 evaluates to some i . We assume all array
 1537 accesses are in bounds, so this is sufficient to use the **S-Arr-Assn** rule to step.

- Case **T-Arr-Assn-Shared**

1539 Same as previous case.

- Case **T-Free**

1541 Trivial via the **S-Free** rule.

- Case **T-Partition**

1543 Trivial via the **S-Partition** rule.

- Case **T-Claim**

1545 Trivial via the **S-Claim** rule.

- Case **T-Lower**

1547 Trivial via the **S-Lower** rule.

- Case **T-Alloc**

1549 We assume that l is not Shared, so we can use the **S-Alloc-Local** or **S-Alloc-Global** rule,
 1550 depending on whether l is Local or Global.

- Case **T-Alloc-Shared**

1552 Trivial via the **S-Alloc-Shared** rule.

1553 *Control Rules.*

- Case **T-Skip**

1556 Trivial

- Case **T-While**

1558 Trivial, all while loops step via the **S-While** rule

- Case **T-If**

1560 By our proof of expression type safety, the expression e steps to either the boolean value true
 1561 or false. We can thus use either the **S-If-True** or **S-If-False** rules to step.

- Case **T-Seq**

1563 In this case we know by our IH that s_1 is either skip or can step. In the former case we use
 1564 the **S-Seq-Done** rule and in the latter we use the **S-Seq-First** rule.

- Case **T-Function-Call**

1566 In this case we know by our expression safety lemma that each of the arguments will evaluate to
 1567 a well-typed value. We also have by assumption that f has a function type, which by inversion

on the **V-Function** rule tells us that it is a closure type. Additionally our assumption that $\Gamma \vdash \eta, \sigma, \Sigma$ tell us that Σ contains f at the same type that Γ does. These premises are sufficient to use the **S-Function-Call** rule.

□

LEMMA A.8. (Statement Preservation) *If $\Gamma \vdash \eta, \sigma, \Sigma$ and $\Gamma \vdash_m^\pi s$ and $\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_{m'}^\pi s \rightsquigarrow s' \dashv_{m''} \eta', \sigma', \Sigma, \Psi', \Phi'$ and $\pi \vdash p$ and $m \geq m'$ and $\Gamma \vdash \Phi$, then there is some Γ' such that $\Gamma \subseteq \Gamma'$ and $\Gamma' \vdash_m^\pi s'$ and $\Gamma' \vdash \eta', \sigma', \Sigma'$ and $m \geq m''$ and $\Gamma' \vdash \Phi'$.*

PROOF. We proceed by induction on the derivation of $\Gamma \vdash_m^\pi s$.

Privilege Management Rules.

- Case **T-Split**

In this case we have by assumption that $(h, n) \vdash p$, $\Gamma \vdash \eta, \sigma, \Sigma$, and $\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_{m'}^\pi \text{split}(n_1, n_2)\{s_1\}\{s_2\} \rightsquigarrow s' \dashv_{m''} \eta', \sigma', \Sigma, \Psi', \Phi'$. Our inductive hypotheses give us that if for any p , if $(h, n_1) \vdash p$ and s_1 steps with partition id p , or if $(h, n_2) \vdash p$ and s_2 steps with partition id p then their results are well typed.

By inversion on our semantic derivation, we are in one of 5 cases.

In the **S-Split-Left** case $p < n_1$ and s_1 steps to s'_1 . This is sufficient to tell us that s'_1 is well-typed and the output environments of that relation η', σ' , and Σ' are all well-typed by $\Gamma' \supseteq \Gamma$, and that the memory is properly bounded by the typing rules.

We can thus use the **T-Split-Left** rule to conclude that the result of this case is well-typed.

The **S-Split-Left-Done** rule is trivial via the **T-Skip** rule.

The **Right** cases are symmetric, with the observation that when $p \geq n_1$ and $p < n_1 + n_2$ then $p - n_1 < n_2$.

The last **S-Split-None** rule is trivial via the **T-Skip** rule.

- Case **T-Destruct**

In this case we have by assumption that $\downarrow \pi$ is defined, so π is either **(Block,1)** or **(Grid,1)**. We also assume that $\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_{m'}^\pi \text{destruct}$ in $s \rightsquigarrow s' \dashv_{m''} \eta', \sigma', \Sigma, \Psi', \Phi'$. We also have by our inductive hypothesis that for any p such that $\downarrow \pi \vdash p$ and s'' such that s steps to s'' at p , then that step preserved well-typedness.

By inversion on the step relation, we are in one of three cases.

If the rule used **S-Destruct-Block**, then we know that s steps to s'' and p is $t \bmod T$. $(\text{Thread}, T) \vdash t \bmod T$ for any t , so we can use our inductive hypothesis to conclude that the s'' stepped to by s is well-typed, as are its environments and memory usage. The **T-Destruct** rule then gives us our desired goal.

The **S-Destruct-Grid** case proceeds similarly, while the **S-Destruct-Done** case is trivial.

- Case **T-Group**

In this case we have by assumption that

$\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_{m'}^{(h, q \cdot n)} \text{group } q \ s \rightsquigarrow s' \dashv_{m''} \eta', \sigma', \Sigma, \Psi', \Phi'$. We have by our inductive hypothesis that s steps to s'' at some partition id p' and $(h, n) \vdash p'$, then s'' is well typed, as are the other outputs of that step.

By inversion on the step relation, we are in one of two cases. The **S-Group-Done** case is trivial, so we shall focus on the **S-Group** case. In this case we have that s steps to s'' at partition id $p \bmod n$. It is always the case that $(h, n) \vdash p \bmod n$ for any p , so we can use our inductive hypothesis to conclude that s'' is well-typed, as are its output environments and memory usage. From there, it is a simple application of the **T-Group** rule to conclude that $\text{group } q \ s'$ is well-typed, and to finish the case.

1618 **Thread Synchronization Rules.** These rules are all trivial: with one exception all thread synchronization
 1619 primitives step to skip without changing environment or memory, and are thus obviously well-typed.
 1620 **S-Sync-Wait-Spin** does not produce skip, but it steps to the same statement as we already assumed
 1621 typechecks in the premise of the lemma, so is straightforward nonetheless.

1622 If we wanted to say something about deadlock freedom we'd have more work here, but we aren't
 1623 doing that, so these rules are easy.

1624 *Asynchrony Rules.*

- Case **T-Async-Partition** In this case we have that s is well-typed in a context where x has been renamed into y , with $(\text{Thread},1)$ privilege. We also have that $\Gamma, x:(^{\text{Thread},1}) \tau[]^l \vdash \eta, \sigma, \Sigma$ and $\Gamma, x:(^{\text{Thread},1}) \tau[]^l \Phi$.

1629 By inversion, we are in one of three cases.

1630 In the **S-Async-Partition-Done** case, we are done.

1631 In the **S-Async-Partition-Congr** case, our inductive hypothesis gives us that there is some
 1632 Γ' such that $\Gamma, y:(^{\text{Thread},1}) \text{async } \tau[]^l \subseteq \Gamma'$ and $\Gamma' \vdash \Phi'$ and $\Gamma' \vdash \text{rename}(\eta, \sigma, \Sigma, x, y, (\text{Thread},1))$
 1633 via our well-typed renaming lemma. This lets us use the **T-Async-Partition** rule to check
 1634 this case, with a choice of Γ' as $\Gamma, x:(^{\text{Thread},1}) \tau[]^l$.

1635 In the **S-Async-Partition-Unwind** case, our assumption that Φ is well-typed tells us that
 1636 $\Gamma, y:(^{\text{Thread},1}) \vdash \{ (\text{Thread},1)_m s$. Thus, we can use the **T-Async-Partition** rule to type this case.

- Case **T-Async-Memcpy**

1638 In this case the statement and environment typing are trivial, we need only to show that the
 1639 async stack remains well typed.

1640 In this case we have that x and y have the same type at $(\text{Thread},1)$. This is sufficient for us
 1641 to check $x=y$ at $(\text{Thread},1)$, meaning that adding that instruction to the stack maintains its
 1642 well-typedness.

- Case **T-Memcpy**

1644 Immediate via use of the **S-Memcpy** rule. We just need to show that the environment remains
 1645 well typed, which we know via our lemmas about **update** and **get**.

1646 *Memory Rules.*

- Case **T-Decl**

1649 By inversion, we are using **S-Decl** rule for evaluation. Our well-typed expression lemma gives
 1650 us that v is well-typed, so it follows from our assumptions and our lemmas about extending
 1651 environments that the extended η and s are well-typed by $\Gamma, x:\pi \tau$.

- Case **T-Free**

1653 Trivial.

- Case **T-Alloc**

1655 By inversion we are either in the **S-Alloc-Local** or **S-Alloc-Global** rules. In either case, we
 1656 assume that $\Gamma, x:\pi \tau[]^l$ checks s , meaning we can use our extended environment lemmas and
 1657 the **T-Seq** and **T-Free** rules to check these cases.

- Case **T-Alloc-Shared**

1659 Same as previous case.

- Case **T-Partition**

1661 In this case we have by assumption that $\Gamma, y:(^{(h,n/c)} \tau[]^l s$ and l is not local and c divides n . By
 1662 inversion on our step relation we must be in the **S-Partition** case, so we have $\eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi$
 1663 $\text{partition}/xycs \sim \text{init}_\psi; s'; \text{dec}_\psi; \text{wait}_\psi \dashv_m \eta', \sigma', \Sigma', \Psi, \Phi$ where $s' = s[(y + c \cdot p)/y]$ and
 1664 $(\eta', \sigma', \Sigma') = \text{rename}(\eta, \sigma, \Sigma, x, y, \pi/c)$. Via our well-typed renaming lemma contexts we know
 1665 that we can check the renamed environments in the extended environment $\Gamma, y:\pi/c \tau[]^l, x:\pi$

1667 $\tau[]^l$, and this context is also sufficient to check s' (via a substitution-preserves-typing lemma
 1668 that is obvious). Using the **T-Skip** rule this is exactly what we need to show to complete this
 1669 case, as the thread sync primitives check trivially via their typing rules.

- 1670 • Case **T-Claim**

1671 Essentially the same as **T-Partition**.

- 1672 • Case **T-Lower**

1673 Essentially the same as **T-Partition**.

1674 *Control Rules.*

- 1675 • Case **T-If**

1676 We have by assumption that **if** e then s_1 else s_2 steps to some s' , and by inversion we
 1677 know that either e evaluates to **true** and s' is s_1 , or e evaluates to **false** and s' is s_2 .

1678 In either case, our inductive hypotheses is sufficient to tell us that these are well-typed. In
 1679 particular, in both cases our IHs tell us that the amount of memory used by stepping each
 1680 branch of the **if** is less than the amount of memory computed by the type system for each
 1681 branch. Because the whole **if** expression checks using the greater of the memory usage of m_1
 1682 or m_2 (i.e., the memory usage on each branch), the resulting usage for the whole conditional
 1683 is also bounded by the type system.

- 1684 • Case **T-Skip**

1685 Trivial: **skip** does not step

- 1686 • Case **T-Seq**

1687 In this case we have that s_1 and s_2 are both well-typed (with m_1 memory and m_2 memory
 1688 respectively), and $m = \mathbf{max}(m_1, m_2)$. We also have by inversion that s_1 either steps to **skip** or
 1689 s'_1 , and our inductive hypothesis tells us that s'_1 is well-typed.

1690 In the former case we can use the **S-Seq-Done** rule to trivially finish the case. In the latter,
 1691 our IH allows us to finish the case, since m_1 is always $\leq \mathbf{max}(m_1, m_2)$

- 1692 • Case **T-While**

1693 By inversion, we have that

$$\begin{aligned} & \eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi \text{while } e \text{ do } s \\ & \rightsquigarrow \text{if } e \text{ then } (s; \text{while } e \text{ do } s) \text{ else skip} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi. \end{aligned}$$

1694 We also have by assumption that e and s are well-typed. With this information, through use
 1695 of the **T-If**, **T-Seq**, **T-While**, and **T-Skip** rules, we can conclude that the result of this rule is
 1696 also well-typed.

- 1697 • Case **T-Function-Call**

1698 By inversion we have that

$$\begin{aligned} & \eta, \sigma, \Sigma, t, b, p, \Psi, \Phi \vdash_m^\pi f(e_1, \dots, e_n) \\ & \rightsquigarrow \text{call } s \text{ with } (x_i : \pi_i \tau_i \mapsto v_i) @ \{\eta, \sigma, \Sigma, m'\} \dashv_m \eta, \sigma, \Sigma, \Psi, \Phi, \end{aligned}$$

1699 and also that $\Sigma(f) = \{[x_1 : \tau_1, \dots, x_n : \tau_n], s\}$, and $\eta, \sigma, \Sigma, t, b, p \vdash^\pi e_i \Downarrow v_i$, and $m' \leq m$.

1700 We also have from the premises of our case that $f : \pi \mathbf{Fun}(x_1 : \pi_1 \tau_1, \dots, x_n : \pi_n \tau_n, \pi, m') \in \Gamma$, and
 1701 $\Gamma \vdash^\pi e_i : \tau_i$, and $m' \leq m$.

1702 Our lemma for expression well-typedness tells us that that each v_i is a well-typed value,
 1703 and our assumption that $\Gamma \vdash \eta, \sigma, \Sigma$ tells us that $\Sigma(f)$ is a well-typed function and thus that
 1704 $\mathbf{fns} \Gamma, x_i : \pi_i \tau_i \vdash_m^\pi s$. We can take the union of this with Γ to produce $\Gamma, x_i : \pi_i \tau_i \vdash_m^\pi s$, which
 1705 clearly checks s and the output environments.

□